



THE INSIDER'S GUIDE TO THE PHILIPS ARM7-BASED MICROCONTROLLERS

An Engineer's Introduction To The LPC2100 Series

Trevor Martin BSc. (hons.) CEng. MIEE

www.hitex.co.uk/arm

Published by Hitex (UK) Ltd.

First Published February 2005

Hitex (UK) Ltd.

Sir William Lyons Road
University Of Warwick Science Park
Coventry, CV4 7EZ

Credits

Author: Trevor Martin
Illustrator: Sarah Latchford
Editors: Alison Wenlock & Michael Beach
Cover: Michael Beach

Acknowledgements

The author would like to thank Kees van Seventer and Chris Davies of Philips Semiconductors for their assistance in compiling this book

© Hitex (UK) Ltd., 11/03/2005

All rights reserved. No part of this publication may be reproduced, stored in a retrieval system or transmitted in any form or by any means, electronic, mechanical or photocopying, recording or otherwise without the prior written permission of the Publisher.

Introduction	8
Chapter 1: The ARM7 CPU Core	10
Outline	10
The Pipeline	10
Registers	11
Current Program Status Register.....	12
Exception Modes	13
ARM 7 Instruction Set.....	16
Branching	18
Data Processing Instructions	19
Copying Registers	20
Copying Multiple Registers	20
Swap Instruction	21
Modifying The Status Registers.....	21
Software Interrupt	22
MAC Unit.....	23
THUMB Instruction Set	24
Summary	26
Chapter 2: Software Development	27
Outline	27
Which Compiler?	27
uVISION IDE.....	28
Tutorial	28
Startup Code.....	29
Interworking ARM/THUMB Code	31
STDIO Libraries	32
Accessing Peripherals	33
Interrupt Service Routines	34
Exception source Constants table C function prototype	34
Software Interrupt	36
Locating Code In RAM.....	36
Operating System Support	38
Fixing Objects At Absolute Locations	38
Inline Assembler	38
Hardware Debugging Tools.....	39

Important!.....	40
Even More Important	40
Summary	40
Chapter 3: System Peripherals	42
Outline.....	42
Bus Structure	42
Memory Map	44
Register Programming.....	45
Memory Accelerator Module	45
Example MAM Configuration	49
FLASH Memory Programming.....	50
Memory Map Control.....	50
Bootloader	51
Philips ISP Utility.....	52
In-Application Programming	53
External Bus Interface	54
External Memory Interface	54
Using The External Bus Interface	57
Booting From ROM	58
Phase Locked Loop	60
VLSI Peripheral Bus Divider	62
Example Code: PLL And VPB Configuration	63
Power Control.....	64
LPC2000 Interrupt System	66
Pin Connect Block.....	66
External Interrupt Pins	66
Interrupt Structure	67
FIQ interrupt.....	68
Leaving An FIQ Interrupt	68
Example Program: FIQ Interrupt	69
Vectored IRQ	70
Leaving An IRQ Interrupt.....	71
Leaving An IRQ Interrupt.....	72
Example Program: IRQ interrupt	72
Non-Vectored Interrupts	73
Leaving A Non-Vectored IRQ Interrupt	73
Example Program: Non-Vectored Interrupt.....	73
Nested Interrupts	75
Summary	76
Chapter 4: User Peripherals.....	78
Outline.....	78

General Purpose I/O	78
General Purpose Timers	79
PWM Modulator	82
Real Time Clock	85
Watchdog	88
UART.....	90
I2C Interface.....	94
SPI Interface	99
Analog To Digital Converter.....	101
Digital To Analog Converter.....	104
CAN Controller	105
ISO 7 Layer Model.....	105
CAN Node Design	106
CAN Message Objects	107
CAN Bus Arbitration	109
Bit Timing	110
CAN Message Transmission	112
CAN Error Containment	115
CAN Message Reception	118
Acceptance Filtering	119
Configuring The Acceptance Filter.....	120
Summary	122
Chapter 5: Keil Tutorial	124
Installation	124
Using the Keil UVISION IDE	125
Exercise 1: Using the Keil Toolset.....	126
Using The Debugger.....	133
Using The ULINK Hardware Debugger	138
Setting up the ULINK JTAG hardware debugger:.....	138
Exercise 2: Startup code.....	141
Exercise 3: Using THUMB code	142
Exercise 4: Using STDIO libraries.....	144
Exercise 5: Simple interrupt	146
Exercise 6: Software Interrupt.....	148
Exercise 7: Memory Accelerator Module	149
Exercise 8: In-Application programming.....	152
Exercise 9: External Bus Interface	153
Exercise 10 Phase Locked Loop	157
Exercise 11: Fast Interrupt	159
Exercise 12: Vectored Interrupt	160
Exercise 13 : Non Vectored Interrupt	162
Exercise 14: Nested Interrupts	163

Exercise 15: General purpose IO pins.....	164
Exercise 16: Timer Capture.....	165
Exercise 17: Timer Match	167
Exercise 18: Dual-Edge (Symmetrical) PWM Generation.	170
Exercise 19: Real Time Clock.....	172
Exercise 20: USART	173
Exercise 21: I2C interface	174
Exercise 22: SPI	175
Exercise 22: SPI	176
Exercise 23: Analog To Digital Converter.....	177
Exercise 2 : Digital to Analogue Converter	178
Exercise 25: Transmitting CAN Data	179
Exercise 26: Receiving CAN Data.....	180
Chapter 6: Tutorial With GNU Tools	182
Introdution.....	182
GCC Startup Code	182
Interworking ARM/THUMB Code	182
Accessing Peripherals	182
Interrupt Service Routines	182
Software Interrupt	183
Inline Functions.....	183
Exercise 1: Using The Keil Toolset With The GNU Compiler	184
Exercise 2: Startup Code	189
Exercise 3: Using THUMB Code.....	190
Exercise 4: Using The GNU Libraries	192
Exercise 5: Simple Interrupt.....	193
Exercise 6: Software Interrupt.....	195
Appendices	197
Appendix A	197
Bibliography	197
Webliography	197
Reference Sites	197
Tools and Software Development	197
Evaluation Boards And Modules	197

Introduction

This book is intended as a hands-on guide for anyone planning to use the Philips LPC2000 family of microcontrollers in a new design. It is laid out both as a reference book and as a tutorial. It is assumed that you have some experience in programming microcontrollers for embedded systems and are familiar with the C language. The bulk of technical information is spread over the first four chapters, which should be read in order if you are completely new to the LPC2000 and the ARM7 CPU.

The first chapter gives an introduction to the major features of the ARM7 CPU. Reading this chapter will give you enough understanding to be able to program any ARM7 device. If you want to develop your knowledge further, there are a number of excellent books which describe this architecture and some of these are listed in the bibliography. Chapter Two is a description of how to write C programs to run on an ARM7 processor and, as such, describes specific extensions to the ANSI C standard which are necessary for embedded programming. In this book a commercial compiler is used in the main text, however the GCC tools have also been ported to ARM.

Appendix A details the ARM-specific features of the GCC tools. Having read the first two chapters you should understand the processor and its development tools. Chapter Three then introduces the LPC2000 system peripherals. This chapter describes the system architecture of the LPC2000 family and how to set the chip up for its best performance. In Chapter Four we look at the on-chip user peripherals and how to configure them for our application code.

Throughout these chapters various exercises are listed. Each of these exercises are described in detail in Chapter Five, the Tutorial section. The Tutorial contains a worksheet for each exercise which steps you through an important aspect of the LPC200. All of the exercises can be done with the evaluation compiler and simulator which come on the CD provided with this book. A low-cost starter kit is also available which allows you to download the example code on to some real hardware and “prove” that it does in fact work. It is hoped that by reading the book and doing the exercises you will quickly become familiar with the LPC2000.

Chapter 1: The ARM7 CPU Core

Outline

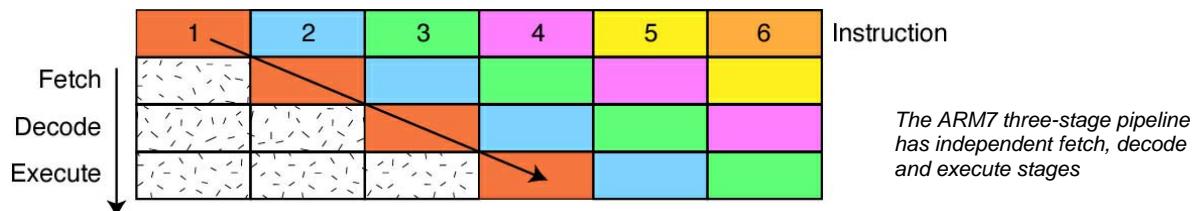
The CPU at the heart of the LPC2000 family is an ARM7. You do not need to be an expert in ARM7 programming to use the LPC2000, as many of the complexities are taken care of by the C compiler. You do need to have a basic understanding of how the CPU is working and its unique features in order to produce a reliable design.

In this chapter we will look at the key features of the ARM7 core along with its programmers' model and we will also discuss the instruction set used to program it. This is intended to give you a good feel for the CPU used in the LPC2000 family. For a more detailed discussion of the ARM processors, please refer to the books listed in the bibliography.

The key philosophy behind the ARM design is simplicity. The ARM7 is a RISC computer with a small instruction set and consequently a small gate count. This makes it ideal for embedded systems. It has high performance, low power consumption and it takes a small amount of the available silicon die area.

The Pipeline

At the heart of the ARM7 CPU is the instruction pipeline. The pipeline is used to process instructions taken from the program store. On the ARM 7 a three-stage pipeline is used.



A three-stage pipeline is the simplest form of pipeline and does not suffer from the kind of hazards such as read-before-write seen in pipelines with more stages. The pipeline has hardware independent stages that execute one instruction while decoding a second and fetching a third. The pipeline speeds up the throughput of CPU instructions so effectively that most ARM instructions can be executed in a single cycle. The pipeline works most efficiently on linear code. As soon as a branch is encountered, the pipeline is flushed and must be refilled before full execution speed can be resumed. As we shall see, the ARM instruction set has some interesting features which help smooth out small jumps in your code in order to get the best flow of code through the pipeline. As the pipeline is part of the CPU, the programmer does not have any exposure to it. However, it is important to remember that the PC is running eight bytes ahead of the current instruction being executed, so care must be taken when calculating offsets used in PC relative addressing.

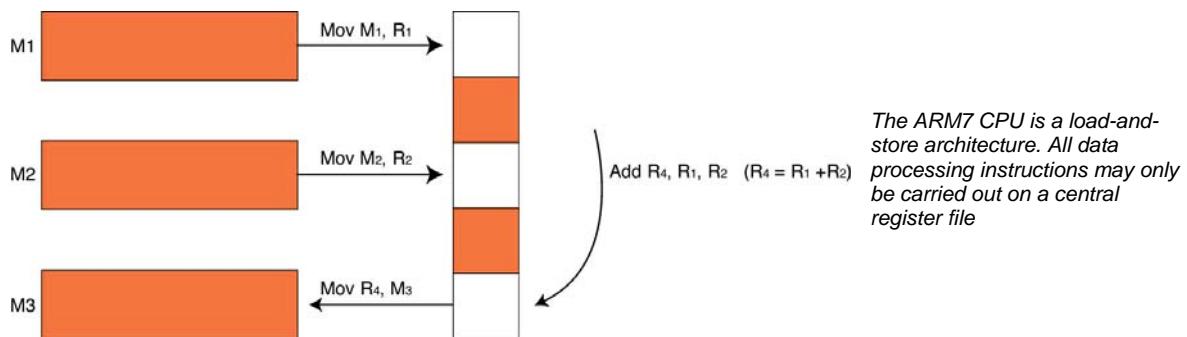
For example, the instruction:

```
0x4000 LDR PC, [PC, #4]
```

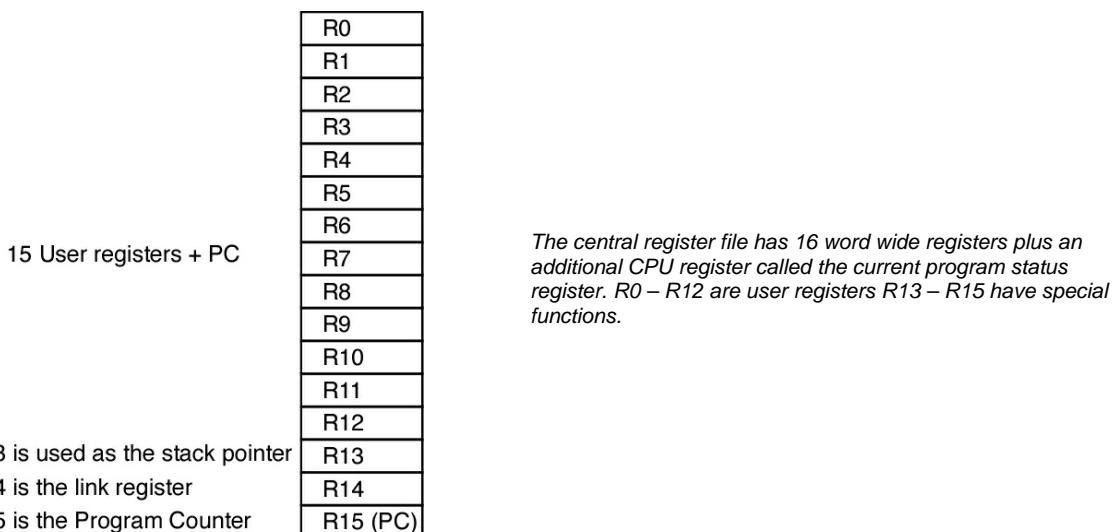
will load the contents of the address PC+4 into the PC. As the PC is running eight bytes ahead then the contents of address 0x400C will be loaded into the PC and not 0x4004 as you might expect on first inspection.

Registers

The ARM7 is a load-and-store architecture, so in order to perform any data processing instructions the data has first to be moved from the memory store into a central set of registers, the data processing instruction has to be executed and then the data is stored back into memory.

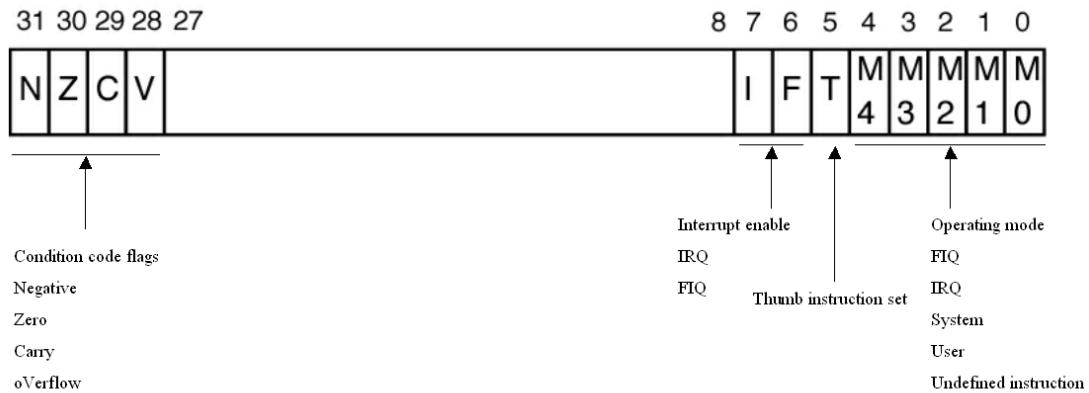


The central set of registers are a bank of 16 user registers R0 – R15. Each of these registers is 32 bits wide and R0 – R12 are user registers in that they do not have any specific other function. The Registers R13 – R15 do have special functions in the CPU. R13 is used as the stack pointer (SP). R14 is called the link register (LR). When a call is made to a function the return address is automatically stored in the link register and is immediately available on return from the function. This allows quick entry and return into a ‘leaf’ function (a function that is not going to call further functions). If the function is part of a branch (i.e. it is going to call other functions) then the link register must be preserved on the stack (R13). Finally R15 is the program counter (PC). Interestingly, many instructions can be performed on R13 – R15 as if they were standard user registers.



Current Program Status Register

In addition to the register bank there is an additional 32 bit wide register called the ‘current program status register’ (CPSR). The CPSR contains a number of flags which report and control the operation of the ARM7 CPU.



The Current Program Status Register contains condition code flags which indicate the result of data processing operations and User flags which set the operating mode and enable interrupts. The T bit is for reference only

The top four bits of the CPSR contain the condition codes which are set by the CPU. The condition codes report the result status of a data processing operation. From the condition codes you can tell if a data processing instruction generated a negative, zero, carry or overflow result. The lowest eight bits in the CPSR contain flags which may be set or cleared by the application code. Bits 7 and 8 are the I and F bits. These bits are used to enable and disable the two interrupt sources which are external to the ARM7 CPU. All of the LPC2000 peripherals are connected to these two interrupt lines as we shall see later. You should be careful when programming these two bits because in order to disable either interrupt source the bit must be set to ‘1’ not ‘0’ as you might expect. Bit 5 is the THUMB bit.

The ARM7 CPU is capable of executing two instruction sets; the ARM instruction set which is 32 bits wide and the THUMB instruction set which is 16 bits wide. Consequently the T bit reports which instruction set is being executed. Your code should not try to set or clear this bit to switch between instruction sets. We will see the correct entry mechanism a bit later. The last five bits are the mode bits. The ARM7 has seven different operating modes. Your application code will normally run in the user mode with access to the register bank R0 – R15 and the CPSR as already discussed. However in response to an exception such as an interrupt, memory error or software interrupt instruction the processor will change modes. When this happens the registers R0 – R12 and R15 remain the same but R13 (LR) and R14 (SP) are replaced by a new pair of registers unique to that mode. This means that each mode has its own stack and link register. In addition the fast interrupt mode (FIQ) has duplicate registers for R7 – R12. This means that you can make a fast entry into an FIQ interrupt without the need to preserve registers onto the stack.

Each of the modes except user mode has an additional register called the “saved program status register”. If your application is running in user mode when an exception occurs the mode will change and the current contents of the CPSR will be saved into the SPSR. The exception code will run and on return from the exception the context of the CPSR will be

restored from the SPSR allowing the application code to resume execution. The operating modes are listed below.

System & User	FIQ	Supervisor	Abort	IRQ	Undefined
R0	R0	R0	R0	R0	R0
R1	R1	R1	R1	R1	R1
R2	R2	R2	R2	R2	R2
R3	R3	R3	R3	R3	R3
R4	R4	R4	R4	R4	R4
R5	R5	R5	R5	R5	R5
R6	R6	R6	R6	R6	R6
R7	R7_fiq	R7	R7	R7	R7
R8	R8_fiq	R8	R8	R8	R8
R9	R9_fiq	R9	R9	R9	R9
R10	R10_fiq	R10	R10	R10	R10
R11	R11_fiq	R11	R11	R11	R11
R12	R12_fiq	R12	R12	R12	R12
R13	R13_fiq	R13_svc	R13_abt	R13_irq	R13_und
R14	R14_fiq	R14_svc	R14_abt	R14_irq	R14_und
R15 (PC)	R15 (PC)	R15 (PC)	R15 (PC)	R15 (PC)	R15 (PC)

CPSR	CPSR SPSR_fiq	CPSR SPSR_svc	CPSR SPSR_abt	CPSR SPSR_irq	CPSR SPSR_und
------	------------------	------------------	------------------	------------------	------------------

The ARM7 CPU has six operating modes which are used to process exceptions. The shaded registers are banked memory that is “switched in” when the operating mode changes. The SPSR register is used to save a copy of the CPSR when the switch occurs

Exception Modes

When an exception occurs, the CPU will change modes and the PC be forced to an exception vector. The vector table starts from address zero with the reset vector and then has an exception vector every four bytes.

Exception	Mode	Address
Reset	Supervisor	0x00000000
Undefined instruction	Undefined	0x00000004
Software interrupt (SWI)	Supervisor	0x00000008
Prefetch Abort (instruction fetch memory abort)	Abort	0x0000000C
Data Abort (data access memory abort)	Abort	0x00000010
IRQ (interrupt)	IRQ	0x00000018
FIQ (fast interrupt)	FIQ	0x0000001C

Each operating mode has an associated interrupt vector. When the processor changes mode the PC will jump to the associated vector.

NB. there is a missing vector at 0x00000014

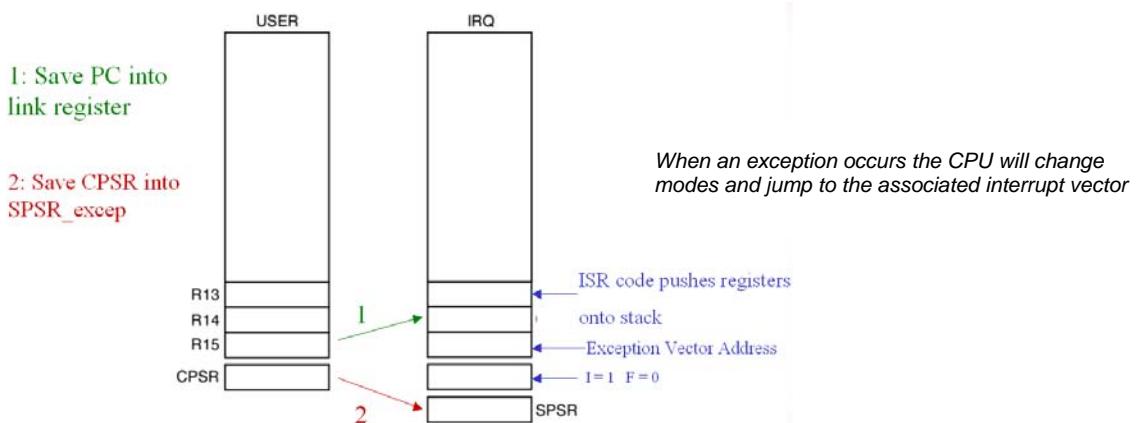
NB: There is a gap in the vector table because there is a missing vector at 0x00000014. This location was used on an earlier ARM architecture and has been preserved on ARM7 to ensure software compatibility between different ARM architectures. However in the LPC2000 family these four bytes are used for a very special purpose as we shall see later.

Priority	Exception
Highest 1	Reset
2	Data Abort
3	FIQ
4	IRQ
5	Prefetch Abort
Lowest 6	Undefined instruction SWI

Each of the exception sources has a fixed priority. The on chip peripherals are served by FIQ and IRQ interrupts. Each peripheral's priority may be assigned within these groups

If multiple exceptions occur then there is a fixed priority as shown below.

When an exception occurs, for example an IRQ exception, the following actions are taken: First the address of the next instruction to be executed (PC + 4) is saved into the link register. Then the CPSR is copied into the SPSR of the exception mode that is about to be entered (i.e. SPSR_irq). The PC is then filled with the address of the exception mode interrupt vector. In the case of the IRQ mode this is 0x00000018. At the same time the mode is changed to IRQ mode, which causes R13 and R14 to be replaced by the IRQ R13 and R14 registers. On entry to the IRQ mode, the I bit in the CPSR is set, causing the IRQ interrupt line to be disabled. If you need to have nested IRQ interrupts, your code must manually re-enable the IRQ interrupt and push the link register onto the stack in order to preserve the original return address. From the exception interrupt vector your code will jump to the exception ISR. The first thing your code must do is to preserve any of the registers R0-R12 that the ISR will use by pushing them onto the IRQ stack. Once this is done you can begin processing the exception.



Once your code has finished processing the exception it must return back to the user mode and continue where it left off. However the ARM instruction set does not contain a “return” or “return from interrupt” instruction so manipulating the PC must be done by regular instructions. The situation is further complicated by there being a number of different return

cases. First of all, consider the SWI instruction. In this case the SWI instruction is executed, the address of the next instruction to be executed is stored in the Link register and the exception is processed. In order to return from the exception all that is necessary is to move the contents of the link register into the PC and processing can continue. However in order to make the CPU switch modes back to user mode, a modified version of the move instruction is used and this is called MOVS (more about this later). Hence for a software interrupt the return instruction is

```
MOVS      R15, R14      ; Move Link register into the PC and switch modes.
```

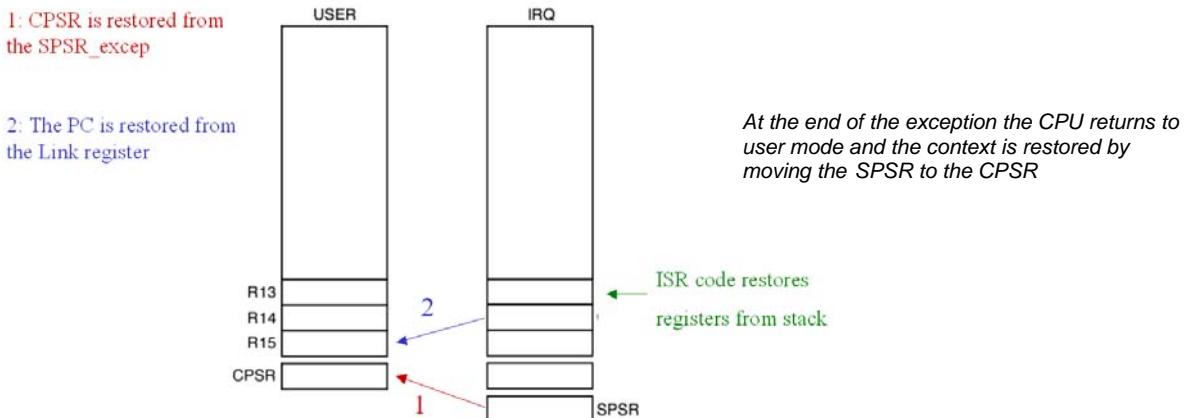
However, in the case of the FIQ and IRQ instructions, when an exception occurs the current instruction being executed is discarded and the exception is entered. When the code returns from the exception the link register contains the address of the discarded instruction plus four. In order to resume processing at the correct point we need to roll back the value in the Link register by four. In this case we use the subtract instruction to deduct four from the link register and store the results in the PC. As with the move instruction, there is a form of the subtract instruction which will also restore the operating mode. For an IRQ, FIQ or **Prog Abort**, the return instruction is:

```
SUBS    R15, R14, #4
```

In the case of a data abort instruction, the exception will occur one instruction after execution of the instruction which caused the exception. In this case we will ideally enter the data abort ISR, sort out the problem with the memory and return to reprocess the instruction that caused the exception. In this case we have to roll back the PC by two instructions i.e. the discarded instruction and the instruction that caused the exception. In other words subtract eight from the link register and store the result in the PC. For a data abort exception the return instruction is

```
SUBS    R15, R14, #8
```

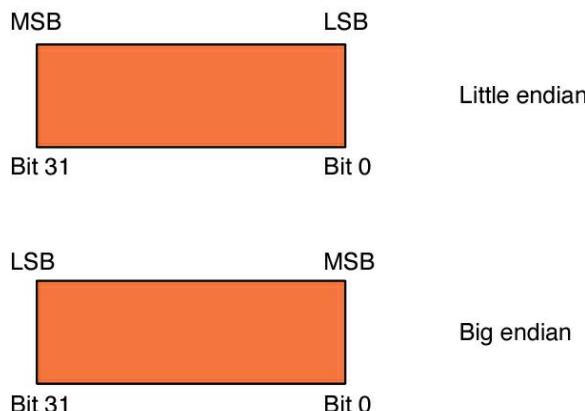
Once the return instruction has been executed, the modified contents of the link register are moved into the PC, the user mode is restored and the SPSR is restored to the CPSR. Also, in the case of the FIQ or IRQ exceptions, the relevant interrupt is enabled. This exits the privileged mode and returns to the user code ready to continue processing.



ARM 7 Instruction Set

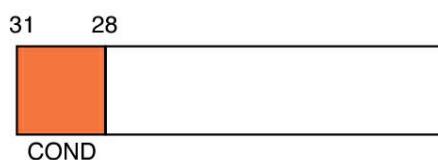
Now that we have an idea of the ARM7 architecture, programmers model and operating modes we need to take a look at its instruction set or rather sets. Since all our programming examples are written in C there is no need to be an expert ARM7 assembly programmer. However an understanding of the underlying machine code is very important in developing efficient programs. Before we start our overview of the ARM7 instructions it is important to set out a few technicalities. The ARM7 CPU has two instruction sets: the ARM instruction set which has 32-bit wide instructions and the THUMB instruction set which has 16-bit wide instructions. In the following section the use of the word ARM means the 32-bit instruction set and ARM7 refers to the CPU.

The ARM7 is designed to operate as a big-endian or little-endian processor. That is, the MSB is located at the high order bit or the low order bit. You may be pleased to hear that the LPC2000 family fixes the endianess of the processor as little endian (i.e. MSB at highest bit address), which does make it a lot easier to work with. However the ARM7 compiler you are working with will be able to compile code as little endian or big endian. You must be sure you have it set correctly or the compiled code will be back to front.



The ARM7 CPU is designed to support code compiler in big endian or little endian format. The Philips silicon is fixed as little endian.

One of the most interesting features of the ARM instruction set is that every instruction may be conditionally executed. In a more traditional microcontroller the only conditional instructions are conditional branches and maybe a few others like bit test and set. However in the ARM instruction set the top four bits of the operand are compared to the condition codes in the CPSR. If they do not match then the instruction is not executed and passes through the pipeline as a NOP (no operation).



Every ARM (32 bit) instruction is conditionally executed. The top four bits are ANDed with the CPSR condition codes. If they do not match the instruction is executed as a NOP

So it is possible to perform a data processing instruction, which affects the condition codes in the CPSR. Then depending on this result, the following instructions may or may not be carried out. The basic assembler instructions such as MOV or ADD can be prefixed with sixteen conditional mnemonics, which define the condition code states to be tested for.

Suffix	Flags	Meaning
EQ	Z set	equal
NE	Z clear	not equal
CS	C set	unsigned higher or same
CC	C clear	unsigned lower
MI	N set	negative
PL	N clear	positive or zero
VS	V set	overflow
VC	V clear	no overflow
HI	C set and Z clear	unsigned higher
LS	C clear and Z set	unsigned lower or same
GE	N equals V	greater or equal
LT	N not equal to V	less than
GT	Z clear AND (N equals V)	greater than
LE	Z set OR (N not equal to V)	less than or equal
AL	(ignored)	always

Each ARM (32-bit) instruction can be prefixed by one of 16 condition codes. Hence each instruction has 16 different variants.

So for example:

```
EQMOV R1, #0x00800000
```

will only move 0x00800000 into the R1 if the last result of the last data processing instruction was equal and consequently set the Z flag in the CPSR. The aim of this conditional execution of instructions is to keep a smooth flow of instructions through the pipeline. Every time there is a branch or jump the pipeline is flushed and must be refilled and this causes a dip in overall performance. In practice there is a break-even point between effectively forcing NOP instructions through the pipeline and a traditional conditional branch and refill of the pipeline. This break-even point is three instructions, so a small branch such as:

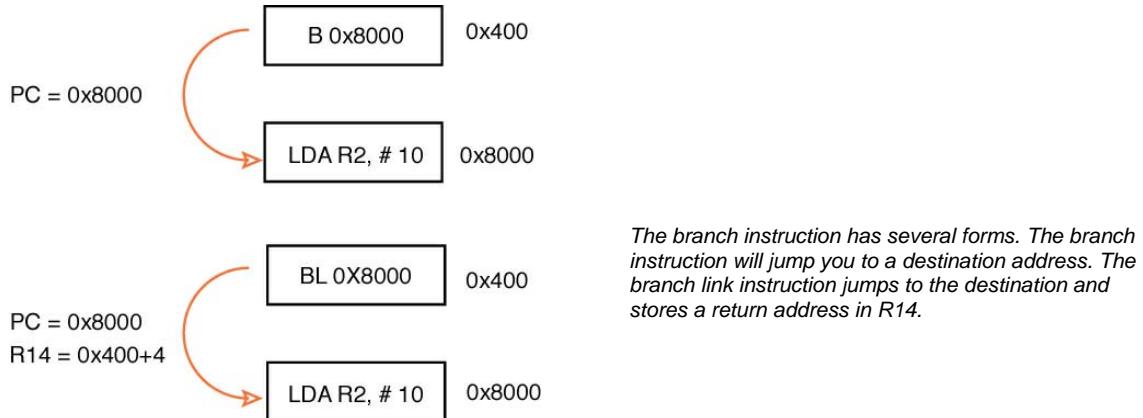
```
if( x<100)
{
    x++;
}
```

would be most efficient when coded using conditional execution of ARM instructions.

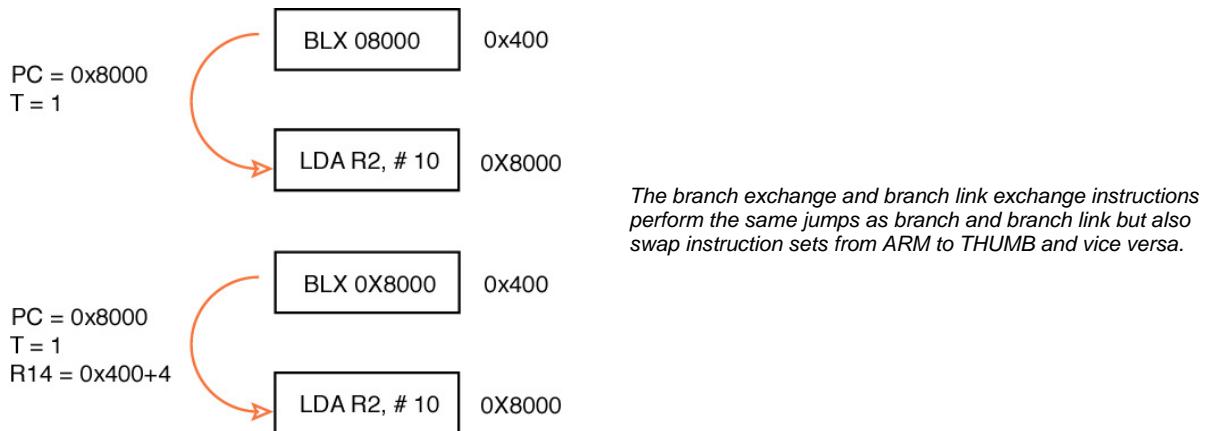
The main instruction groups of the ARM instruction set fall into six different categories, Branching, Data Processing, Data Transfer, Block Transfer, Multiply and Software Interrupt.

Branching

The basic branch instruction (as its name implies) allows a jump forwards or backwards of up to 32 MB. A modified version of the branch instruction, the branch link, allows the same jump but stores the current PC address plus four bytes in the link register.



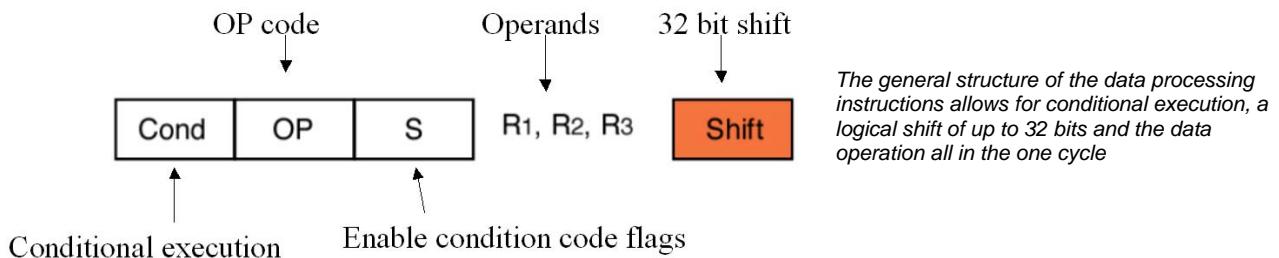
So the branch link instruction is used as a call to a function storing the return address in the link register and the branch instruction can be used to branch on the contents of the link register to make the return at the end of the function. By using the condition codes we can perform conditional branching and conditional calling of functions. The branch instructions have two other variants called “branch exchange” and “branch link exchange”. These two instructions perform the same branch operation but also swap instruction operation from ARM to THUMB and vice versa.



This is the only method you should use to swap instruction sets, as directly manipulating the “T” bit in the CPSR can lead to unpredictable results.

Data Processing Instructions

The general form for all data processing instructions is shown below. Each instruction has a result register and two operands. The first operand must be a register, but the second can be a register or an immediate value.



In addition, the ARM7 core contains a barrel shifter which allows the second operand to be shifted by a full 32-bits within the instruction cycle. The "S" bit is used to control the condition codes. If it is set, the condition codes are modified depending on the result of the instruction. If it is clear, no update is made. If, however, the PC (R15) is specified as the result register and the S flag is set, this will cause the SPSR of the current mode to be copied to the CPSR. This is used at the end of an exception to restore the PC and switch back to the original mode. Do not try this when you are in the USER mode as there is no SPSR and the result would be unpredictable.

Mnemonic	Meaning
AND	Logical bitwise AND
EOR	Logical bitwise exclusive OR
SUB	Subtract
RSB	Reverse Subtract
ADD	Add
ADC	Add with carry
SBC	Subtract with carry
RSC	Reverse Subtract with carry
TST	Test
TEQ	Test Equivalence
CMP	Compare
CMN	Compare negated
ORR	Logical bitwise OR
MOV	Move
BIC	Bit clear
MVN	Move negated

These features give us a rich set of data processing instructions which can be used to build very efficiently-coded programs, or to give a compiler-designer nightmares. An example of a typical ARM instruction is shown below.

```
if (Z ==1)R1 = R2+(R3x4)
Can be compiled to: EQADDS R1,R2,R3,LSL #2
```

Copying Registers

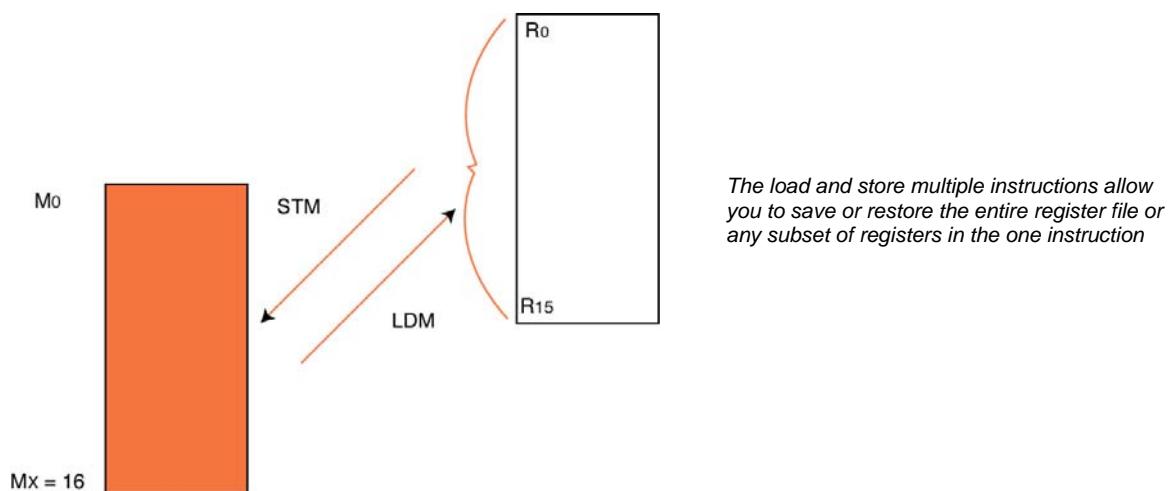
The next group of instructions are the data transfer instructions. The ARM7 CPU has load-and-store register instructions that can move signed and unsigned Word, Half Word and Byte quantities to and from a selected register.

Mnemonic	Meaning
LDR	Load Word
LDRH	Load Half Word
LDRSH	Load Signed Half Word
LDRB	Load Byte
LRDSB	Load Signed Byte
STR	Store Word
STRH	Store Half Word
STRSH	Store Signed Half Word
STRB	Store Byte
STRSB	Store Signed Half Word

Since the register set is fully orthogonal it is possible to load a 32-bit value into the PC, forcing a program jump anywhere within the processor address space. If the target address is beyond the range of a branch instruction, a stored constant can be loaded into the PC.

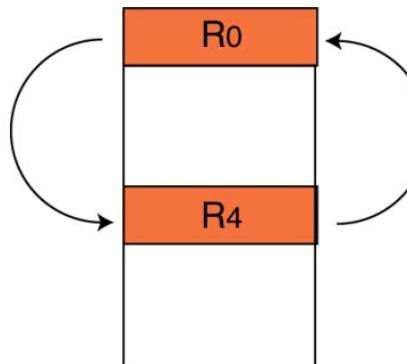
Copying Multiple Registers

In addition to load and storing single register values, the ARM has instructions to load and store multiple registers. So with a single instruction, the whole register bank or a selected subset can be copied to memory and restored with a second instruction



Swap Instruction

The ARM instruction set also provides support for real time semaphores with a swap instruction. The swap instruction exchanges a word between registers and memory as one atomic instruction. This prevents crucial data exchanges from being interrupted by an exception.

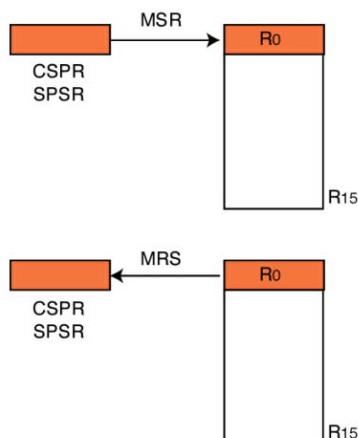


The swap instruction allows you to exchange the contents of two registers. This takes two cycles but is treated as a single atomic instruction so the exchange cannot be corrupted by an interrupt.

This instruction is not reachable from the C language and is supported by intrinsic functions within the compiler library.

Modifying The Status Registers

As noted in the ARM7 architecture section, the CPSR and the SPSR are CPU registers, but are not part of the main register bank. Only two ARM instructions can operate on these registers directly. The MSR and MRS instructions support moving the contents of the CPSR or SPSR to and from a selected register. For example, in order to disable the IRQ interrupts the contents of the CPSR must be moved to a register, the “I” bit must be set by ANDing the contents with 0x00000080 to disable the interrupt and then the CPSR must be reprogrammed with the new value.



The CPSR and SPSR are not memory-mapped or part of the central register file. The only instructions which operate on them are the MSR and MRS instructions. These instructions are disabled when the CPU is in USER mode.

The MSR and MRS instructions will work in all processor modes except the USER mode. So it is only possible to change the operating mode of the process, or to enable or disable

interrupts, from a privileged mode. Once you have entered the USER mode you cannot leave it, except through an exception, reset, FIQ, IRQ or SWI instruction.

Software Interrupt

The Software Interrupt Instruction generates an exception on execution, forces the processor into supervisor mode and jumps the PC to 0x00000008. As with all other ARM instructions, the SWI instruction contains the condition execution codes in the top four bits followed by the op code. The remainder of the instruction is empty. However it is possible to encode a number into these unused bits. On entering the software interrupt, the software interrupt code can examine these bits and decide which code to run. So it is possible to use the SWI instruction to make calls into the protected mode, in order to run privileged code or make operating system calls.

31	28 27	24 23	
Cond	1111		Ordinal

The Software Interrupt Instruction forces the CPU into SUPERVISOR mode and jumps the PC to the SWI vector. Bits 0-23 are unused and user defined numbers can be encoded into this space.

The Assembler Instruction:

SWI #3

Will encode the value 3 into the unused bits of the SWI instruction. In the SWI ISR routine we can examine the SWI instruction with the following code pseudo code:

```
switch( *(R14-4) & 0x00FFFFFF)      // roll back the address stored in link reg
{                                     // by 4 bytes
    // Mask off the top 8 bits and switch
    // on result
    case ( SWI-1)
    ....
```

Depending on your compiler, you may need to implement this yourself, or it may be done for you in the compiler implementation.

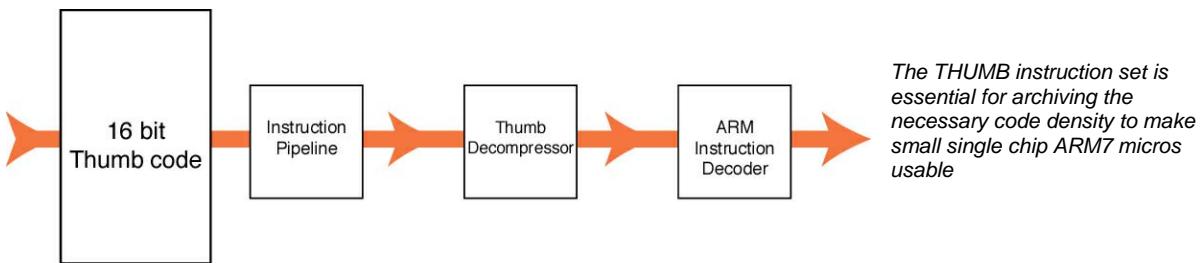
MAC Unit

In addition to the barrel shifter, the ARM7 has a built-in Multiply Accumulate Unit (MAC). The MAC supports integer and long integer multiplication. The integer multiplication instructions support multiplication of two 32-bit registers and place the result in a third 32-bit register (modulo32). A multiply-accumulate instruction will take the same product and add it to a running total. Long integer multiplication allows two 32-bit quantities to be multiplied together and the 64-bit result is placed in two registers. Similarly a long multiply and accumulate is also available.

Mnemonic	Meaning	Resolution
MUL	Multiply	32 bit result
MULA	Multiply accumulate	32 bit result
UMULL	Unsigned multiply	64 bit result
UMLAL	Unsigned multiply accumulate	64 bit result
SMULL	Signed multiply	64 bit result
SMLAL	Signed multiply accumulate	64 bit result

THUMB Instruction Set

Although the ARM7 is a 32-bit processor, it has a second 16-bit instruction set called THUMB. The THUMB instruction set is really a compressed form of the ARM instruction set.



This allows instructions to be stored in a 16-bit format, expanded into ARM instructions and then executed. Although the THUMB instructions will result in lower code performance compared to ARM instructions, they will achieve a much higher code density. So, in order to build a reasonably-sized application that will fit on a small single chip microcontroller, it is vital to compile your code as a mixture of ARM and THUMB functions. This process is called interworking and is easily supported on all ARM compilers. By compiling code in the THUMB instruction set you can get a space saving of 30%, while the same code compiled as ARM code will run 40% faster.

The THUMB instruction set is much more like a traditional microcontroller instruction set. Unlike the ARM instructions THUMB instructions are not conditionally executed (except for conditional branches). The data processing instructions have a two-address format, where the destination register is one of the source registers:

ARM Instruction

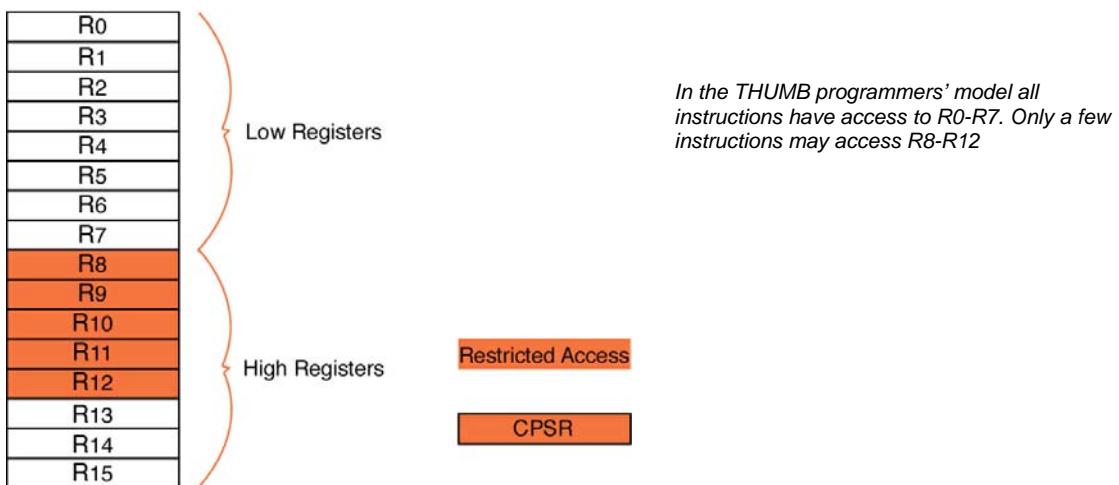
`ADD R0, R0,R1`

THUMB Instruction

`ADD R0 ,R1`

$R0 = R0+R1$

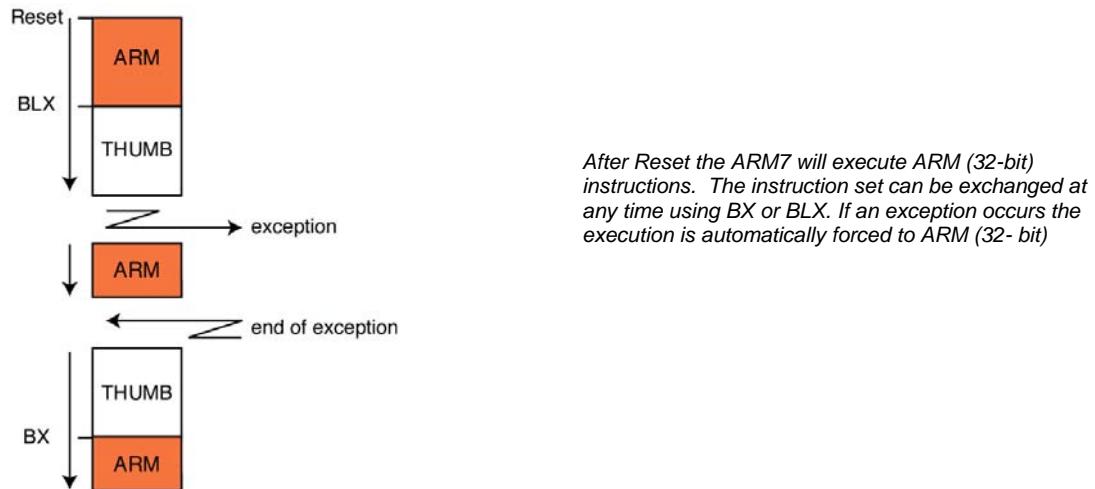
The THUMB instruction set does not have full access to all registers in the register file. All data processing instructions have access to R0 –R7 (these are called the “low registers”).



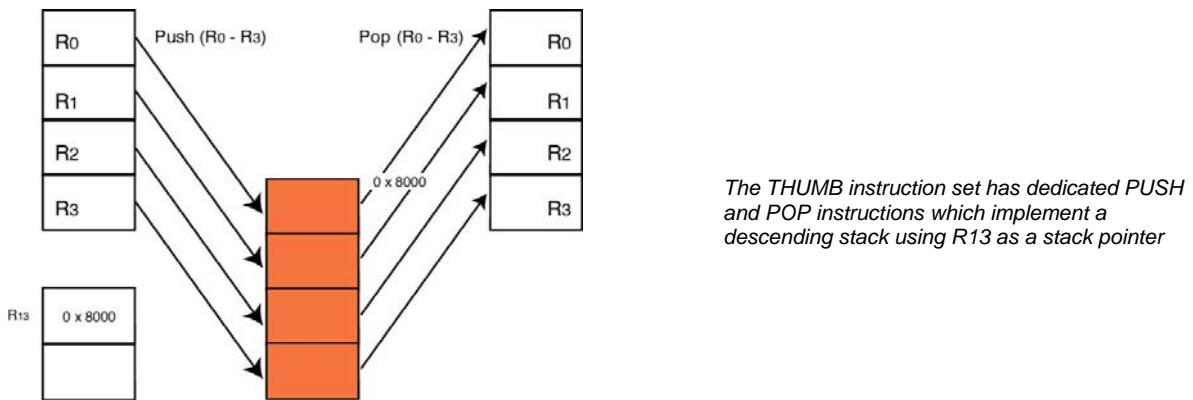
However access to R8-R12 (the “high registers”) is restricted to a few instructions:

MOV, ADD, CMP

The THUMB instruction set does not contain MSR and MRS instructions, so you can only indirectly affect the CPSR and SPSR. If you need to modify any user bits in the CPSR you must change to ARM mode. You can change modes by using the BX and BLX instructions. Also, when you come out of RESET, or enter an exception mode, you will automatically change to ARM mode.



The THUMB instruction set has the more traditional PUSH and POP instructions for stack manipulation. They implement a fully descending stack, hardwired to R13.



Finally, the THUMB instruction set does contain a SWI instruction which works in the same way as in the ARM instruction set, but it only contains 8 unused bits, to give a maximum of 255 SWI calls.

Summary

At the end of this chapter you should have a basic understanding of the ARM7 CPU. Please see the bibliography for a list of books that address the ARM7 in more detail. Also included on the CD is a copy of the ARM7 user manual.

Chapter 2: Software Development

Outline

In this book we will be using an Integrated Development Environment from Keil Electronic. This IDE is called uVISION (pronounced “MicroVision”) and versions already exist for other popular microcontrollers including the 8051 and the Infineon C16X family. uVISION successfully integrates project management, editor, compiler and debugger in one seamless front-end. Although we are concentrating on the LPC2000 family in this book, the Keil ARM tools can be used for any other ARM7 based microcontroller.

Which Compiler?

The uVISION development environment can be used with several different compiler tools. These include the ARM ADS compiler, the GNU compiler and Keil’s own ARM compiler. In this book the examples are based on the Keil CA-ARM compiler. However, a parallel set of examples is also included for the GNU compiler and Appendix A details the differences between the Keil and GNU compilers. This does beg the question of which compiler to use. First of all the GNU compiler is free, can be downloaded from the internet and is also included on the CD which comes with this book. So why use an expensive commercial compiler? Well, before you embark on a full project, it is worth looking at the table of benchmarks comparing some of the most popular C compilers available for the ARM CPU.

Dhystone V2.1

Parameter	Compiler			
	Keil CA BETA	GNU V3.22	ARM ADS V1.2	IAR V4.11A
Execution Speed (µSeconds)	25.4	112.9	16.8	24.4
Dhystones/sec	39,370.1	8,857.4	59,382.4	40,983.6
Total Code Size (bytes)	10,330	36,004	22,266	19,892
Stack Size (bytes)	208	852	608	356
Total Data Size (bytes)	10,256	11,912	10,256	10,269

All tests were performed under identical conditions using the Keil µVision Simulator.
The ARM device used was a Philips LPC2294 running at 60MHz in Thumb Mode.

Whetstone

Parameter	Compiler			
	Keil CA BETA	GNU V3.22	ARM ADS V1.2	IAR V4.11A
Execution Speed (seconds)	0.195308	2.461430	0.268623	0.635633
Whetstone (KWIPS)	5,263	406	3,846	1,587
Total Code Size (bytes)	8,492	44,428	28,516	16,316
Stack Size (bytes)	212	1,552	710	488
Total Data Size (bytes)	72	1,768	76	72

All tests were performed under identical conditions using the Keil µVision Simulator.
The ARM device used was a Philips LPC2294 running at 60MHz in Thumb Mode.

We can see from this simple analysis that the commercial compilers are streets ahead of the GNU tools in terms of code density and speed of execution. The reasons to use each of the given compilers can be summed up as follows: if you want the fastest code and standard tools use the ARM compiler, for best code density use the Keil, if you have no budget or a simple project use the GNU. Since we are writing code for a small single-chip microcontroller with limited on-chip resources, the obvious choice for us is the Keil ARM compiler. When deciding on a toolset it is also important to examine how much support is given to a specific ARM7 implementation. Although a toolset may generate code for an ARM7, it may not understand how the ARM7 is being used in a specific system i.e. LPC2000. Using a “raw” ARM7 will generate code, which will run on the LPC2000, but you will have to spend time writing the start-up code and struggle with a debugger, which will not understand the LPC peripherals. This can lead to “fighting” the development tools, which needless to say can be very frustrating.

uVISION IDE

uVISION also includes two debug tools. Once the code has been compiled and linked, it can be loaded into the uVISION simulator. This debugger simulates the ARM7 core and peripherals of the supported micro. Using the simulator is a very good way of becoming familiar with the LPC2000 devices. Since the simulator gives cycle- accurate simulation of the peripherals, as well as the CPU, it can be a very useful tool for verifying that the chip has been correctly initialised and that the correct values for things such as timer prescaler values have been calculated.

However, the simulator can only take you so far and sooner or later you will need to take some inputs from the real world. This can be done to a certain extent with the simulator scripting language, but eventually you will need to run your code on the real target. The simulator front end can be connected to your hardware by the Keil Ulink interface. The ULINK interface connects to the PC via USB and connects to the development hardware by the LPC2000 JTAG interface. The JTAG interface is a separate peripheral on the ARM7 which supports debug commands from a host. By using the JTAG you can use the uVISION simulator to have basic run control of the LPC2000 device. The JTAG allows you to download code onto the target, to single step and run code at full speed, to set breakpoints and view memory locations.

Tutorial

Included with this book is a demonstration version of the Keil uVISION IDE. The installation comes with two compilers, the Keil ARM compiler and the GNU tools. The tutorial section talks you through example programs illustrating the major features of the LPC2000. These examples can be run on the simulator, or if you have the starter kit they can be downloaded and run on the MCB2100 evaluation board. There are two sets of examples on the CD, one for the Keil compiler and one for the GNU. The main text concentrates on the Keil compiler. However, Appendix A describes how to use the GNU compiler and also

describes the GNU version of the exercises up to exercise 6. After exercise 6 you can use the exercise descriptions in the main text.

As you read through the rest of the book, at the end of each section there will be an exercise described in the tutorial section which illustrates what has been discussed. The best way to use this book is to read each section, then jump to the tutorial and do the exercise. This way, by the time you have worked through the book you will have a firm grasp of the ARM7, its tools and the LPC2000 microcontroller.

Exercise 1: Configuring A New Project

The first exercise covers installing the uVISION software and setting up a first project.

Startup Code

In our example project we have a number of source files. In practice the .c files are your source code, but the file startup.s is an assembler module provided by Keil. As its name implies, the startup code is located to run from the reset vector. It provides the exception vector table as well as initialising the stack pointer for the different operating modes. It also initialises some of the on-chip system peripherals and the on-chip RAM before it jumps to the main function in your C code. The startup code will vary, depending on which ARM7 device you are using and which compiler you are using, so for your own project it is important to make sure that you are using the correct file. The startup code for the Keil compiler may be found in C:\keil\ARM\startup and for the GNU use the files in C:\keil\GNU\startup.

First of all the startup code provides the exception vector table as shown below

Declare the external C exception routines. →

```
EXTERN CODE32 (Undef_Handler?A)
EXTERN CODE32 (SWI_Handler?A)
EXTERN CODE32 (PAbt_Handler?A)
EXTERN CODE32 (DAbt_Handler?A)
EXTERN CODE32 (IRQ_Handler?A)
EXTERN CODE32 (FIQ_Handler?A)
```

The suffix ?A denotes an ARM routine

Unused vector, this will become → important later !

```
Vectors: LDR PC, Reset_Addr
          LDR PC, Undef_Addr
          LDR PC, SWI_Addr
          LDR PC, PAbt_Addr
          LDR PC, DAbt_Addr
          NOP             /* Reserved Vector */
          LDR PC, IRQ_Addr
          LDR PC, FIQ_Addr
```

Constants table for ISR address →

```
Reset_Addr: DD  Reset_Handler
Undef_Addr:  DD  Undef_Handler?A
SWI_Addr:   DD  SWI_Handler?A
PAbt_Addr:  DD  PAbt_Handler?A
DAbt_Addr:  DD  DAbt_Handler?A
          DD  0           /* Reserved Address */
IRQ_Addr:   DD  IRQ_Handler?A
FIQ_Addr:   DD  FIQ_Handler?A
```

The vector table is located at 0x00000000 and provides a jump to interrupt service routines (ISR) on each vector. To ensure that the full address range of the processor is available, the LDR (Load Register) instruction is used. This loads a constant from a table stored immediately above the vector table. The vector table and the constants table take up the first 64 bytes of memory. On the LPC2000 this first 64 bytes can be mapped from several

sources, depending on the operating mode of the LPC2000. (This is discussed more fully later on.) The NOP instruction is used to pad out the vector table at location 0x000000014 which is the location of the ‘missing’ vector. Again this location is used by the LPC2000 bootloader (discussed again later.) You are responsible for managing the vector table in the startup code as it is not done automatically by the compiler.

The startup code is also responsible for configuring the stack pointers for each of the operating modes.

```

// Setup Stack for each mode
    LDR R0, =Top_Stack
// Enter Undefined Instruction Mode and set its Stack
// Pointer
    → MSR CPSR_c
    → MOV SP, R0
    → SUB R0, R0, : The six on chip stack pointers (R13)
                      are initialised at the top of on chip
                      memory. Care must be taken to
                      allocate enough memory for the
                      maximum size of each stack
// Enter Abort Mode and s
    MSR CPSR_c
    MOV SP, R0
    SUB R0, R0, :
.....FIQ,IRQ and supervisor stacks

// Enter User Mode and set its Stack Pointer
    → MSR CPSR_c, #Mode_USR
    → MOV SP, R0
// Enter the C code
    LDR R0,=C?INIT
    TST R0,#1 ; Bit-0 set: INIT is Thumb
    LDREQ LR,=exit?A ; ARM Mode
    LDRNE LR,=exit?T ; Thumb Mode
    BX R0

```

Finally switch to USER mode and enable interrupts

Check for ARM or Thumb mode.

Load an Exit address into the link register

Jump to the C init routine

Since each operating mode has a unique R13 there are effectively six stacks in the ARM7. The strategy used by the compiler is to locate user variables from the start of the on-chip RAM and grow upwards. The stacks are located at the top of memory and grow downwards. The startup code enters each different mode of the ARM7 and loads each R13 with the starting address of the stack

Stack Configuration (Stack Sizes in Bytes)	
Undefined Mode	0x0000 0004
Supervisor Mode	0x0000 0004
Abort Mode	0x0000 0004
Fast Interrupt Mode	0x0000 0004
Interrupt Mode	0x0000 0080
User/System Mode	0x0000 0400
PLL Setup	<input checked="" type="checkbox"/>
MAM Setup	<input checked="" type="checkbox"/>
External Memory Controller (EMC)	<input checked="" type="checkbox"/>

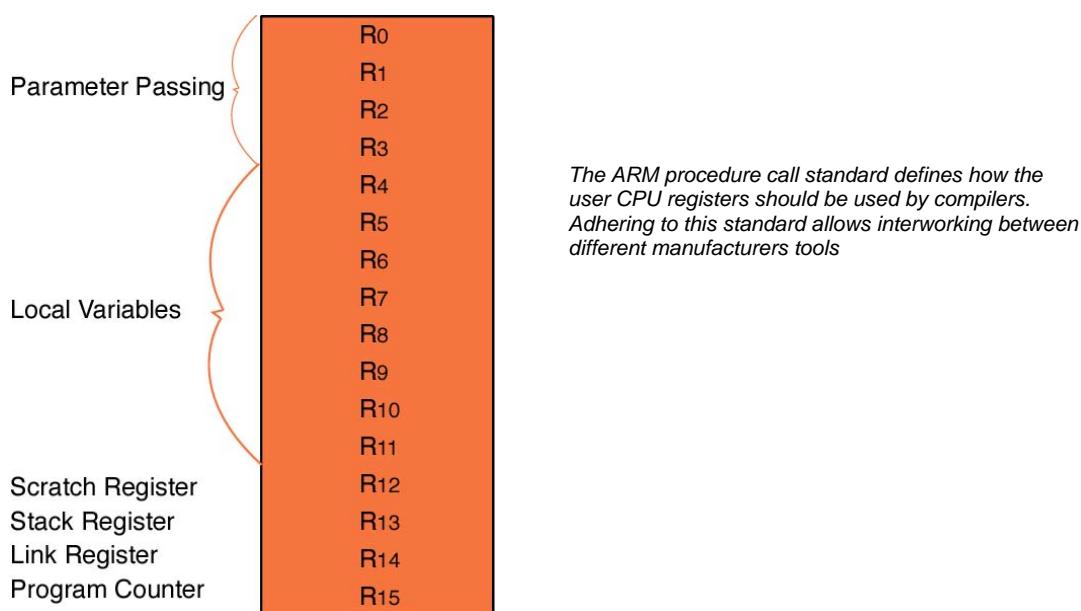
Like the vector table you are responsible for configuring the stack size. This can be done by editing the startup code directly, however Keil provide a graphical editor that allows you to more easily configure the stack spaces. In addition the graphical editor allows you to configure some of the LPC2000 system peripherals. We will see these in more detail later but remember that they can be configured directly in the startup code.

Exercise 2: Startup code

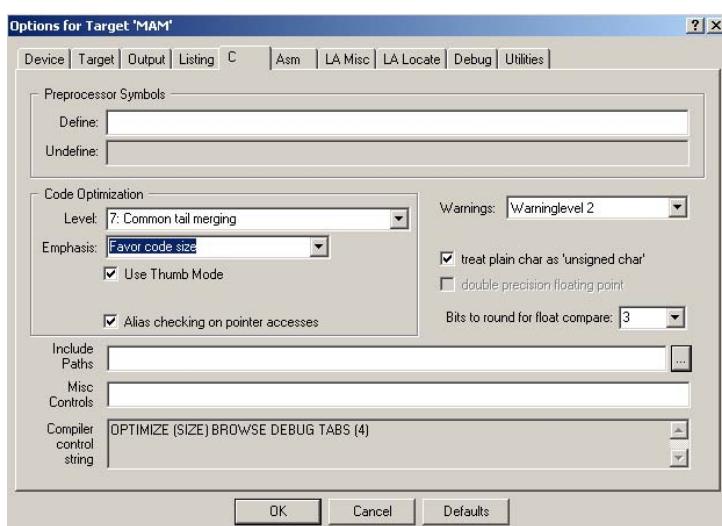
The second exercise in the tutorial takes you through allocating space for each processor stack and examines the vector table.

Interworking ARM/THUMB Code

One of the most important things that we need to do in our application code is to interwork the ARM and THUMB instruction sets. In order to allow this interoperability, ARM have defined a standard called the ARM THUMB Procedure Call Standard (ATPCS). The ATPCS defines among other things how functions call one another, how parameters are passed and how stacks are handled. The APCS adds a veneer of assembler code to support various compiler features. The more you use, the larger these veneers get. In theory the APCS allows code built in different toolsets to work together so that you can take a library compiled by a different compiler and use it with the Keil toolset.



The APCS splits the register file into a number of regions: R0 to R3 are used for parameter passing between functions. If you need to pass more than 16 bytes then spilled parameters are passed via the stack. Local variables are allocated R4 – R11 and R12 is reserved as a memory location for the intra-call veneer code. In the Keil compiler all code is built for interworking and the global instruction set is the THUMB, so all code will be compiled as THUMB instructions (except for interrupt code which defaults to ARM.) This global default can be changed in the “Options for Target” menu. In the CC tab uncheck the “use THUMB code” box and the default instruction set will be ARM.



In addition the programmer can force a given function to be compiled as ARM or THUMB code. This is done with the two programming directives #Pragma ARM and #pragma THUMB as shown below. The main function is compiled as ARM code and calls a function called THUMB_function, (No prizes for guessing that this function is compiled in the 16 bit instruction set.)

```
#pragma ARM // Switch to ARM instructions

int main(void)
{
while(1)
{
THUMB_function(); //Call THUMB function
}
}

#pragma THUMB //Switch to THUMB instructions

void THUMB_function(void)
{
unsigned long i,delay;

for (i = 0x00010000;i < 0x01000000 ;i = i<<1) //LED FLASHER
{
for (delay = 0;delay<0x00010000;delay++) //simple delay loop
{
;
}
IOSET1 = i; //Set the next led
}
}
```

It is also possible to declare individual functions as either ARM or THUMB functions by using the following declarations on the function prototype:

```
int ARM_FUNCTION ( int my_var) __THUMB
{
...
}

int THUMB_FUNCTION ( int my_var) __THUMB
{
...
}
```

Exercise 3: Interworking

The next exercise demonstrates setting up a project which interworks ARM and THUMB code.

STDIO Libraries

The high-level, formatted IO functions in the STDIO library, such as printf and scanf, are directed at UART0 on the LPC2000. It is up to the programmer to initialise the UART to the correct BAUD rate. Once this is done it is possible to use these high- level functions to stream data to a terminal program on a PC for example. The STDIO functions use two low- level drivers to send and receive a single character to the conio, the UART in this case. The two functions are called putchar and getchar and the source for them is available in serial.c in the Keil lib directory. By adding this file to your project the default library version is

ignored and the code in serial.c is used in its place. So, by rewriting the putchar and getchar routines, the high level printf and scanf function can be redirected to any IO device you want to use, such as an LCD and keypad. Bear in mind that the high level STDIO functions are quite bulky and should only be used if your application is very I/O driven.

Exercise 4: STDIO

This exercise demonstrates the low-level routines used by printf and scanf and configures them to read and write to the on-chip UART.

Accessing Peripherals

Once we have built some code and got it running on an LPC2000 device, it will at some point be necessary to access the special function registers (SFR) in the peripherals. As all the peripherals are memory-mapped, they can be accessed as normal memory locations. Each SFR location can be accessed by ‘hardwiring’ a volatile pointer to its memory location as shown below.

```
#define SFR    (*((volatile unsigned long *) 0xFFFFF000))
```

The Keil compiler comes with a set of include files which define all the SFR’s in the different LPC2000 variants. Just include the correct file and you can directly access the peripheral SFR’s from your C code. The names of the include files are:

```
LPC21xx.h  
LPC22xx.h  
LPC210x.h
```

Interrupt Service Routines

In addition to accessing the on-chip peripherals, your C code will have to service interrupt requests. It is possible to convert a standard function into an ISR, as shown below:

```
void fiqint (void) __fiq
{
    IOSET1      = 0x00FF0000; // Set the LED pins
    EXTINT     = 0x00000002; // Clear the peripheral interrupt flag
}
```

The keyword `__fiq` defines the function as a fast interrupt request service routine and so will use the correct return mechanism. Other types of interrupt are supported by the keywords `__IRQ`, `__SWI`, `__ABORT`.

As well as declaring a C function as an interrupt routine, you must link the interrupt vector to the function.

Vectors:	LDR	PC,Reset_Addr
	LDR	PC,Undef_Addr
	LDR	PC,SWI_Addr
	LDR	PC,PAbt_Addr
	LDR	PC,DAbt_Addr
	NOP	/* Reserved Vector */
;	LDR	PC,IRQ_Addr
	LDR	PC,[PC, #-0x0FF0] /* Vector from VicVectAddr */
	LDR	PC,FIQ_Addr
Reset_Addr:	DD	Reset_Handler
Undef_Addr:	DD	Undef_Handler?A
SWI_Addr:	DD	SWI_Handler?A
PAbt_Addr:	DD	PAbt_Handler?A
DAbt_Addr:	DD	DAbt_Handler?A
	DD	0 /* Reserved Address */
IRQ_Addr:	DD	IRQ_Handler?A
FIQ_Addr:	DD	FIQ_Handler?A

The vector table is in two parts. First there is the physical vector table, which has a Load Register Instruction (LDR) on each vector. This loads the contents of a 32-bit wide memory location into the PC, forcing a jump to any location within the processor's address space. These values are held in the second half of the vector table, or the constants table which follows immediately after the vector table. This means that the complete vector table takes the first 64 bytes of memory. The Keil startup code contains predefined names for the Interrupt Service Routines (ISR). You can link your ISR functions to each interrupt vector by using the same name as your C function name. The table below shows the constants table symbols and the corresponding C function prototypes which should be used.

Exception source	Constants table	C function prototype
Undefined Instruction	Undef_Handler?A	void Undef_Handler (void) __abort
Prefetch Abort	PAbt_Handler?A	void PAbt_Handler (void) __abort
Data Abort	DAbt_Handler?A	void DAbt_Handler (void) __abort
Fast Interrupt	FIQ_Handler?A	void FIQ_Handler (void) __fiq

The SWI and IRQ exceptions are special cases, as we will see later. The ?A is used to tell the linker that the corresponding function should be compiled with the ARM instruction set ?T is used for the THUMB instruction set. Only the IRQ and FIQ interrupt sources can be disabled. The protection exceptions (Undefined instruction, Prefetch Abort, and Data abort)

are always enabled. Consequently these exceptions must always be trapped. If you do not declare a corresponding C function for these interrupt sources, then the compiler will default to using a tight loop to trap any entry to these exceptions.

Pabt_Handler:

B Pabt Handler

*Default handling of exceptions for
which no C function has been
declared*

Exercise 5: Exception Handling

In this exercise we configure a C routine to be a simple interrupt and see it working in the debugger. Later on we will see how the LPC2000 hardware is configured to service interrupts.

Software Interrupt

The Software Interrupt exception is a special case. As we have seen, it is possible to encode an integer into the unused portion of the SWI opcode.

```
#define SWIcall2      asm{ swi#2}
```

However, in the Keil CA ARM compiler, there is a more elegant method of handling software interrupts. A function can be defined as a software interrupt by using the following non ANSI keyword adjacent to the function prototype:

```
int Syscall2 (int pattern) __swi(2)
{
    .....
}
```

In addition the assembler file SWI_VEC.S must be included as part of the project.

Now when a call is made to the function an SWI instruction is used, causing the processor to enter the supervisor privileged mode and execute the code in the SWI_VEC.S file. This code determines which function has been called and handles the necessary parameter passing. This mechanism makes it very easy to take advantage of the exception structure of the ARM7 processor and to partition code which is non-critical code running in user mode, or privileged code such as a BIOS or operating system. In the tutorial section we will take a closer look at how this works.

Exercise 6: Software Interrupt

The SWI support in the Keil compiler is demonstrated in this example. You can easily partition code to run in either the user mode or in supervisor mode.

Locating Code In RAM

As we shall see later, the main performance bottleneck for the ARM7 CPU is fetching the instructions to execute from the FLASH memory. The LPC2000 has special hardware to solve this problem for the on-chip FLASH. However if you are running from external FLASH you are stuck with the access time of the external FLASH. One trick is to boot the executable code into fast RAM and then run from this RAM. This means that you need to compile position-independent code which can be copied into the RAM, or compile code so that it runs in the RAM and is loaded by a separate bootloader program. Both of these solutions will work, but require extra effort to develop. Fortunately the Keil compiler has a directive which defines a function as a RAM function. The startup code will copy the function into RAM and the linker will resolve all calls to it as being located in the defined RAM area. The function declaration is shown below

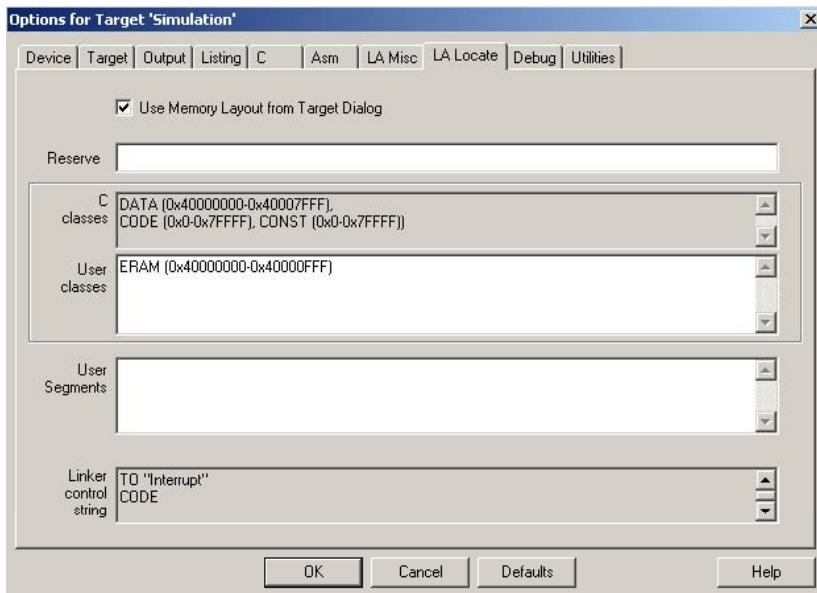
```
int      RAM_FUNCTION (int my_VAR) __ram
{
    ....
}
```

It is also necessary to define which section of memory will be used to hold these functions. This is done by declaring a section of the RAM as executable RAM or ERAM. This declaration makes use of the classes directive to allocate a region of RAM to contain all the executable RAM functions.

The basic syntax is shown below:

ERAM (0x40000000 – 0x40000FFF)

This entry should be made in the LA Locate dialogue of the options for target menu.



The compiler does not check if your RAM function is calling functions or library functions which are not also stored in the RAM. So if your “fast” RAM function makes calls to a maths routine stored in the FLASH memory, you may not get the performance you were expecting. This method of locating functions in RAM is not only simple and easy to use, it has the added advantage that the linker knows where the function will finally end up and can place the debug symbols at the correct address. This will give you not only a ROMable image which will run standalone, but also an image which can be debugged.

Inline Functions

It is also possible to increase the performance of your code by inlining your functions. The inline keyword can be applied to any function as shown below

```
void NoSubroutine (void) __inline
{
    ...
}
```

When the inline keyword is used the function will not be coded as a subroutine, but the function code will be inserted at the point where the function is called, each time it is called. This removes the prologue and epilogue code which is necessary for a subroutine, making its execution time faster. However, you are duplicating the function every time it is called, so it is expensive in terms of your FLASH memory.

Operating System Support

If you are using an operating system for the LPC2000, the OS is likely to take care of the system stacks and context switching. To avoid duplicating this by the compiler, it is possible to declare a function as a task within the operating system. This causes the compiler to just translate the code within the function and not to add the normal prologue and epilogue code which saves and restores registers to the stack. A function may be declared as a task as shown below

```
void  AnalogueSample(void)      __task
{
    ...
}
```

Fixing Objects At Absolute Locations

The compiler also allows you to fix any C object, such as a variable or a function at any absolute memory location. The compiler has an extension to the C language as shown below

```
int  checksum  __at  0x40000000;
```

Variables declared using this keyword cannot be initialised by the startup code. You must also be careful to fix variables on the correct boundaries, or you will get a memory abort. (For example if an integer is located at an uneven memory address.)

Inline Assembler

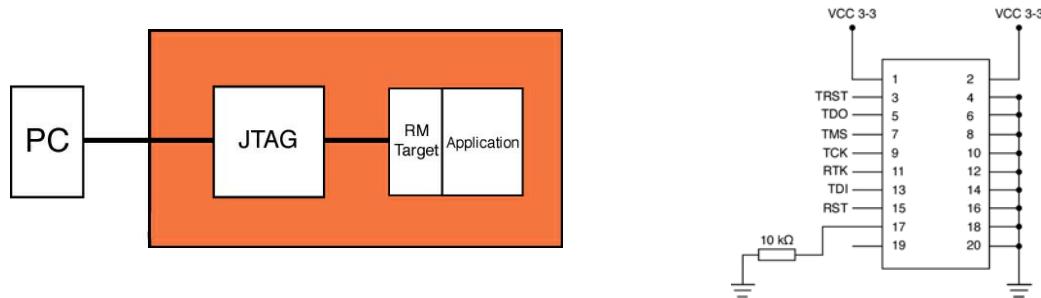
The compiler also allows you to use ARM or THUMB Assembler instructions within a C file. This can be done as shown below:

```
__asm { mov r15,r2; }
```

This can be useful if you need to use features which are not supported by the C language, for example the MRS and MSR instructions.

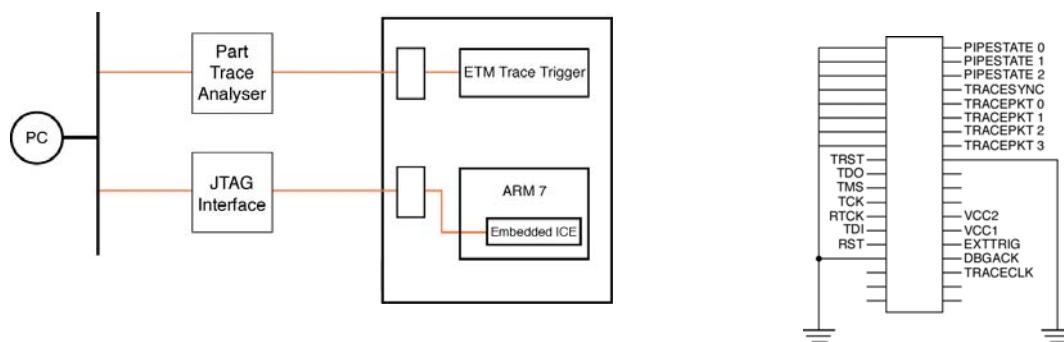
Hardware Debugging Tools

Philips have designed the LPC2000 to have the maximum on-chip debug support. There are several levels of support. The simplest is a JTAG debug port. This port allows you to connect to the LPC2000 from the PC for a debug session. The JTAG interface allows you to have basic run control of the chip. That is, you can single step lines of code, run halt and set breakpoints and also view variables and memory locations once the code is halted.



Debug support on the LPC2000 includes a JTAG port for Flash programming and basic run control debugging.

In addition, Philips has included the ARM embedded trace module. The embedded trace module provides much more powerful debugging options and real time trace, code coverage, triggering and performance analysis toolsets. In addition to more advanced debug tools, the ETM allows extensive code verification and software testing which is just not possible with a simple JTAG interface. If you are designing for safety critical applications, this is a very important consideration.



In addition to the JTAG port Philips have included the ARM ETM module for high end debugging tools

The final on-chip debug feature is the Real Time Monitor. This is a kernel of code which is resident in a reserved area of memory. During a debug session the debugger can start the real monitor via the JTAG port. The real monitor can be used to provide “on the fly” updates as your code is running. This process is pseudo real time in that the real monitor code interrupts your code and uses some processor time to read and communicate debug information to the PC.

Important!

The JTAG and ETM tools simply provide a fairly “dumb” serial debug connection to the ARM7 core. A generic ARM JTAG tool does not have any understanding of the overall LPC2000 architecture. This means that a generic tool will always enter the bootloader after reset because it does not write the “program signature” into the FLASH (this feature is discussed later) and consequently will never run your code. If you are new to the LPC2000 this is likely to catch you out and be very frustrating. Since the Keil tools are developed for ARM7 based general purpose microcontrollers MicroVision (“uVISION”) understands the LPC2000 memory architecture and will debug the device seamlessly.

Even More Important

As mentioned above, the JTAG port is a simple serial debug connection to the ARM7 device. It is very important to understand its behaviour during reset. If the ARM7 CPU is reset, all of the peripherals including the JTAG are reset. When this happens the ULINK debugger loses control of the chip and has to re-establish control once the LPC2000 device comes out of reset. This will take a finite number of clock cycles. While this is happening, any code which is on the chip will be run as normal. Once the ULINK gets back control of the chip, it performs a soft reset by forcing the PC back to address zero. However, the on-chip peripherals are no longer in the reset condition ie peripherals will be initialised, interrupt enabled etc. You must bear this in mind if the application you are developing could be adversely affected by this. A quick solution is to place a simple delay loop in the startup code or at the beginning of main(). After a reset occurs, the CPU will be trapped in this loop until the ULINK regains control of the chip. None of the application code will have run, leaving the LPC2000 in its initialised condition.

Summary

So, by the end of this section you should be able to set up a project in the Keil uVISION IDE, select the compiler and LPC2000 variant you want to use, configure the startup code, be able to interwork the ARM and THUMB instruction sets, access the LPC2000 peripherals and write C functions to handle exceptions. With this grounding we can now have a look at the LPC2000 system peripherals.

Chapter 3: System Peripherals

Outline

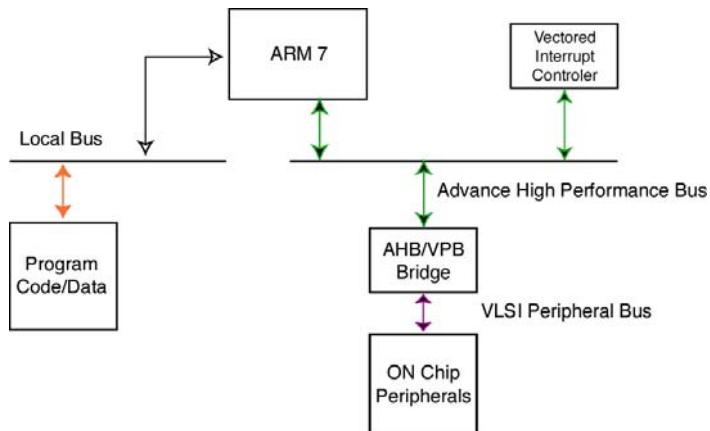
Now that we have some familiarity with the ARM7 core and the necessary development tools, we can begin to look at the LPC2000 devices themselves. In this section we will concentrate on the system peripherals, that is to say the features which are used to control the performance and functional features of the device. This includes the on-chip flash and SRAM memory, the external bus interface which is present on the LPC22xx devices, the phase locked loop which is used to multiply the external oscillator in order to provide a maximum of 60MHz processor clock and the power control features. Finally, we will take a look at the simplest user interrupt source, the external interrupt pins, before going on to look at the exception system in detail in the next section.

Bus Structure

To the programmer, the memory of all LPC2100 devices is one contiguous 32 bit address range. However, the device itself is made up of a number of buses. The ARM7 core is connected to the Advanced High performance Bus (AHB) defined by ARM. As its name implies, this is the fastest way of connecting peripheral devices to the ARM7 core. Connected to the AHB is the vector interrupt controller and a bridge to a second bus called the VLSI peripheral bus (VPB). Since the Interrupt vector controller is responsible for managing all the device interrupt sources, it is connected to the ARM7 core by the fastest bus.

All the remaining user peripherals are connected to the VPB. The VBP bridge contains a clock divider, so the VPB bus can be run at a slower speed than the ARM7 core and the AHB. This is useful for two reasons. Firstly, we can run the user peripherals at a slower clock rate than the main processor to conserve power. Secondly it gives Philips the option of adding a slower peripheral to the LPC2000 family without it becoming a bottleneck on the AHB bus. Currently all the on-chip peripherals are capable of running at 60MHz so the VPB bus can be set to the same speed as the AHB bus. It is important to note that after reset the VPB divider is set to divide down the AHB clock by four, so all the on-chip peripherals will be running at $\frac{1}{4}$ the CPU clock frequency.

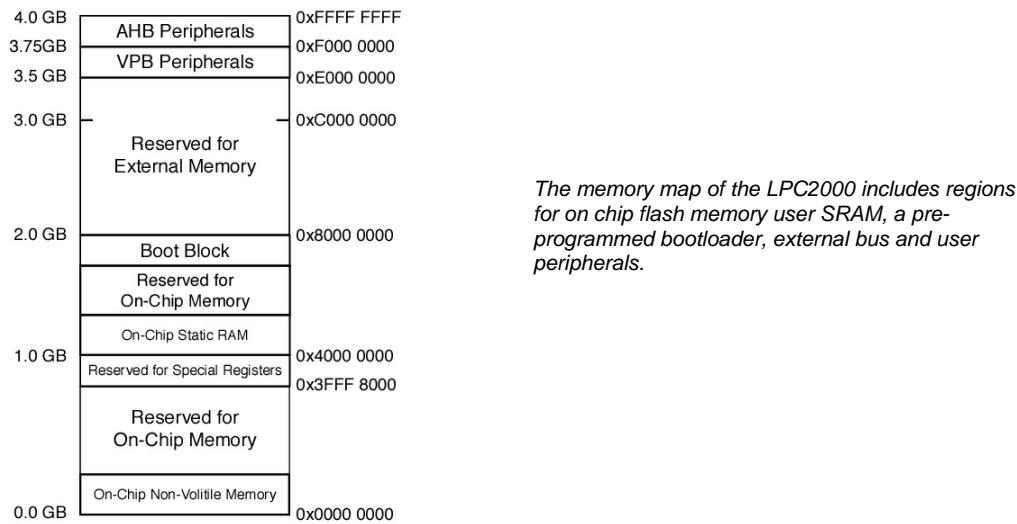
Finally, there is a third local bus which is used to connect the on-chip Flash and RAM to the CPU. Connection of the program code and data store to the ARM7 CPU via the AHB bus is possible, but this introduces some execution stalls because of contention on the bus. Using a separate local bus removes the possibility of these stalls to give the best processor performance.



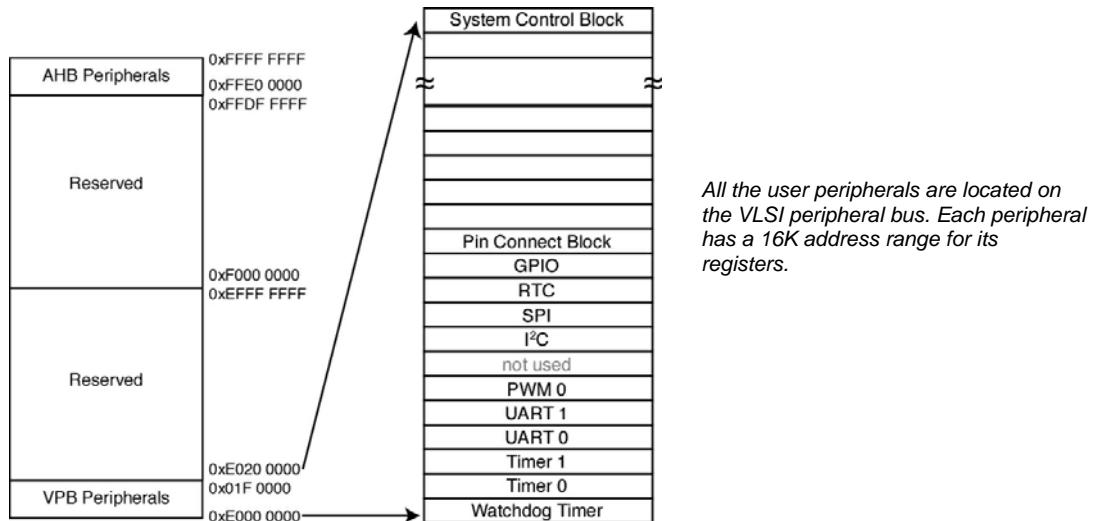
Although the LPC2000 has a linear address space there are several internal buses. It is important to be aware of the difference between them and how the performance of the processor is affected

Memory Map

Despite the number of internal buses, the LPC2000 has a completely linear memory map. The general layout is shown below.



The on-chip flash is fixed at 0x00000000 upwards with the user RAM fixed at 0x4000000 upwards. The LPC2000 is pre-programmed at manufacture with a FLASH bootloader and the ARM real monitor debug program. These programs are placed in the region 0x7FFFFFFF – 0x8000000. The region between 0x8000000 and 0xE000000 is reserved for external memory. Currently the LPC22xx devices are capable of addressing external memory via four chipselects each with a 16 Mbyte page.

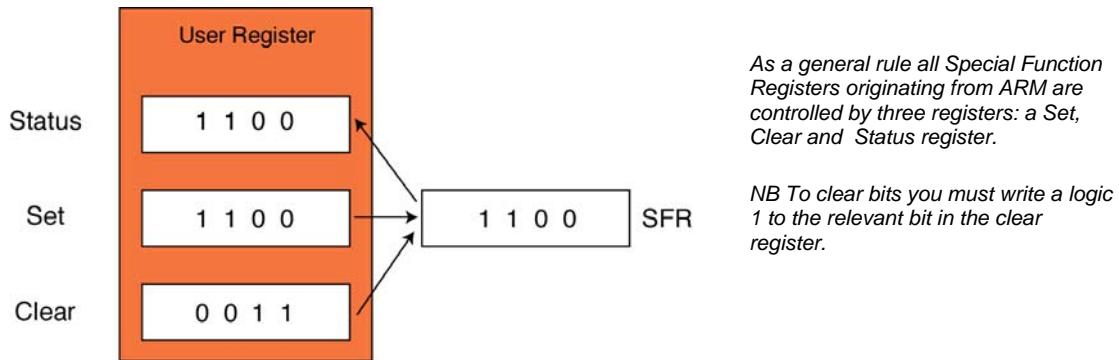


The user peripherals located on the VPB are all mapped into the region between 0xE000000 and 0xE020000 and each peripheral is allocated a 16K memory page. Finally the Vector Interrupt Unit is located at the top of the address range at 0xFFFFF000.

If your user code tries to access memory outside these regions, or non-existent memory within them, an abort exception will be produced by the CPU. This mechanism is hardwired into the design of the processor and cannot be changed or switched off.

Register Programming

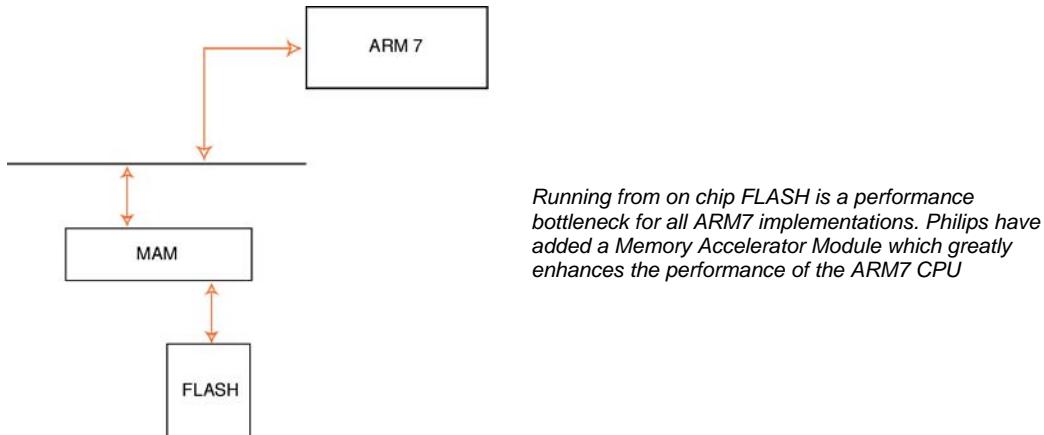
Before we start our tour through the system block, it is worth noting how Special Function Registers (SFR) are programmed on ARM7 chips.



Each underlying SFR is controlled by three user registers. A Set register which is used to set bits, a Clear register which is used to clear bits by writing a logic 1 to the bits you wish to clear and a Status register which is used to read the current contents of the register. The most common mistake made when new to the LPC2100 is to write zero into the Clear register which has no effect.

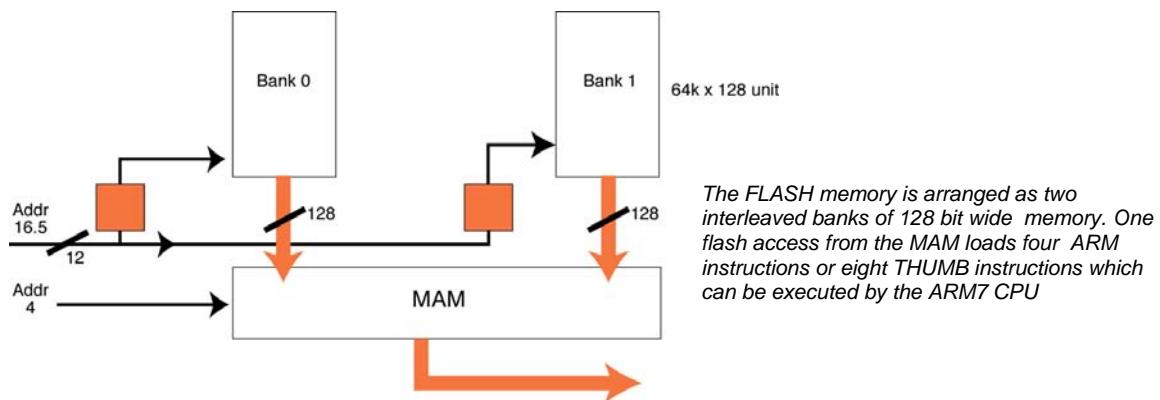
Memory Accelerator Module

The Memory Accelerator Module (MAM) is the key to the high instruction execution rate of the LPC2100 family. The MAM is present on the local bus and sits between the FLASH memory and the ARM7 CPU.



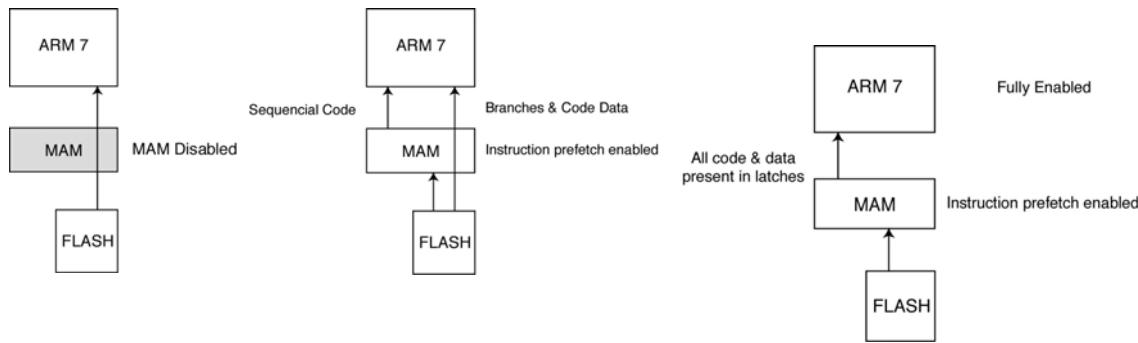
One of the main constraints in designing a high performance, single-chip microcontroller based on the ARM7 is the access time to the on-chip FLASH memory. The ARM CPU is capable of running up to 80MHz, however the on-chip FLASH has an access time of 50ns. Consequently, just running out of the FLASH would limit the execution speed to 20MHz (a quarter of the possible clock rate of the processor.) There are a number of ways round this problem. The simplest is to load the critical sections of your program into RAM and run out of RAM. As the RAM has a much faster access time, our overall performance will be greatly increased. The down side is that on-chip RAM is a finite and precious resource. Using it to hold program instructions greatly limits the size of application code which we could run. Another approach would be to have an on-chip cache. A cache is a small region of memory placed between the processor and memory store, which stores regions of recently referenced main memory. In a well-designed cache, the processor will use the cache memory whenever possible, thus reducing the bottleneck imposed by slow memory. However, a full cache is a complex peripheral which demands a high number of gates and consequently a large portion of the LPC200 die area. This flies in the face of the ARM7 design, which has simplicity as its watchword. Another downside of a full cache is that the runtime of code using the cache is no longer deterministic and could not be used by any application which required predictability and repeatability.

The Memory Accelerator Module is a compromise between the complexity of a full cache and the simplicity of allowing the processor to directly access the FLASH memory.



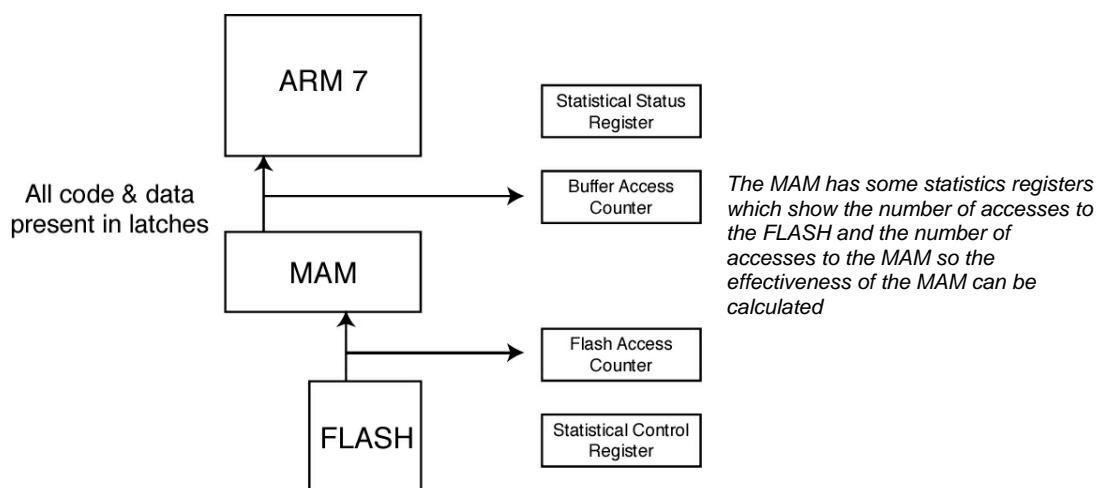
Like a cache, the MAM attempts to have the next ARM instruction in its local memory in time for the CPU to execute. First of all the FLASH memory is split into two banks which are 128 bits wide and can be independently accessed. This means that a single FLASH access can load four ARM instructions or eight THUMB instructions. User code is interleaved between the two banks, so during sequential code execution the code fetched from one bank into the MAM is being executed, while the next 128 bits of instructions from the second bank is being fetched. This ensures that it will be ready for execution once the last 128 bits has been executed. This technique works particularly well with the ARM instructions, which can use the condition codes to iron out small branches in order to keep the code-flow largely linear. In the case of small loops and jumps the MAM has branch and trail buffers that hold recently loaded instructions which can be re-executed if required.

The complexities of the MAM are transparent to the user and it is configured by two registers, the timing register and the control register. There are some additional registers to provide runtime information on the effectiveness of the MAM. The timing register is used to control the relationship between the CPU clock and the FLASH access time. By writing to the first three bits of the timing register you can specify the number of CPU clock cycles required by the MAM to access the FLASH. As the FLASH has an access time of 20 MHz and the CPU clock can be set to a maximum of 60MHz, the number of cycles required to access the FLASH is 3. So, for each three CPU cycles, we can load four instructions which keep the MAM ahead of the game. The MAM configuration register is used to define the operating mode of the MAM.



On reset the MAM is disabled and all access to code and constant data are made directly to the FLASH. It is possible to partially enable the MAM so that all sequential code is fetched from it, but branches and constant data stored in the FLASH are accessed directly from the FLASH. Finally, the MAM may be fully enabled so that it fetches all FLASH memory accesses from the MAM. The reason for these modes is that, like a cache code, running from the MAM is not deterministic, so we have the option to switch it off or reduce its impact if we need to guarantee the run time of our application code. However, even in its full operating mode the impact of the MAM is not as great as a cache. It is possible to predict runtime performance particularly with the ‘use performance analysis’ features in development tools.

To help with this analysis and also to gauge the effectiveness of the MAM, there are a group of statistical registers which can be used to measure the MAM’s performance.



The Statistics registers are based around two counters which record the accesses made to the FLASH and the accesses made to the MAM buffers. The statistical control register can further refine the type of access which will cause the counters to increment. By configuring the statistical control register we can differentiate between code constant and instruction fetches, so it is possible to determine the instruction or data hit rate or the combined instruction and data hit rate. These metrics can give us some information on the efficacy of the MAM with our application. On the CD there is a simple example which demonstrates the use of the MAM, its statistical registers and demonstrates how vital it is to the overall performance of the LPC2000 family.

Example MAM Configuration

The example code shown below starts the LPC2000 with the PLL set to 60MHz and the MAM disabled. The code FLASHes each LED in sequence with a delay loop between each increment. An A/D conversion is also done and if the result is above 0x00000080, the code enables the MAM for maximum execution speed. The effect of the MAM can be seen on the update rate of the LEDs. In the next section we will look at burning the code into the FLASH to observe its operation.

```

int main(void)
{
    unsigned int delay;
    unsigned int FLASHer = 0x00010000;           // define locals

    IODIR1 = 0x00FF0000;                      // set all ports to output
    VPBDIV = 0x02;
    ADCR   = 0x00270601;                      // Setup A/D: 10-bit AIN0 @ 3MHz
    ADCR |= 0x01000000;                        // Start A/D Conversion

    while(1)
    {
        do
        {
            val = ADDR;                         // Read A/D Data Register
        }

        while ((val & 0x80000000) == 0);
        val = ((val >> 6) & 0x03FF);

        if (val <0x80)
        {
            MAMCR = 0;
            MAMTIM = 0x03;
            MAMCR = 0x02;
        }
        else
        {
            MAMCR = 0x0;
        }
        for(delay = 0;delay<0x100000;delay++) //simple delay loop
        {
            ;
        }

        ChangeGPIOPinState(FLASHer);          //set the state of the ports
        FLASHer = FLASHer <<1;              //shift the active led

        if(FLASHer&0x01000000)
        {
            FLASHer = 0x00010000;           //Increment FLASHer led and test for
                                              // overflow
        }
    }

    void ChangeGPIOPinState(unsigned int state)
    {

        IOCLR1 = ~state;      //clear output pins
        IOSET1 = state;       //set output pins
    }
}

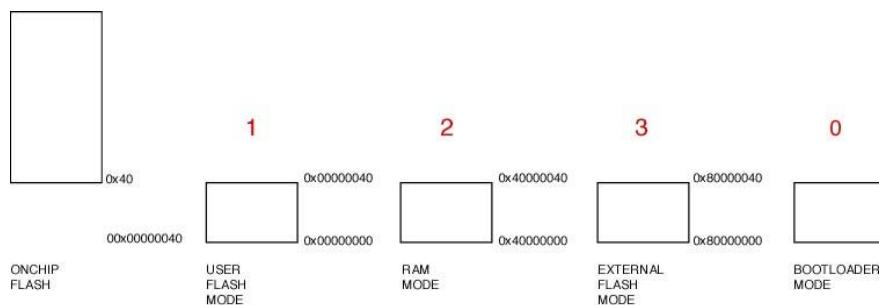
```

FLASH Memory Programming

Although the internal FLASH is arranged as two interleaved banks, you will be relieved to know that, to the user, it can be treated as one contiguous memory space and no special tools are required to prepare the code prior to programming the chip. In terms of programming the FLASH, to the user it appears as a series of 8K sectors which can be individually erased and programmed. There are several methods which can be used to program the on-chip FLASH. The easiest is by the built-in bootloader which allows your code to be downloaded via UART 0 into RAM and then be programmed into the FLASH. It is also possible to use a JTAG development tool to program the memory. This is useful during development because it can be done from the debugging environment without the need to keep switching between debugger and bootloader. Also, the JTAG connection can be very fast, up to 400Kbytes/sec download, so in large applications, particularly those using external FLASH memory, it can be the best method of production programming. Finally it is also possible to reprogram sections of the FLASH memory under command of the application already on the chip. This, in application programming, can use any method to load the new code onto the chip (SPI CAN I2C) and then load it into a given section of FLASH. So there is an easy to use mechanism which allows field updates to your application.

Memory Map Control

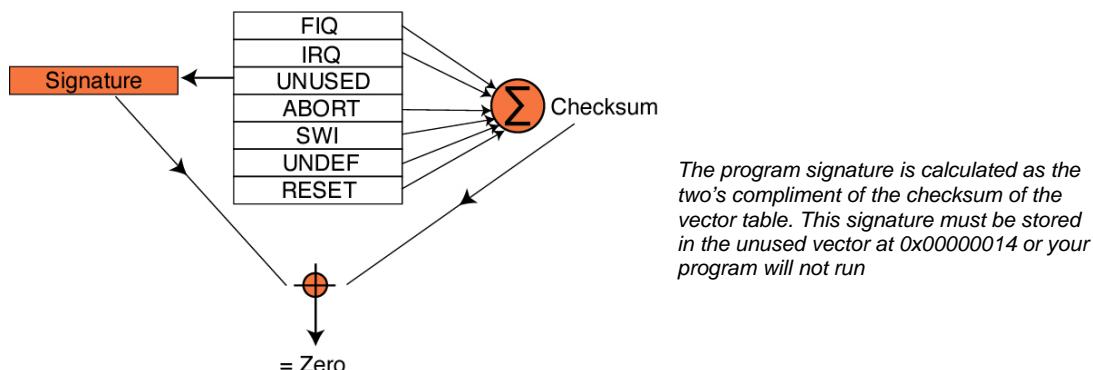
Before looking at the operation of the bootloader we must first understand the different memory modes available on the LPC2100. As we have seen, the ARM7 interrupt vector table and its constants table take up the first 64 bytes of memory. In the LPC2000 these first 64 bytes may be mapped from a number of locations, depending on the mode set in the MEMMAP register. It is important to note that these modes have nothing to do with the ARM7 operating modes. The MEMMAP register allows you to select between boot mode, FLASH mode, RAM mode and External memory mode. When selected, a new vector table will be mapped into the first 64 bytes of memory. So for the RAM mode the contents of 0x4000000- 0x400003F will be mapped to the start of memory. This allows a program to be loaded into RAM starting at 0x4000000 and the vector table can then be redirected, thus allowing the program and its interrupts to run in RAM. This mode is normally only used for debugging small programs. FLASH mode leaves the first 64 bytes of user FLASH unchanged and is the normal mode for user applications. Boot mode replaces the first 64 bytes of FLASH with the vector table for the bootloader and places a jump to the on-chip bootloader on the reset vector.



Bootloader

Every time the LPC2000 comes out of reset its memory map will be in boot mode, so the instruction on the reset vector will cause it to jump into the bootloader code entry point at 0x7FFFFFFF. This can be the bane of new users if they load their code into FLASH with a JTAG, reset and single step the first instruction only to find that the program counter is at some wild high address. If this happens, you need to program the MEMMAP register to 0x00000002, to force the chip into FLASH mode and return the user vector table.

Once the bootloader code has been entered, it will perform a number of checks to see if the FLASH needs to be programmed. First the watchdog is checked to see if the processor has had a hard reset or a soft reset. If it is a hard reset, the logic level on pin1.4 will be tested. If it is low, then the bootloader command handler will be entered. If it is a soft reset (ie watchdog timeout) or pin 1.4 is high, then there is no external request to reprogram the FLASH. However, before handing over to the user application, the bootloader will check to see if there is a valid user program in FLASH. In order to detect if a valid program is present, every user program must have a program signature. This signature is a word-wide number that is stored in the unused location in the ARM7 vector table at 0x00000014. The program signature is the two's compliment of the checksum of the ARM7 vector table



When this value is summed with the program signature the result will be zero for a valid program. If a valid program is detected, the memory operating mode is switched to FLASH, which restores the user vector table, the program counter is forced to zero and the user application starts execution. If there is no valid program, then the bootloader enters its command handler. So, without the program signature your code will never run! The program signature can be added to your startup code as shown below:

```

LDR      PC, Reset_Addr
        LDR      PC, Undefined_Addr
        LDR      PC, SWI_Addr
        LDR      PC, Prefetch_Addr
        LDR      PC, Abort_Addr

        .long    0XB8A06F58      /* Program signature */
        LDR      PC, IRQ_Addr
        LDR      PC, FIQ_Addr

```

Philips ISP Utility

If there is a valid program signature, or pin 0.14 is held low after reset, the LPC2000 will start the bootloader. Before handing over to the command handler it enters an auto-baud routine. This routine listens on UART 0 for a synchronisation character. When this is sent by the host, the LPC2000 measures the bit period and adjusts the Uart 0 baud rate generator to match the host. Once this is done some further handshaking and configuration takes place and then control is passed to the command handler.

The Bootloader command handler takes commands from UART0 in ASCII format. The command set is shown below and allows you full programming control of the FLASH. In addition the GO command is a simple debugging command which can be used to start execution of code loaded into RAM. A full description of the bootloader communication protocol is given in the LPC2000 datasheet.

ISP Command	Usage
Unlock	U <Unlock Code>
Set Baud Rate	B <Baud Rate> <stop bit>
Echo	A <Setting>
Write to RAM	W <start address> <number of bytes>
Read Memory	R <address> <number of bytes>
Prepare sector(s) for write operation	P <start sector number> <end sector number>
Copy RAM to Flash	C <Flash address> <RAM address> <number of bytes>
Go	G <address> <Mode>
Erase sector(s)	E <start sector number> <end sector number>
Blank check sector(s)	I <start sector number> <end sector number>
Read Part ID	J
Read Boot code version	K
Compare	M <address> <address2> <number of bytes>

When the bootloader has been entered it will accept commands as ASCII characters on serial port 0

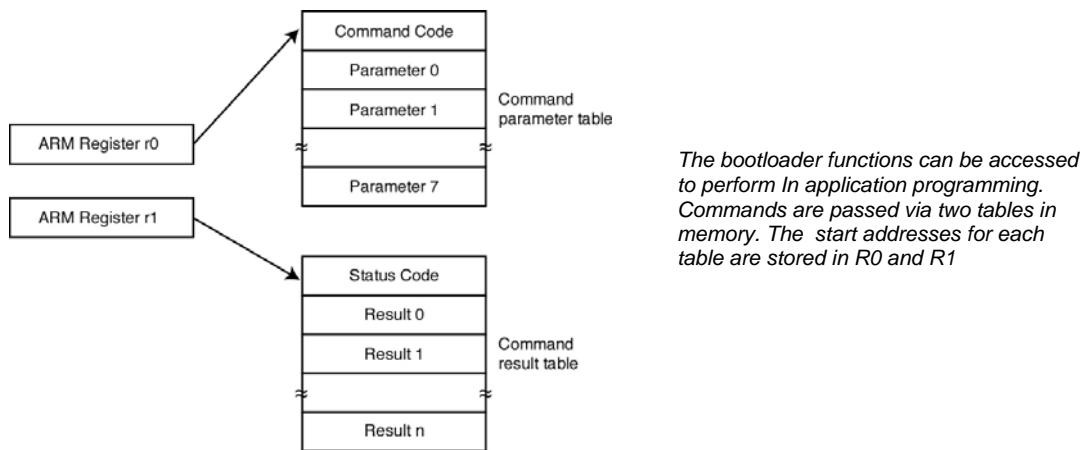
Philips provide a ready made FLASH In System Programming utility for the PC which can be used to program the development board. This tool automatically calculates and adds the program signature to your code, to ensure that your program will run. If you are using this tool to program the FLASH, your code should have a NOP instruction on the unused vector for this tool to work correctly.

Exercise 7: Memory Accelerator Module and Flash programming Utility

This exercise describes the use of the Philips Flash programming tool to load a simple program into the LPC2000. This program runs without the MAM switched on. By adjusting the A/D value the MAM is enabled so we can see the performance increase caused by this important peripheral.

In-Application Programming

It is also possible to reprogram the FLASH memory from within your program. All of the bootloader commands are available as an on-chip API and can be called by your code. To access the bootloader functions you must set up a table in RAM which contains a command code for the function you want to use followed by its parameters. The start address of this table is stored in R0. The start address of a second table which contains the status code and function results is stored in R1.



The IAP entry point is at 0x7FFFFFFF0 if you wish to call the functions from a THUMB function or at 0x7FFFFFFF1 if you wish to enter from an ARM function. The return address is expected to be stored in the link register. This convention is designed to work within the ARM procedure call standard. A method of calling the IAP routines through function pointers is detailed in the datasheet. An alternative method is shown below and both methods are used in the example program. If you are short of program space you can experiment with both methods to see which is the most efficient in your compiler.

If we define a THUMB function with three parameters as shown below.

```
void iap (unsigned *cmd, unsigned *rslt, unsigned entry)
{
    asm("mov r15,r2");
}
```

We can pass the start address of a command and result array and by the APCS convention these values will be stored in R0 and R1. We can also store the address of the entry point to the IAP routines in the next available parameter register R2. In THUMB mode we cannot program the high registers directly, but we can move low registers to high registers, hence we can move the contents of R2 directly into the program counter and initiate the requested In Application Programming routine. When the IAP routine has finished, it will return to your application code using the value stored in the link register, which is the next instruction in the function which called our void IAP (...) function. You should also note that the In Application functions return in ARM mode not THUMB. The IAP functions require the top 32 bytes of on-chip RAM, so you must either locate the stacks to start below this region so it

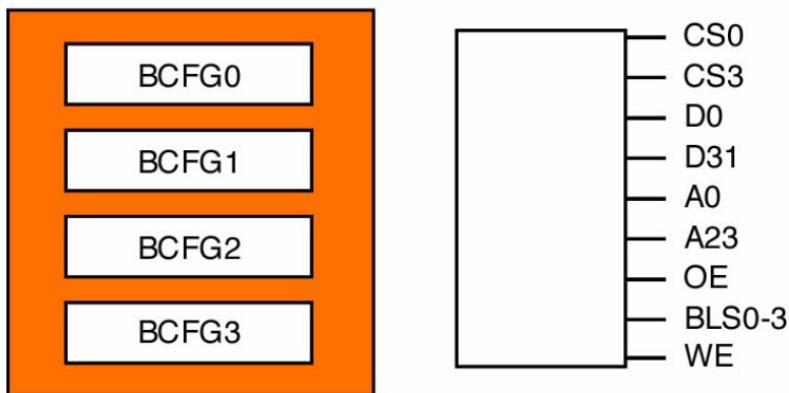
is unused, or, if you need all the RAM, place the IRQ stack at the top of memory and disable interrupts before you enter the IAP routines. Using a pointer you can now copy the top 32 bytes of on-chip SRAM into a temporary array and then restore them once you return from the IAP functions. This way you will not risk corrupting any stacked data.

External Bus Interface

The LPC22xx variants have an External Memory Controller (EMC). When enabled, the EMC provides four chipselects from 0x80000000. Each chipselect has a fixed 16Mbyte address range and a programmable wait state generation and can be programmed as an 8,16 or 32-bit wide bus. As well as allowing additional memory and peripheral devices to be interfaced to the LPC22xx devices, it is possible to boot the chip from external FLASH memory located on chip select zero.

External Memory Interface

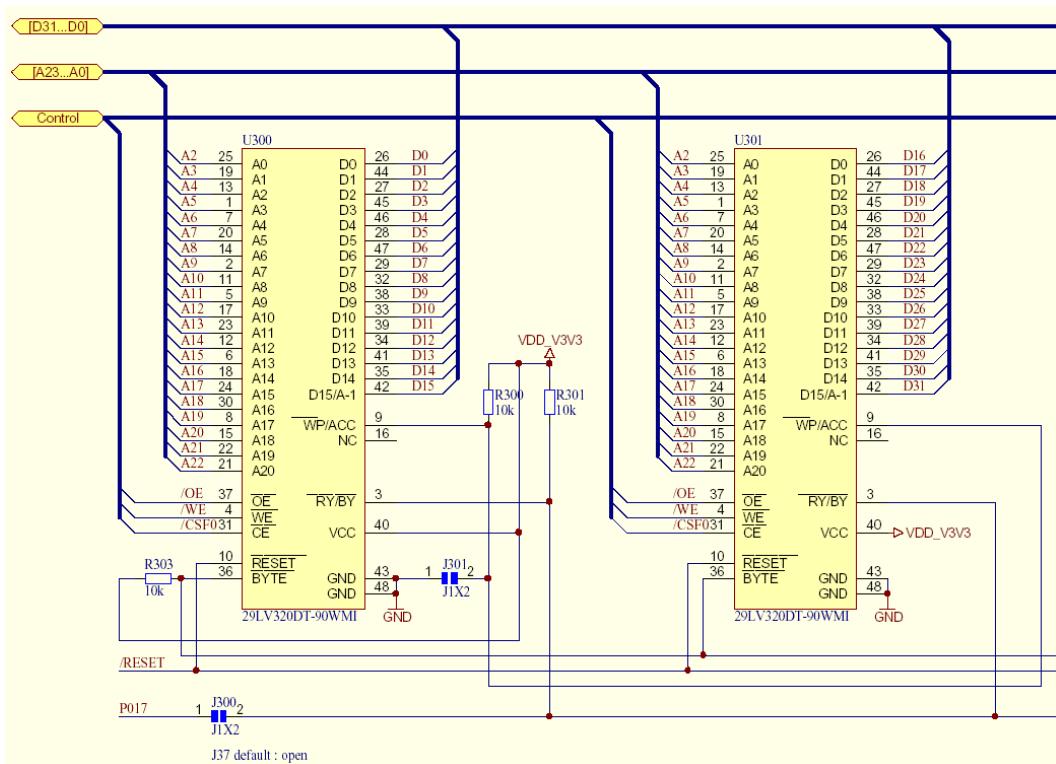
The External Memory Interface of the LPC22xx devices is shown below.



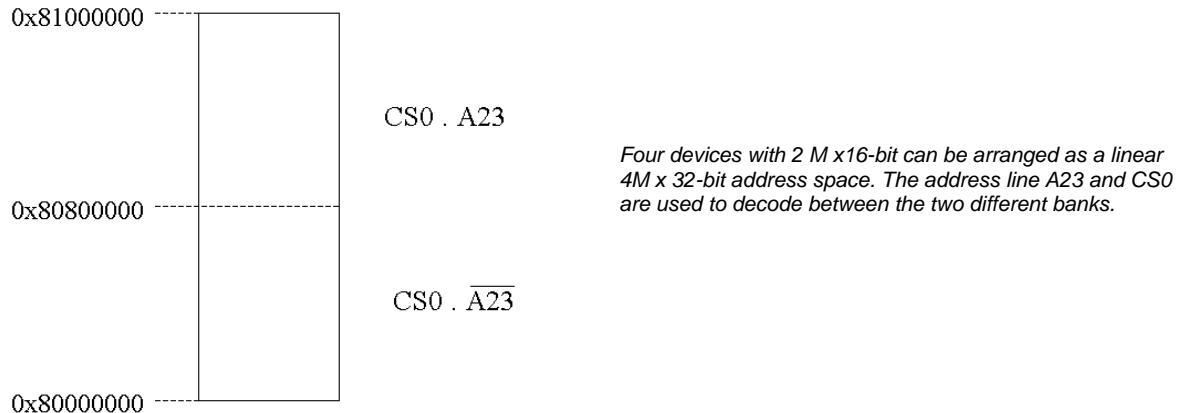
The data bus uses port 2 GPIO pins 2.0 – 2.31 and the address bus uses Port 3 GPIO pins 3.0 – 3.23. The remainder of port 3 is used for the Chipselects 1 – 3, the bytelane select pins and the write enable signal. The remaining signal Chipselect 0 and output enable are on port 1. The two boot pins are multiplexed with the databus pins D26 and D27. Depending on the state of these pins at reset, the LPC22xx variants can boot from internal FLASH or any width of memory connected to Chipselect zero. The table below shows the states the pins should be held in to boot from a particular device. These two pins are fitted with weak internal pull up resistors which ensure the device will boot from internal FLASH in its default condition.

Boot 1:0	CS0 Buswidth
00	8 Bit
01	16 Bit
10	32 Bit
11	Internal FLASH

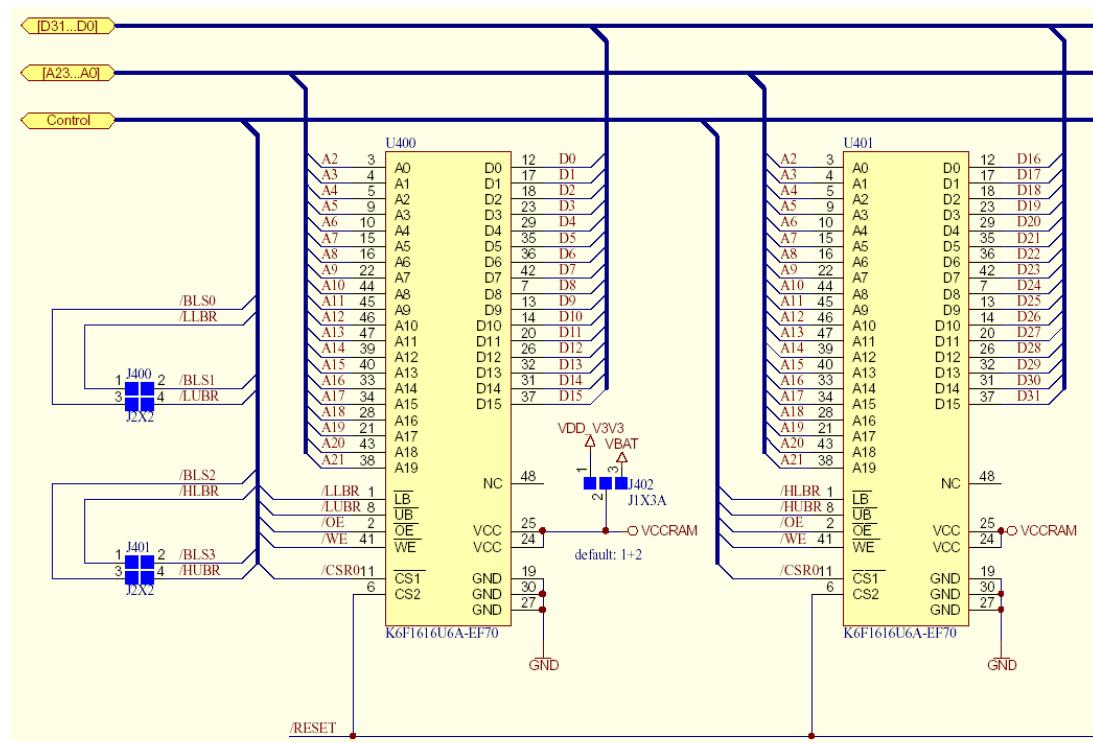
The LPC22xx datasheet shows basic schematics for the most common memory interfacing options. However, we will consider a practical example of interfacing external FLASH and static RAM onto a 32-bit bus. The FLASH memory we will use is the AMD AM29LV320DT. This is a 32 megabit FLASH memory which can be arranged as 4M by 8 bits, or 2M by 16 bits. For the RAM we will use a K6F1616U6A which is a 1M by 16 bit static RAM. Both these devices are designed for low power applications and the programming algorithm is supported by the ULINK JTAG interface. The FLASH is connected to Chipselect 0 and the RAM is connected to Chipselect 1. The schematic for each Chipselect is shown below.



Two of the 29LV320DT devices are arranged as 16-bit wide memories to give a 2M page of 32-bit wide FLASH memory. The byte# pin is pulled high on each device to enable the 16-bit mode. The FLASH device is designed to be a boot sector device and consequently has an option to protect the top and bottom sectors so that they cannot be corrupted. This feature is enabled by pulling the |WP/ACC pin low. Since we do not want this feature, the pin is pulled high allowing us to reprogram any sector of the FLASH memory. We are also not using the Ready/Busy output, so this is also tied high. The remaining control signals reset, Output enable (OE), Write Enable (WE) and Chip Enable are connected directly to the processor. As the memory is to be arranged word-wide (32 bits) we need to be able to address it every quad bytes, hence A0 and A1 are not used. If it is necessary to add more memory onto this chipselect the 29LV320 can be replaced with a XXX to give a 4M page of word-wide memory. To access the full 16 Mbyte address range, a duplicate pair of devices can be added and the chipselect gated with A23 to provide a chipselect for each half of the memory page.



The RAM is interfaced to the address bus in a similar fashion using Chipselect1 except the devices are 1 Mbyte in size so we are using A2 to A21. Further devices can be mapped in by multiplexing A22 and A23 with the chipselect line. As this is a RAM device and we may want to access it word, half-word or byte wide we can use the byte lane pins to allow access to the upper and lower bytes in each device



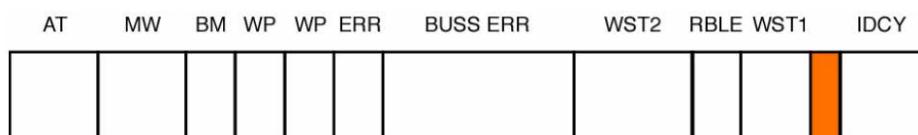
Finally the boot pins D26 and D27 must be pulled low if we want to boot from the external device.

Using The External Bus Interface

Each chipselect has a fixed address range and has a dedicated bus configuration register BCFG0 – BCFG3. The address range of each chipselect is shown below.

Chip Select	Address	Range
CS0	8000 0000	80FFF FFFF
CS1	8100 0000	81FFF FFFF
CS2	8200 0000	82FFF FFFF
CS3	8300 0000	83FFF FFFF

In our hardware example above we have mapped the FLASH onto chip select 0 at 0x80000000 and the ram onto chipselect 1 at 0x81000000. Before we can use the external memory we must setup the chipselect configuration registers.



- IDCY Controls the minimum number of idle cycles between read and write operations
- WST1 Controls the length of read accesses
- RBLE 0 for byte wide devices (drives BLS signals high)
 1 for 16/32 bit devices (drives BLS signals low)
- WST2 Controls the length of write accesses
- WPERR Set if software attempts to write to a memory bank
- WP Write protect, once set write protects a bank
- BM Enables memory bank as a Burst ROM
- MW Defines the width of the data bus
- AT, BUSS, ERR - Not used

Each of the chipselects in use must be programmed with the correct parameters to match the external device connected on to it. In the case of the FLASH memory, it has a 90ns read cycle so at 60MHz with a cycle time of 16 ns we need 6 Cclk read waitstates with one idle cycle. The FLASH is accessed word-wide, so RBLE is set to zero to disable the byte laning.

MW	Buswidth
0 0	8 Bit
0 1	16 Bit
1 0	32 Bit
1 1	Reserved

Each Chipselect may be configured with a buswidth of 8, 16 or 32 bits

During normal operation the FLASH will not be written to, so WST2 is set to zero. Also, the write protect may be set to detect accidental writes to the FLASH bank, but during development it may be wise to set it to zero and disable write protect in case it interferes with the FLASH programming algorithm of the ULINK. Finally the bus width is set to 32 bits. This gives a configuration value for Chipselect zero of 0x20000060.

In the case of the RAM it has a 70ns read and write time. Consequently at 60MHz the read and write waitstate (WST1 and WST2) should be set to 5 Cclk cycles with IDCY set to one cycle. As the RAM is a byte-partitioned device, the byte lane control must be enabled by setting RBLE to one. And again the bus width must be set to 32 bits. This gives us a chipselect configuration value of 0x20001440.

These values can be configured with the graphical editor in the Keil startup code.

<input type="checkbox"/>	Stack Configuration (Stack Sizes in Bytes)	
<input type="checkbox"/>	PLL Setup	<input checked="" type="checkbox"/>
<input type="checkbox"/>	MAM Setup	<input checked="" type="checkbox"/>
<input type="checkbox"/>	External Memory Controller (EMC)	<input checked="" type="checkbox"/>
<input type="checkbox"/>	Bank Configuration 0 (BCFG0)	<input checked="" type="checkbox"/>
	IDCY: Idle Cycles	1
	WST1: Wait States 1	6
	WST2: Wait States 2	0
	RBLE: Read Byte Lane Enable	<input type="checkbox"/>
	WP: Write Protect	<input type="checkbox"/>
	BM: Burst ROM	<input type="checkbox"/>
	MW: Memory Width	32-bit
<input type="checkbox"/>	Bank Configuration 1 (BCFG1)	<input checked="" type="checkbox"/>
	IDCY: Idle Cycles	1
	WST1: Wait States 1	5
	WST2: Wait States 2	5
	RBLE: Read Byte Lane Enable	<input checked="" type="checkbox"/>
	WP: Write Protect	<input type="checkbox"/>
	BM: Burst ROM	<input type="checkbox"/>
	MW: Memory Width	32-bit
<input type="checkbox"/>	Bank Configuration 2 (BCFG2)	<input type="checkbox"/>
<input type="checkbox"/>	Bank Configuration 3 (BCFG3)	<input type="checkbox"/>

Booting From ROM

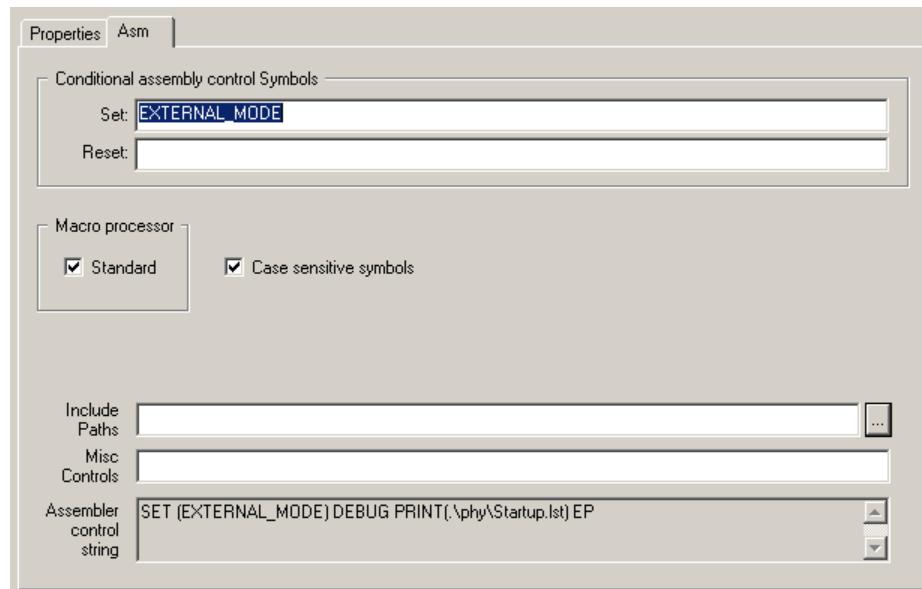
By default the LPC22xx devices will boot from their internal FLASH memory and can access the external memory once the chipselects are configured. However, if the external bootpins are pulled low, the chip will boot from external memory. In this case Chipselect zero will be enabled in the bus width selected by the boot pins. Its waitstate parameters will default to 34 Cclk cycles for WST1 and WST2 and 16 Cclk cycles for the IDCY. This ensures that the accesses on Chipselect zero will be slow enough to interface with any external device. When booting from an external device is selected, the value in the MEMMAP register will be set to 0x3 (boot from external FLASH) and the first 64 bytes of external memory on Chipselect 0 will appear at Zero. This means that you must build your

code so that the interrupt vector table and the constants table are located from address 0x80000000. In practice this means changing the start address to 0x80000000 instead of 0x00000000. In the Keil startup code this is done by an assembler directive, which is used to relocate the CODE_BASE segment containing the vector table.

```
$IF (EXTERNAL_MODE)
    CODE_BASE      EQU      0x80000000
$ELSE
    CODE_BASE      EQU      0x00000000
$ENDIF

AREA  STARTUPCODE, CODE, AT CODE_BASE // READONLY, ALIGN=4
PUBLIC_startup
```

The define EXTERNAL_MODE is declared in the assembler local options menu as shown below:



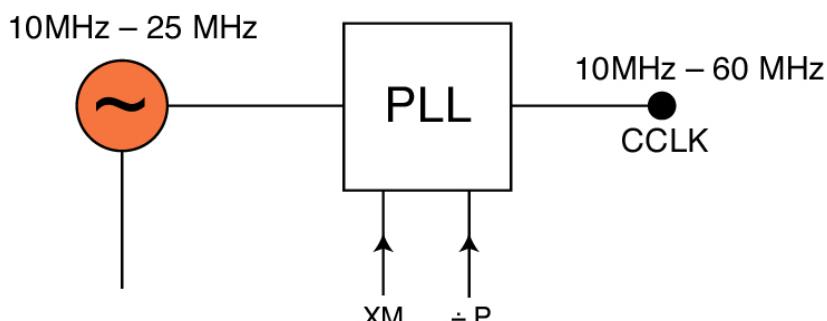
Once we have our program ready to run from external FLASH, there is a slight chicken and egg situation. In order to be able to program the external FLASH the chipselect must be configured, but to do this we must have code running on the chip. One solution would be to place a configuration program into the on-chip FLASH, boot from this and use it to configure the chipselects. However, some LPC variants are available without on-chip FLASH. Fortunately the Ulink JTAG can run a script file to setup the chipselects as required and then program the external memory.

Exercise 9: External Bus Interface

This exercise shows the necessary changes to the project we set up in exercise 1 so that it will boot and run from external memory.

Phase Locked Loop

The Phase Locked Loop is used to take an external oscillator frequency from between 10 MHz – 25MHz from a fundamental crystal and multiply this frequency up to a maximum of 60MHz to provide the on-chip clocks for the ARM7 CPU and peripherals. This allows the LPC2000 to run at its maximum frequency with a low value external oscillator, thus minimising the EMC emissions of the LPC2000. The PLL output frequency can also be changed dynamically, allowing the device to throttle back its execution speed in order to conserve power when it is idling.



The PLL is used to multiply the external crystal frequency up to the maximum 60 MHz. It is controlled by the constants M and P

Within the PLL are two constants which must be programmed in order to determine the clock for the CPU and AHB. This clock is called Cclk. The first constant is a straightforward multiplier of the external crystal. The output frequency of the PLL is given by

$$Cclk = M \times Osc$$

In the feedback path of the PLL is a current-controlled oscillator which must operate in the range 156MHz – 320 MHz.. The second constant acts as a programmable divider which ensures that the CCO is kept in specification. The operating frequency of the CCO is defined as

$$F_{CCO} = Cclk \times 2 \times P$$

On our development board there is a 12MHz oscillator so to reach the maximum CPU frequency of 60MHz

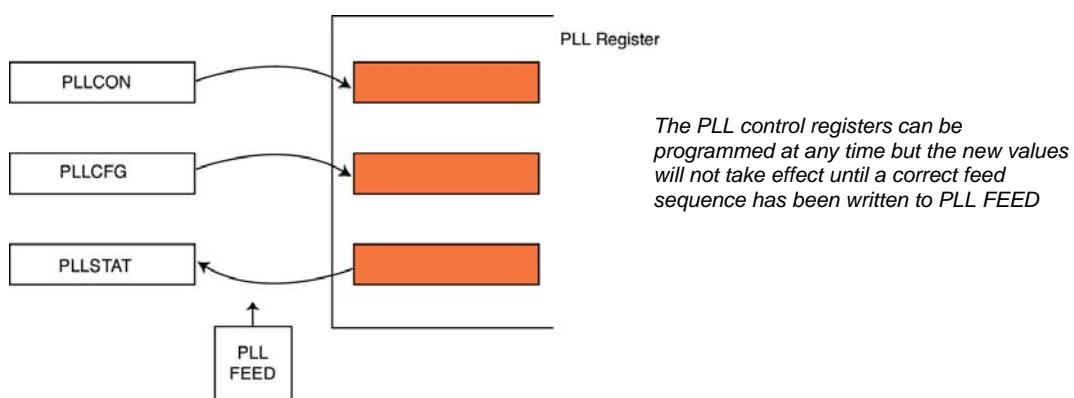
$$M = Cclk/Osc = 60/12 = 5$$

And then for P

$$156 < F_{CCO} < 320 = 60 \times 2 \times P$$

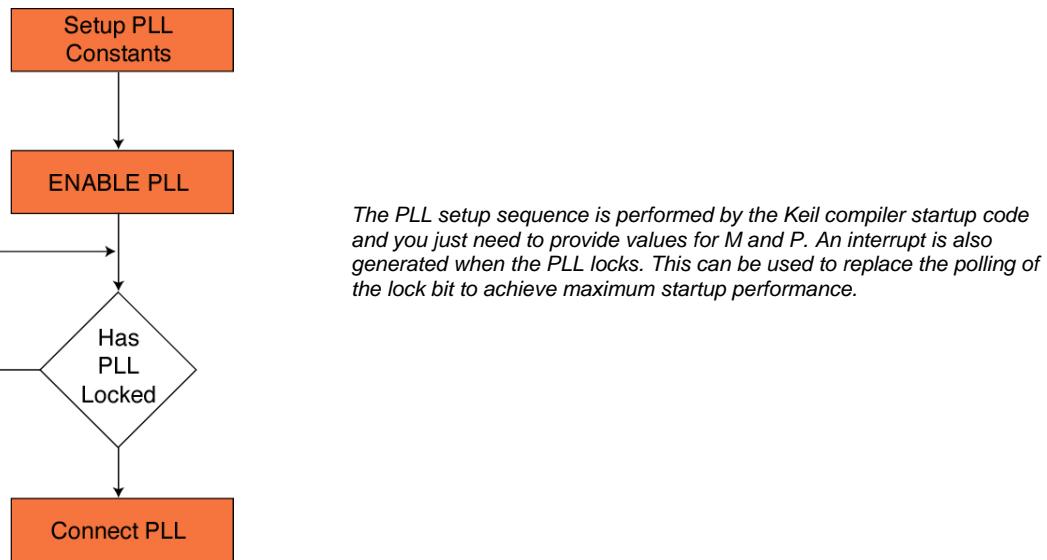
By inspection, P = 2

The programming interface for the PLL is shown below.



The values written to the user SFRs are not transferred to the internal PLL registers until a feed sequence is written to the PLL feed register. Once you have updated the PLLCON and PLLCFG registers, you must write 0x000000AA followed by 0x00000055 (the PLLFEED register). These values must be written on consecutive cycles. If you program the PLL with interrupts enabled, it is conceivable that an interrupt could occur after the first word of the sequence is written and the new PLL settings would not become effective. To set up the PLL you must write the values for P and M to the PLLCFG register. Then, using the PLLCON register, the PLL is enabled. This starts up the PLL but there is a finite startup time before it is stable enough to be used as the Cclk source. The startup of the PLL can be monitored by reading the LOCK bit in the PLLSTATUS register. Once the lock bit is set, the PLL can be used as the main clock source. Alternatively an interrupt can be generated when the PLL

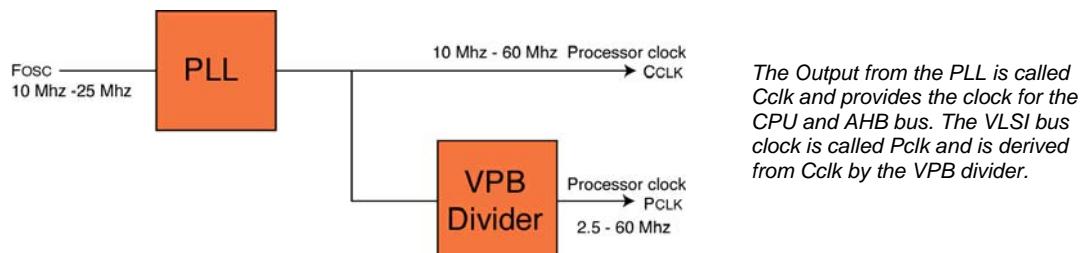
locks, so that you can carry out other tasks while the PLL starts. Once the PLL has locked as a stable clock source, it can replace the external oscillator as the source for Cclk. This is done via the PLLCON register.



Care should be taken with the values stored for the constants in the PLLCFG register. The values written to the register for the constants are P-1 and M-1, which ensures that the values of P and M in the PLL are never zero. Also the value for M is 5 bits long, so the value for P is not on an even nibble boundary. If you make a simple mistake setting up the PLL the whole chip may be running out of specification. If the chip enters power down mode, the PLL is switched off and disconnected. A wakeup from power down does not restore the PLL so the same startup sequence must be followed each time the chip exits the power down setting.

VLSI Peripheral Bus Divider

The external oscillator or the output of the PLL is used as the source for the Cclk which is the clock source for the ARM7 CPU and the AHB bus. The peripherals are on the separate VPB bus.



The clock on the VPB bus is called Pclk. This clock is derived from Cclk via the VPB bridge. The VPB bridge contains a divider which can divide down the Cclk by a factor of 1,2 or 4. The VPB divider register can be programmed by your application at any time. At reset it is set to the maximum value of four, so the Pclk is running at a quarter of the Cclk value at startup. Currently all the peripherals on the LPC2000 derivatives can run at the full 60MHz, so the VPB divider is principally used for power-saving by running the VPB clock at the slowest speed acceptable for your application.

Example Code: PLL And VPB Configuration

The code below demonstrates how to configure the PLL to give 60MHz Cclk and 30 MHz Pclk with an external crystal of 12MHz.

```
void init_PLL(void)
{
    PLLCFG = 0x00000024;           // Set multiplier and divider of PLL to
                                    // give 60.00 MHz
    PLLCON = 0x00000001;          // Enable the PLL

    PLLFEED = 0x000000AA;         // Update PLL registers with feed sequence
    PLLFEED = 0x00000055;

    while (!(PLLSTAT & 0x00000400)); // test Lock bit

    PLLCON = 0x00000003;          // Connect the PLL

    PLLFEED = 0x000000AA;         //Update PLL registers
    PLLFEED = 0x00000055;

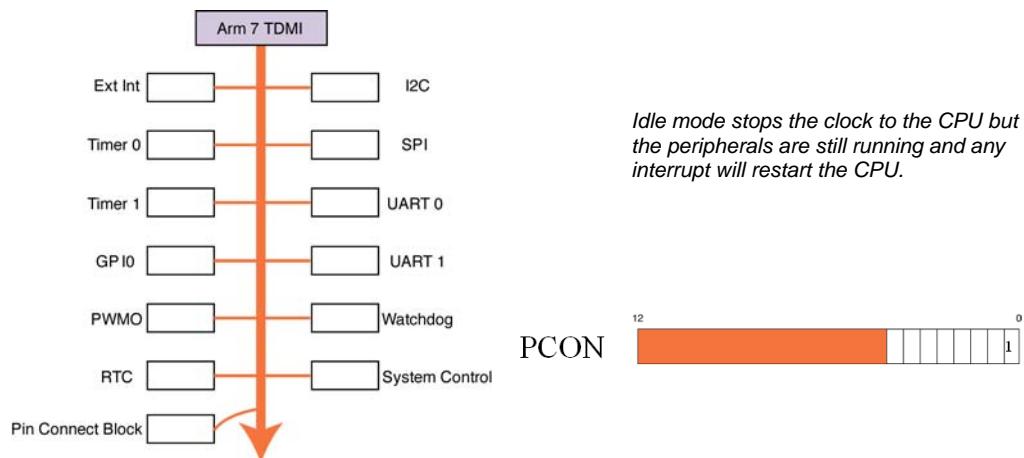
    VPBDIV = 0x00000002;          //Set the VLSI peripheral bus to 30.000MHz
}
```

Exercise 10: Phase Locked Loop

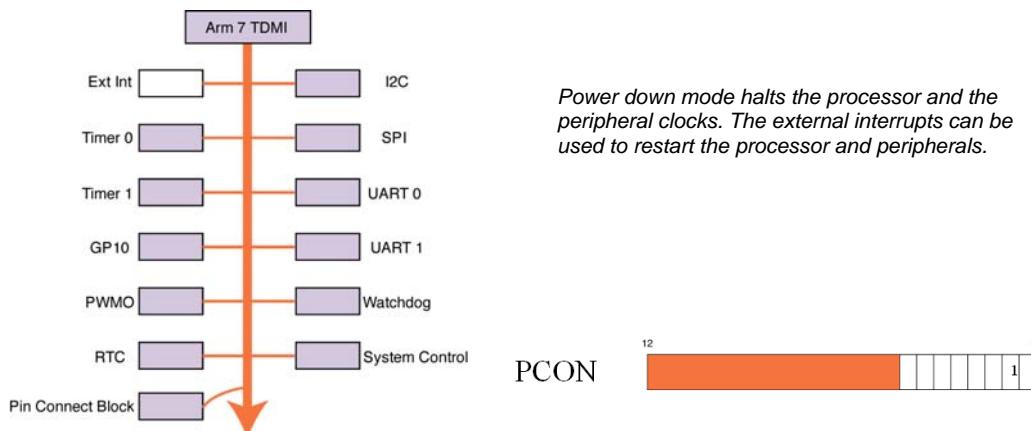
In this exercise we configure the PLL to generate a Cclk of 60MHz and a Pclk of 30MHz

Power Control

Power consumption on all (well-designed) microcontrollers is a direct relationship with the number of gates and the switching speed. The LPC2000 is no exception, The simplicity and low gate count of the ARM7 core contributes to its low power consumption. Intelligent use of the PLL and VPB divider can contribute to reducing the runtime switching speed. In addition, the LPC2000 has additional dedicated power control features. The ARM7 CPU has two power down modes controlled by the first two bits of the PCON register. The CPU may be placed into Idle mode where the CPU is halted, but the peripherals are still operational. Any interrupt from a peripheral will wake up the CPU and processing will resume.



The ARM7 can also be placed into a power down mode which halts both the CPU and the peripherals. In this mode only a reset or an interrupt generated by the external interrupt pins will cause the chip to wake up. In power down mode the oscillator is shut down. All the internal states of the processor registers and on-chip SRAM are preserved, as are the static logic levels on the I/O pins. On wake up from power down the clock source is the external oscillator and the PLL must be reconfigured.



The LPC2000 has an internal wake up timer which ensures the external oscillator is stable and the on chip memory and peripherals have initialised before the CPU starts to execute instructions. From wake up the oscillator will start to resonate. When its cycles become strong enough to drive the chip, the wake up timer will count 4096 cycles before initialising the FLASH memory and resuming program execution. This ensures the minimum restart delay after a power down or chip reset. It is also possible to power down an individual peripheral if it is not being used via its power control bit in the PCONP register. A few peripherals cannot be powered down: these are the Watchdog, GPIO, pin connect block and the system control block. Your code can optimise the configuration of the LPC2000 for minimum power consumption for a given application. Some unofficial power consumption figures are given below.

LPC2106	@60MHz	30mA
	Power down	10 – 15uA
LPC2129	@60MHz	55mA
	@60MHz VPBDIV = 4	40mA

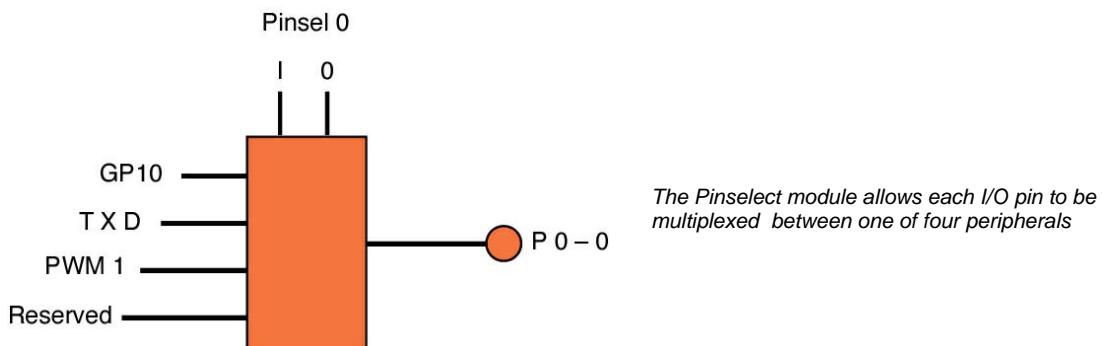
During development it is likely you will be using a JTAG development tool connected to the ARM7 via a dedicated serial link. If you place the CPU into Idle or Power Down mode no further debugging will be possible until the CPU is woken up.

LPC2000 Interrupt System

In the C code section we saw how to deal with ARM7 exceptions for an undefined instruction, a memory abort and a SWI instruction. In this section we will look at the remaining two exception sources: the General Purpose Interrupt (IRQ) and Fast Interrupt (FIQ). These two exceptions are used to handle all the interrupt sources external to the ARM7 CPU. In the case of the LPC2100 these are the user peripherals. In order to examine the LPC interrupt structure, we need a simple interrupt source. For this we can use the external interrupt pins which are the easiest peripheral to configure and EINT1 is connected to a switch on the development board which allows us to trigger an interrupt at will and observe the results with the debugger.

Pin Connect Block

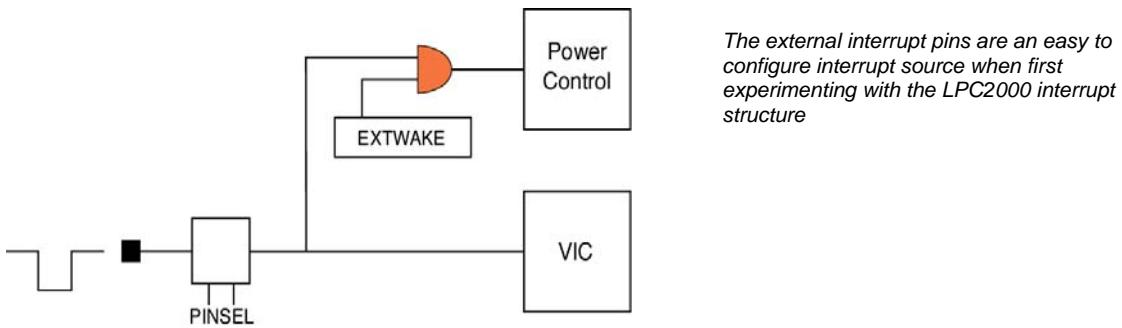
All of the I/O pins on the LPC2000 are connected to a number of internal functions via a multiplexer called the pin select block. The pin select block allows a user to configure a pin as GPIO or select up to three other functions.



On reset all the I/O pins are configured as GPIO. The secondary functions are selected through the PINSEL registers. The EINT1 interrupt line shares the same I/O pin as GPIO 0.14 and a UART1 control line. So, in order to use EINT1 we must configure the pin select register to switch from the GPIO function to EINT1.

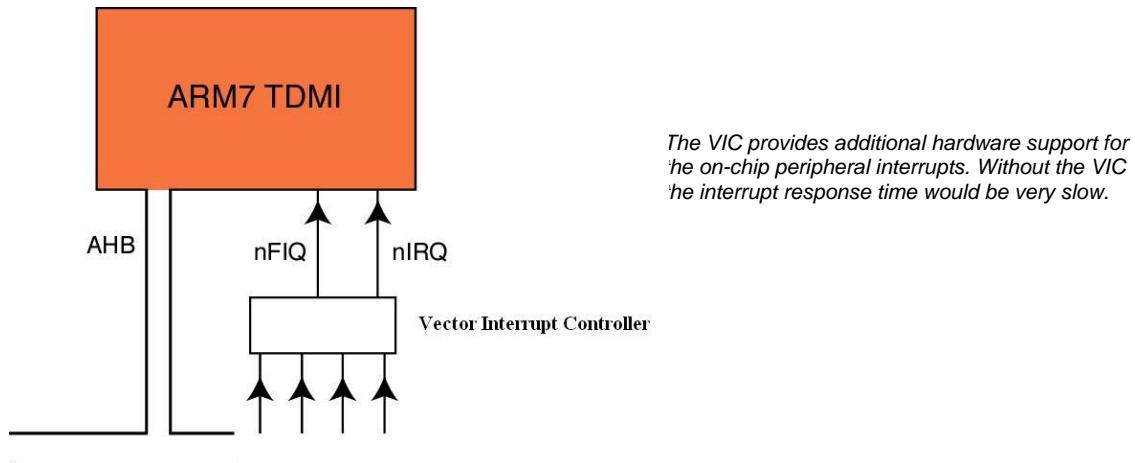
External Interrupt Pins

The external interrupts are controlled by the four registers shown below. The EXMODE register can select whether the interrupt is level or edge sensitive. If an external interrupt is configured as edge sensitive, the EXPOL register is used to qualify whether the interrupt is triggered on the rising or falling edge. In the case of level-sensitive triggering, the external interrupts can only trigger on a logic zero level. If the power down mode is being used, the EXWAKE register can enable an interrupt to wake up the CPU. So to set up a simple interrupt source program configure the EINT1 interrupt to be level sensitive and then connect it to the processor pin via the pinsel0 register.

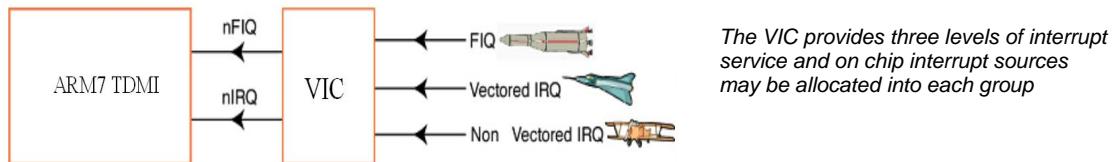


Interrupt Structure

The ARM7 CPU has two external interrupt lines for the fast interrupt request (FIQ) and general purpose interrupt IRQ request modes. As a generalisation, in an ARM7 system there should only be one interrupt source which generates an FIQ interrupt so that the processor can enter this mode and start processing the interrupt as fast as possible. This means that all the other interrupt sources must be connected to the IRQ interrupt. In a simple system they could be connected through a large OR gate. This would mean that when an interrupt was asserted the CPU would have to check each peripheral in order to determine the source of the interrupt. This could take many cycles. Clearly a more sophisticated approach is required. In order to handle the external interrupts efficiently an on-chip module called the Vector Interrupt Controller (VIC) has been added.



The VIC is a component from the ARM prime cell range of modules and as such is a highly optimised interrupt controller. The VIC is used to handle all the on-chip interrupt sources from peripherals. Each of the on-chip interrupt sources is connected to the VIC on a fixed channel: your application software can connect each of these channels to the CPU interrupt lines (FIQ, IRQ) in one of three ways. The VIC allows each interrupt to be handled as an FIQ interrupt, a vectored IRQ interrupt, or a non vectored IRQ interrupt. The interrupt response time varies between these three handling methods. FIQ is the fastest followed by vectored IRQ with non-vectored IRQ being the slowest. We will look at each of these interrupt handling methods in turn.

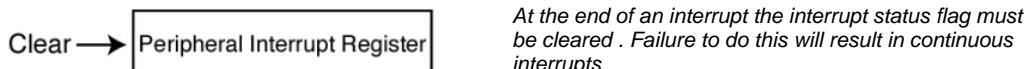


FIQ interrupt

Any interrupt source may be assigned as the FIQ interrupt. The VIC interrupt select register has a unique bit for each interrupt. Setting this bit connects the selected channel to the FIQ interrupt. In an ideal system we will only have one FIQ interrupt. However setting multiple bits in the Interrupt Select Register will enable multiple FIQ interrupt sources. If this is the case, on entry the interrupt source can be determined by examining the VIC FIQ Status register and the appropriate code executed. Clearly having several FIQ sources slows entry into the ISR code. Once you have selected an FIQ source the interrupt can be enabled in the VIC interrupt enable register. As well as configuring the VIC, the peripheral generating the interrupt must be configured and its own interrupt registers enabled. Once an FIQ interrupt is generated, the processor will change to FIQ mode and vector to 0x0000001C, the FIQ vector. You must place a jump to your ISR routine at this location in order to serve the interrupt.

Leaving An FIQ Interrupt

As we have seen, declaring a C function as an FIQ interrupt will make the compiler use the correct return instructions to resume execution of the background code at the point at which it was interrupted. However, before you exit the ISR code you must make sure that any interrupt status flags in the peripheral have been cleared. If this is not done you will get continuous interrupts until the flag is cleared. Again, be careful, as to clear the flag you will have to write a logic 1 not a logic 0.



Example Program: FIQ Interrupt

This function sets up the external interrupt as an FIQ interrupt then sits in a loop.

```
void main (void)
{
    IODIR1 = 0x00FF0000; //Set the LED pins as outputs
    PINSEL0 = 0x20000000; //Select the EINT1 function in the pin connect block
    VICIntSelect = 0x00008000; //Enable a Vic Channel as FIQ
    VICIntEnable = 0x00008000; //Enable the EINT1 interrupt in the VIC

    IOCLR1 = 0x00FF0000; // Clear the LED's

    while(1); //Loop here forever
}
```

In the startup code the FIQ interrupt routine must be added to the vector table. The address of the FIQ interrupt routine is suffixed with ?A to denote the routine as an ARM (32 Bit instruction set) routine.

```
EXTERN CODE32 (fiqint?A)
__startup PROC CODE32

Vectors:      LDR    PC,=Reset_Addr
              LDR    PC,Undef_Addr
              LDR    PC,SWI_Addr
              LDR    PC,PAbt_Addr
              LDR    PC,DAbt_Addr
              NOP
                           /* Reserved
Vector */
              LDR    PC,[PC, #-0xFF0]
              LDR    PC,FIQ_Addr           //Load the address of the FIQ
routine into thePC from the constants table

Reset_Addr:   DD     Reset_Handler
Undef_Addr:   DD     Undef_Handler
SWI_Addr:    DD     SWI_Handler
PAbt_Addr:   DD     PAbt_Handler
DAbt_Addr:   DD     DAbt_Handler
              DD     0                  /* Reserved Address */
IRQ_Addr:    DD     IRQ_Handler
FIQ_Addr:    DD     fiqint?A           // The address of the FIQ routine is
stored here
```

When the INT1 button is pressed on the MCB2100 the FIQ interrupt is generated and the code will vector to the fiqint routine. The routine is declared as an interrupt routine by using the __fiq language extension. Before exiting the ISR the peripheral flag is cleared.

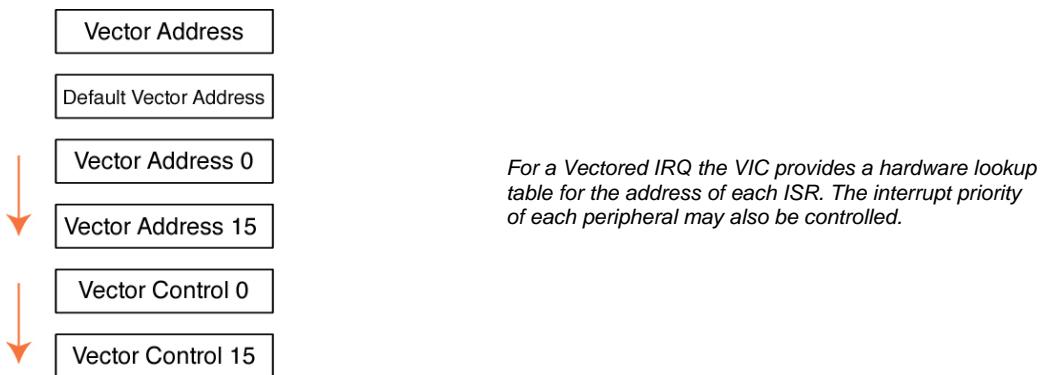
```
void fiqint (void) __fiq
{
    IOSET1 = 0x00FF0000; //Set the LED pins
    EXTINT = 0x00000002; //Clear the peripheral interrupt flag
}
```

Exercise 11: FIQ interrupt

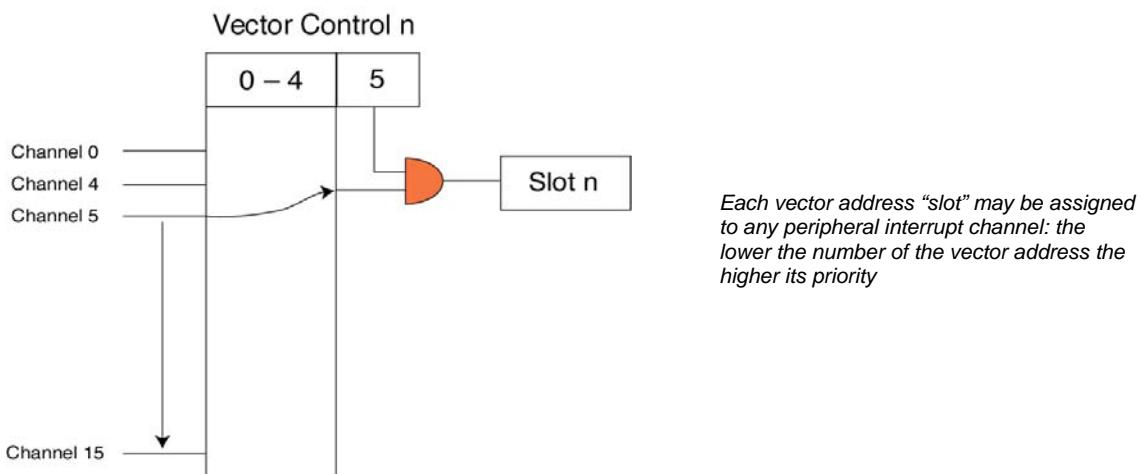
This exercise sets up the VIC to respond to an external interrupt line as a FIQ exception

Vectored IRQ

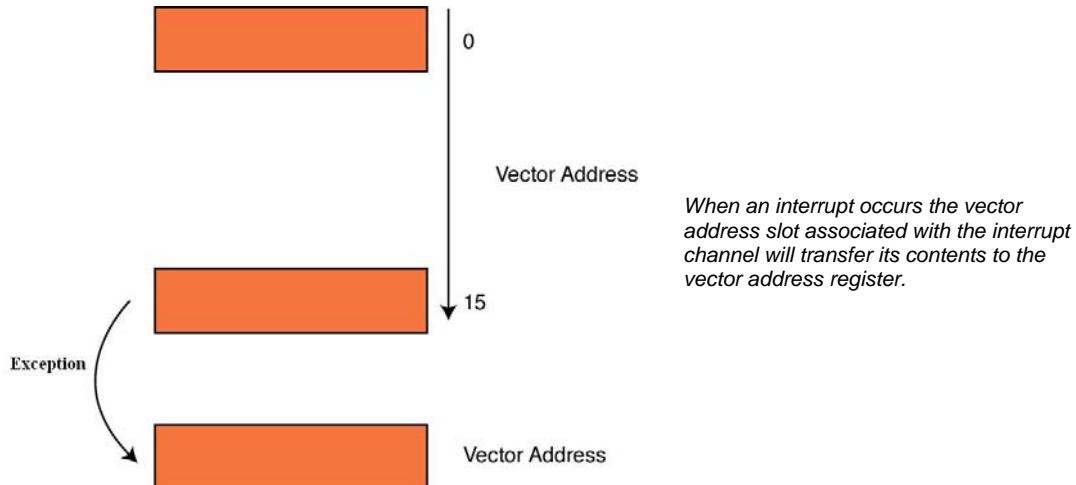
If we have one interrupt source defined as an FIQ interrupt all the remaining interrupt sources must be connected to the remaining IRQ line. To ensure efficient and timely processing of these interrupts, the VIC provides a programmable hardware lookup table which delivers the address of the C function to run for a given interrupt source. The VIC contains 16 slots for vectored addressing. Each slot contains a vector address register and a vector control register.



The Vector Control Register contains two fields: a channel field and an enable bit. By programming the channel field, any interrupt channel may be connected to any given slot and then activated using the enable bit. The priority of a vectored interrupt is given by its slot number, the lower the slot number, the more important the interrupt.



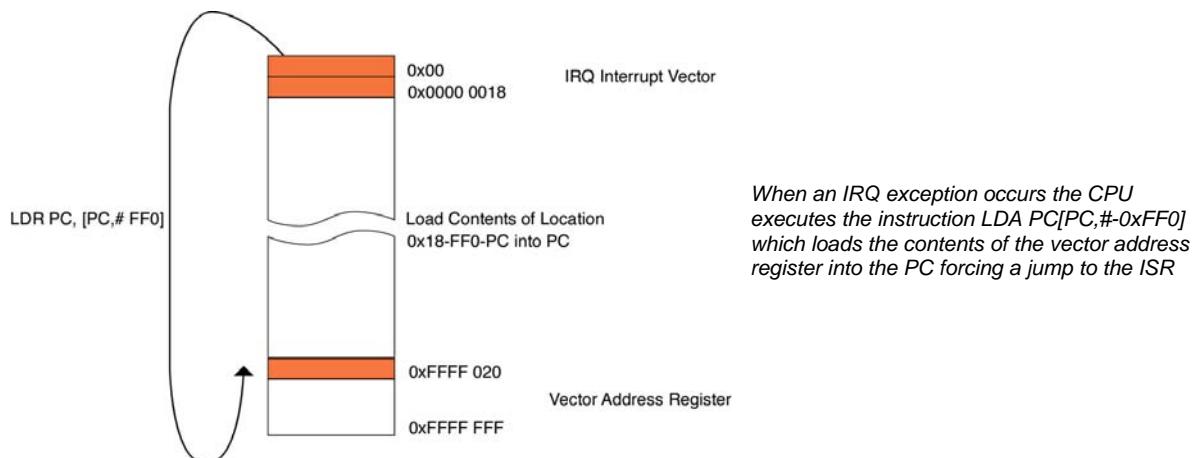
The other register in the VIC slot is the Vector Address Register. As its name suggests, this register must be initialised with the address of the appropriate C function to run when the interrupt associated with the slot occurs. In practice, when a vectored interrupt is generated the interrupt channel is routed to a specific slot and the address of the ISR in the slot's Vector Address Register is loaded into a new register called the Vector Address Register. So whenever an interrupt configured as a vectored interrupt is generated, the address of its ISR will be loaded into a fixed memory location called the Vector Address Register.



While this is happening in the VIC unit, the ARM7 CPU is going through its normal entry into the IRQ mode and will vector the 0x00000018 the IRQ interrupt vector. In order to enter the appropriate ISR, the address in the VIC Vector Address Register must be loaded into the PC. The assembly instruction shown below does this in a single cycle.

```
LDA PC, [PC, #-0xFF0]
```

As we are on the IRQ we know the address is 0x00000018 + 8 (for the pipeline). If we deduct 0xFF0 from this, it wraps the address round the top of the 32-bit address space and loads the contents of address 0xFFFFF020 (the Vector Address Register.)

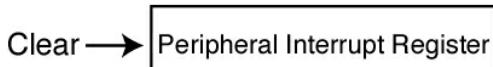


Leaving An IRQ Interrupt

As in the FIQ interrupt, you must ensure that the interrupt status flags are cleared in the peripheral which generated the request. In addition, at the end of the interrupt you must do a dummy write to the Vector Address Register. This signals the end of the interrupt to the VIC and any pending IRQ interrupt will be asserted.



At the end of a vectored IRQ interrupt you must make a dummy write to the Vector Address Register in addition to clearing the peripheral flag to clear the interrupt.



Example Program: IRQ interrupt

This example is a repeat of the FIQ example but demonstrates how to set up the VIC for a vectored IRQ interrupt.

The vector table should contain the instruction to read the VIC vector address as follows:

```
Vectors:      LDR      PC,Reset_Addr
                  LDR      PC,Undef_Addr
                  LDR      PC,SWI_Addr
                  LDR      PC,PAbt_Addr
                  LDR      PC,DAbt_Addr
                  NOP
                  LDR      PC,[PC, #-0x0FF0]      /* Vector from
VicVectAddr */
                  LDR      PC,FIQ_Addr
```

The C routines to enable the VIC and sever the interrupt are shown below:

```
void main (void)
{
    IODIR1 = 0x000FF000;           //Set the LED pins as outputs
    PINSEL0 = 0x20000000;          //Enable the EXTINT1 interrupt
    VICVectCntl0 = 0x00000002F;    //select a priority slot for a
                                  // given interrupt
    VICVectAddr0 = (unsigned)EXTINTVectoredIRQ; //pass the address
                                                // of the IRQ into
                                                // the VIC slot
    VICIntEnable = 0x00008000; //enable interrupt

    while(1);
}

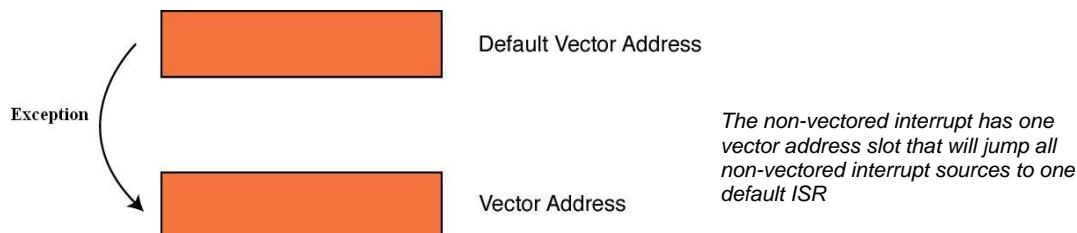
void EXTINTVectoredIRQ (void) __irq
{
    IOSET1 = 0x000FF000;           //Set the LED pins
    EXTINT = 0x00000002;           //Clear the peripheral interrupt flag
    VICVectAddr = 0x00000000;      //Dummy write to signal end
                                  //of interrupt
}
```

Exercise 12: Vectored interrupt

This exercise uses the same interrupt source as in exercise 11 but this time the VIC is configured to respond to it as a vectored IRQ exception.

Non-Vectored Interrupts

The VIC is capable of handling 16 peripherals as vectored interrupts and at least one as an FIQ interrupt. If there are more than 17 interrupt sources on the chip, any extra interrupts can be serviced as non-vectored interrupts. The non-vectored interrupt sources are served by a single ISR. The address of this ISR is stored in an additional vector address register called the default vector address register. If an interrupt is enabled in the VIC and is not configured as an FIQ or does not have a vectored interrupt slot associated with it, then it will act as a non-vectored interrupt. When such an interrupt is asserted the address in the default vector address is loaded into the vector address register, causing the processor to jump to this routine. On entry the CPU must read the IRQ status register to see which of the non-vectored interrupt sources has generated the exception.

**Leaving A Non-Vectored IRQ Interrupt**

As with the vectored IRQ interrupt, you must clear the peripheral flag and write to the vector address register.

Example Program: Non-Vectored Interrupt

```
void main (void)
{
    IODIR1 = 0x000FF000;           //Set the LED pins as outputs
    PINSEL0 = 0x20000000;         //Enable the EXTINT0 interrupt
    VICDefVectAddr = (unsigned)NonVectoredIRQ; //pass the address of the IRQ
                                                //into the VIC slot
    VICIntEnable = 0x8000;        //Enable EXTINT0 in the VIC
    while(1);
}

Vectors:      LDR      PC,Reset_Addr
              LDR      PC,Undef_Addr
              LDR      PC,SWI_Addr
              LDR      PC,PAbt_Addr
LDR      PC,DAbt_Addr
NOP
              LDR      PC,[PC, #-0xFF0]          /* Vector from VicVectAddr */
              LDR      PC,FIQ_Addr
```

```

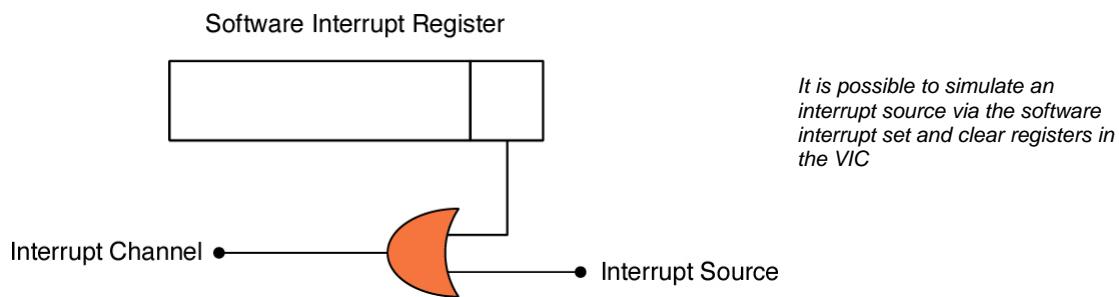
void NonVectoredIRQ (void) __irq
{
    if(VICIRQStatus&0x00008000)          //Test for the interrupt source
    {
        IOSET1      = 0x00FF0000;        //Set the LED pins
        EXTINT     = 0x00000002;        //Clear the peripheral interrupt flag
        update++;
    }
    VICVectAddr = 0x00000000; //Dummy write to signal end of interrupt
}

```

Exercise 13 : Non Vectored Interrupt

This final exercise with the VIC demonstrates how to handle a non-vectored interrupt. It is included for completeness since this mode will not normally be required.

Within the VIC it is possible for the application software to generate an interrupt on any given channel through the VIC software interrupt registers. These registers are nothing to do with the software interrupt instruction (SWI), but allow interrupt sources to be tested either for power-on testing or for simulation during development.



In addition the VIC has a protected mode which prevents any of the VIC registers from being accessed in USER mode. If the application code wishes to access the VIC, it has to enter a privileged mode. This can be in an FIQ or IRQ interrupt, or by running a SWI instruction.

Typical latencies for interrupt sources using the VIC are shown below. In the case of the non-vectored interrupts use the latency for the vectored interrupt plus the time taken to read the IRQstatus register and decide which routine to run.

- FIQ

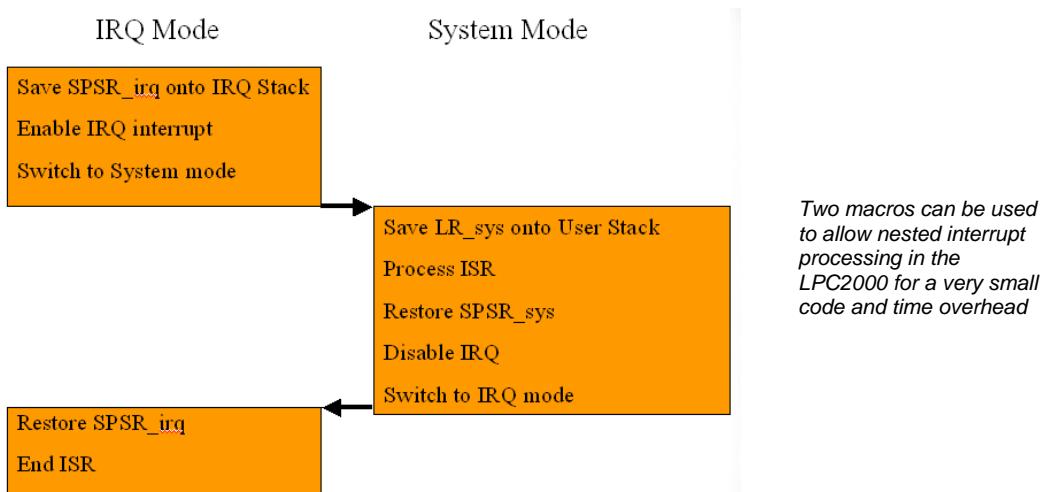
- Interrupt Sync
- + Worst Case Instruction execution
- + Entry to first Instruction
- = FIQ Latency = 12 cycles = 200 nS @ 60MHz

- IRQ

- Interrupt sync
- + worst case instruction execution
- + Entry to first instruction
- + Nesting
- = IRQ Latency = 25 cycles = 416nS @ 60MHz

Nested Interrupts

The interrupt structure within the ARM7 CPU and the VIC does not support nested interrupts. If your application requires interrupts to be able to interrupt ISRs then you must provide support for this in software. Fortunately this is easy to do with a couple of macros. Before discussing how nested interrupts work, it is important to remember that the IRQ interrupt is disabled when the ARM7 CPU responds to an external interrupt. Also, on entry to a C function that has been declared as an IRQ interrupt routine, the LR_isr is pushed onto the stack.



Once the processor has entered the IRQ interrupt routine, we need to execute a few instructions to enable nested interrupt handling. First of all the SPSR_irq must be preserved by placing it on the stack. This allows us to restore the CPSR correctly when we return to user mode. Next we must enable the IRQ interrupt to allow further interrupts and switch to the system mode (remember system mode is user mode but the MSR and MRS instructions work). In system mode the new link register must again be preserved because it may have values which are being used by the background (user mode) code so this register is pushed onto the system stack (also the user stack). Once this is done we can run the ISR code and then execute a second macro that reverses this process. The second macro restores the state of the link register, Disables the IRQ interrupts and switches back to IRQ mode finally restores the SPSR_irq and then the interrupt can be ended. The two macros that perform these operations are shown below.

```

#define IENABLE          /* Nested Interrupts Entry */
__asm { MRS      LR, SPSR      } /* Copy SPSR_irq to LR */
__asm { STMFD   SP!, {LR}     } /* Save SPSR_irq */
__asm { MSR      CPSR_c, #0x1F } /* Enable IRQ (Sys Mode) */
__asm { STMFD   SP!, {LR}     } /* Save LR */

#define IDISABLE         /* Nested Interrupts Exit */
__asm { LDMFD   SP!, {LR}     } /* Restore LR */
__asm { MSR      CPSR_c, #0x92 } /* Disable IRQ (IRQ Mode) */
__asm { LDMFD   SP!, {LR}     } /* Restore SPSR_irq to LR */
__asm { MSR      SPSR_cxsf, LR } /* Copy LR to SPSR_irq */

```

The total code overhead is 8 instructions or 32 Bytes for ARM code and execution of both macros takes a total of 230 nSec. This scheme allows any interrupt to interrupt any other interrupt. If you need to prioritise interrupt nesting then the macros would need to block low priority interrupts by disabling the lower priority interrupt sources in the VIC.

Exercise 14: Nested Interrupts

OK, one last interrupt exercise. This exercise demonstrates setting a timer to generate a regular periodic interrupt which must run. It also configures an interrupt which is triggered by Eint1. The external interrupt uses the above technique to allow the timer interrupt to run even if the external interrupt routine is active.

Summary

This is the most important chapter in this book as it describes the system architecture of the LPC2000 family. You must be familiar with all the topics in this chapter in order to be able to successfully configure the LPC2000 for its best performance and to avoid many of the common pitfalls which trap people new to this family of devices.

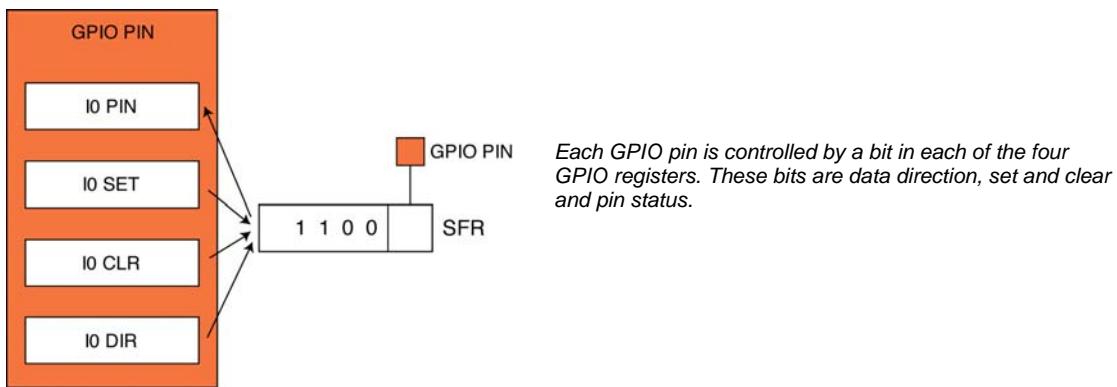
Chapter 4: User Peripherals

Outline

This chapter presents each of the user peripherals in turn. The examples show how to configure and operate each peripheral. Once you are familiar with how the peripherals work the example code can be used as the basis for a set of low-level drivers.

General Purpose I/O

On reset the pin connect block configures all the peripheral pins to be general purpose I/O (GPIO) input pins. The GPIO pins are controlled by four registers, as shown below.



The IODIR pin allows each pin to be individually configured as an input (0) or an output (1). If the pin is an output the IOSET and IOCLR registers allow you to control the state of the pin. Writing a ‘1’ to these registers will set or clear the corresponding pin. Remember you write a ‘1’ to the IOCLR register to clear a pin not a ‘0’. The state of the GPIO pin can be read at any time by reading the contents of the IOPIN register. A simple program to flash the LED on the evaluation board is shown below.

```
int main(void)
{
    unsigned int delay;
    unsigned int flasher = 0x00010000;           // define locals

    IODIR1 = 0x00FF0000;                         // set all ports to output

    while(1)
    {
        for(delay = 0;delay<0x10000;delay++)      //simple delay loop
        {
            ;
        }
        IOCLR1 = ~flasher;                      //clear output pins
        IOSET1 = flasher;                        //set the state of the ports

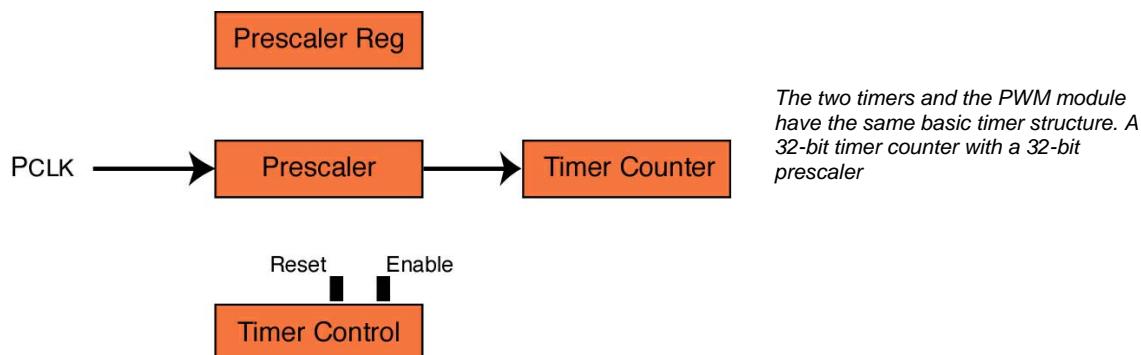
        flasher = flasher <<1;                  //shift the active led
        if(flasher&0x01000000) flasher = 0x00010000; //Increment flasher
                                                //led and test for
                                                //overflow
    }
}
```

Exercise 15: GPIO

This simple exercise demonstrates using the GPIO as an LED chaser program.

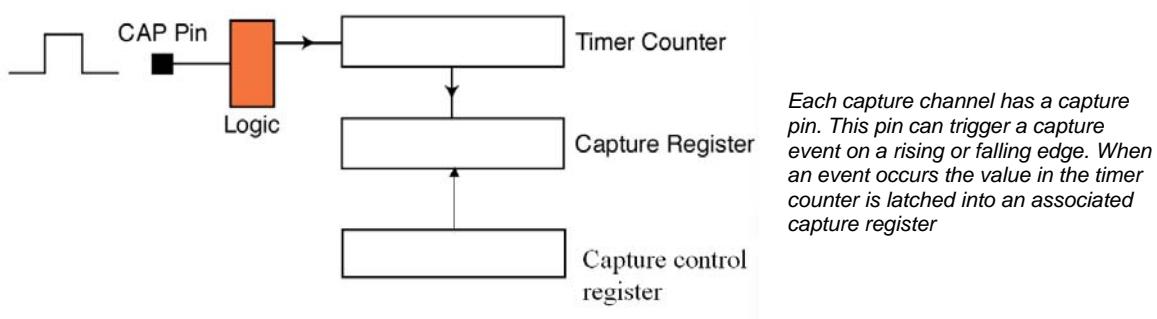
General Purpose Timers

The LPC2000 have a number of general purpose timers. The exact number will vary depending on the variant, but there are at least two timers. All of the general purpose timers are identical in structure but vary slightly in the number of features supported. The timers are based around a 32-bit timer counter with a 32-bit prescaler. The clock source for all of the timers is the VLSI peripheral clock PCLK



The tick rate of the timer is controlled by the value stored in the prescaler register. The prescale counter will increment on each tick of Pclk until it reaches the value stored in the prescaler register. When it hits the prescale value the timer counter is incremented by one and the prescale counter resets to zero and starts counting again. The Timer control register contains only two bits which are used to enable/disable the timer and reset its count.

In addition to the basic counter each timer has up to four capture channels. The capture channels allow you to capture the value of the timer counter when an input signal makes a transition.



Each capture channel has an associated capture pin which can be enabled via the pin connect block. The Capture control register can configure if a rising or falling edge, or both, on this

pin will trigger a capture event. When the capture event occurs, the current value in the timer counter will be transferred into the associated capture register and if necessary an interrupt can be generated. The code below demonstrates how to configure a capture channel. This example sets up a capture event on a rising edge on pin 0.2 (Capture 0.0) and generates an interrupt.

```
int main(void)
{
    VPBDIV      = 0x00000002;           //Set pclk to 30 MHz
    PINSEL0     = 0x00000020;          //Enable pin 0.2 as capture channel0
    T0PR        = 0x00007530;          //Load prescaler for 1 Msec tick
    T0TCR       = 0x00000002;          //Reset counter and prescaler
    T0CCR       = 0x00000005;          //Capture on rising edge of channel0
    T0TCR       = 0x00000001;          //enable timer

    VICVectAddr4 = (unsigned)T0isr;      //Set the timer ISR vector address
    VICVectCntl4 = 0x00000024;          //Set channel
    VICIntEnable = 0x00000010;          //Enable the interrupt

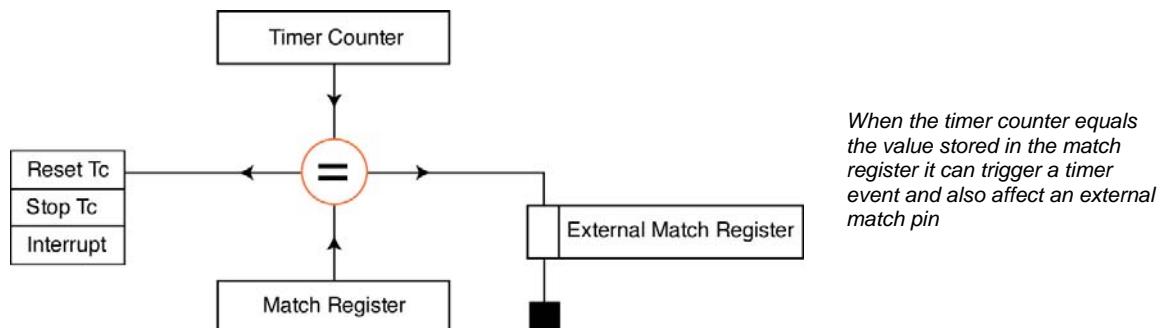
    while(1);
}

void T0isr (void)  __irq
{
    static int value;
    value        = T0CR0;             // read the capture value
    T0IR        |= 0x00000001;         //Clear match 0 interrupt
    VICVectAddr4 = 0x00000000;         //Dummy write to signal end of interrupt
}
```

Exercise 16: Timer Capture.

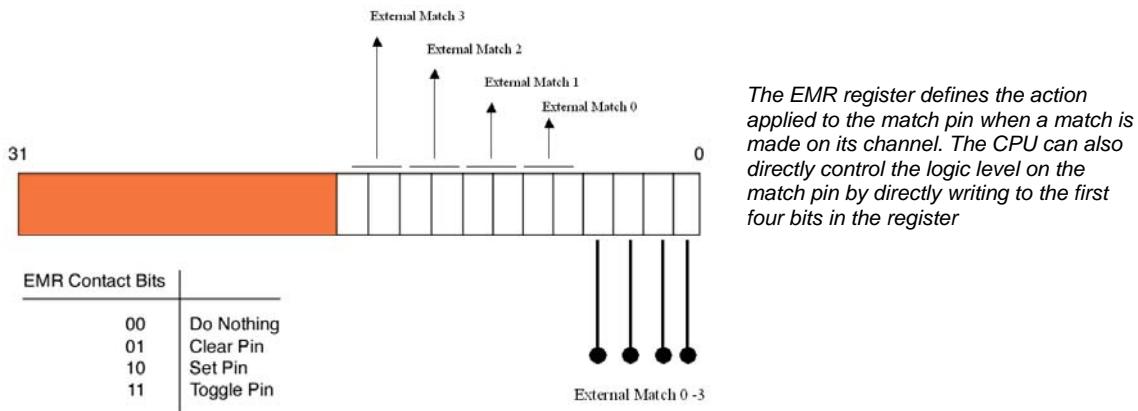
This exercise configures a general purpose timer with a capture event to measure the width of a pulse applied to a capture pin.

Each timer also has up to four match channels. Each match channel has a match register which stores a 32-bit number. The current value of the timer counter is compared against the match register. When the values match an event is triggered. This event can perform an action to the timer (reset, stop or generate interrupt) and also affect an external pin (set, clear,toggle).



To configure the timer for a match event, load the match register with the desired value. The internal match event can now be configured through the Match Control Register. In this register each channel has a group of bits which can be used to enable the following actions on a match event: generate a timer interrupt, reset the timer or stop the timer. Any

combination of these events may be enabled. In addition, each match channel has an associated match pin which can be modified when a match event occurs. As with the capture pins, you must first use the pin connect block to connect the external pin to the match channel. The match pins are then controlled by the first four bits in the external match register.



The external match register contains a configuration field for each match channel. Programming this field decides the action to be carried out on the match pin when a match event occurs. In addition, each match pin has a bit that can be directly programmed to change the logic level on the pin.

The example below demonstrates how to perform simple pulse width modulation using two match channels. Match channel zero is used to generate the period of the PWM signal. When the match event occurs the timer is reset and an interrupt is generated. The interrupt is used to set the Match 1 pin high. Match channel 1 is used to control the duty cycle. When the match 1 event occurs the Match 1 pin is cleared to zero. So by changing the value in the Match 1 register it is possible to modulate the PWM signal

```

int main(void)
{
    VPBDIV = 0x00000002;           // Configure the VPB divi
    PINSEL0 |= 0x00000800;         // Match1 as output
    T0PR  = 0x0000001E;           //Load presaler
    T0TCR = 0x00000002;           //Reset counter and presale
    T0MCR = 0x00000003;           //On match reset the counter and generate an
                                  //interrupt
    T0MR0 = 0x00000010;           //Set the cycle time
    T0MR1 = 0x00000008;           // Set 50% duty cycle
    T0EMR = 0x00000042;           //On match clear MAT1 and set MAT1 pin high for
                                  // first cycle
    T0TCR = 0x00000001;           //Enable timer
    VICVectAddr4 = (unsigned)T0isr; //Set the timer ISR vector address
    VICVectCnt14 = 0x00000024;     //Set channel
    VICIntEnable |= 0x00000010;     //Enable the interrupt

    while(1);
}

```

```

void T0isr (void) __irq
{
    TOEMR |= 0x00000002;           // Set MAT1 high for beginning of the cycle
    TOIR |= 0x00000001;           // Clear match 0 interrupt
    VICVectAddr = 0x00000000;    // Dummy write to signal end of interrupt
}

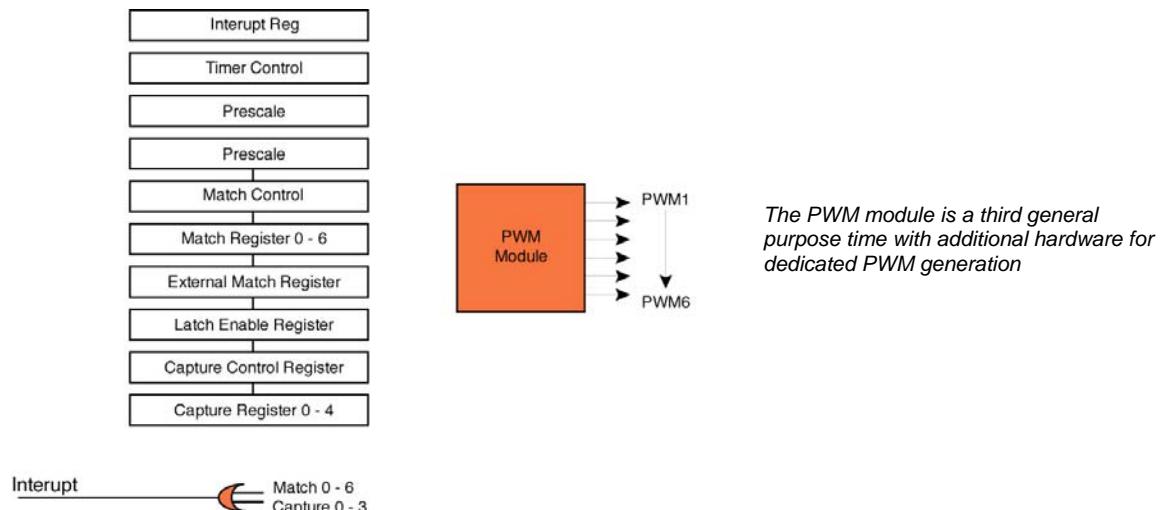
```

Exercise 17: Timer Match

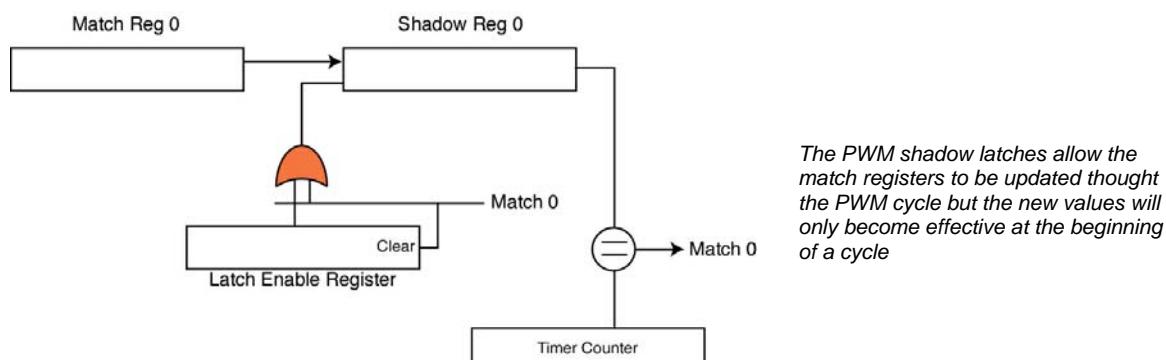
This second timer exercise uses two match channels to generate a PWM signal, there is some CPU overhead in the timer interrupt routine.

PWM Modulator

At first sight the PWM modulator looks a lot more complicated than the general purpose timers. However it is really an extra general purpose timer with some additional hardware. The PWM modulator is capable of producing six channels of single edge controlled PWM or three channels of dual edge controlled PWM.

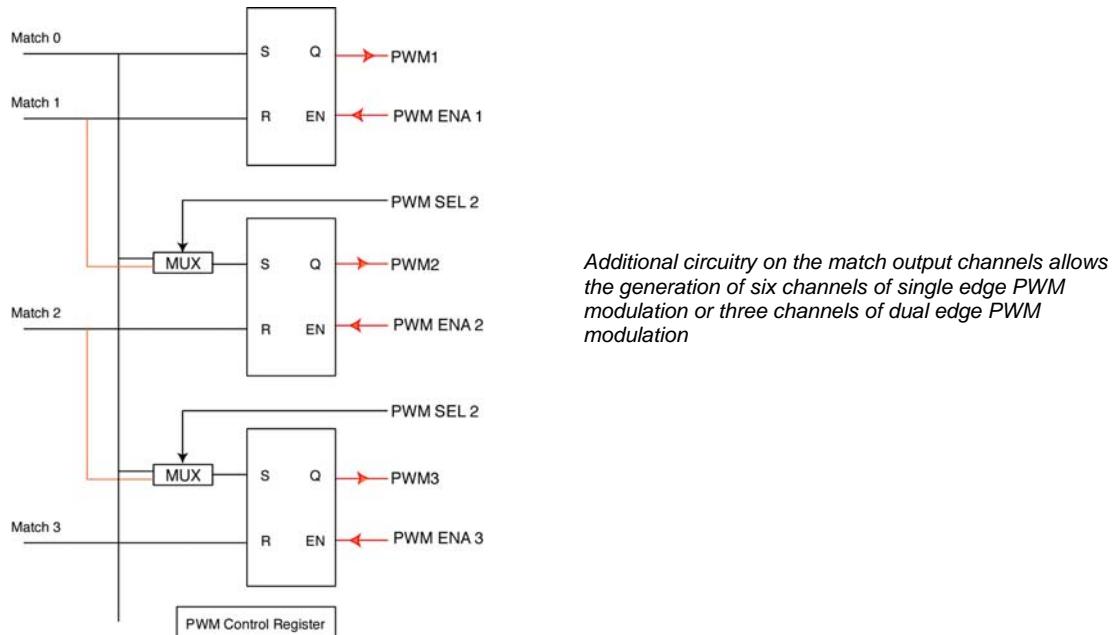


In the general purpose timers when a new value is written to a match register the new match value becomes effective immediately. Unless care is taken in your software this may be part way through a PWM cycle. If you are updating several channels, the new PWM values will take effect at different points in the cycle and may cause unexpected results. The PWM modulator has an additional shadow latch mechanism which allows the PWM values to be updated on the fly, but the new values will only take effect simultaneously at the beginning of a new cycle.

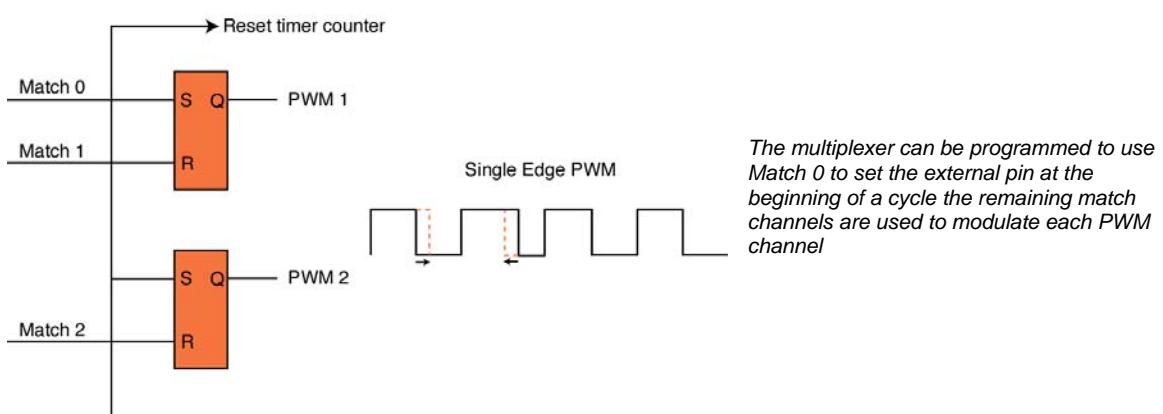


The value in a given match register may be updated at any time but it will not become effective until the bit corresponding to the match channel is set in the Latch Enable register (LER). Once the LER is set, the value in the match register will be transferred to the shadow register at the beginning of the next cycle. This ensures that all updates are done simultaneously at the beginning of a cycle. Apart from the shadow latches the PWM modulator match channels function in the same way as the timer match registers.

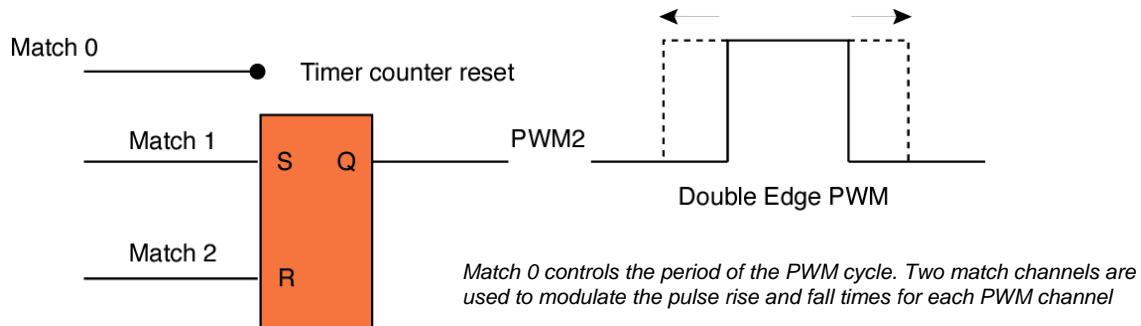
The second hardware addition to the PWM modulator over the basic timers is in the output to the device pins. In place of the match channels directly controlling the match output pin are a series of SR flip-flops



This arrangement of SR flip-flop and multiplexers allows the PWM modulator to produce either single edge or dual edge controlled PWM channels. The multiplexer is controlled by the PWMSEL register and can configure the output stage in one of two configurations. The first arrangement is for single edge modulation



Here the multiplexer is connecting Match 0 to the S input of each flip-flop and each of the remaining channels are connected to the R input. With this scheme Match 0 is set up to count to total cycle period. At the end of the cycle it will reset the counter and set match 0 high. This causes all the flip-flops to be set at the beginning of the cycle. The output Q goes high raising all the output pins high. Modulation of the PWM signal is done with the remaining match channels. Each PWM channel has an associated match channel which is connected to the R input of the flip-flop. When the match is made the flip-flop is reset and the PWM pin is set low. This allows modulation of the PWM signal by changing the value of the dedicated match channel.



By reprogramming the multiplexer the output stage of the PWM modulator can be configured to dual edge controlled modulation. In this configuration Match 0 is not connected to any output and is used solely to reset the timer at the end of each PWM period. In this configuration the S and R inputs to each flip-flop have a dedicated Match channel. At the beginning of a cycle the PWM output is low. The rising edge of the pulse is controlled by the Match channel connected to the S input and the falling edge is controlled by the Match channel connected to the R input. The example below illustrates how to configure the PWM module for dual edge PWM .

```
void main(void)
{
    PINSEL0 |= 0x00028000;      //Enable pin 0.7 as PWM2
    PWMPMR = 0x00000001;        //Load prescaler

    PWMPCR = 0x0000404;        //PWM channel 2 double edge control, output enabled
    PWMMCR = 0x00000003;        //On match with timer reset the counter
    PWMMR0 = 0x00000010;        //set cycle rate to sixteen ticks
    PWMMR1 = 0x00000002;        //set rising edge of PWM2 to 2 ticks
    PWMMR2 = 0x00000008;        //set falling edge of PWM2 to 8 ticks
    PWMILER = 0x00000007;        //enable shadow latch for match 0 - 2
    PWMEMR = 0x00000280;        //Match 1 and Match 2 outputs set high
    PWMTCR = 0x00000002;        //Reset counter and prescaler
    PWMTCR = 0x00000009;        //enable counter and PWM, release counter from reset

    while(1)                  // main loop
    {
        //.....      //Modulate PWMMR1 and PWMMR2
    }
}
```

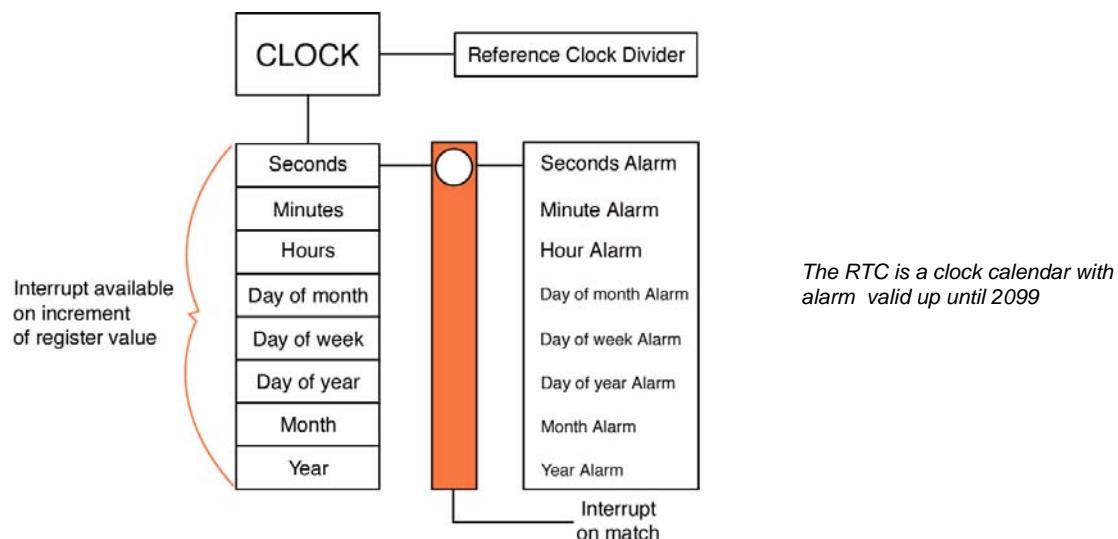
One important line to note is that the PWMEMR register is used to ensure the output of the match channel is logic 1 when the match occurs. If this register is not programmed correctly the PWM scheme will not work. Also the PWM modulator does not require any interrupt to make it work unlike the basic timers.

Exercise 18: Centre-Aligned PWM

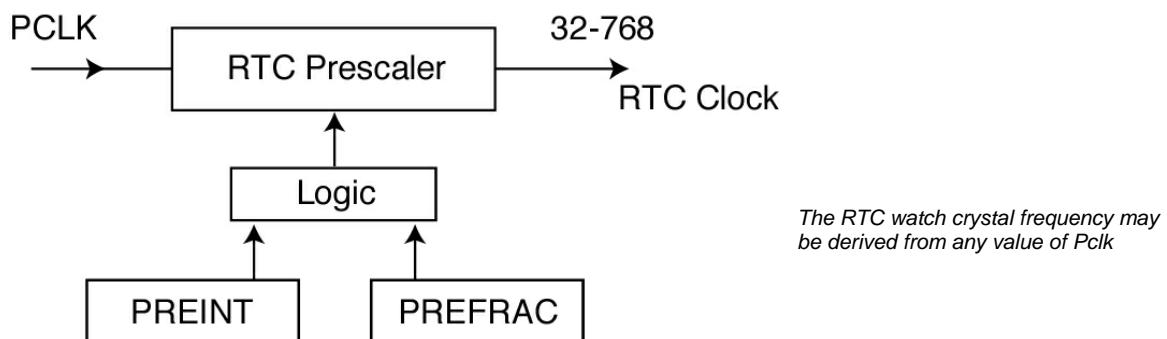
This exercise configures the PWM unit to produce a centre aligned PWM signal without any CPU overhead.

Real Time Clock

The LPC2xxx Real time clock (RTC) is a clock calendar accurate up to 2099. Like all the other peripherals the RTC runs off the PCLK so an additional external oscillator is not required. The RTC is designed to be an ultra low power peripheral and through use of the LPC2xxx low power modes is suitable for running off batteries. As well as providing a clock calendar, the RTC has a set of alarm registers that can be used to trigger a particular date and time or on a specific value held in a time-count register.



The RTC clock runs on a standard 32.7KHz clock crystal frequency. In order to derive this frequency the Pclk is connected to the reference clock divider. In effect this is a prescaler which can accurately divide any Pclk frequency to produce the required 32KHz frequency.



To ensure that the RTC clock can be accurately derived from any Pclk the prescaler is more complicated than the general purpose timer prescalers. The prescaler is programmed by two

registers called PREINT and PREFRAC. As their name implies, these hold integer and fractional divisor values. The equations used to calculate the load values for these registers are as follows:

```
PREINT = (int)(pclk/32768)-1
PREFRAC = pclk - ((PREINT+1) x 32768)
```

So for a 30MHz Pclk:

```
PREINT = (int)( 30,000,000/32768)-1 = 914
```

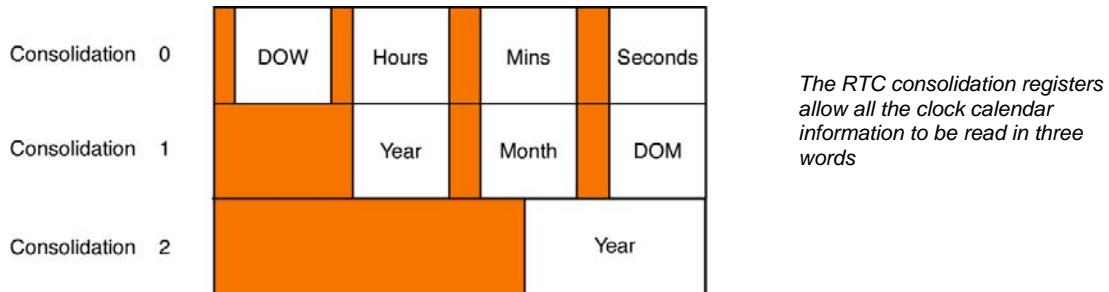
Then

```
PREFRAC = 30,000,000 - ((914+1) x 32768) = 17280
```

These values can be programmed directly into the RTC prescaler registers and the RTC is then ready to run. Just enable the clock in the clock control register and the time counters will start.

```
PREINT      = 0x00000392;          //Set RTC prescaler for 30.000 MHz Pclk
PREFRAC    = 0x00004380;
CCR        = 0x00000001;          //Start the RTC
```

There are eight time-counter registers, each of which contains a single time quantity which can be read at any time. In addition there are a set of consolidation registers which present the same time quantities in three words, allowing all the time information to be read in just three operations.



As well as maintaining a clock, the RTC can also generate alarm events as interrupts. There are two interrupt mechanisms. You can program the RTC to generate an interrupt when any time-counter register is incremented, so you could generate an interrupt every second when the second counter is updated, or once a year when the year counter is incremented. The counter increment interrupt register allows you to enable an increment interrupt for each of the eight time-counter registers.

The second method for generating an RTC interrupt is with the alarm registers. Each time-counter register has a matching Alarm register. If the matching Alarm register is unmasked it is compared to the time counter register. If all the unmasked alarm registers match the time counter registers then an interrupt is generated. So it is possible to set an alarm between now and 2099 with one second's accuracy. The Alarm Mask register controls which alarm registers are used in the compare. As both the increment and alarm events can generate an RTC interrupt it is necessary to distinguish between them from within the interrupt. The Interrupt location register provides two flags which can be interrogated to see what caused

the RTC interrupt. Again, remember that these flags must be cleared to cancel the interrupt. An RTC program which sets the clock and uses both styles of interrupt is shown below.

```

int main(void)
{
    VPBDIV = 0x00000002;
    IODIR1 = 0x00FF0000; // set LED ports to output
    IOSET1 = 0x00020000;
    PREINT = 0x00000392;      //Set RTC prescaler for 30MHz Pclk
    PREFRAC = 0x00004380;
    CIIR = 0x00000001;        //Enable seconds counter interrupt
    ALSEC = 0x00000003;        //Set alarm register for 3 seconds
    AMR = 0x000000FE;         //Enable seconds Alarm
    CCR = 0x00000001;         //Start the RTC

    VICVectAddr13 = (unsigned)RTC_isr; //Set the timer ISR vector address
    VICVectCntl13 = 0x0000002D;           //Set channel
    VICIntEnable = 0x00002000;           //Enable the interrupt

    while(1);
}

void RTC_isr(void)
{
    unsigned led;

    if(ILR&0x00000001)                  //Test for RTC counter interrupt
    {
        led = IOPIN1;                  //read the current state of the IO pins
        IOCLR1 = led&0x00030000; //Clear the illuminated LED
        IOSET1 = ~led&0x00030000; //Set the idle LED
        ILR = 0x00000001;            //Clear the interrupt register
    }

    if(ILR & 0x00000002)
    {
        IOSET1 = 0x00100000;        //Set LED 0.7
        ILR = 0x00000002;           //clear the interrupt register
    }
    VICVectAddr = 0x00000000;           //Dummy write to signal end of interrupt
}

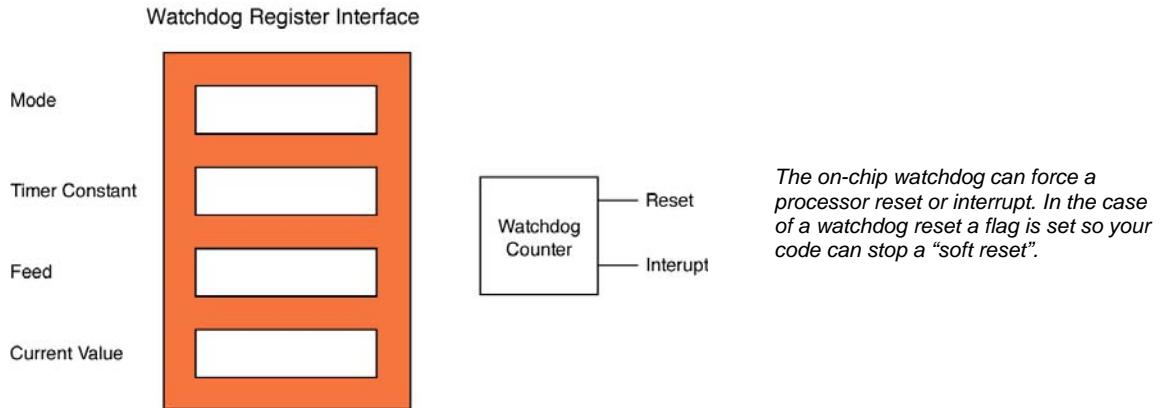
```

Exercise 19: Real Time Clock

This exercise configures the RTC and demonstrates both the alarm and increment interrupts.

Watchdog

In common with many microcontrollers the LPC2xxx family has a watchdog system to provide a method of recovering control of a program that has crashed.

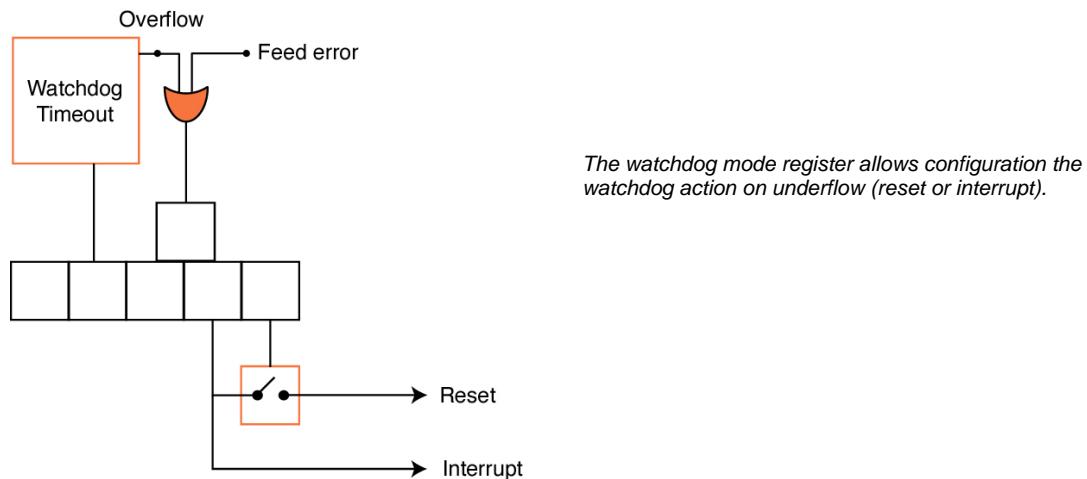


The watchdog has four registers as shown above. The watchdog timeout period is set by a value programmed into the Watchdog Constant Register (WDTCR). The timeout period is determined by the following formula

$$Wdperiod = Pclk \times WDTC \times 4$$

The minimum value for WDTC is 256 and the maximum is 2^{32} . Hence the minimum watchdog period at 60MHz is 17.066us and the maximum is just under 5 minutes.

Once the watchdog constant is programmed the operating mode of the watchdog can be configured. The Watchdog mode register contains three enable bits controlling: whether the watchdog generates an interrupt, whether it generates a reset and a final bit which is used to enable operation of the watchdog.



The Mode register also contains two flags, the WDTOF is set when the watchdog times out and is only cleared after an external hard reset. This allows your startup code to detect if the reset event was a power on reset or a reset due to a program error. The Mode register also contains the watchdog interrupt flag. This flag is read-only, but it must be read in order to clear the watchdog interrupt. If you need to debug code with the watchdog active you should

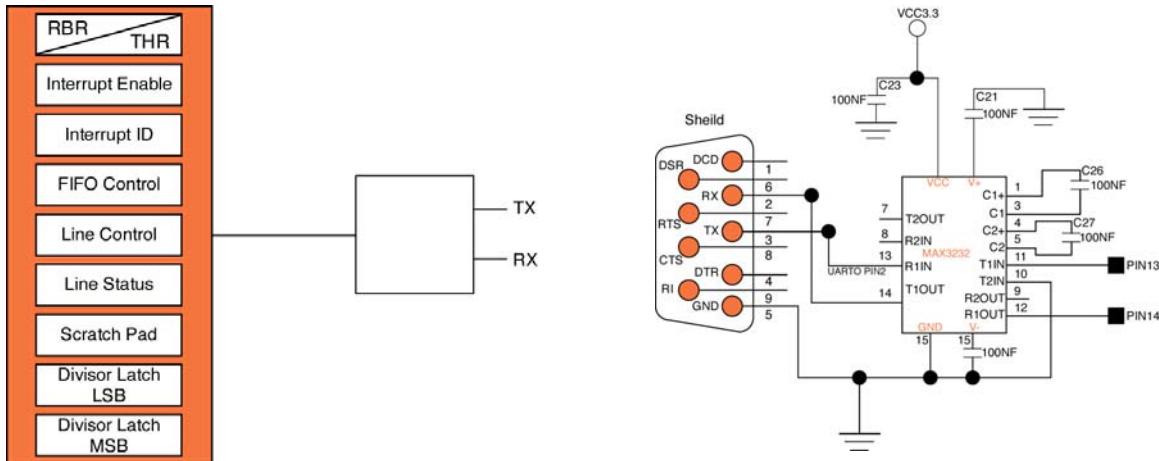
not enable the reset option as this will trip up the JTAG debugger when the watchdog times out.

Once the watchdog timer constant and mode registers have been configured, the watchdog can be kicked into action by writing to the feed register. This needs a feed sequence similar to the PLL. To feed the watchdog you must write 0xAA followed by 0x55. If this sequence is not followed, a watchdog feed error occurs and a watchdog timeout event is generated with its resulting interrupt/reset. It is also important to note that although the watchdog may be enabled via the watchdog mode register, it does not start running until the first correct watchdog feed sequence is encountered. Once fully started the watchdog must receive regular feed sequences in order to stop the watchdog counter reaching zero and timing out.

The final Watchdog register is the Watchdog Timer Value Register which allows you to read the current value of the watchdog timer.

UART

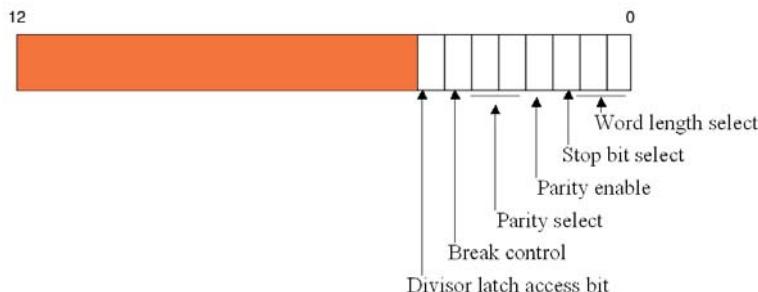
The LPC2xxx devices currently have two on-chip UARTS. They are both identical to use, except UART1 has additional modem support. Both peripherals conform to the “550 industry standard” specification. Both have a built-in baud rate generator and 16 byte transmit and receive FIFOs.



Initialisation of the UART0 is shown below:

```
void init_serial (void)                                /* Initialize Serial Interface      */
{
    PINSEL0 = 0x00050000;                            /* Enable RxD1 and TxD1           */
/* */
    U1LCR     = 0x00000083;                          /* 8 bits, no Parity, 1 Stop bit */
/* */
    U1DLL     = 0x000000C2;                          /* 9600 Baud Rate @ 30MHz VPB Clock
/* */
    U1LCR     = 0x00000003;                            /* DLAB = 0
/* */
}
```

First the pinselect block must be programmed to switch the processor pins from GPIO to the UART functions. Next the UART line control register is used to configure the format of the transmitter data character.



UART Line control register: The LCR configures the format of transmitted data. Setting the DLAB bit allows programming of the BAUD rate generators

In our example the character format is set to 8 bits, no parity and one stop bit. In the LCR there is an additional bit called DLAB which is the divisor latch access bit. In order to be able to program the baud rate generator this bit must be set.

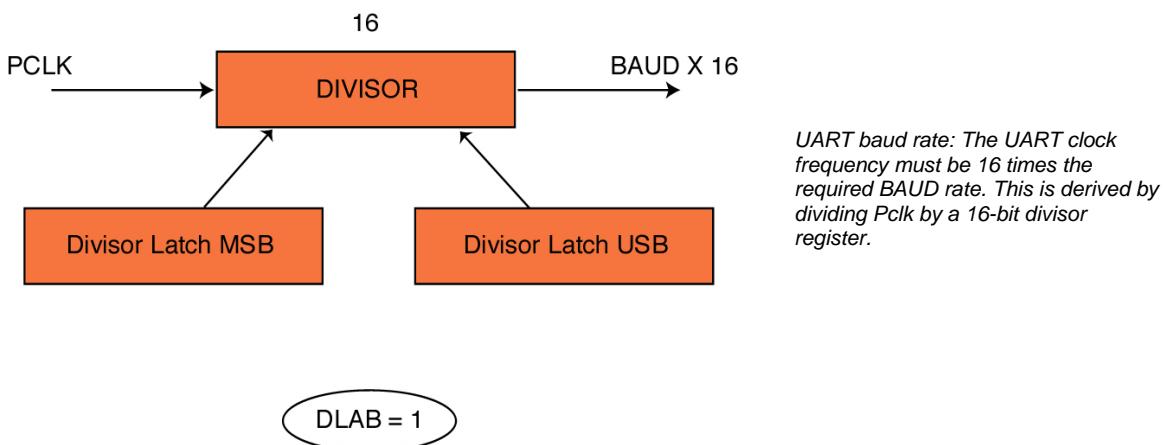
The baud rate generator is a sixteen bit prescaler which divides down Pclk to generate the UART clock which must run at 16 times the baud rate. Hence the formula used to calculate the UART baud rate is:

$$\text{Divisor} = \text{Pclk}/16 \times \text{BAUD}$$

In our case at 30MHz:

$$\text{Divisor} = 30,000,000/16 \times 9600 = (\text{approx}) 194 \text{ or } 0xC2$$

This gives a true baud rate of 9665. Often it is not possible to get an exact baud rate for the UARTs however they will work with up to around a 5% error in the bit timing. So you have some leeway with the UART timings if you need to adjust the Pclk to get exact timings on other peripherals such as the CAN bit timings. The divisor value is held in two registers, Divisor latch MSB (DLM) and Divisor latch LSB (DLL). The first eight bits of both registers holds each half of the divisor as shown below. Finally the DLAB bit in the LCR register must be set back to zero to protect the contents of the divisor registers.



Once the UART is initialised, characters can be transmitted by writing to the Transmit Holding Register. Similarly, characters may be received by reading from the Receive Buffer Register. In fact both these registers occupy the same memory location, writing a character places the character in the transmit FIFO and reading from this location loads a character from the Receive FIFO. The two routines shown below demonstrate handling of transmit and receive characters.

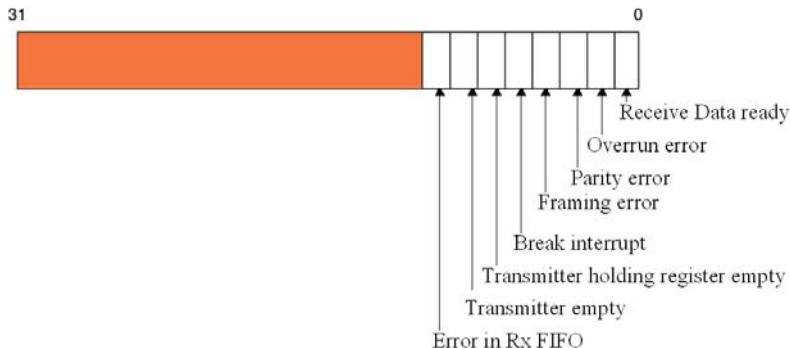
```

int putchar (int ch)           /* Write character to Serial Port */
{
    if (ch == '\n') {
        while (!(U1LSR & 0x20));
        U1THR = CR;           /* output CR */
    }
    while (!(U1LSR & 0x20));
    return (U1THR = ch);
}

int getchar (void)            /* Read character from Serial Port */
{
    while (!(U1LSR & 0x01));
    return (U1RBR);
}

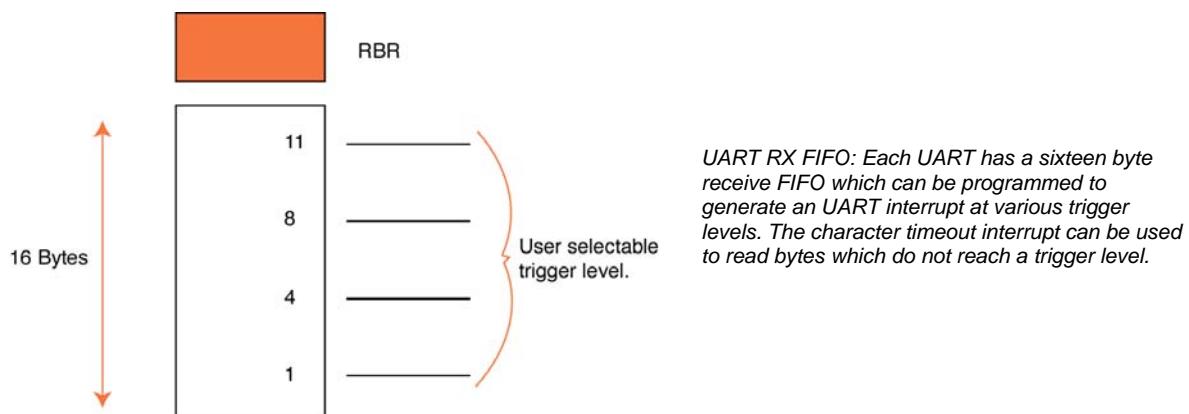
```

The putchar() and getchar functions are used to read/write a single character to the UART. These low level drivers are called by the Keil STDIO functions such as printf() and scanf(). So, if you want to redirect the standard I/O from the UART to say an LCD display and a keypad, rewrite these functions to support sending and receiving a single character to your desired I/O devices. Both the putchar() and getchar() functions read the Link Status Register (LSR) to check on UART error conditions and to check the status of the receive and transmit FIFOs.



UART Line Status Register: The LSR contains flags which indicate events within the UART. It may be polled or should be read after a UART interrupt is generated.

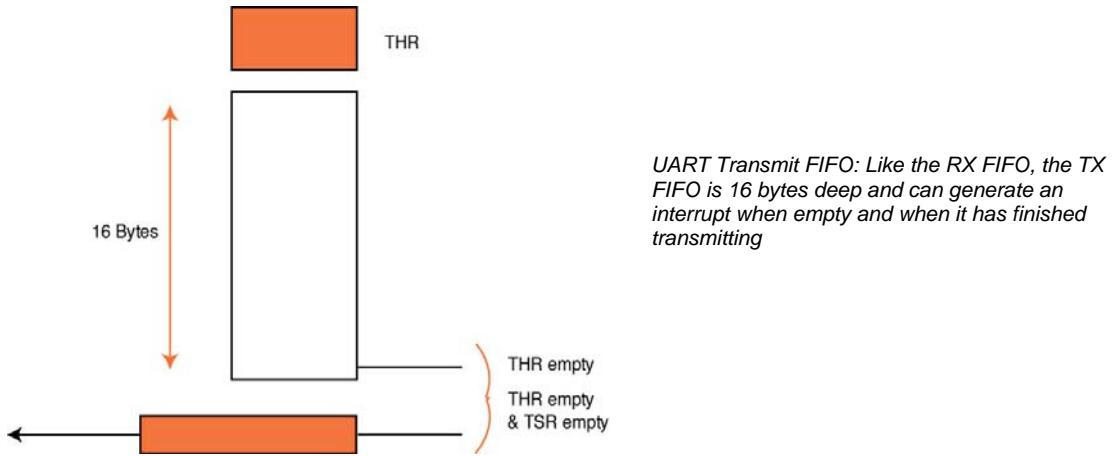
The UART has a single interrupt channel to the VIC, but three sources of interrupt. UART interrupts can be generated on a change in the Receive line status. So, if an error condition occurs, an interrupt is generated and the LSR can be read to see what is the cause of the error. The remaining two interrupt sources are receive and transmit interrupts. The receive interrupt is triggered by characters being received into the RX FIFO. The depth at which the interrupt is triggered is set in the UART FIFO control register.



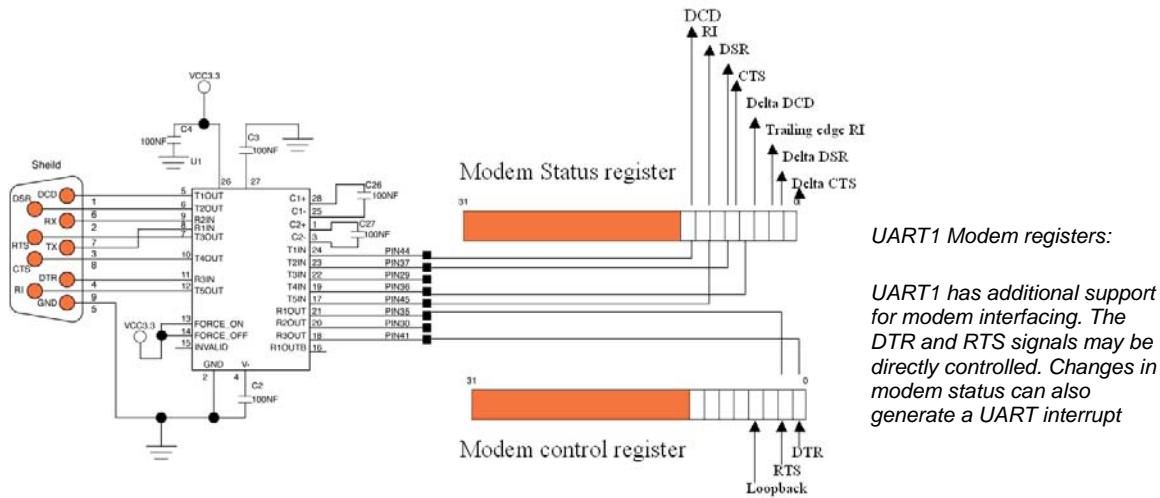
UART RX FIFO: Each UART has a sixteen byte receive FIFO which can be programmed to generate an UART interrupt at various trigger levels. The character timeout interrupt can be used to read bytes which do not reach a trigger level.

The receive interrupt can be set to trigger after it has received 1,4,8 or 14 characters. So, if the interrupt is set to trigger when eight characters are in the buffer and a total of 34 characters are sent, then four interrupts will be generated with two characters left in the FIFO. These remaining characters will cause a “character time out indication” (CTI) interrupt. The CTI interrupt occurs when there are one or more characters in the FIFO and no FIFO activity has occurred for 3.5- 4.5 character times.

The transmit FIFO will also generate interrupts when the transmit holding register is empty and when the transmit shift register is empty.



UART1 has the same basic structure as UART0, however it has additional support for modem control. This consists of additional external pins to support the full modem interface (CTS, DCD, DSR, DTR, RI, RTS), there are two additional registers the modem control register and the modem status register and an additional interrupt source to provide a modem status interrupt.



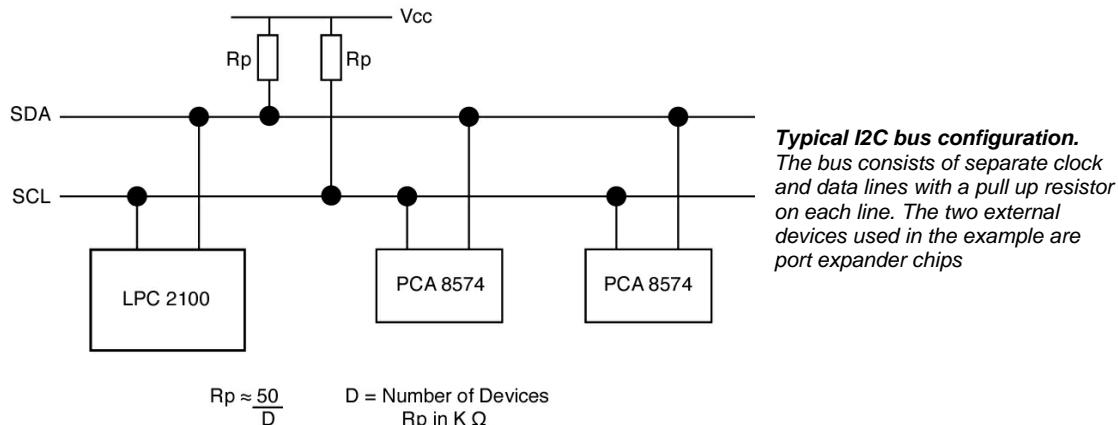
These additional features allow optimal connection to a modem with an interrupt generated each time there is a change in the modem status register.

Exercise 20: UART

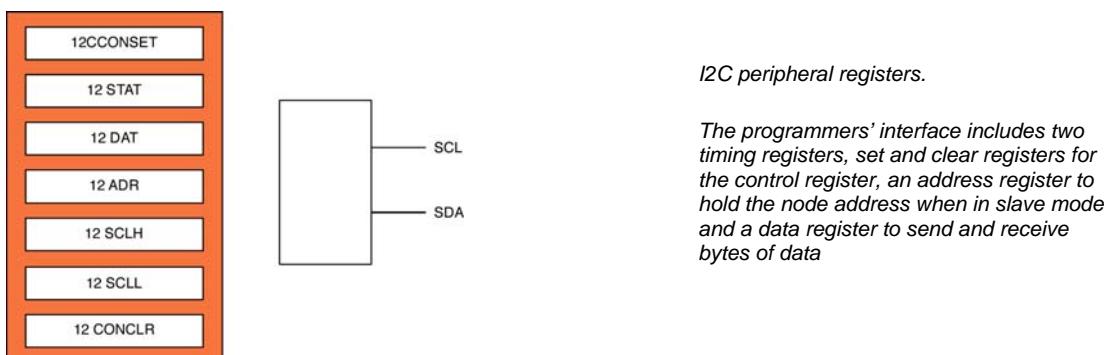
In Exercise 4 we saw how to use the STUDIO library with the UARTs. In this example we look at how the UARTs are initialised to run at a specific baud rate.

I2C Interface

As Philips were the original inventors of the I2C bus standard, it is not surprising to find the LPC2000 equipped with a fully featured I2C interface. The I2C interface can operate in master or slave mode up to 400K bits per second and in master mode it will automatically arbitrate in a multi-master system.



A typical I2C system is shown above where the LPC2000 is connected to two external port expander chips. As with the other peripherals the Serial Clock (SCL) and Data (SDA) lines must be converted from GPIO pins to I2C pins via the pin connect block.



The I2C peripheral interface is composed of seven registers. The control register has two separate registers which are used to set and clear bits in the control register (I2CONSET, I2CONCLR). The bit rate is also determined by two registers (I2SCLH, I2SCLL). The status register returns control codes which relate to different events on the bus. The data register is used to supply each byte to be transmitted, or as data is received it will be transferred to this register. Finally, when the LPC2000 is configured as a slave device its network address is set by programming the I2ADR register.

In order to initialise the I2C interface we need to run the following lines of code:

```
VICVectCnt11 = 0x00000029;           //select a priority slot for a given interrupt
VICVectAddr1 = (unsigned)I2CISR    //pass the address of the IRQ into the VIC slot
VICIntEnable = 0x00000200;          //enable interrupt

PINSEL0 = 0x50;                     //Switch GPIO to I2C pins
I2SCLH      = 0x08;    //Set bit rate to 57.6KHz
I2SCLL      = 0x08;
```

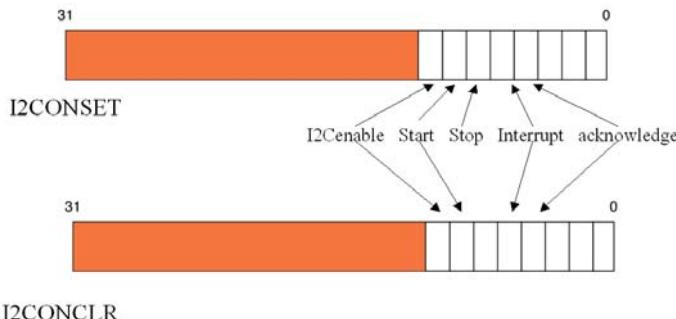
The I2C peripheral must be programmed to respond to each event which occurs on the bus. This makes it a very interrupt-driven peripheral. Consequently the first thing we must do is to configure the VIC to respond to an I2C interrupt. Next the pinselect block is configured to connect the I2C data and clock lines to the external pins. Lastly we must set the bit rate by programming I2SCLH and I2CSLL. In both of these registers only the first 16 bits are used to hold the timing values. The formula for the I2C bit rate is given as:

$$\text{Bit Rate} = \text{Pclk}/(\text{I2SCLH}+\text{I2CSLL})$$

In the above example the PLL is not enabled and the external crystal is 14.7456MHz. Hence the I2C bit rate is:

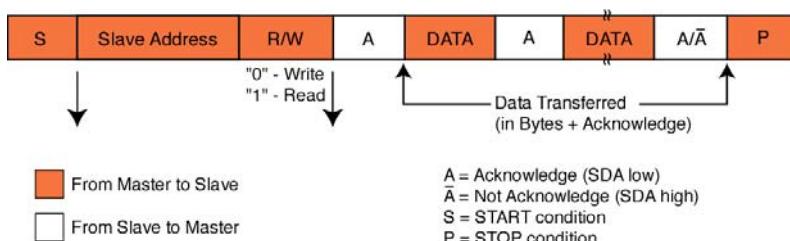
$$\text{Bit Rate} = 14.7456/\text{B} (8 + 8) = 937500$$

Once configured, the LPC2100 can initiate communication with other bus devices to read and write data as a bus master, or receive and reply to requests from a bus master. The contents of the I2C control register are shown below. Remember this register is controlled by the CONSET and CONCLR registers.



I2C control registers:
The control registers are used to enable the I2C peripheral and interrupt as well as controlling the I2C bus start, stop and ack conditions.

We will first look at the bus master mode. To enter this mode the I2C peripheral must be enabled and the acknowledge bit must be set to zero. This prevents the I2C peripheral acknowledging any potential master and entering the slave mode. In the master mode the LPC2000 device is responsible for initiating any communication. During a I2C bus transfer a number of bus events must occur.



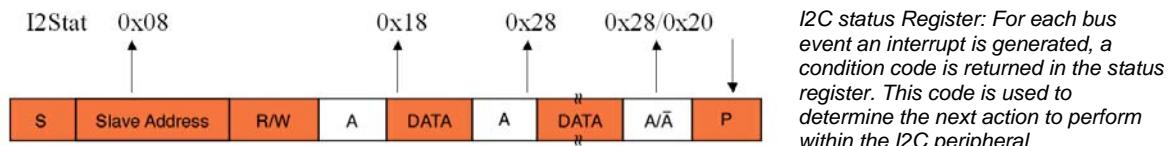
Typical I2C transaction :
A I2C bus transaction is characterised by a start condition, slave address data exchange and stop condition with acknowledge handshaking

The bus master must first signal a start condition. To do this the I2C clock line is pulled high and the data is pulled low. The address of the slave which the master wants to talk to is then written onto the bus, followed by a bit which states if a read or write is being requested. If the slave has received this preamble correctly, it will reply with an acknowledge. Then data can be transferred as a series of bytes and acknowledges, until the master terminates the transaction with a stop condition. The I2C peripheral on the LPC2000 series is really a I2C engine. It controls all the bus events but has no intelligence. This means that the ARM7 CPU has to micro-manage the I2C bus for each transaction. Fortunately this is easy to do

and is centred around the I2C interrupt. Once the I2C peripheral is initialised in master mode we can start a write data transfer as follows:

```
void I2CTransferByte(unsigned Addr,unsigned Data)
{
    I2CAddress = Addr;           //Place address and data in Globals to be used by
                                //the interrupt
    I2CData = Data;
    I2CONCLR = 0x000000FF;      //Clear all I2C settings
    I2CONSET = 0x00000040;      //Enable the I2C interface
    I2CONSET = 0x00000020;      //Start condition
}
```

The slave address and data to be sent are placed in global variables so that they can be used by the I2C interrupt routine. The address is a seven-bit address with the LSB set for write and cleared for read. The routine next clears the I2C control flags, enables the I2C peripheral and asserts a start condition. Once the start condition has been written onto the bus an interrupt is generated and a result code can be read from the I2C status register.



If the start condition has been successful, this code will be 0x08. Next the application software must write the slave address and the R/W bit into the I2Cdata register. This will be written on to the bus and will be acknowledged by the slave. When the acknowledge is received, another interrupt is generated and the status register will contain the code 0x18 if the transfer was successful. Now that the slave has been addressed and is ready to receive data, we can write a string of bytes into the I2C data register. As each byte is written it will be transmitted and acknowledged. When it is acknowledged an interrupt is generated and 0x28 will be in the status register if the transfer was successful. If it failed and had a NACK the code will be 0x20 and the byte must be sent again. So, as each byte is transferred an interrupt is generated, the status code can be checked and the next byte can be sent. Once all the bytes have been sent the stop condition can be asserted by writing to the I2C control register and the transaction is finished. The I2C interrupt is really a state machine which examines the status register on each interrupt and performs the necessary action. This is easy to implement as a switch statement as shown below.

```

void I2CISR (void)           //I2C interrupt routine
{
    switch (I2STAT)          //Read result code and switch to next action
    {
        case ( 0x08):          //Start bit
            I2CONCLR = 0x20;   //Clear start bit
            I2DAT = I2CAddress; //Send address and
            //write bit
            break;

        case (0x18):           //Slave address+W, ACK
            I2DAT = I2Cdata;    //Write data to tx register
            break;

        case (0x20):           //Slave address +W, Not ACK
            I2DAT = I2CAddress; //Resend address and write bit
            break;

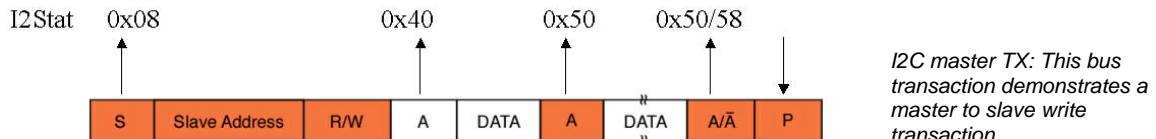
        case (0x28):           //Data sent, Ack
            I2CONSET = 0x10;    //Stop condition
            break;

        default :
            break;
    }

    I2CONCLR      = 0x08;      //Clear I2C interrupt flag
    VICVectAddr = 0x00000000; //Clear interrupt in
}

```

This example sends a single byte but could be easily modified to send multiple bytes. Additional case statements may be added to handle a master request for data.



In the case of a master receive, the start condition will be the same but this time the address written on to the bus will have the R/W bit cleared. When the acknowledge is received after the slave address is sent, it will be followed by the first byte of data from the slave so the master does not have to do anything. However, in the case statement we can set the acknowledge bit so that an ACK is generated as soon as the byte has been transferred. As each byte is transferred, the data can be read from I2CDAT. When all the bytes have been received, the stop condition can be asserted and the transaction ends.

The same I2CtransferByte() function can be used to start a read transaction and the additional case statements required in the interrupt are shown below.

```
case (0x40) :           //Slave Address +R, ACK
    I2CONSET = 0x04;    //Enable ACK for data byte
break;

case (0x48) :           //Slave Address +R, Not Ack
    I2CONSET = 0x20;    //Resend Start condition
break;

case (0x50) :           //Data Received, ACK
    message = I2DAT;
    I2CONSET = 0x10;    //Stop condition
lock = 0;                //Signal end of I2C activity

break;

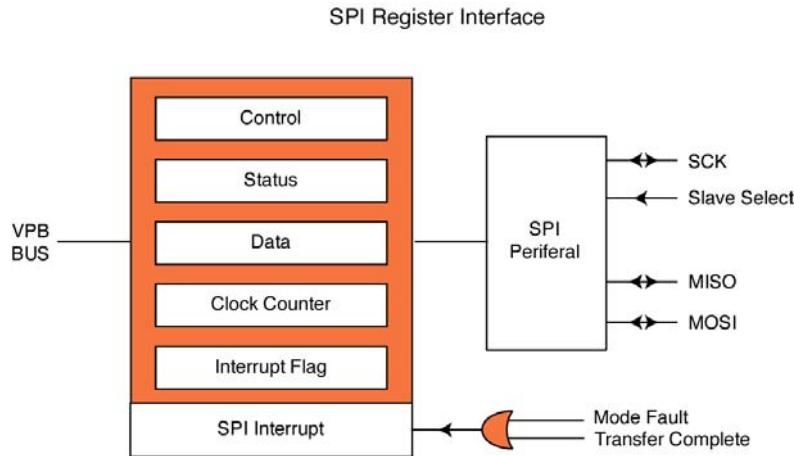
case (0x58):            //Data Received, Not Ack
    I2CONSET = 0x20;    // Resend Start condition
break;
```

Exercise 21: I2C

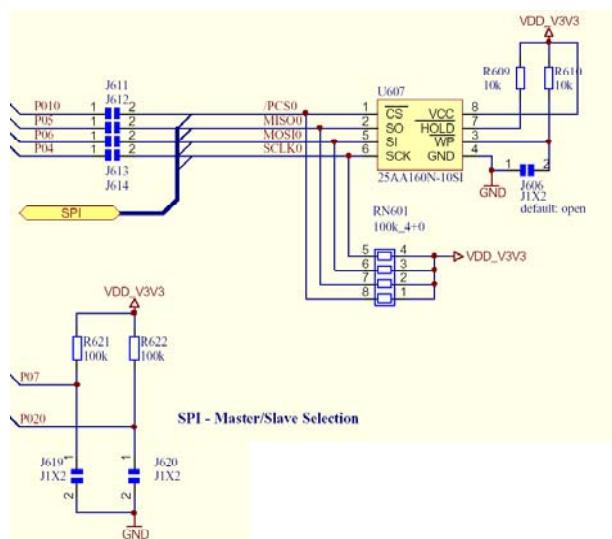
This exercise demonstrates how to use the I2C interface to communicate to an I2C EEPROM.

SPI Interface

Like the I2C interface the SPI interface is a simple peripheral “engine” which can write and read data to the SPI bus, but is not intelligent enough to manage the bus. It is up to your code to initialise the SPI interface and then manage the bus transfers.



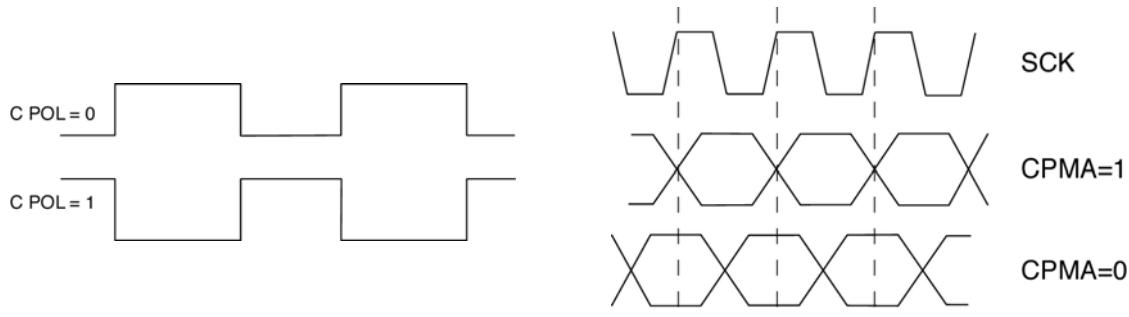
The SPI peripheral has four external pins: a serial clock pin, slave select pin and two data pins master in/slave out and master out/slave in. The serial clock pin provides a clock source of up to 400Kbits/sec when in master mode, or will accept an external clock source when in slave mode. The SPI bus is purely a serial data connection for high-speed data transfer and unlike I2C does not have any addressing scheme built into the serial transfer. An external peripheral is selected by a slave select pin which is a separate pin. Typically, if the LPC2000 is acting in master mode, it could use a GPIO pin to act as slave select (chip enable) for the desired SPI peripheral. When the SPI peripheral is in slave mode, it has its own slave select input which must be pulled low to allow an SPI master to communicate with it. The two data transfer pins master in / slave out and master out / slave in are connected to the remote SPI device and their orientation depends on whether the device is operating in master or slave mode. The diagram below shows a typical configuration for connecting to an EEROM device.



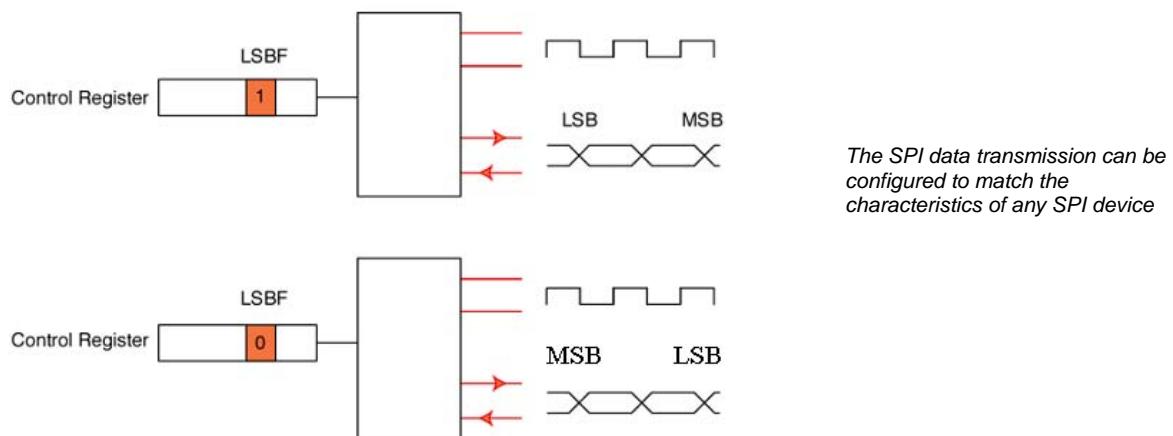
SPI EEROM peripheral:
This diagram shows how to interface an external EEROM onto the SPI bus of the LPC2000. It should be noted that pins P0.7 and P0.20 must be pulled high to enable the SPI peripheral as a master

The programmers’ interface for the SPI peripheral has five registers. The clock counter register determines the baud rate. Pclk is simply divided by the value in the clock counter to give the SPI bit rate. This register must hold a minimum value of eight. The control register

is used to configure the operation of the SPI bus. Because of the simple nature of the SPI data transfer and the wide range of SPI peripherals available, the SPI clock and data lines can be configured to operate in several different configurations. Firstly the polarity and phase of the clock must be defined. The polarity can be active high or active low as shown below and the clock phase can be edge or centre aligned.



Finally the data orientation may also be defined as the most significant bit transferred first or the least significant bit transferred first.



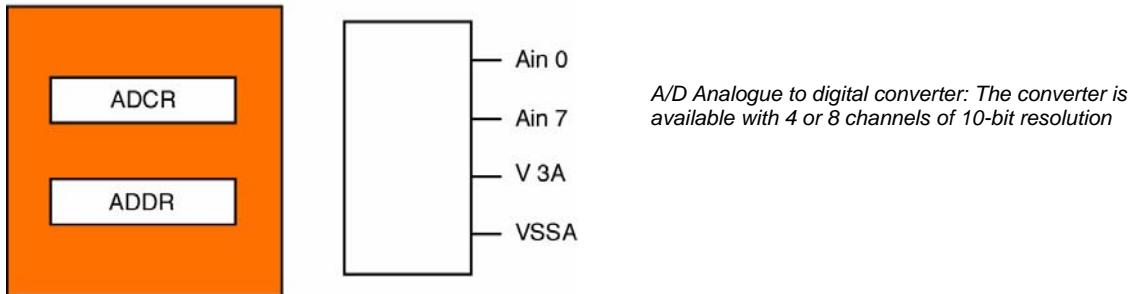
Each of these configuration features has a configuration bit in the control register and you must program these bits to match the SPI peripheral you are trying to communicate with. Once the bit rate has been set and the control register configured, then communication can begin. To communicate with the SPI memory shown above, first set the GPIO pin to enable the memory for communication. Then writing to the SPI data register will send a byte of data and reading from the register will collect any data sent from the external peripheral. The actual data format used in the transaction will depend on the SPI device you are trying to communicate with.

Exercise 22: SPI

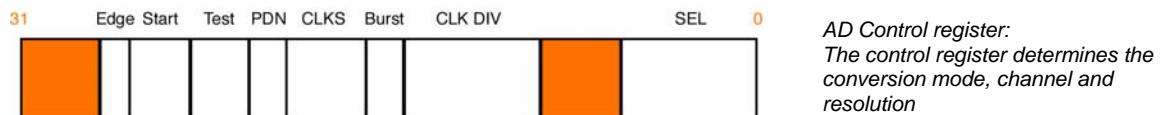
This exercise demonstrates how to configure the SPI peripheral and communicate with an external EEPROM on the SPI bus

Analog To Digital Converter

The A/D converter present on some LPC2000 variants is a 10-bit successive approximation converter, with a conversion time of 2.44 uSec or just shy of 410 KSps. The A/D converter has either 4 or 8 multiplexed inputs depending on the variant. The programming interface for the A/D converter is shown below.



The A/D control register establishes the configuration of the converter and controls the start of conversion. The first step in configuring the converter is to set up the peripheral clock. As with all the other peripherals, the A/D clock is derived from the PCLK. This PCLK must be divided down to equal 4.5MHz. This is a maximum value and if PCLK cannot be divided down to equal 4.5MHz then the nearest value below 4.5MHz which can be achieved should be selected.



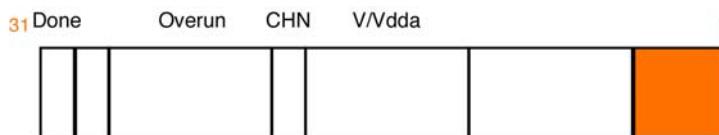
PCLK is divided by the value stored in the CLKDIV field plus one. Hence the equation for the A/D clock is as follows:

$$\text{CLKDIV} = (\text{PCLK}/\text{Adclk}) - 1$$

As well as being able to stop the clock to the A/D converter in the peripheral power down register, the A/D has the ability to fully power down. This reduces the overall power consumption and the on-chip noise created by the A/D. On reset, the A/D is in power down mode, so as well as setting the clock rate the A/D must be switched on. This is controlled by the PDN bit in ADCR. Logic one in this field enables the converter. Unlike other peripherals the A/D converter can make measurements of the external pins when they are configured as GPIO pins. However, by using the pinselect block to make the external pins dedicated to the A/D converter the overall conversion accuracy is increased.

Prior to a conversion the resolution of the result may be defined by programming the CLKS field. The A/D has a maximum resolution of 10 bits but can be programmed to give any resolution down to 3 bits. The conversion resolution is equal to the number of clock cycles per conversion minus one. Hence for a 10-bit result the A/D requires 11 ADCLK cycles and four for a 3-bit result. Once you have configured the A/D resolution, a conversion can be made. The A/D has two conversion modes, hardware and software. The hardware mode

allows you to select a number of channels and then set the A/D running. In this mode a conversion is made for each channel in turn until the converter is stopped. At the end of each conversion the result is available in the A/D data register.



AD data register: The data register contains the conversion result, channel overrun error and conversion done flag.

At the end of a conversion the Done bit is set and an interrupt may also be generated. The conversion result is stored in the V/Vdda field as a ratio of the voltage on the analogue channel divided by the voltage on the analogue power supply pin. The number of the channel for which the conversion was made is also stored alongside the result. This value is stored in the CHN field. Finally, if the result of a conversion is not read before the next result is due, it will be overwritten by the fresh result and the OVERUN bit is set to one. The example below demonstrates use of the A/D converter in hardware mode.

```
int main(void)
{
    VPBDIV = 0x00000002;           //Set the Pclk to 30 MHz
    IODIR1 = 0x00FF0000;          // P1.16..23 defined as Outputs
    ADCR   = 0x00270607;          // Setup A/D: 10-bit AIN0 @ 3MHz

    VICVectCntl0 = 0x00000032;    //connect A/D to slot 0
    VICVectAddr0 = (unsigned)AD_ISR; //pass the address of the IRQ into the VIC
                                    //slot
    VICIntEnable = 0x00040000;     //enable interrupt

    while(1)
    {
        ;
    }
}

void AD_ISR (void)
{
    unsigned val,chan;
    static unsigned result[4];

    val = ADCR;
    val = ((val >> 6) & 0x03FF);      //Extract the A/D result
    chan = ((ADCR >>0x18) & 0x07);
    result[chan] = val;
}
```

The A/D has a second software conversion mode. In this case, a channel is selected for conversion using the SEL bits and the conversion is started under software control by writing 0x01 to the START field. This causes the A/D to perform a single conversion and store the results in the ADDR in the same fashion as the hardware mode. The end of conversion can be signalled by an interrupt, or by polling the done bit in the ADDR. In the software conversion mode it is possible to start a conversion when a match event occurs on timer zero or timer one. Or when a selected edge occurs on P0.16 or P0.22, the edge can be rising or falling, as selected by the EDGE field in the ADCR.

Start Bits	Conversion Trigger
0 0 0	Halted
0 0 1	Start Now
0 1 0	P0 - 16
0 1 1	P0 - 22
1 0 0	MAT 0-1
1 0 1	MAT 0-3
1 1 0	MATT 1-0
1 1 1	MAT 1-1

The A/D may be started by a software event or it may be started by several hardware triggers

The simplest method of using the A/D converter is shown below.

```

VPBDIV = 0x02;           //Set the Pclk to 30 MHz
IODIR1 = 0x00FF0000;     // P1.16..23 defined as Outputs
ADCR   = 0x00270601;     // Setup A/D: 10-bit AIN0 @ 3MHz
ADCR |= 0x01000000;      // Start A/D Conversion

while(1)
{
    do
    {
        val = ADDR;          // Read A/D Data Register
    }
}

```

Exercise 23 : Analog To Digital Converter

This exercise uses the A/D to convert an external voltage source and modulate a bank of LEDs with the result.

Digital To Analog Converter

The LPC2132/2138 variants have a 10-bit Digital to Analogue converter. This is an easy-to-use peripheral as it only has a single register.

The DAC is enabled by writing to bits 18 and 19 of PINSEL1 and converting pin 0.25 from GPIO to the AOUT function. It should also be noted that a channel of the analogue to digital converter also shares this pin.

BIAS 15	Value 6	0	

The DAC is controlled by a single register.
The value to be converted is written here along with the bias value

31

Once enabled a conversion can be started by writing to the VALUE bits in the control register. The conversion time is dependant on the value of the BIAS bit. If it is set to one the conversion time is 2.5uSec but it can drive 700 uA. If it is zero the conversion time is 1 uSec but it is only able to deliver 350 uA. However, the total settling time is also dependent on the external impedance. Figures for the impedance of the DAC have not yet been released.

Exercise 24: Digital to Analog converter

This exercise simulates a sine wave which is sampled by the Analogue to digital converter. These values are loaded straight into the Digital to Analogue converter to regenerate the sine wave. The two sine waves can be compared in the logic analyser window.

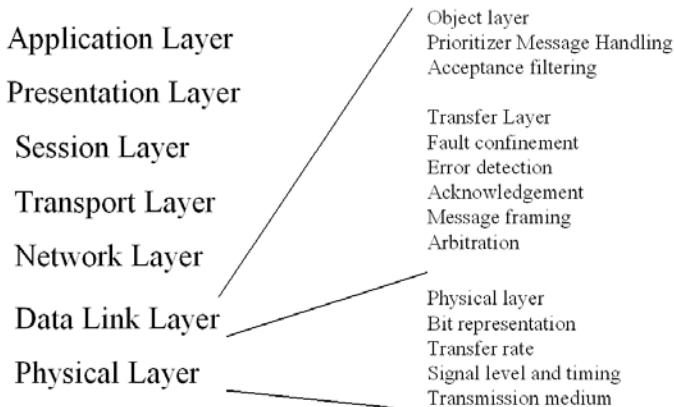
CAN Controller

Variants of the LPC2000 are available with up to 4 independent CAN controllers on board the chip. The CAN controllers are one of the more complicated peripherals on the LPC2000. In this section we will have a look at the CAN protocol and the LPC2000 CAN peripherals.

The Controller Area Network (CAN) Protocol was developed by Robert Bosch for Automotive Networking in 1982. Over the last 22 Years CAN has become a standard for Automotive Networking and has had a wide uptake in non-automotive systems where it is required to network together a few embedded nodes. CAN has many attractive features for the embedded developer. It is a low-cost, easy-to-implement, peer to peer network with powerful error checking and a high transmission rate of up to 1 Mbit/sec. Each CAN packet is quite short and may hold a maximum of eight bytes of data. This makes CAN suitable for small embedded networks which have to reliably transfer small amounts of critical data between nodes.

ISO 7 Layer Model

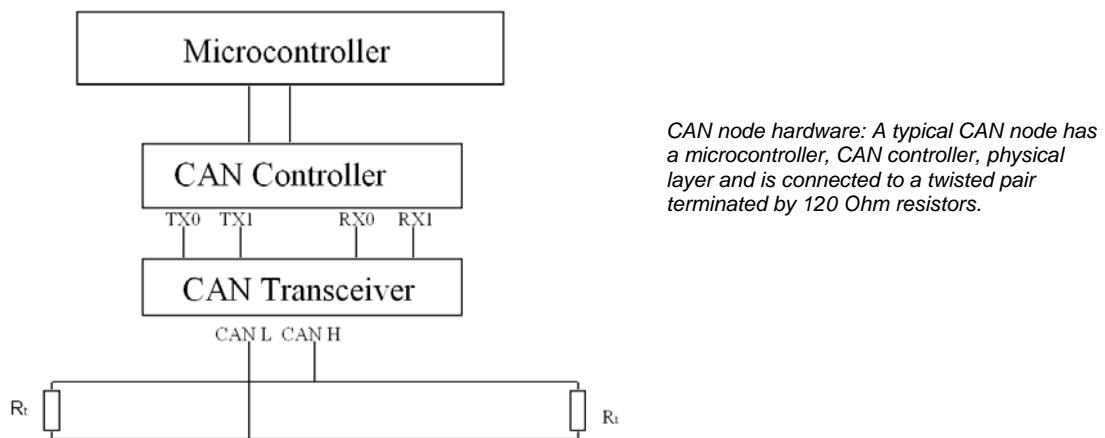
In the ISO seven layer model the CAN protocol covers the layer two ‘data link layer’, i.e. forming the message packet, error containment, acknowledgement and arbitration.



CAN does not rigidly define the layer 1 ‘Physical layer’ so CAN messages may be run over many different physical mediums. However, the most common physical layer is a twisted pair and standard line drivers are available. The other layers in the IOS model are effectively empty and the application code directly addresses the registers of the CAN peripheral. In effect, the CAN peripheral can be used as a glorified UART without the need for an expensive and complex protocol stack. Since CAN is also used in Industrial Automation there are a number of software standards that define how the CAN messages are used to transfer data between different manufacturers’ equipment. The most popular of these application layer standards are CANopen and Device net. The sole purpose of these standards is to provide interoperability between different OEM equipment. If you are developing your own closed system you do not need these application layer protocols and are free to implement your own proprietary protocol, which is what most people do.

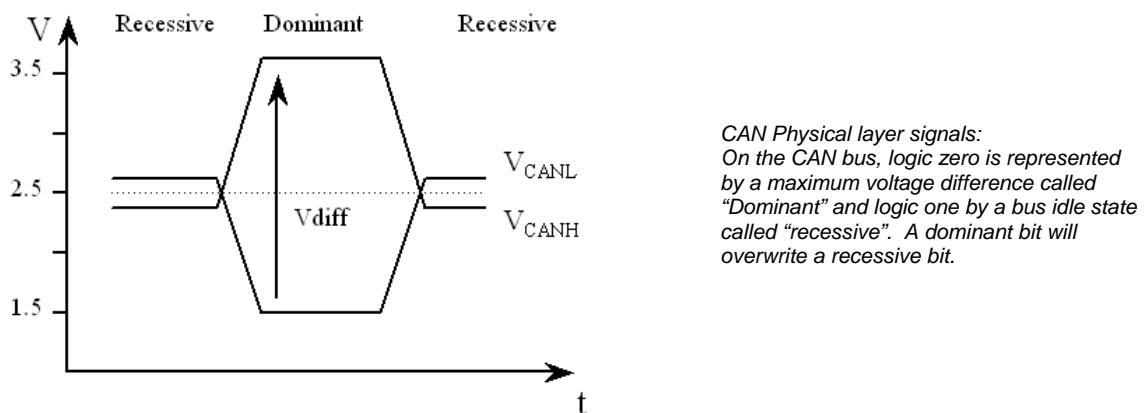
CAN Node Design

A typical CAN node is shown below. Each node consists of a microcontroller and a separate CAN controller. The CAN controller may, as in the case of the LPC2000, be fabricated on the same silicon as the microcontroller or it may be a stand-alone controller in a separate chip to the microcontroller. The CAN controller is interfaced to the twisted pair by a line driver and the twisted pair is terminated at either end by a 120 Ohm resistor. The most common mistake with a first CAN network is to forget the terminating resistors and then nothing works.



One important feature about the CAN node design is that the CAN controller has separate transmit and receive paths to and from the physical layer device. So, as the node is writing on to the bus it is also listening back at the same time. This is the basis of the message arbitration and for some of the error detection.

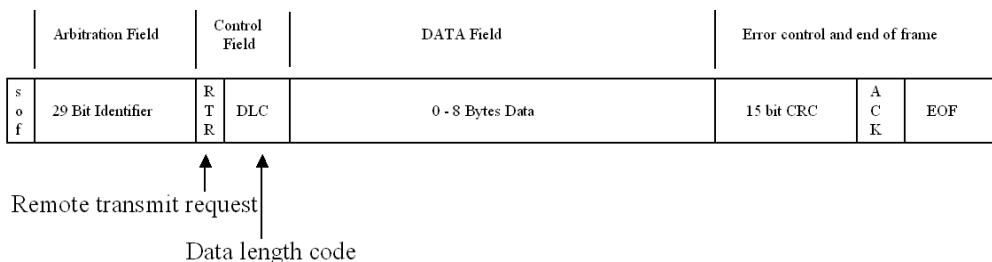
The two logic levels are written onto the twisted pair as follows, a logic one is represented by bus idle with both wires held half way between 0 and Vcc. A logic Zero is represented by both wires being differentially driven.



In “CAN speak” a logic one is called a recessive bit and a logic zero is called a dominant bit. In all cases a dominant bit will overwrite a recessive bit. So, if ten nodes write recessive and one writes dominant, then each node will read back a dominant bit. The CAN bus can achieve bit rates up to a maximum of 1 Mbit/sec. Typically this can be achieved over about 40 metres of cable. By dropping the bit rate, longer cable runs may be achieved. In practice you can get at least 1500 metres with the standard drivers at 10 Kbit/sec.

CAN Message Objects

The CAN bus has two message objects which may be generated by the application software. The message object is used to transfer data around the network. The message packet is shown below.



CAN message packet : The message packet is formed by the CAN controller, the application software provides the data bytes, the message identifier and the RTR bit

The message packet starts with a dominant bit to mark the start of frame. Next comes the message identifier which may be up to 29 bits long. The message identifier is used to label the data being sent in the message packet. CAN is a producer / consumer protocol. A given message is produced from one unique node and then may be consumed by any number of nodes on the network simultaneously. It is also possible to do point-to-point communication by making only one node interested in a given identifier. Then a message can be sent from the producer node to one given consumer node on the network. In the message packet the RTR bit is always set to zero. (This field will be discussed shortly.) The DLC field is the data length code and contains an integer between 0 and 8 which indicates the number of data bytes being sent in this message packet.

So, although you can send a maximum of 8 bytes in the message payload it is possible to truncate the message packet in order to save bandwidth on the CAN bus. After the 8 bytes of data there is a 15-bit cyclic redundancy check. This provides error detection and correction from the start of frame up to the beginning of the CRC field. After the CRC there is an acknowledge slot. The transmitting node expects the receiving nodes to assert an acknowledge in this slot within the transmitting CAN packet. In practice the transmitter sends a recessive bit and any node which has received the CAN message up to this point will assert a dominant bit on the bus, thus generating the acknowledge. This means that the transmitter will be happy if just one node acknowledges its message, or if 100 nodes generate the acknowledge. So when developing your application layer care must be taken to treat the acknowledge as a weak acknowledge, rather than confirmation that the message has

reached all its destination nodes. After the acknowledge slot there is an end of frame message delimiter.

It is also possible to operate the CAN bus in a master / slave mode. A CAN node may make a remote request onto the network by sending a message packet which contains no data, but has the RTR bit set. The remote frame is requesting a message packet to be transmitted with a matching identifier. On receiving a remote frame, the node which generates the matching message will transmit the corresponding message frame.

S O F	Identifier	RTR	DLC	CRC	ACK	EOF
-------------	------------	-----	-----	-----	-----	-----

Remote Transmit request: The RTR frame is used to request message packets from the network as a master / slave transaction

As previously mentioned, the CAN message identifier can be up to 29 bits long. There are two standards of CAN protocol, the only difference being the length of the message identifier.

2.0A Has an 11-bit identifier

2.0B Passive Has an 11-bit identifier

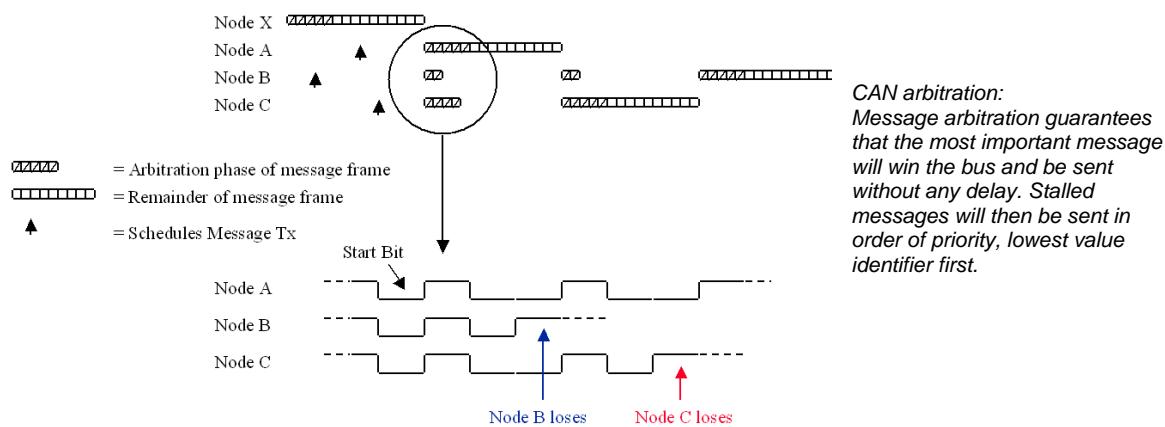
2.0B Active Has a 29-bit identifier

It is possible to mix the two protocol standards on the same bus but you must not send a 29-bit message to an 2.0A device

	Frame with 11 bit ID	Frame with 29 bit ID
V2.0B Active CAN Module	Tx/Rx OK	Tx/Rx OK
V2.0B Passive CAN Module	Tx/Rx OK	Ignored
V2.0A CAN Module	Tx/Rx OK	<u>Bus ERROR</u>

CAN Bus Arbitration

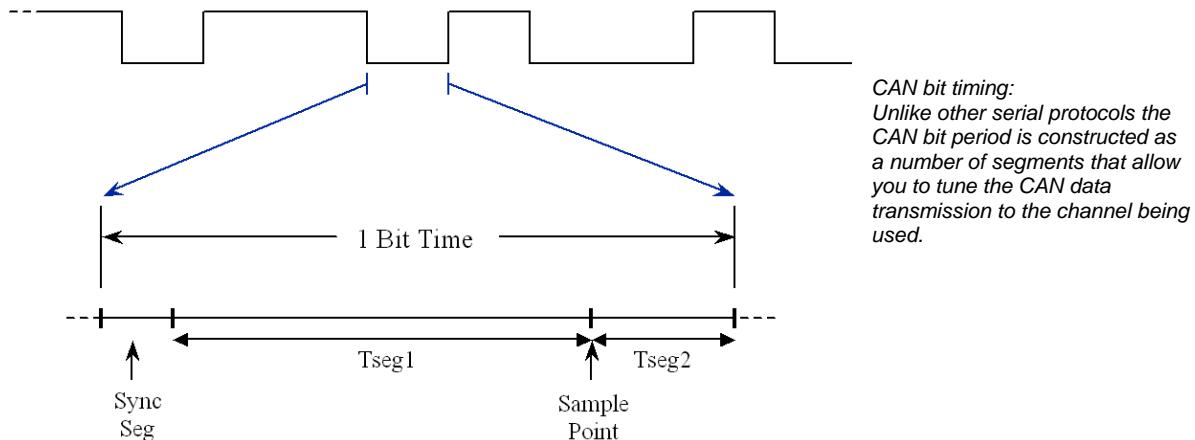
If a message is scheduled to be transmitted on to the bus and the bus is idle, it will be transmitted and may be picked up by any interested node. If a message is scheduled and the bus is active, it will have to wait until the bus is idle before it can be transmitted. If several messages are scheduled while the bus is active, they will start transmission simultaneously once the bus becomes idle, being synchronised by the start of frame bit. When this happens, the CAN bus arbitration will take place to determine which message wins the bus and is transmitted.



CAN arbitrates its messages by a method called “non-destructive bit wise arbitration”. In the diagram above, three messages are pending transmission. Once the bus is idle and they are synchronised by the start bit, they will start to write their identifiers onto the bus. For the first two bits, all three messages write the same logic and hence read back the same logic so each node continues transmission. However on the third bit, node A and C write dominant bits and node B writes recessive. At this point, node B wrote recessive but reads back dominant. In this case it will back off the bus and start listening. Node A and C will continue transmission until node C writes recessive and node A writes dominant. Now node C stops transmission and starts listening. Now node A has won the bus and will send its message. Once A has finished, nodes B and C will transmit and node C will win and send its message. Finally node B will send its message. If node A is scheduled again, it will win the bus even though the node B and C messages have been waiting. In practice the CAN bus will transmit the message with the lowest value identifier.

Bit Timing

Unlike many other serial protocols, the CAN bit rate is not just defined by a baud rate prescaler. The CAN peripheral contains a Baud rate prescaler but it is used to generate a time quanta i.e. a time slice. A number of these time quanta are added together to get the overall bit timing.



The bit period is split into three segments. First is the sync segment, which is fixed at one time quanta long. The next two segments are Tseg1 and Tseg2 where the user defines the number of time quanta in each region. The minimum number of time quanta in a bit period is 8 and the maximum is 25. The receiving sample point is at the end of Tseg1 so changing the ratio of Tseg1 to Tseg2 adjusts the sample point. This allows the CAN protocol to be tuned to the transmission channel. If you are using long transmission lines, the sample point can be moved backwards. If you have drifting oscillators you can bring the sample point forward. In addition, the receivers can adjust their bit rate to lock onto the transmitter. This allows the receivers to compensate for small variations in the transmitter bit rate. The amount that each bit can be adjusted is called the “synchronous jump width” and may be set to between 1 – 4 time quanta and is again user definable.

To calculate the bit timing, the formula is given by

$$\text{Bit rate} = \frac{\text{Pclk}}{(\text{BRP} \times (1 + \text{Tseg1} + \text{Tseg2}))}$$

Where: BRP = Baud rate prescaler

This calculation has a lot of unknowns. If we assume that we want to reach a bit rate of 125K with a 60 MHz Pclk and a sample point of about 70%, here is how the BRP calculation is performed.

The total number of time quanta in a bit period is given by $(1 + \text{Tseg1} + \text{Tseg2})$. If we call this term QUANTA and rearrange the equation in terms of the baud rate prescaler:

$$\text{BRP} = \frac{\text{Pclk}}{(\text{Bit rate} \times \text{QUANTA})}$$

Using our known values:

$$\text{BRP} = 60 \text{ MHz} / (125K \times \text{QUANTA})$$

Now we know that we can have between 8 and 25 time quanta in the bit period, so using a spreadsheet we can substitute in integer values between 8 and 25 for QUANTA until we get an integer value for BRP.

In this case when QUANTA = 16 BRP = 30;

$$\text{Then } 16 = \text{Quanta} = (1 + \text{Tseg1} + \text{Tseg2})$$

So we can adjust the ratio between Tseg1 and Tseg2 to give us the desired sample point.

$$\text{Sample point} = (\text{QUANTA} \times 70) / 100$$

Hence $16 * 0.7 = 11.2$. This gives Tseg 1 = 10, Tseg2 = 5 and the sample point = 68.8%

The value for the synchronous jump width may be calculated via the following rule of thumb.

Tseg2 >= 5 Tq then program SJW to 4
Tseg2 < 5 Tq then program SJW to (Tseg2 - 1) Tq

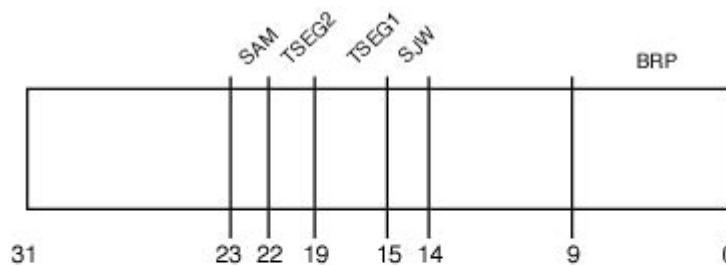
In this case SJW = 4.

CAN Message Transmission

In the LPC2000, each CAN controller has a number of status and control registers plus three transmit buffers and a receive buffer.

CAN Control Register	Name	Description
Rx Buffer	CAN MOD	Controls Operating Mode
Tx1 Buffer	CAN CMR	Command Register
Tx2 Buffer	CAN GSR	Global Status Register
Tx3 Buffer	CAN ICR	Interrupt Status
	CAN IER	Interrupt Enable
	CAN BTR	Bit Timing Register
	CAN EWL	Error Warning Limit
	CAN SR	Status Register

In order to configure CAN controller we must program the bit timing register. However the bit timing register is a protected register and may only be written to when the CAN controller is in reset. Bit zero of the mode register is used to place the CAN controller into reset.



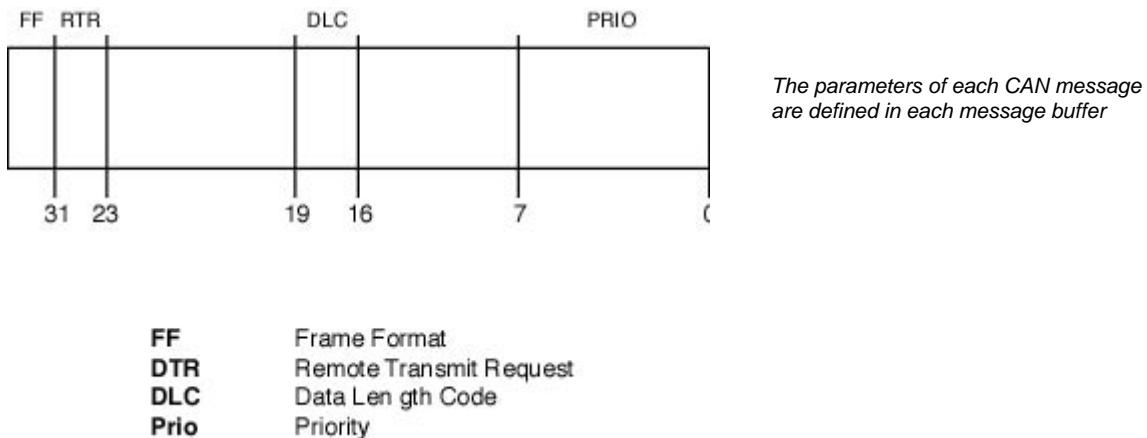
The CAN bit timing is defined by 5 separate parameters

SAM	Sampling
TSEG2	Timing Segment 2
TSEG1	Timing Segment 1
SJW	Synchronous Jump Width
BRP	Baudrate Prescaler

We can use the values calculated above to initialise one of the CAN controllers to 125Kbit/sec. It is important to note that the values stored in the register are the calculated values minus 1. This ensures that no timing segment is set to zero. Once the CAN controller has been initialised, it is possible to transmit a message by writing to a transmit buffer. Each transmit buffer is made up of four words.

Name	Description
CANT Fix	Transmit Frame Info
CANT iDx	Transmit Identifier
CAN TDAx	Transmit Data 1-4
CAN TDBx	Transmit Data 5-8

Two words are used to hold the 8 bytes of data and one word holds the message identifier. The final register is the frame information register.



This register holds the values of the DLC and the RTR bit. In addition, there is a frame format (FF) bit that defines whether the message has an 11-bit or 29-bit identifier. As there are three TX buffers it is possible to define an internal priority for each TX buffer. If several buffers are scheduled simultaneously, the CAN controller will use internal arbitration to decide which is transmitted first. This can be done in one of two ways; if the TPM bit in the MODE register is Zero, the transmit buffer with the lowest value identifier will be sent first. If TPM is high, then arbitration will use the values stored in the PRIO field in the Tx Frame Information register and the buffer with the lowest PRIO value is sent first. Once the buffer has been filled with a message, transmission can be started by setting the Transmit request bit (TR) in the COMMAND register. The code below shows some code fragments to initialise the CAN peripheral and transmit a message.

```

C2MOD = 0x00000001;           //Set CAN controller into reset
C2BTR = 0x001C001D;           //Set bit timing to 125k
C2MOD = 0x00000000;           //Release CAN controller

if(C2SR & 0x00000004)        //See if Tx Buffer 1 is free
{
    C2TFI1 = 0x00040000;      //Set DLC to 4 bytes
    C2TID1 = 0x00000022;      //Set address to 0x22 Standard Frame
    C2TDA1 = NetworkData;     //Copy some data into first four bytes
    C2CMR = 0x00000001;       //Transmit the message
}

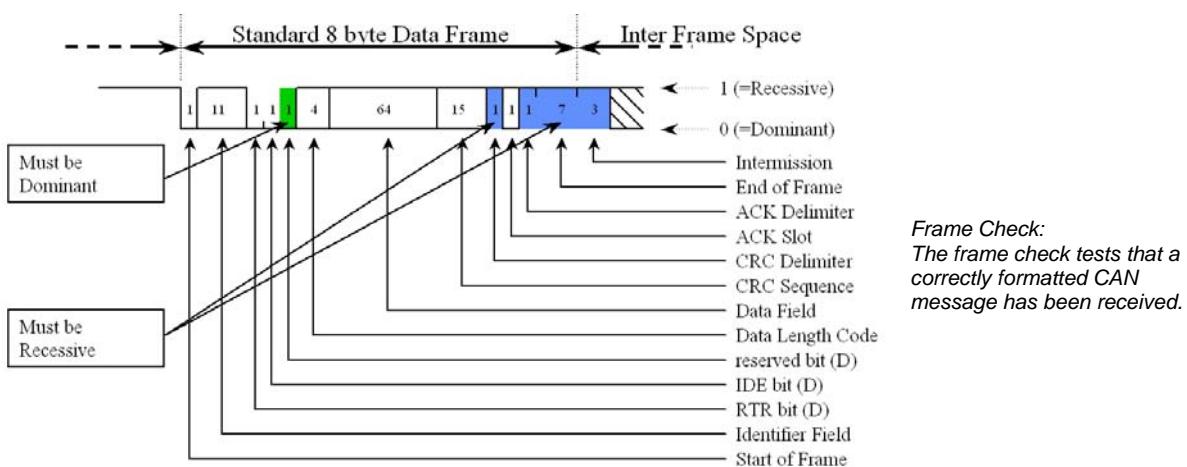
```

Exercise 25: CAN Transmit

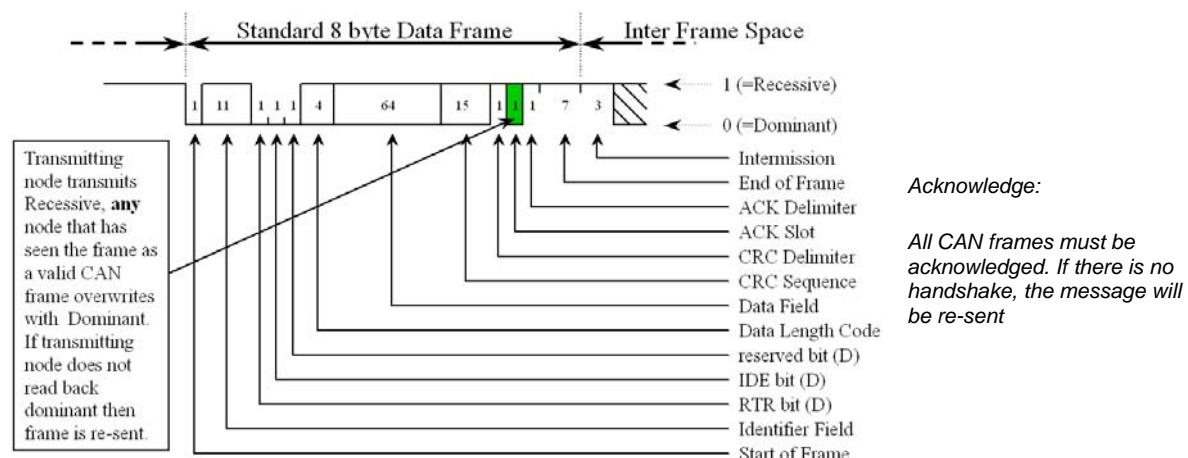
This exercise configures the second CAN channel for 125K bits\second and repeatedly transmits a CAN message frame.

CAN Error Containment

The CAN protocol has five methods of error containment built into the silicon. If any error is detected, it will cause the transmitter to resend the message so the CPU does not need to intervene unless there is a gross error on the bus. There are three error detection methods at the packet level; form check, CRC, and acknowledge plus two at the bit level; bit check error and bit stuffing error. Within the CAN message there are a number of fields that are added to the basic message. On reception, the message telegram is checked to see if all these fields are present. If not, the message is rejected and an error frame is generated. This ensures that a full, correctly formatted message has been received.

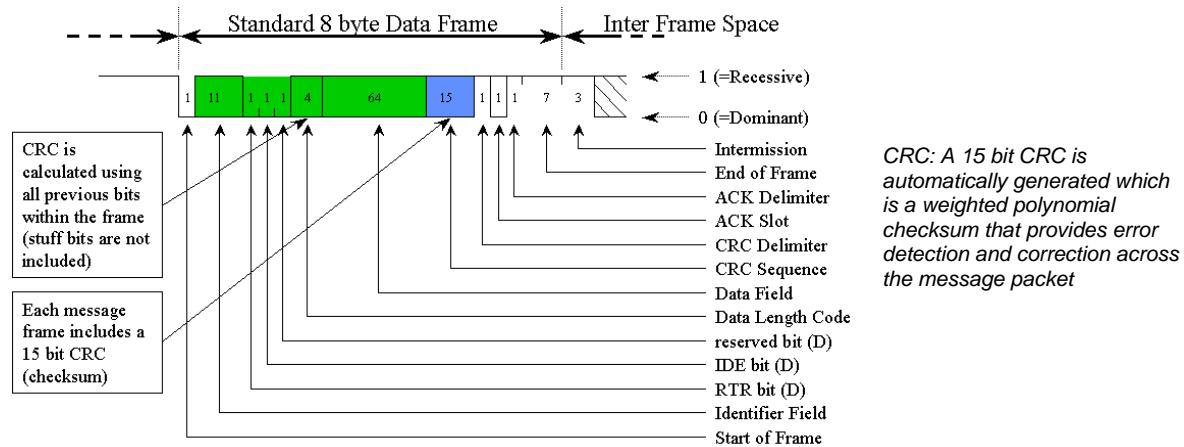


Each message must be acknowledged by having a dominant bit inserted in the acknowledge field. If no acknowledge is received, the transmitter will continue to send the message until an acknowledge is received.

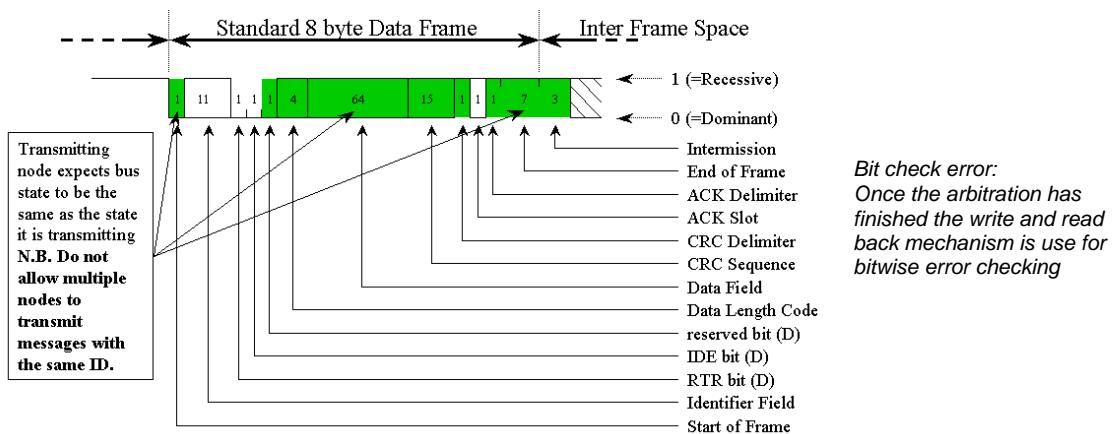


The CAN message packet also contains a 15 bit CRC which is automatically generated by the transmitter and checked by the receiver. This CRC can detect and correct 4 bits of error

in the region from the start-of-frame to the beginning of the CRC field. If the CRC fails and the message is rejected, an error frame is placed onto the bus.

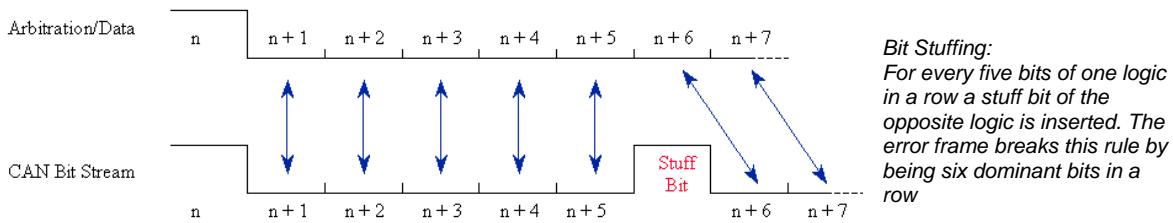


Once a node has won arbitration it will start to write its message onto the bus. As during arbitration as each bit is written onto the bus, the CAN controller is reading back the level written onto the bus. As the node has won arbitration nothing else should be transmitting so each bit level written onto the bus must match the level read back. If the wrong level is read back, the transmitter generates an error frame and reschedules the message. The message is sent in the next message slot but must still go through the arbitration process with any other scheduled message.



This leads to one of the golden rules in developing a CAN network. In a CAN network, every identifier must be uniquely generated. So you must not have the same identifier sent from two different nodes. If this happens, it is possible that two messages with the same ID are scheduled together, both messages will fight for arbitration and both will win as they have the same ID. Once they have won arbitration they will both start to write their data onto the bus. At some point this data will be different and this will cause a bit check error. Both messages will be rescheduled, win arbitration and go into error again. Potentially this ‘deadly embrace’ can lock up the network, so beware!

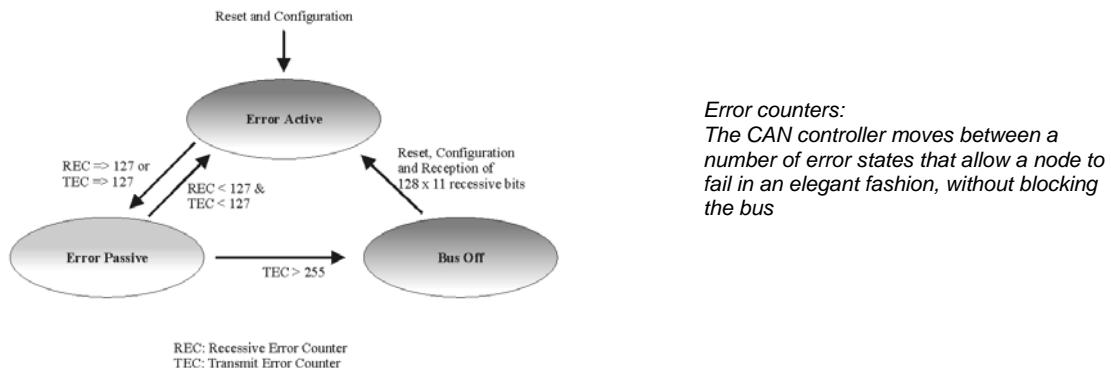
At the bit level, CAN also implements a bit stuffing scheme. For every five dominant bits in a row, a recessive bit is inserted.



Bit Stuffing:

For every five bits of one logic in a row a stuff bit of the opposite logic is inserted. The error frame breaks this rule by being six dominant bits in a row

This helps to break up DC levels on the bus and provides plenty of edges in the bit stream which are used for resynchronisation. An error frame in the CAN protocol is simply six dominant bits in a row. This allows any CAN controller to assert an error onto the bus as soon as the error is detected, without having to wait until the end of a message. Internally each CAN controller has two counters.



Error counters:

The CAN controller moves between a number of error states that allow a node to fail in an elegant fashion, without blocking the bus

These are a receive error counter and a transmit error counter. These counters will count up when receiving or transmitting an error frame. If either counter reaches 128, then the CAN controller will enter an ‘error passive’ mode. In this mode it still responds to error frames but if it generates an error frame, it writes recessive bits in place of dominant bits. If the transmit error counter reaches 255 then the CAN controller will go into a bus-off condition and take no further part in CAN communication. To restart communication, the CPU must intervene to reinitialise the controller and put it back onto the bus. Both these mechanisms are to ensure that if a node goes faulty, it will fail gracefully and not block the bus by continually generating error frames.

The LPC2000 CAN controllers have a number of error detection mechanisms. First of all, the current count of the transmit and receive error counters can be read in the Global Status Register.

Also in this register are two error flags, the Bus Status flag will be set when the maximum error count is reached and the CAN controller is removed from the bus. The second error flag is the Error Status flag, which is set when the CAN error counters reach a warning limit. This warning limit is an arbitrary value that is set by writing a value into the Error Warning limit register. The default value in this register is 96. Like the bit timing registers, the EWL

register may only be modified when the CAN controller is in reset. In addition, the Interrupt Capture Register provides extensive diagnostics for managing events on the CAN bus.

The CAN controller has the following interrupt sources,

1. Transmit interrupt (one for each buffer)
2. Receive interrupt
3. Error Warning
4. Data overrun
5. Wake up
6. Error Passive
7. Arbitration lost
8. Bus error
9. ID ready

CAN Message Reception

Once initialised, the CAN controller is able to receive messages into its receive buffer. This is similar in layout to the transmit buffers

Name	Description
CAN RFS	Reciever Frame Status
CAN RID	Recieved Identifier
CAN RDA	Recieved Data 1-4
CAN RDB	Recieved Data 5-8

The Rx Frame Status register is analogous to the Tx Frame information register. However it has two additional values. These are the ID Index and the BP bit and these will be explained in the next section.

The code below demonstrates how to receive a CAN message:

```
int main(void)
{
    VPBDIV = 0x00000001;           //Set PClk to 60MHz
    IODIR1 = 0x00FF0000;          // set all ports to output
    PINSEL1 |= 0x00040000;        //Enable Pin 0.25 as CAN1 RX
    C1MOD = 0x00000001;          //Set CAN controller into reset
    C1BTR = 0x001C001D;          //Set bit timing to 125k
    C1IER = 0x00000001;          //Enable the Receive interrupt
    VICVectCntl0 = 0x0000003A;   //select a priority slot for a given interrupt
    VICVectAddr0 = (unsigned)CAN1IRQ; //pass the address of the IRQ
                                    //into the VIC slot
    VICIntEnable = 0x04000000;    //enable interrupt
    AFMR = 0x00000001;           //Disable the Acceptance filters
    C1MOD = 0x00000000;          //Release CAN controller

    while(1){;}
}
```

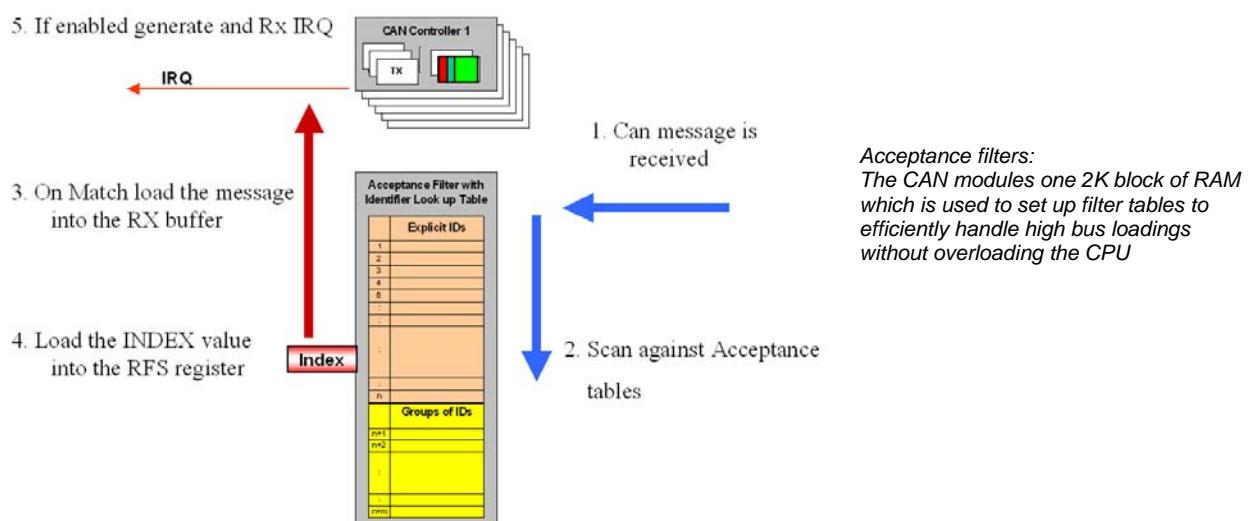
```

void CAN1IRQ (void) __irq
{
    IOCLR1 = ~C1RDA;      //clear output pins
    IOSET1 = C1RDA;       //set output pins
    C1CMR = 0x00000004;   //release the receive buffer
    VICVectAddr = 0x00000000; //Signal the end of interrupt
}

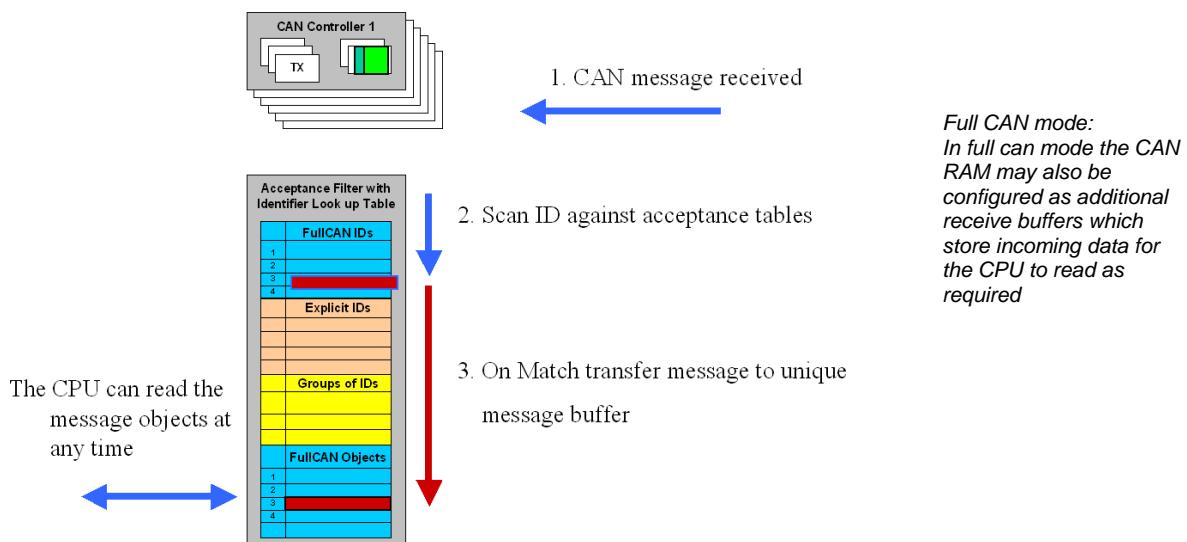
```

Acceptance Filtering

While the receive example shown above will work perfectly well, it suffers from two problems. Firstly, it receives every message transmitted on the bus. In a fully loaded CAN bus this could mean a message would be received every 72us. As the LPC2000 has up to 4 CAN controllers, the CPU would have to spend a lot of time just managing the CAN busses. Secondly, once the message has been received the CAN controller would have to read and decode the message identifier in order to decide what to do with the message. In order to overcome these problems, the LPC2000 CAN controllers have a sophisticated acceptance filtering scheme. The acceptance filter is used to screen messages as they come in from the CAN bus. The acceptance filter can be programmed to pass or block message identifiers before they enter the CAN controller for processing. This prevents unwanted messages entering the CAN receive buffer and consequently greatly reduces the overhead on the CPU. The acceptance filter has 2K of RAM (512 x 32), which may be allocated into tables of identifiers. This allows ranges of messages and individual messages to be able to enter into the CAN receive buffer.



As a message passes through the acceptance filter, it is assigned an ID Index. This is an integer number that relates to the message ID's offset in the acceptance filter table. This number is stored in the RX Frame Status register. So rather than decode the raw message ID, it is easier and faster to use the index value to decide what message has been received.

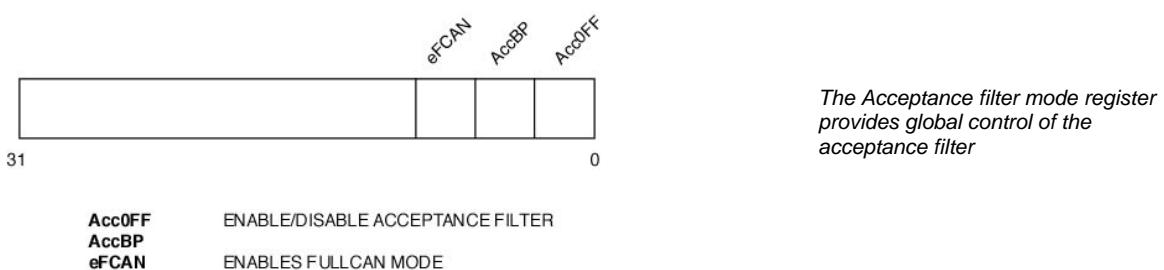


The acceptance filter also has a full CAN mode. In this mode the messages are received and scanned against the table of permissible identifiers. If a match is made, the message is stored not in the CAN controller receive buffer but in a dedicated message buffer within the acceptance filter memory. In this mode, each message has its own unique message buffer at a fixed location, making all the CAN data easily accessible from the CPU.

Configuring The Acceptance Filter

The acceptance filter is configured by seven registers. Control of the filter is via the mode register. The various ID tables are configured by the next five registers and the seventh register is an error reporting register.

Before configuration of the acceptance filter can start it must be disabled. This is done by setting the AccOff bit and clearing the AccBP bit in the acceptance filter mode register. If the CAN controller is run with this configuration, then all messages on the bus will be received.



Once the acceptance filter is disabled, each of the four filter tables may be configured. The four tables are as follows:

Individual standard identifiers	(11 bit ID)
Groups of standard identifiers	(11 bit ID)
Individual Extended identifiers	(29 bit ID)
Groups of extended identifiers	(29 bit ID)

The acceptance filter RAM starts at 0xE0038000. Each of the tables must be defined and fixed at absolute locations in the filter RAM. The start address of each table should then be written into the relevant acceptance filter register. The tables should start at the beginning of RAM and use the memory contiguously. Finally, the address of the last used location of RAM should be written into the End of Table register. To enable the Acceptance filter, set the ACCOff bit to logic one and AccBP bits to zero.

Each of the tables is constructed as follows;

31 15	29 13	26 10	16 0
Controller #	Dis able	not used	Identifier

The Individual Standard identifier table allows you to define individual 11-bit identifiers that will pass through the acceptance filter. Each definition takes two bytes, the first 11 bits contains the message identifier to be passed. This is followed by a bit to dynamically enable or disable this filter entry. Finally, the top three bits associates this filter entry with a particular CAN controller.

31	29	26	16	10	0		
Controller #	dis able	not used	Lower Identifier Bound	Controller #	dis able	not used	Upper Identifier Bound

The group standard identifier table uses the same format but two entries are used to define the upper and lower identifier address range for messages that are allowed to pass through the acceptance filter

31	29	28	0
Controller #	Identifier		

The individual extended identifier table uses four bytes per entry, as shown above. The first 29 bits define the message identifier to be passed through the acceptance filter and the top three bits associates the filter entry with a particular CAN controller. The group extended identifier table uses two words in the same format as the individual extended table to build up a start and end identifier values in the same fashion as the standard message group table

The following code shows how the acceptance filters may be configured for the basic CAN mode.

```

unsigned int StandardFilter[2] _at_ 0xE0038000; //Declare the standard
                                                 //acceptance filter table
unsigned int GroupStdFilter[2] _at_ 0xE0038008; //Next the standard Group
                                                 //filter table
unsigned int IndividualExtFilter[2] _at_ 0xE0038010; //Now the extended filter
                                                 //table
unsigned int GroupExtFilter[2] _at_ 0xE0038018; //Finally the Group extended
                                                 //filter table

AFMR = 0x00000001;                         //Disable the Acceptance filters
StandardFilter[0] = 0x20012002;               //Setup the standard filter table
StandardFilter[1] = 0x20032004;               //Allow Ids 1,2,3 & 4
SFF_sa = 0x00000000;                        //Set start address of Standard table
SFF_GRP_sa = 0x00000008;                     //Set start address of Standard group table
EFF_sa = 0x00000008;                        //Set start address of Extended table
EFF_GRP_sa = 0x00000008;                     //Set start address of Extended group table
ENDofTable = 0x00000008;                     //Set end of table address
AFMR = 0x00000000;                          //Enable Acceptance filters
C1MOD = 0x00000000;                         //Release CAN controller

```

Exercise 26 CAN Receive

Like the last exercise this example configures the CAN peripheral for 125Kbits/sec and sets the acceptance filters to receive one of three message frames.

Summary

This chapter is a bit of a moving target! The LPC2000 is a rapidly growing family with new variants being released on a regular basis. Check the CD that came with this book for a .PDF update to this chapter or keep an eye on the web at <http://www.hitex.co.uk/arm/lpcbook>

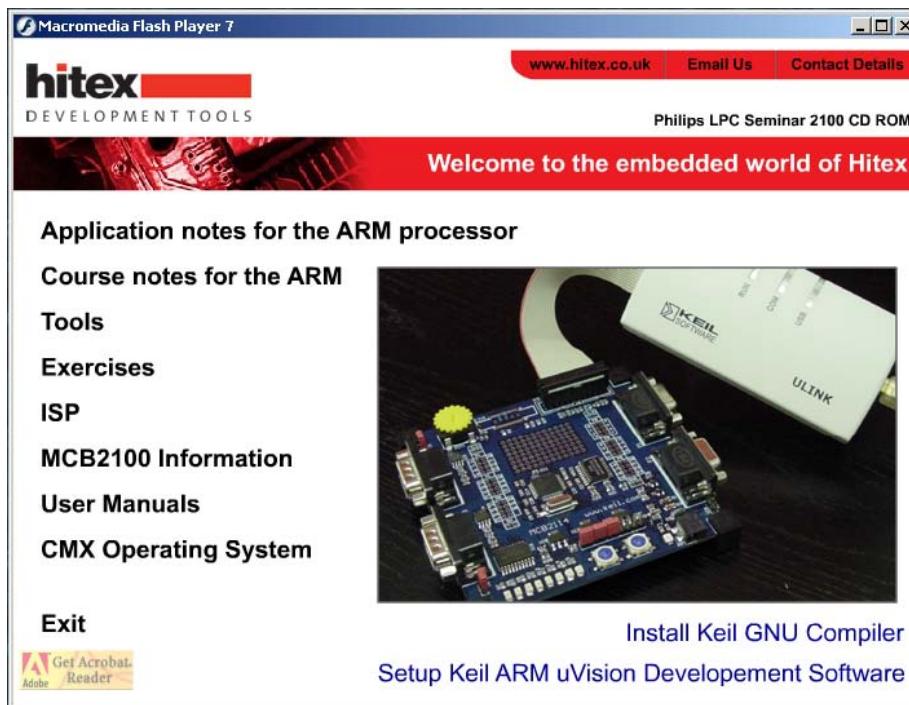
If you have worked through this and the proceeding chapters, you should now have a firm grasp of the LPC2000 family the ARM7 CPU and the necessary development tools. Appendix B lists further reading and web resources for the ARM7 and the LPC2000 in particular.

Chapter 5: Keil Tutorial

This chapter contains worksheets for the practical examples that are available on the CD. There are two sets of examples one for the Keil compiler and one for the GNU compiler. This chapter is written for the Keil compiler. If you want to use the GNU compiler you should start with Appendix A which details the non-ANSI additions to the GCC compiler. Appendix A also contains example worksheets for the first six exercises that deal with the specifically with GNU tools. After exercise six you can rejoin this chapter and use either the GNU or Keil examples for the remaining examples.

Installation

All the necessary software for the practical examples is on the CD that comes with this book. If you place the CD in the drive on your PC the following window will appear.



1. First it is necessary to install the Keil uVISION software, this will also install the Keil compiler.
2. If you wish to use the GNU compiler you will need to install this separately once uVISION is installed.

Next install the example set for the compiler you plan to use

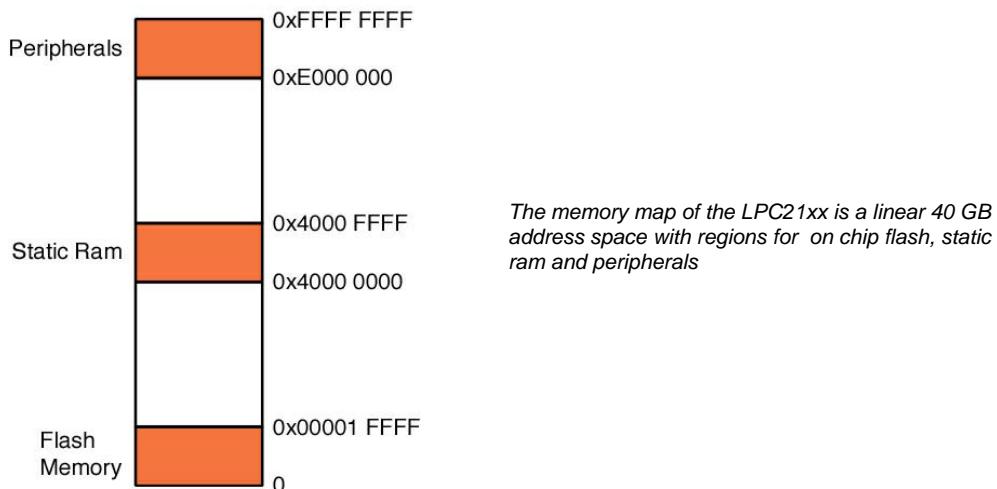
Finally install the Philips ISP flash programming tool

Once the software has been installed you are ready to start the tutorial exercises

Using the Keil UVISION IDE

This section will cover the development tools that can be used to develop code for the LPC2000. In this book all the example are written for the Keil ARM toolset. The Keil toolset comprises of the UVISION IDE which contains an editor and project manager, an ARM7 Compiler and linker and a Software simulator. The simulator will simulate the ARM7 core and the LPC2000 peripherals so it is possible to see the full operation of the chip by just using the simulator. An evaluation version of the toolset is available free from the Keil website at www.keil.com or on the CD supplied with this book. It is also possible to purchase a starter kit which contains an evaluation board and JTAG debugger that allows you to develop code on a real target device.

The Keil ARM compiler allows us to write in the C language and compile code to run on the LPC2000 devices. In order to cope with the microcontroller architecture the compiler has a number of non ANSI extensions that allow us to handle features such as interrupts ARM/Thumb interworking and accessing device peripherals. Before we start its worth looking at a simplified memory map of an LPC2000 so we can understand how to build a project.



Since the reset and exception vector table are located from zero upwards the on chip flash memory is located from zero for up to 256K. The on chip SRAM starts at 0x40000000 for up to 64K and the on chip peripherals are mapped from 0xE0000000 to the top of memory.

Before we begin to look at the compiler in detail I will run through a step by step tutorial on how to set up a UVISION project, compile the code and run the debugger. This does not cover all the features of uVISION but once you have a basic understanding of the IDE feel free to explore. If you copy the example from the CD onto your hard disk the source files referred to below can be found in C:\examples\ex1-first project. Ok lets build our first project.

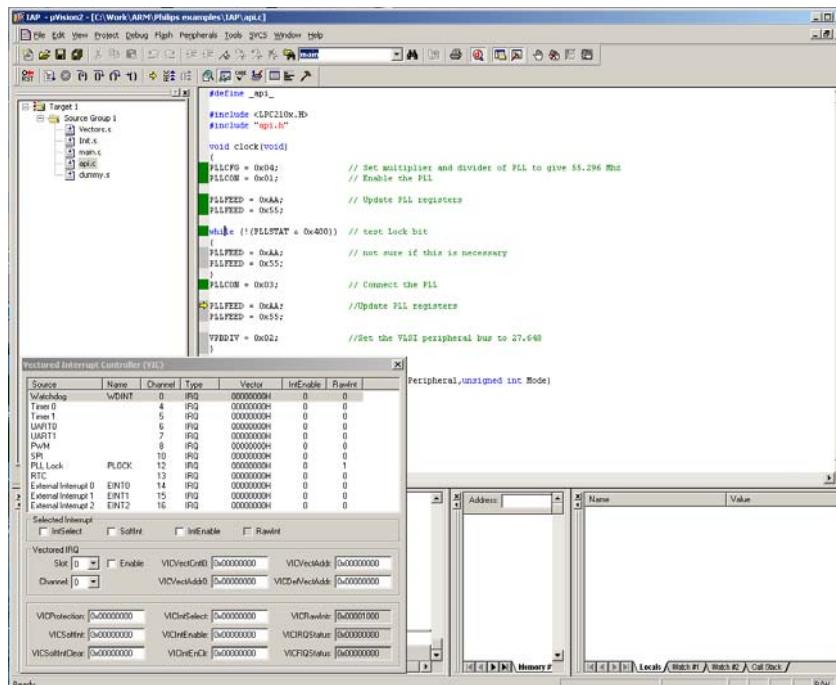
Exercise 1: Using the Keil Toolset

This example is based on the source code that can be found in

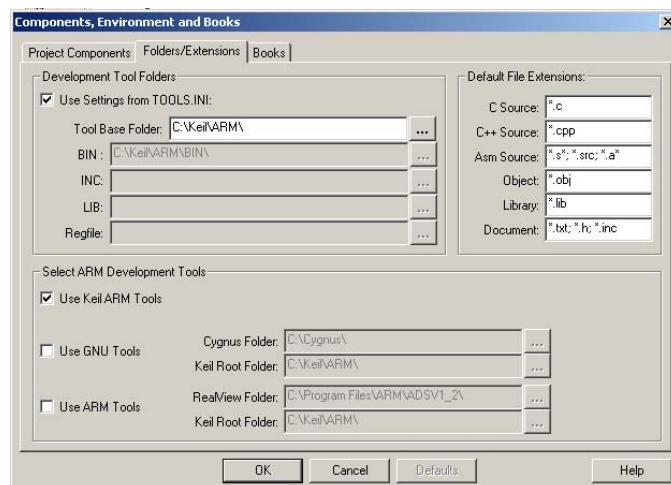
C:\Work\EX1 first program

In this first exercise we will spend some time defining a first project, building the code and downloading it into the simulator for debugging. We will then cover the basic debugging functions of the Keil simulator.

Double click on the Keil Uvision3 icon to start the IDE.

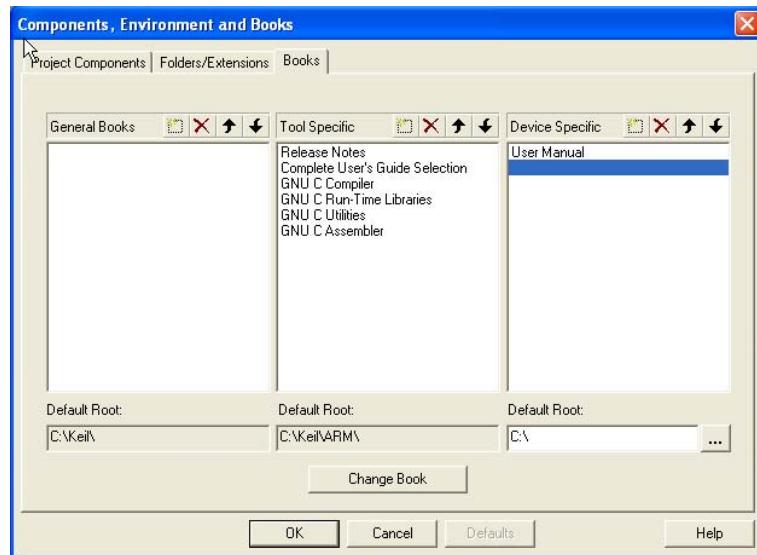


The Keil UVISION (a.k.a “uVISION”) IDE is designed to support several compilers, the Gnu C compiler, The ARM development suite and the Keil ARM compiler. Before compiling make sure you have the GNU compiler selected. This is done by activating the project workspace, right clicking and selecting manage components. In this dialog select the

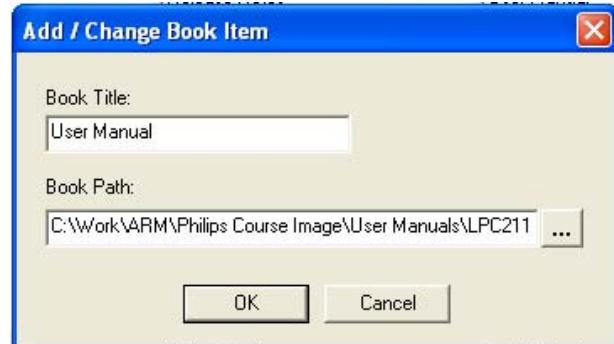


Folders/extensions tab and make sure the “Keil ARM tools” box is selected.

Next click on the “Books” tab.

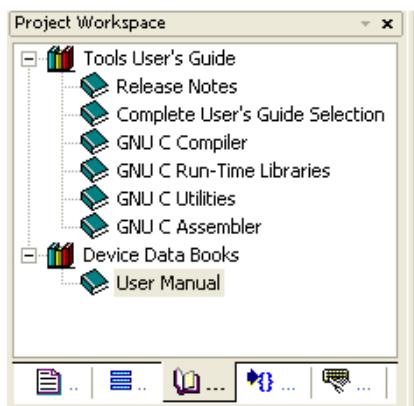


Highlight the UserManual entry and press the “Change Book” button.



Now change the path to point to the LPC2000 user manual, which can be found on the CD in the “User Manuals” directory. You can now add any other documentation you wish through these menus.

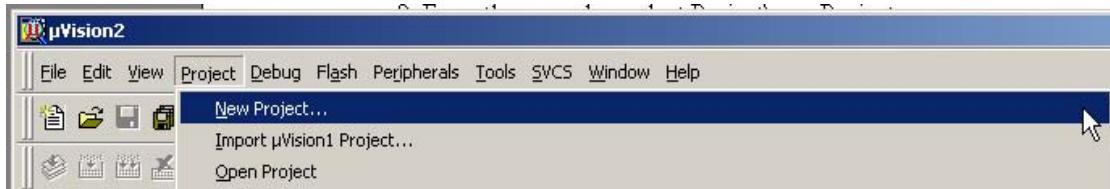
Once the project has been fully configured the on-line documentation can be accessed through the Books tab in the project workspace window. The full Keil documentation for



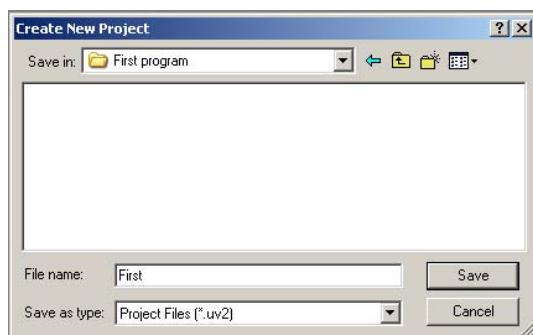
uVISION and the CARM compiler is found under “Complete Users Guide Selection”

Once you have added the datasheet click the OK button to continue defining the project.

From the menu bar select Project\new Project.

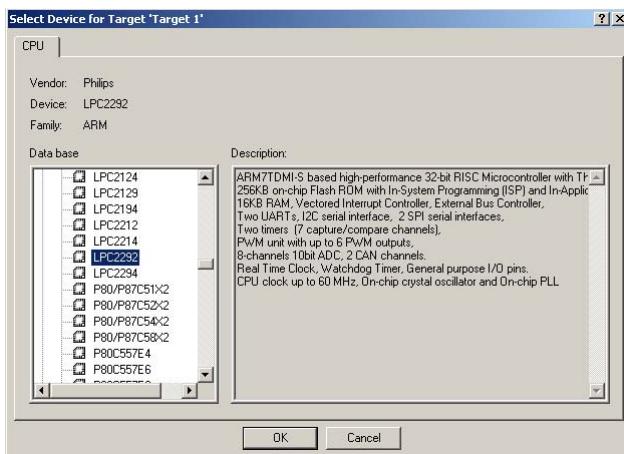


In the New project dialog navigate to your desired project directory.

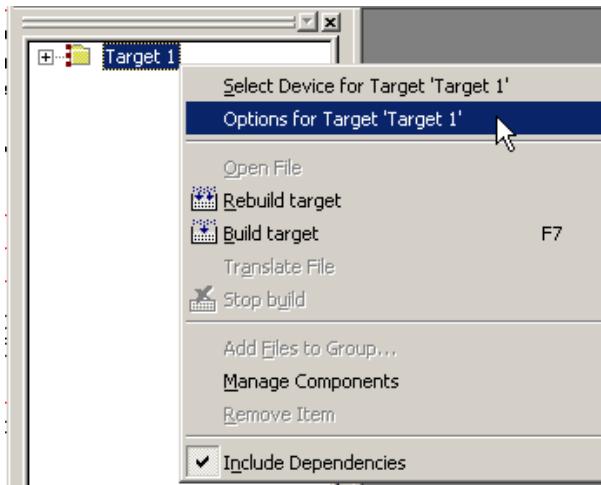


In the new project dialog name the project “first.uv2” and select Save.

A ‘select new device for target’ dialog will appear. Navigate through the device database and select the Philips\LPC2129 folder and then OK.



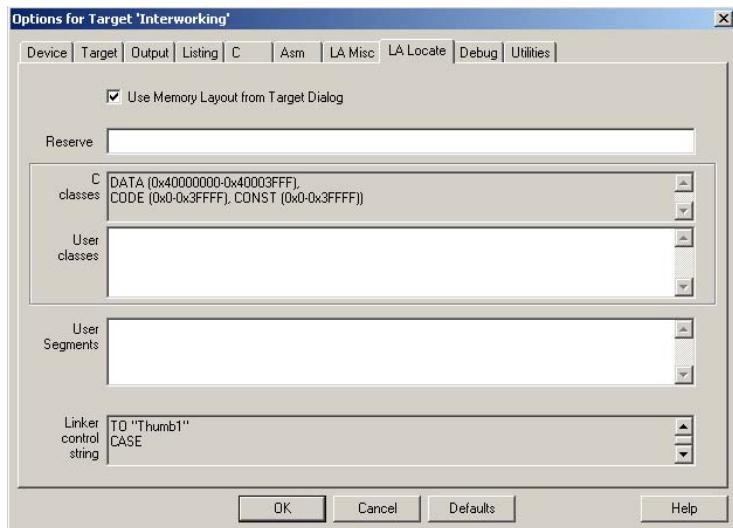
In the project browser highlight the ‘Target1’ root folder and select the local menu by pressing the right mouse button. In this menu select ‘options for target’.



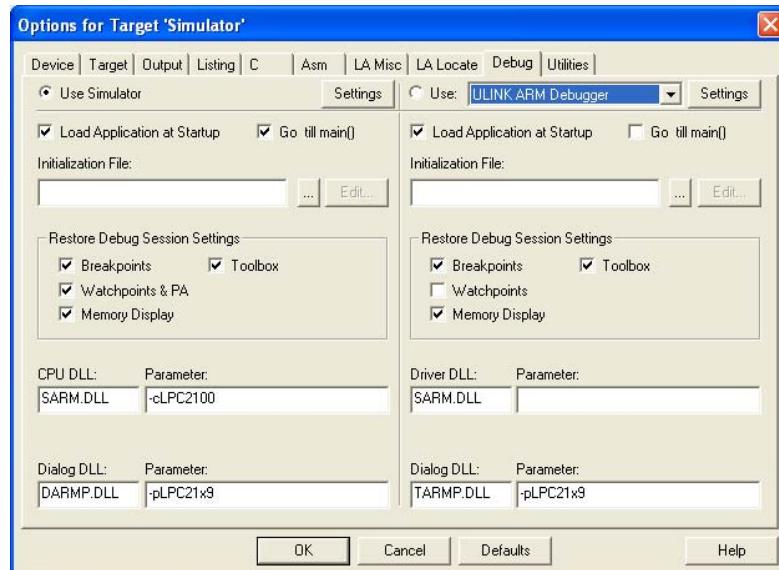
In the ‘Target’ tab, set the simulation frequency to 12.000 MHz. Also make sure the “Use on chip Rom” and “Use On chip Ram” boxes are ticked.



In the LA Locate tab, make sure that the “Use memory layout from target dialog” box is ticked.



In the debug tab select make sure the use simulator radio button is checked along with the “Load application at startup” and “Go till main”.

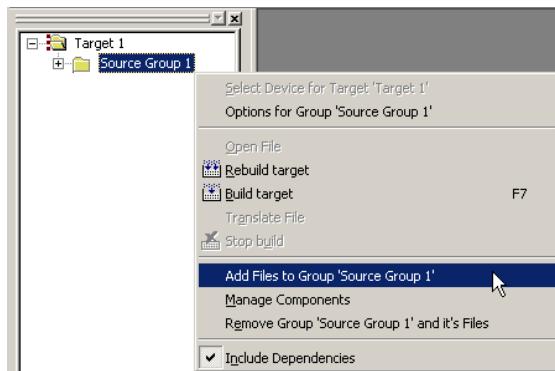


Select OK to complete the target options.

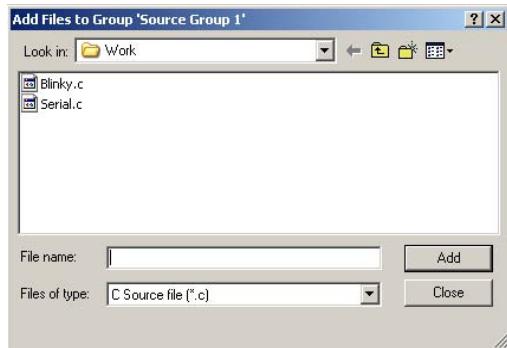
In the project browser expand the ‘target1’ root node to show the source group 1 folder.



Highlight the ‘source group 1’ folder, open the local menu with a right click and select ‘Add files to source group source group1’.



In the ‘Add files to group’ dialog add the file blinky.c and serial.c.



Change the ‘Type of file’ filter to ASM and add the file startup.s

These are all the source files necessary for the project so select “Close” .

Notes:

- (i) You can view the source code contained in a file by double clicking on the file name in the project browser window.
- (ii) The “manage components/project components” option also allows you to customise your project by adding extra source groups and different build options such as build for RAM debugging or Flash, debug in the simulator or with the JTAG.

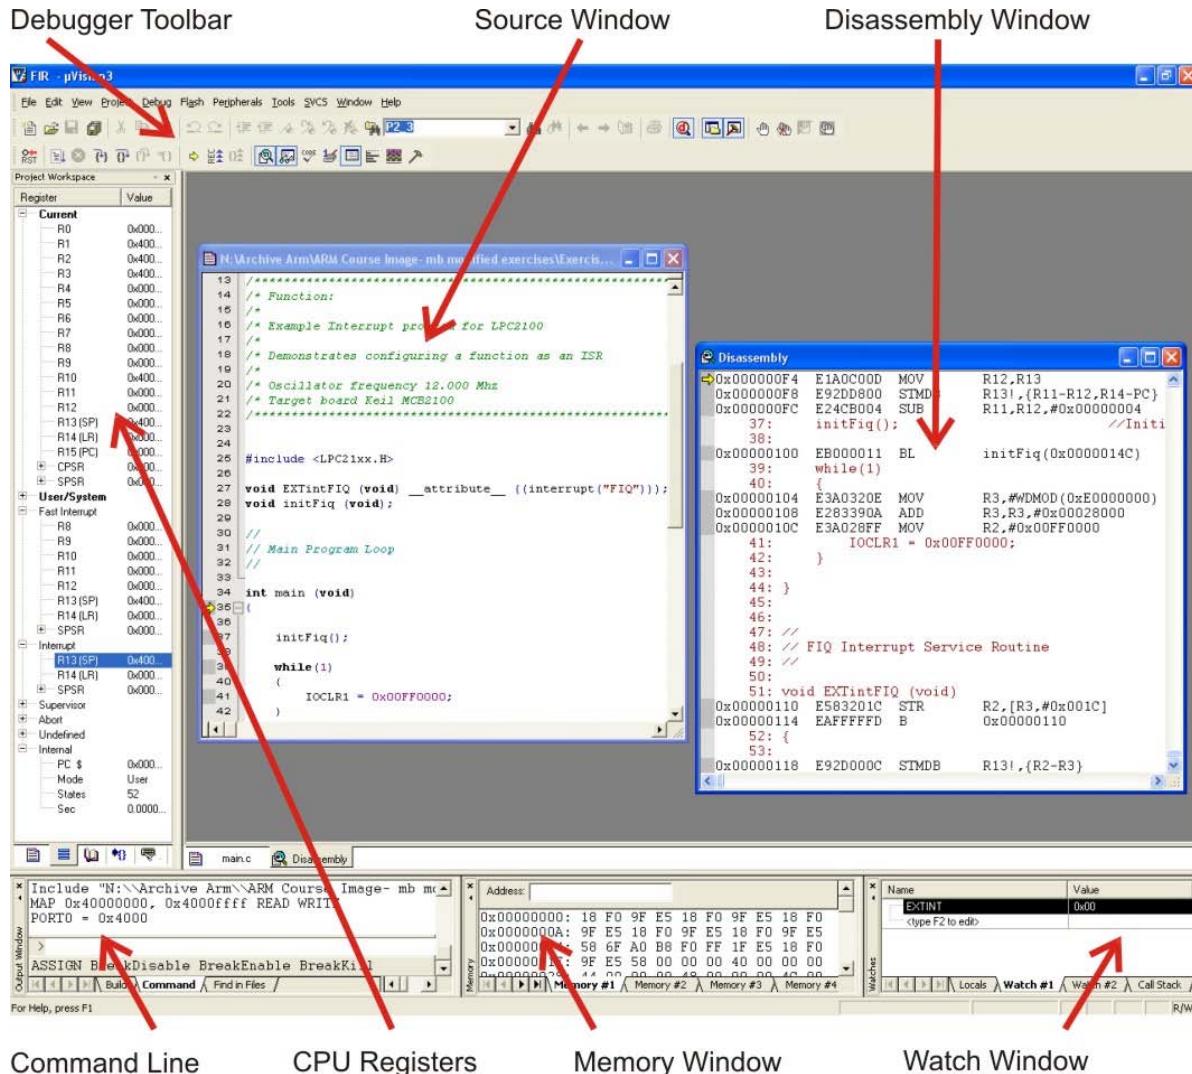
Build the code by selecting the Project\build target menu or the F7 key. Build Icons are also available on the toolbar.



For the rest of the tutorial, the projects will be defined but relevant bits of code will be missing. A complete copy of the exercise can be found in the solution directory. All the example code is included on your CD.

Using The Debugger

Launch the debugger by selecting the ‘debugger\start/stop debugger session menu or the button on the toolbar.



The code will be loaded into the simulator and executed from the reset vector until it reaches `main()`.

The project browser is replaced by a register window that allows you to view the contents of the CPU registers. Here you can:

- View the registers in each of the different operating modes
- Open the SPSR and CPSR registers to view the flags
- View the internal mode to the simulated cycle count and timestamp
- Change the contents of a register by triple clicking on its value and then entering a new value.

From main(), single step the code:

Use F11 to step a line of code

Use F10 to block step lines of code and functions

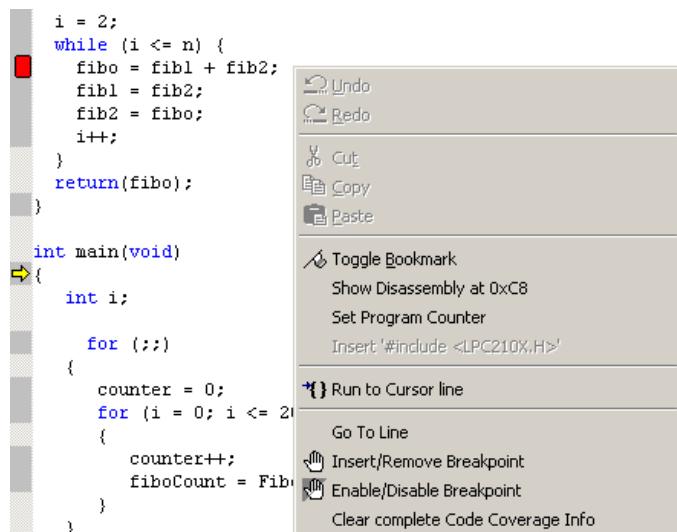
Use F5 to run the code at full speed

Use Esc to halt the code

Note: For the single step commands to work, a source code window must be the active window

To set a breakpoint:

Select a line of code, right click for the local menu and select “insert/remove breakpoint”. Now press F5 to run to the breakpoint.



Run to a point in the code:

This is a quick way of getting to an arbitrary pointing your code.

Select a line of code, open the local menu, “select run to cursor”. This will execute code until this line is reached. You can also select “Set program counter”. This forces the PC to the current position without running any code.

Select view/disassembly to see the underlying assembler code:

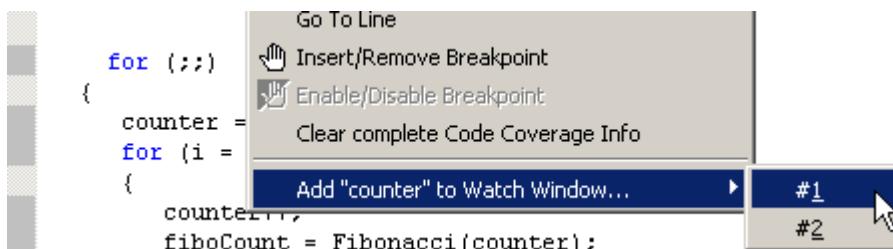
```

Disassembly
26: {
27:     int i;
28:
29:     for (;;)
30:     {
31:         counter = 0;
32:         for (i = 0; i <= 20; i++)
33:         {
34:             counter++;
35:             fiboCount = Fibonacci(counter);
36:         }
37:     }
38:     BL      Fibonacci(0x000000B0)
39:     LDR    R3,[PC,#0x0014]
40:     STR    R0,[R3]

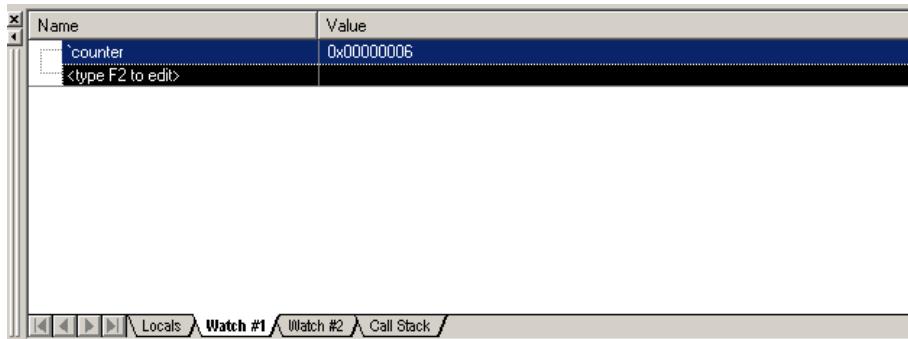
```

Examine program variables:

In the C code window place the cursor on the “counter” variable. Open the local window and select “Add counter to watch window” in the sub menu select #1



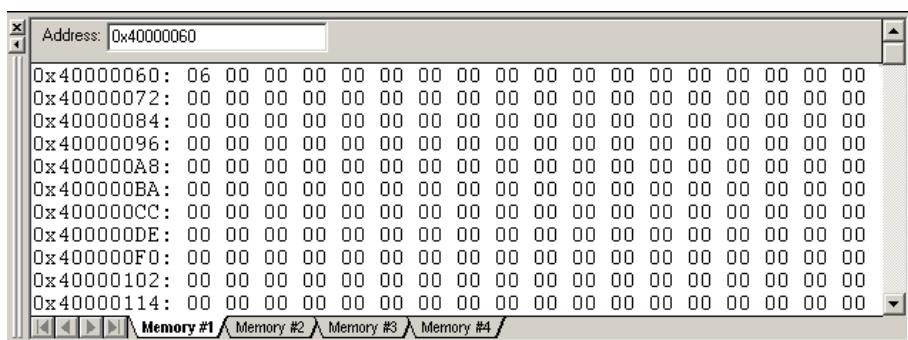
In the “Watch and call stack window”, select the Watch #1 tab to view the contents of counter



Examine memory locations:

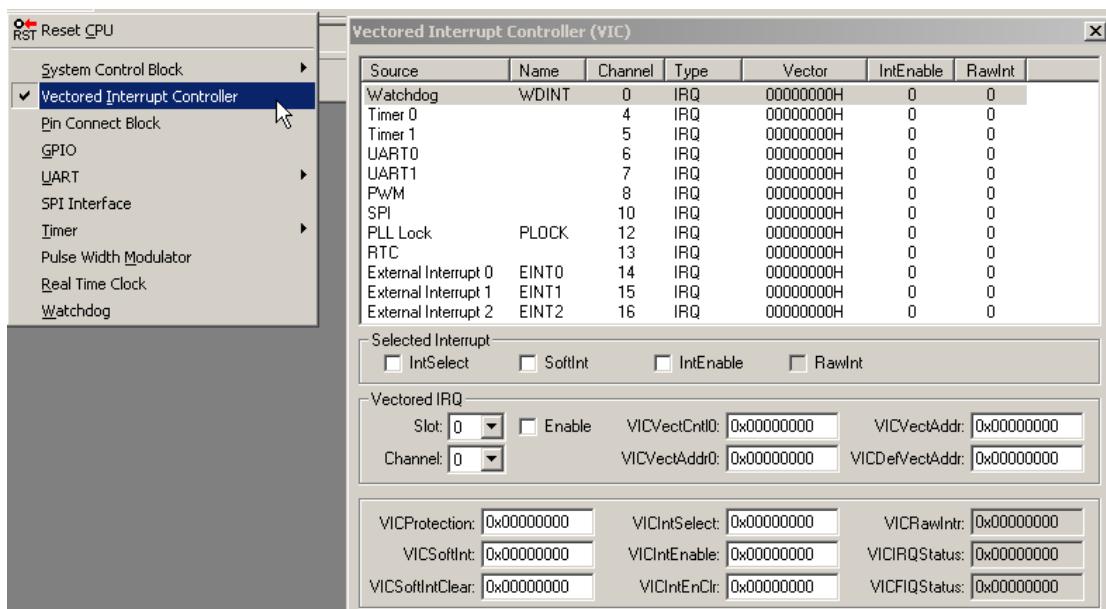
Open the memory window with “view\memory window”

Set the start of the memory window to 0x40000060, the address of the counter variable.



View the device peripherals:

Open some of the windows under the Peripherals menu.



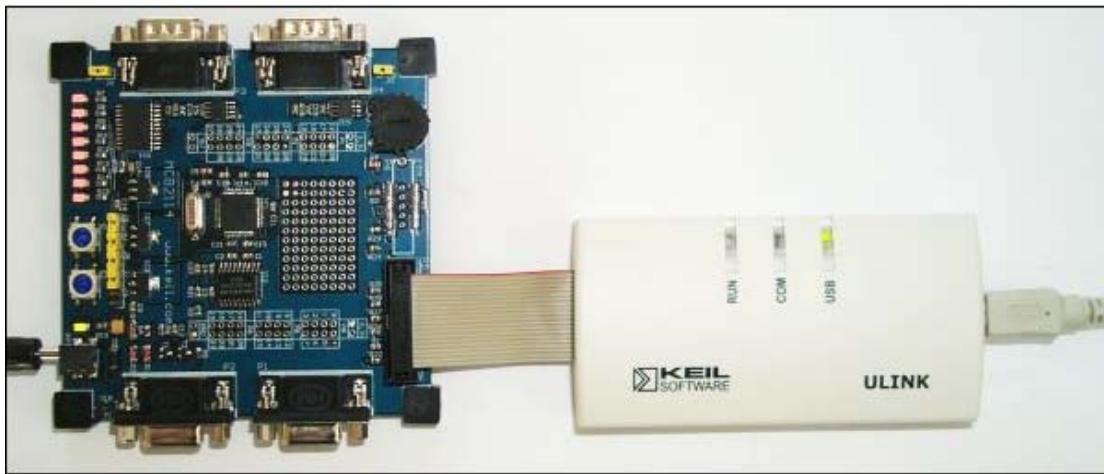
The debugger has many mode functions but the above debugging tools will allow you to run the course exercises. The simulator can also be replaced by a JTAG debugger and the same front end can be used to debug real hardware.

Using The ULINK Hardware Debugger

The JTAG debugger included with the starter kit is the Keil ULINK. This connects to the JTAG port on MCB2100 (P5) and then connects to the PC via USB. To switch from using the simulator to using the ULINK, follow the steps below.

Setting up the ULINK JTAG hardware debugger:

Connect the ULINK to the MCB2100 as shown below and plug the USB connection into the PC. Power should also be connected to the MCB2100 (6.5V).



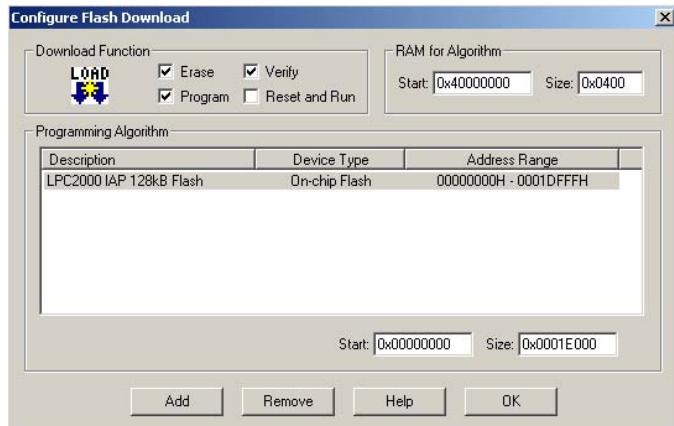
Configure UVISION to use the ULINK in place of the simulator:

First open the utilities menu in the “Options for target” dialogue . Select the “use target device for Flash programming” radio button and select the Ulink ARM7 debugger from the drop down menu. Also tick the update target before debugging box



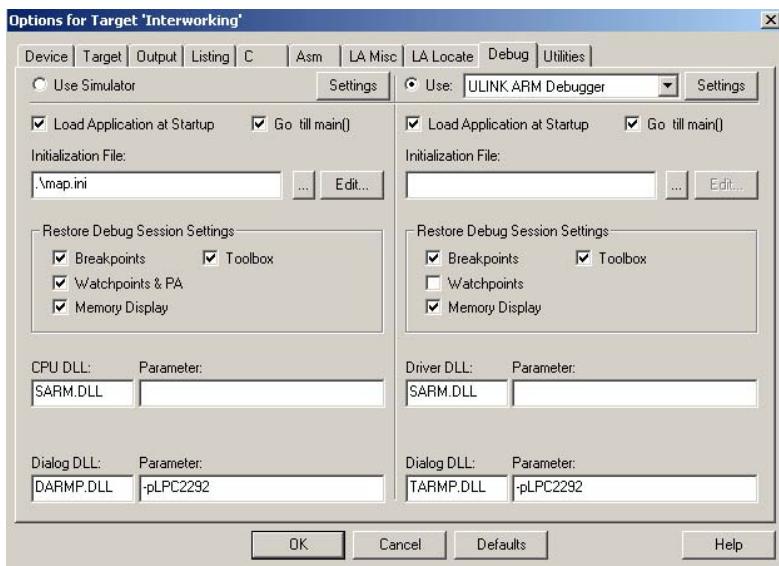
Configure the flash algorithm:

Next click the setting button. Set the start address for the RAM at 0x40000000 with a size of 0x400. Click the add button and select the flash algorithm for the device you are using then click OK to quit.



Switch from the simulator to the JTAG debugger:

Again open the options for target dialogue and select the debugger menu. On the right hand side of the menu select the ULINK ARM7 Debugger from the dropdown menu and tick the Use radio button.

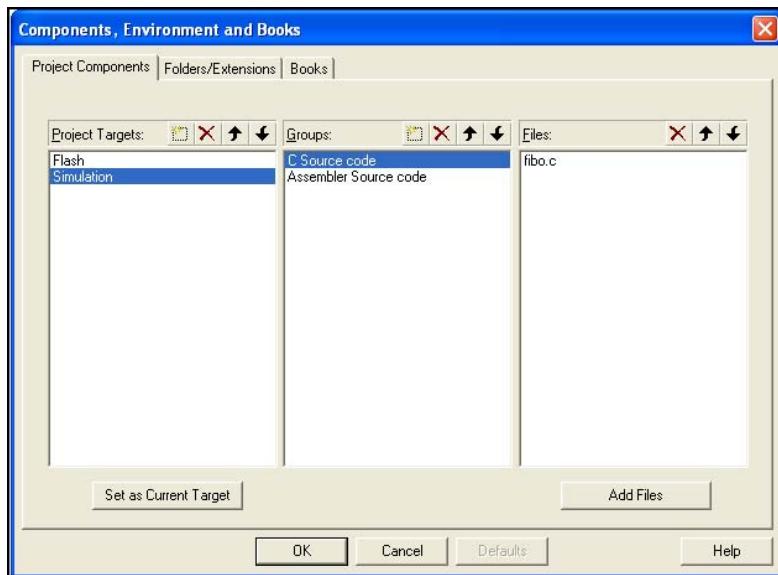


UVISION is now ready to use the ULINK JTAG in place of the simulator.

If you are in the JTAG debugger or simulator you must halt any running code and quit the debugger before you can rebuild the code or quit the project.

Important Note for the following exercises:

All the following exercises have builds to be debugged in either the Keil Simulator or to be downloaded onto the MCB2100 target hardware and debugged via the Ulink JTAG debugger. The root folder of the project will be named “Simulator” for the simulation version or “Flash” for the version built to be debugged on the hardware. Only the debug options in the project are changed. To switch between the two versions make the “project workspace” the active window right click and select manage components.



In the project components tab select the project target you want (either Simulation or Flash) and click the “Set as current target” button.

In a real project this feature of uVISION allows you to setup several different builds of a project or have several different programs in a project which are part of one larger application.

Exercise 2: Startup code

In this exercise we will configure the compiler startup code to configure the stack for each operating mode of the Arm7 we will also ensure that the interrupts are switched on and that our program is correctly located on the interrupt vector.

Open the project in C:\work\EX2 startup

Open the file Startup.s and using the graphical editor configure the operating mode stacks as follows.

Stack Configuration	
Top of Stack Address	0x4000 4000
Stack Sizes (in Bytes)	
Undefined Mode	0x0000 0080
Supervisor Mode	0x0000 0080
Abort Mode	0x0000 0080
Fast Interrupt Mode	0x0000 0080
Interrupt Mode	0x0000 0080
User/System Mode	0x0000 0400
PLL Setup	<input checked="" type="checkbox"/>

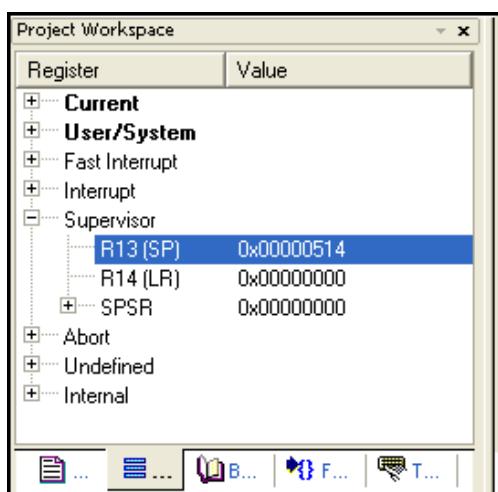
Now:

Compile the code

Start the simulator and when the PC reaches main examine the contents of each R13 register.

Examine the flags in the CPSR to determine the operating mode and instruction set being used

At last entry in the register window is the “Internal Mode” this gives additional information including timing information (if you are in the simulator).



Each stack is allocated a space of 0x80. The user stack is 0x400 bytes so user data will start at 0x40003d80 – 0x400

Exercise 3: Using THUMB code

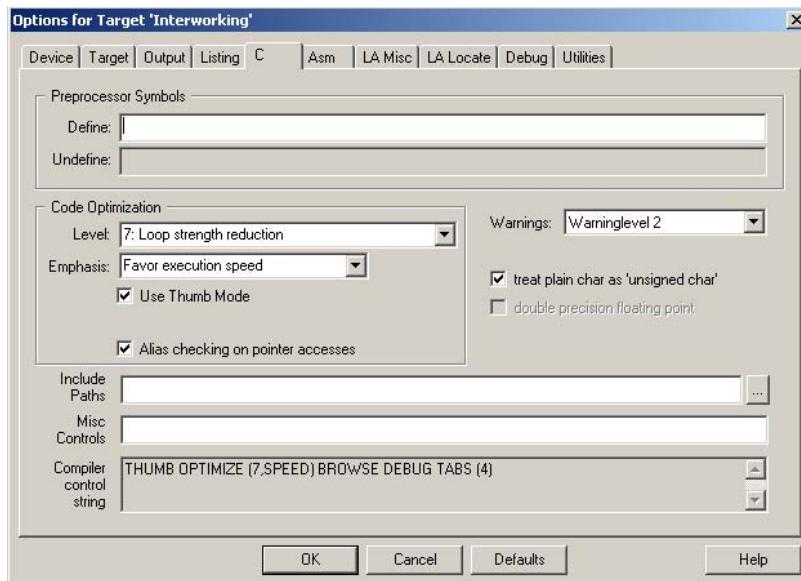
In this example we will build a very simple program to run in the ARM 32-bit instruction set and call a 16-bit thumb function and then return to the 32-bit ARM mode.

Open the project in C:\ work\ EX3 Thumb code

In main.c complete the function declarations for the main function and thumb function as follows.

```
void main (void) __arm  
void thumb_function(void) __thumb //Nb use double underscore
```

Again in the file browser select the root target (Flash) and in the local menu “options for target”. In the C tab you can select the global instruction set by checking or unchecking the “use thumb mode” box. It is also possible to define the instructions set to be used for a C module by checking or unchecking the same flag in the local C options menu for the C module.



Compile and download the code into the debugger

Open the disassembly window and single step through the code using the F11 key

Observe the switch from 32-bit to 16-bit code and the THUMB flag in the CPSR

32-bit ARM code

```

CPSR      0x60000010  0x00000013C  E59FC000  LDR      R12,[PC]
N         0           0x000000140  E12FFF1C  BX       R12
Z         1           0x000000144  000000149  DD       0x000000149
C         1           46: for (i = 0x00010000;i < 0x010000 ;i = i<<1)
V         0           47: {
I         0           48:
F         0           0x000000148  4806    LDR      R0,[PC,#0x0018]
T         0           0x00000014A  E007    B        0x00000015C
M         0x10

```

The processor is running in ARM (32-bit) mode, the T bit is clear and the instructions are 4 bytes long. A call to the THUMB function is made which executes a BX instruction forcing the processor into THUMB mode (16-bit).

Exercise 4: Using STDO libraries

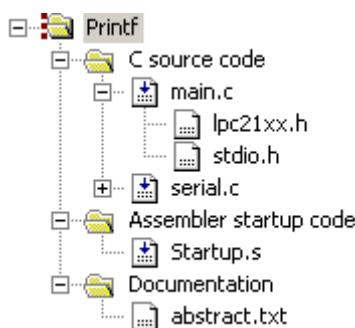
In this exercise we will look at tailoring the `Printf` function to work with the LPC2100 UART. We will look at the registers of the UART's in more detail later.

Open the project in EX4 printf\work

In main.c add a message for transmission to the `printf` statement

```
while(1)
{
    printf("Your Message Here \n"); //Call the printf function
}
```

Add the file serial.c in the work directory to the project



In serial.c complete the `putchar()` function so it writes a single character to the serial port.

```
int putchar (int ch)
{
    if (ch == '\n')
    {
        while (!(U0LSR & 0x20));
        U0THR = CR;
    }

    while (!(U0LSR & 0x20));

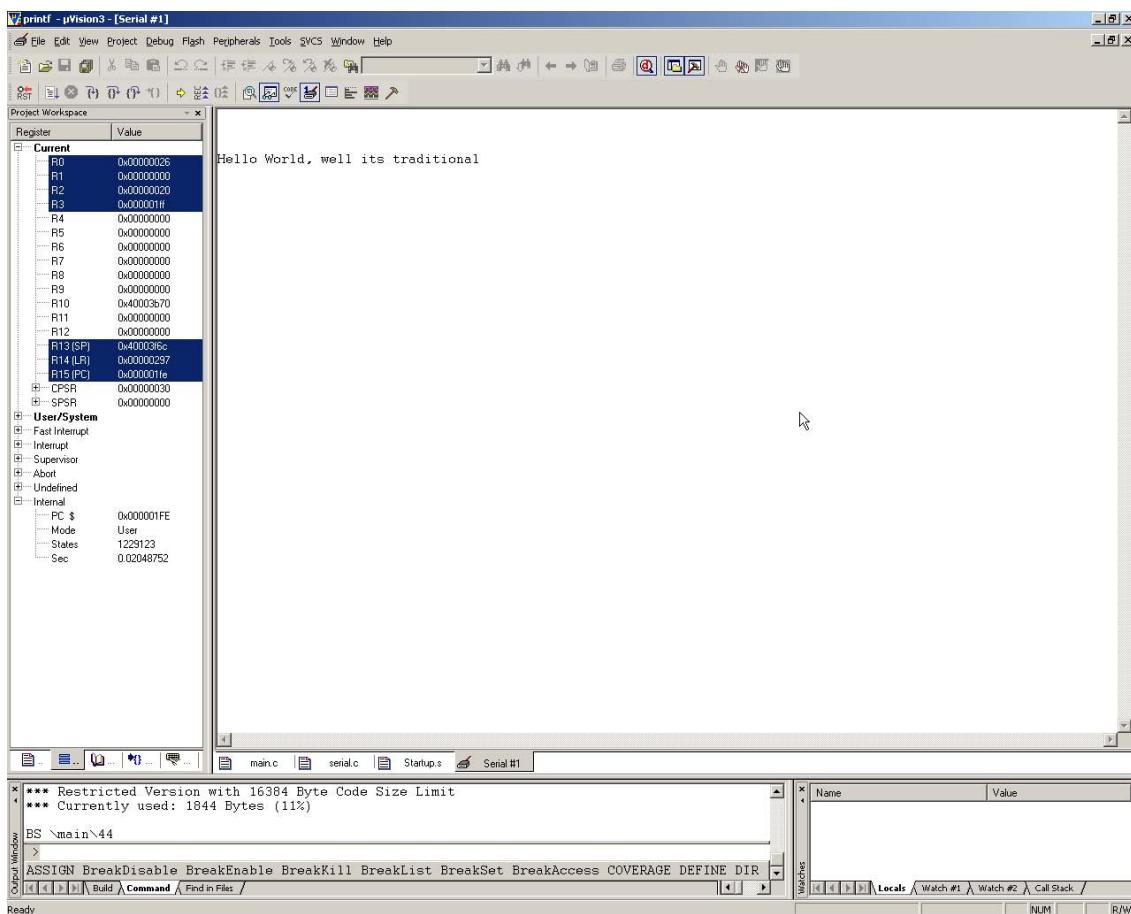
    return (U0THR = ch);
}
```

Compile the code and download it to the development board or simulator

Connect COM0 on the board to PC comm. Port 1 and start hyperterminal with the configuration file in the exercise directory.

If you are using the simulator select view/serial window #1. This opens a terminal window within the simulator that displays the UART0 output

Run the code and check that the message appears on the terminal window.



Exercise 5: Simple interrupt

In this exercise we will setup a basic FIQ interrupt and see it serviced.

Open the project in C:\work\EX5-Interrupt

In main.c complete the definition of the FIQ_Handler function to define it as the FIQ interrupt service routine

```
void FIQ_Handler (void) __fiq
```

In startup.s complete the vector constants table to define EXTintFIQ as the FIQ ISR.

__startup	PROC	CODE32	
Vectors:	LDR	PC, =Reset_Addr	
	LDR	PC, Undef_Addr	
	LDR	PC, SWI_Addr	
	LDR	PC, PAbt_Addr	
	LDR	PC, DAbt_Addr	
	NOP		/* Reserved Vector */
	LDR	PC, [PC, #--0xFF0]	
	LDR	PC, FIQ_Addr	
Reset_Addr:	DD	Reset_Handler	The FIQ interrupt vector, the instruction loads the address of the c routine into the PC
Undef_Addr:	DD	Undef_Handler	
SWI_Addr:	DD	SWI_Handler	
PAbt_Addr:	DD	PAbt_Handler	
DAbt_Addr:	DD	DAbt_Handler	
	DD	0	/* Reserved Address */
IRQ_Addr:	DD	IRQ_Handler	
FIQ_Addr:	DD	FIQ_Handler?A	The constants table holds the address of the FIQ C routine

Compile the code and download it onto the board.

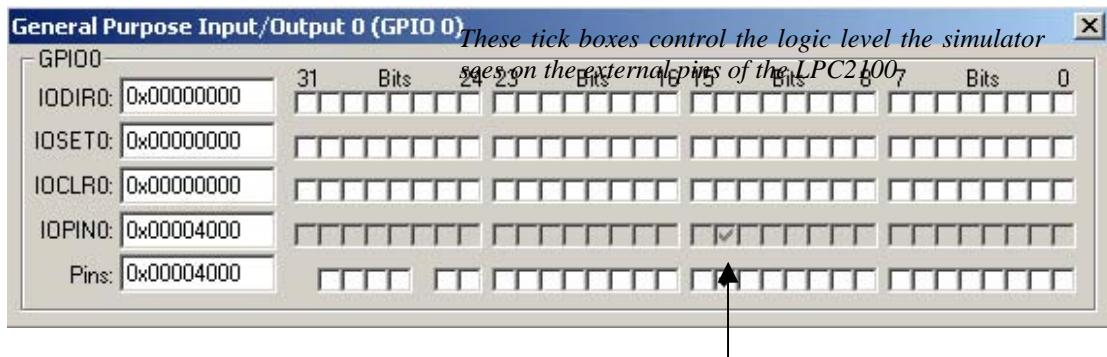
Step through the code until you reach the while loop

Set a breakpoint in the FIQ_Handler function

Press F5 to set the program running

On the MCB2100 board press the INT button to generate the interrupt.

If you want to see the entry and exit mechanisms to the exception it is best to use the simulator and single step in the disassembly window. This way you can watch the program flow and the actions on the CPU registers. To control the interrupt in the simulator open the peripherals/GPIO port 0 window. If you set the program running unchecking the Pin1.4 box will generate the interrupt. You must raise the pin high again to stop interrupts.



Alternatively in the toolbox there is a “Generate EINT1” button. This button will generate a simulated pulse onto the interrupt pin.

Toolbox button



Toolbox with user configurable scripts



Within uVISION there is a full scripting language that allows you to simulate external events. These scripts are based on the C language and are stored in text files. The script used to simulate the pulse is shown below.

```
signal void Toggle(void)
{
    PORT0 = (PORT0 ^ 0x4000);
    twatch(200);
    PORT0 = (PORT0 ^ 0x4000);
}

KILL BUTTON *
DEFINE BUTTON "GenerateEINT1", "Toggle()"
```

This script is stored in the file signal.ini and is added to the project in the debug window. For more details on the scripting language see the uVISION documentation.

Exercise 6: Software Interrupt

In this exercise we will define an inline assembler function to call a software interrupt and place the value 0x02 in the calling instruction. In the Software interrupt SWI we will decode the instruction to see which SWI function has been called and then use a case statement to run the appropriate code.

Open the project in C:\work\EX6 SWI\

Add the file SWI_VEC.S to the project

In Main.c declare the two function as SWI functions as follows

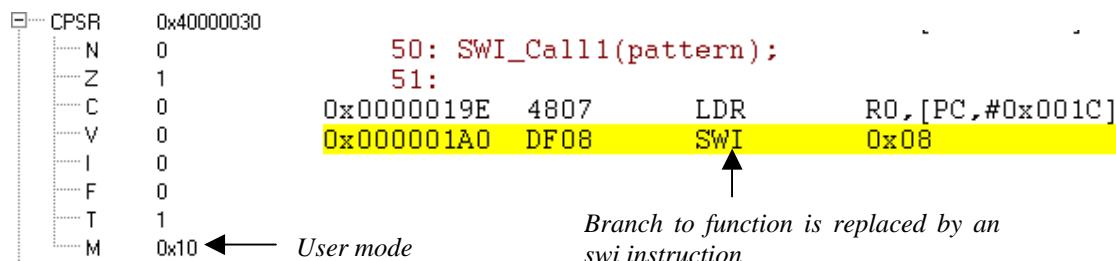
```
void SWI_Call1(int pattern) __swi(8)
{
    ...
}

void SWI_Call2(void)    __swi(9)
{
    ...
}
```

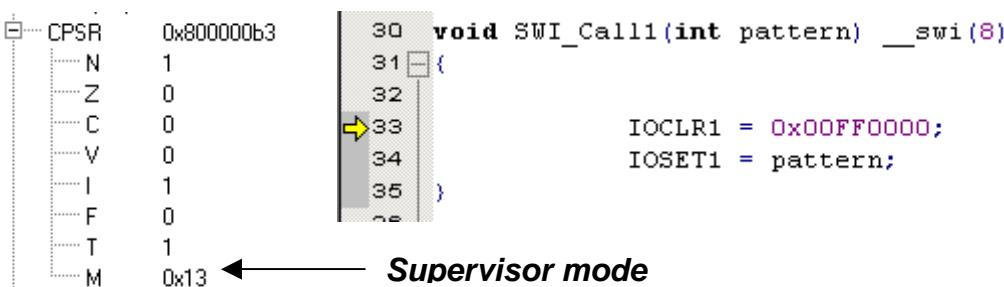
Compile and download the code into the debugger

Step the code and observe the SWI being serviced

In the disassembly window we can see the call to the first function. This passes the parameter and generates an SWI which is packed with the integer 8. In the register window we can read the CPSR which shows we are in user mode.



Once we have generated the software interrupt the chip will change modes run the code in vectors.s and then enter our number 8 SWI function .



Exercise 7: Memory Accelerator Module

This exercise demonstrates the importance of the memory accelerator module. Initially the PLL is set to 60MHz operation but the MAM is disabled. A simple LED flashing routine is used to illuminate the LEDs on the target board in sequence. This shows the sort of performance you can expect from an ARM7 running directly from on chip FLASH memory. When the value of the potentiometer is changed the MAM is enabled and the code will run faster making the LEDs flash faster. This increase in performance caused solely by the MAM which is why it is so important to this kind of small single chip microcontroller. In this example we will use the bootloader to load the code into the flash in place of the JTAG.

Open the project in exercises/EX7-MAM

In main.c complete the code to enable the MAM

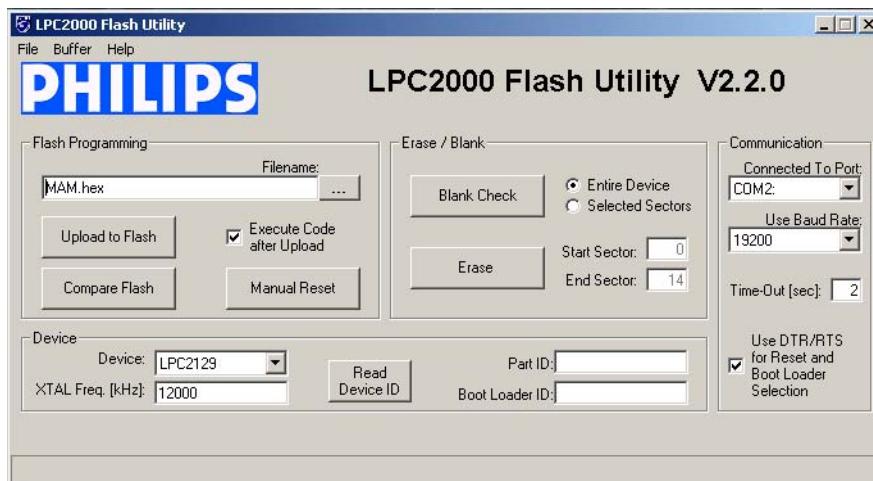
In Options for target/output tick the generate hex box

Build the code

Connect the PC serial port to COM0 on the target board

Apply power to the board.

Start the Philips ISP Utility to get the screen shown below



Make sure the “Use DTR/RTS” box is ticked.

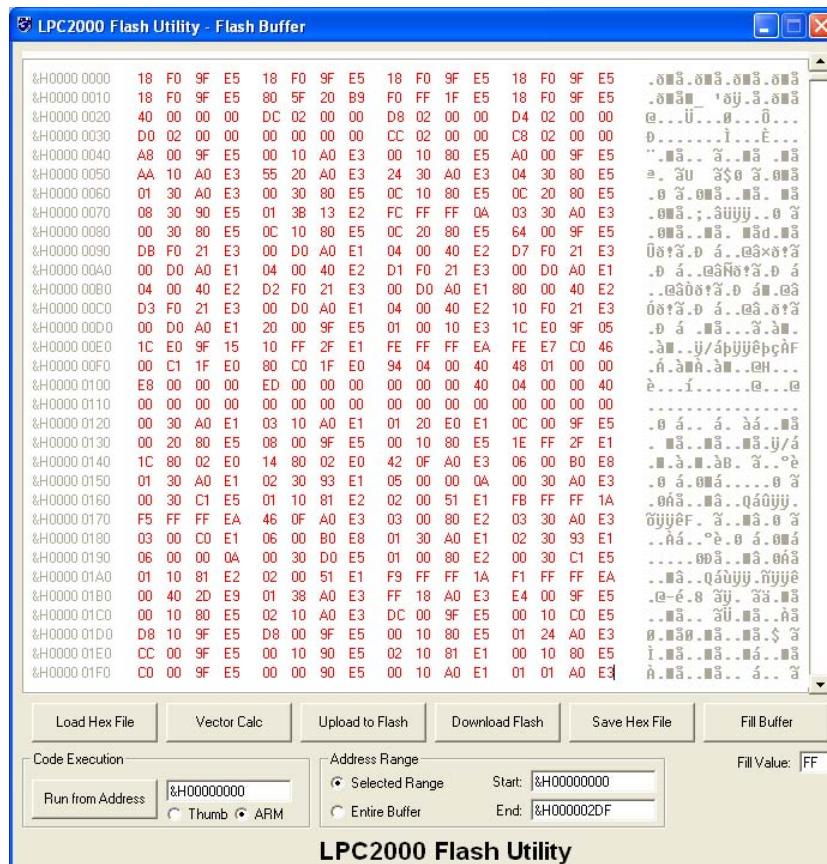
Press the “Read Device ID” button, If the board is connected ok the part ID number and bootload version will be displayed.

Make a note of these numbers as we will use them in the next exercise.

Next select the MAM.hex file from the project directory and press “Upload to Flash” this will program the target LPC2100.

You can also use the “Compare” button to verify that the flash has programmed correctly.

If you select the buffer option the same operations can be performed along with calculation of the program signature and limited debugging options.



Once the Target LPC2100 has been programmed the chip will automatically be rebooted and start to run your code.

Turn the potentiometer fully clockwise

Reset the code and the LED's will start to sequence

If you turn the potentiometer anticlockwise the mam will be enabled and you can see the LPC2000 “turbo” kick in.

In the ISP utility under the buffer option you can view a HEX dump of your program. In this view the calculated program signature is also shown.

If you reconnect the JTAG and start the debugger without downloading the program you can examine the interrupt vector table. As we have programmed the flash with the Philips ISP tool the program signature has been added in location 0x00000014 which exists as a NOP in the startup code.

Exercise 8: In-Application programming

This example demonstrates how to call the bootloader API from within your application programs. The code makes a call into the bootloader functions to run the “Read\Part ID” function. This example has only been built to work with the MCB2100 evaluation board.

Open the project in C:\work\ EX8 IAP

Add the following lines of code to main.c

```
command[ 0 ] = 0x36;  
iap(command,result,0x7FFFFFFF0);
```

Add the following lines to API.c

```
#pragma asm  
  
    mov r15,r2;           //move entry address into PC  
  
#pragma endasm
```

Now make sure that main is compiled as ARM code and API is built as Thumb code

Build the code and start the debugger

Run the code up to the first call to the IAP function.

Set a breakpoint on the line of code after this function

Step into the IAP function

Look at the contents of R0-R2, These registers should contain the addresses of the command table, the results table and the entry address of the bootloader functions. Also R14 should contain the return address to the main function.

Step the line of assembly code which will jump you to 0x7FFFFFFF. This is the entry to the bootloader code.

Now run the code at full speed, it will execute the bootloader function called and return to the address held in R14 (the next line in main.c).

When the code hits the breakpoint examine the CPSR to see which instruction set is now being used.

Take a look at the result table and read the PartID number returned by the bootloader. Compare this to the value returned in the previous example.

Run the next call to the Bootloader API and check that the correct bootloader version is returned.

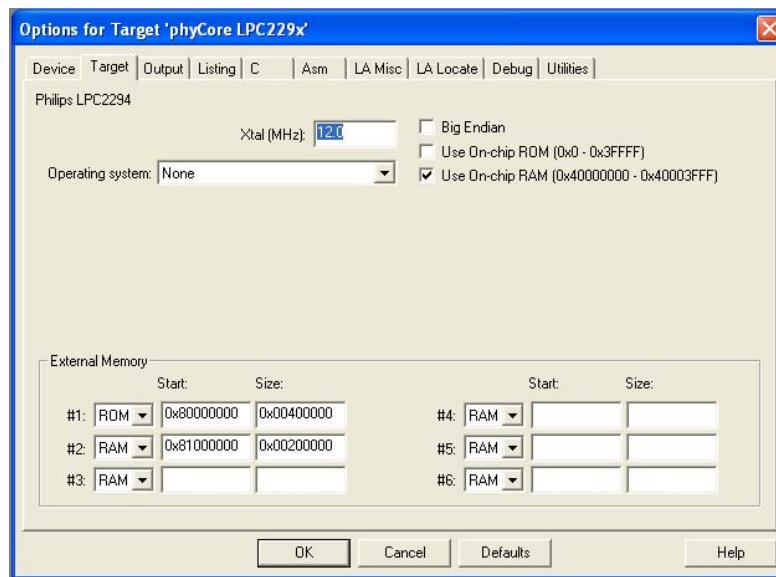
Exercise 9: External Bus Interface

This exercise demonstrates how to build a project to use the external bus. It is necessary to link the code so it is located in an external flash device, configure any chipselect that is used. Because the MCB2100 evaluation board uses a true single chip device this example uses the simulator to demonstrate the external bus interface. The project is built to place code in flash at 0x80000000 which is the boot chipselect and RAM at 0x81000000 which is chipselect one. On real hardware the boot method is determined by the state of the two external boot pins, in the simulator we use a script file to set the correct condition. However we will also look at configuring the JTAG to program external flash memory so this example can be used on real hardware.

Open the project in C:\work\ebi\

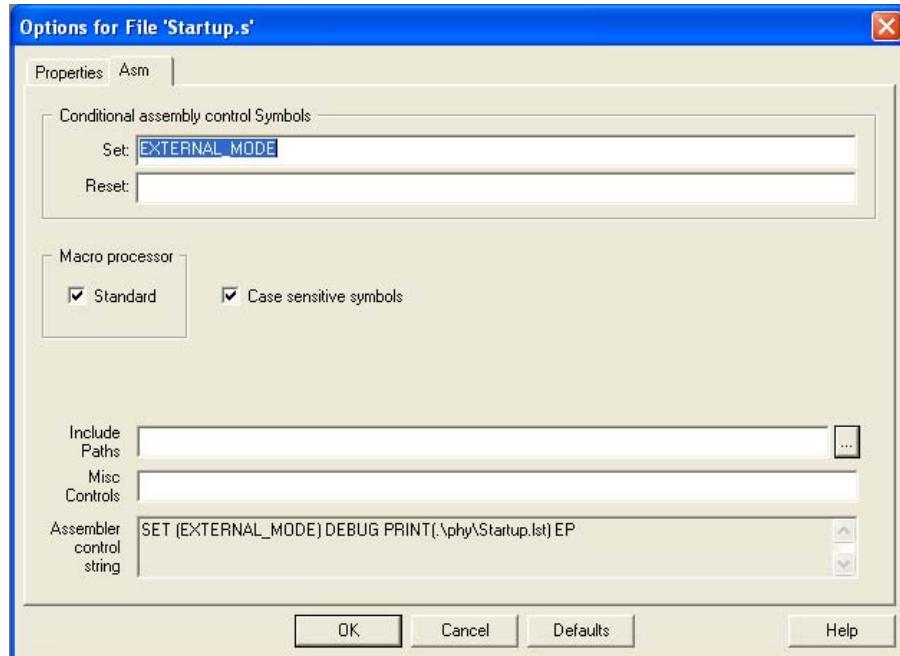
In the “Options for target” menu select the device tab and select the LPC2294.

Next, select the Target tab and fill in two regions of external memory and uncheck the “Use on chip ROM” as shown below.



Next highlight the startup.s file right click and select “Options for file” and select the ASM tab.

Enter the “EXTERNAL_MODE” symbol in the control directives box as shown below.



This will ensure the startup code is correctly built for an external boot.

Next in the graphical display of the startup code add the parameters necessary for the chipselect configuration

External Memory Controller (EMC)	<input checked="" type="checkbox"/>
Bank Configuration 0 (BCFG0)	<input checked="" type="checkbox"/>
IDCY: Idle Cycles	3
WST1: Wait States 1	7
WST2: Wait States 2	7
RBLE: Read Byte Lane Enable	<input type="checkbox"/>
WP: Write Protect	<input type="checkbox"/>
BM: Burst ROM	<input type="checkbox"/>
MW: Memory Width	32-bit
Bank Configuration 1 (BCFG1)	<input checked="" type="checkbox"/>
IDCY: Idle Cycles	3
WST1: Wait States 1	7
WST2: Wait States 2	7
RBLE: Read Byte Lane Enable	<input checked="" type="checkbox"/>
WP: Write Protect	<input type="checkbox"/>
BM: Burst ROM	<input type="checkbox"/>
MW: Memory Width	32-bit

Build the code and start the simulator

In the disassembly window the PC is set to address 0x00000000 and no meaningful code has been loaded.

0x00000000	00000000	ANDEQ	R0,R0,R0
0x00000004	00000000	ANDEQ	R0,R0,R0
0x00000008	00000000	ANDEQ	R0,R0,R0
0x0000000C	00000000	ANDEQ	R0,R0,R0
0x00000010	00000000	ANDEQ	R0,R0,R0
0x00000014	00000000	ANDEQ	R0,R0,R0
0x00000018	00000000	ANDEQ	R0,R0,R0
0x0000001C	00000000	ANDEQ	R0,R0,R0
0x00000020	00000000	ANDEQ	R0,R0,R0
0x00000024	00000000	ANDEQ	R0,R0,R0
0x00000028	00000000	ANDEQ	R0,R0,R0
0x0000002C	00000000	ANDEQ	R0,R0,R0
0x00000030	00000000	ANDEQ	R0,R0,R0
0x00000034	00000000	ANDEQ	R0,R0,R0
0x00000038	00000000	ANDEQ	R0,R0,R0
0x0000003C	00000000	ANDEQ	R0,R0,R0
0x00000040	00000000	ANDEQ	R0,R0,R0

Open the peripherals/system control block/ memory mapping control



Switch from “User Flash Mode” to “Ext Flash Mode”

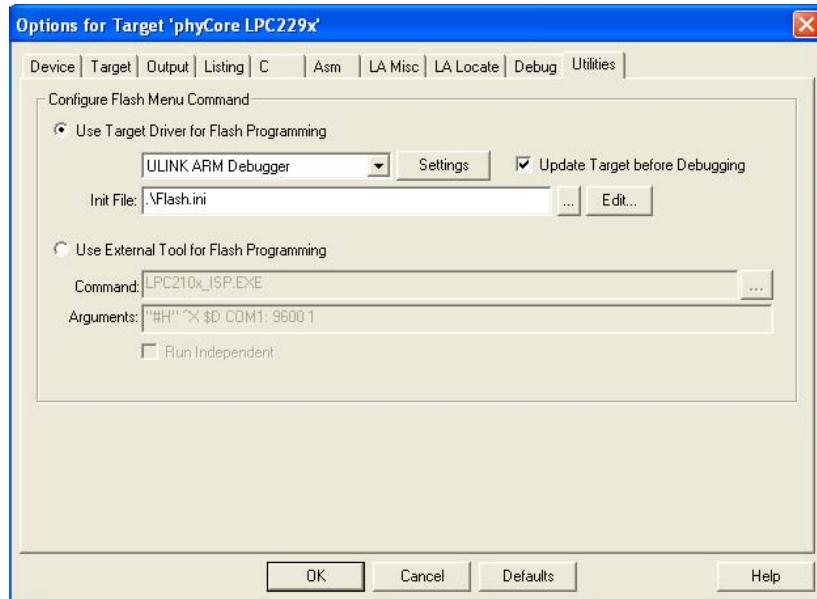
This will map the first 64 bytes of chipselect Zero into address 0x00000000 upwards. This is our vector table that is located in the external flash.

0x00000000	E59FF018	LDR	PC,[PC,#0x0018]
0x00000004	E59FF018	LDR	PC,[PC,#0x0018]
0x00000008	E59FF018	LDR	PC,[PC,#0x0018]
0x0000000C	E59FF018	LDR	PC,[PC,#0x0018]
0x00000010	E59FF018	LDR	PC,[PC,#0x0018]
0x00000014	E1A00000	NOP	
0x00000018	E51FFFF0	LDR	PC,[PC,#-0x0FF0]
0x0000001C	E59FF018	LDR	PC,[PC,#0x0018]
0x00000020	80000040	DD	0x80000040
0x00000024	80000324	DD	0x80000324
0x00000028	80000320	DD	0x80000320
0x0000002C	8000031C	DD	0x8000031C
0x00000030	80000318	DD	0x80000318
0x00000034	00000000	DD	0x00000000
0x00000038	80000314	DD	0x80000314
0x0000003C	80000310	DD	0x80000310

If you single step off the reset vector you will jump to 0x80000040 and start running code from the external Flash device.

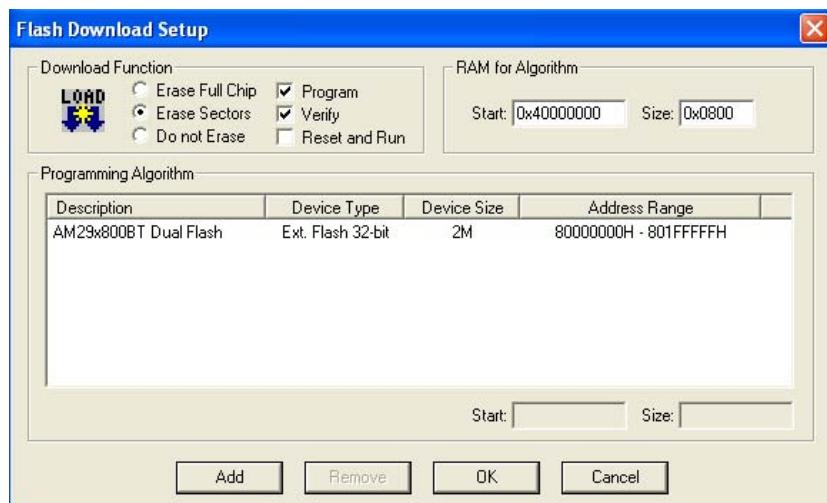
The Ulink JTAG can also be used to program external flash devices. This can be achieved as follows.

Open the options for target/utilities window and add the file flash.ini as shown below



This is a script file that is used by the Ulink JTAG to configure the necessary chipselects to allow code to be programmed into the external flash and as such should reflect the values used in the startup code.

Next select the settings button and you can add the flash algorithm to match the external device used.



With this configuration using the debugger with the JTAG and external flash becomes seamless.

Exercise 10 Phase Locked Loop

In this exercise we will configure the operation of the PLL to give maximum speed of operation for the ARM 7 core for a 12.00MHz oscillator. We will also configure the VLSI bus to run at half the speed of the ARM7 core.

Using $C_{clk} = M \times OSC$ calculate the maximum integer value for M given that $OSC = 12.00$ MHz and $C_{clk} \leq 60$ MHz

Using $F_{cco} = c_{clk} \times 2 \times P$ calculate a suitable value for P using the result for C_{clk} calculated above and were 156 MHz $< F_{cco} < 320$ MHz

Using the results for M and P calculate the value for PLLCFG

(answer M = 5 and P = 2)

Calculate the value for VPBDIV to set Pclk at half the frequency of C_{clk}

Open the project in C:\work\ EX7 PLL

Complete the code in main as follows

Set multiplier and divider of PLL in PLLCFG to give 60.00 MHz

```
PLLCFG = 0x00000024;
```

Enable the PLL in PLLCON

```
PLLCON = 0x00000001;
```

Update the internal PLL registers with the feed sequence

```
PLLFEED = 0x000000AA;
PLLFEED = 0x00000055;
```

Test the Lock bit in PLLSTAT until the PLL is stable

```
while (!(PLLSTAT & 0x00000400))
```

Connect the PLL as the main clock source

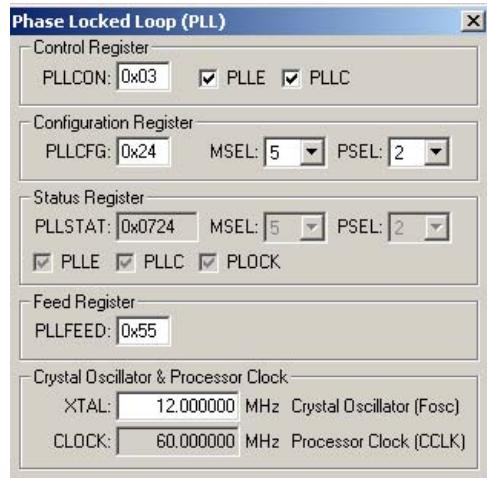
```
PLLCON = 0x00000003;
```

Set the VLSI peripheral bus to 30.00MHz

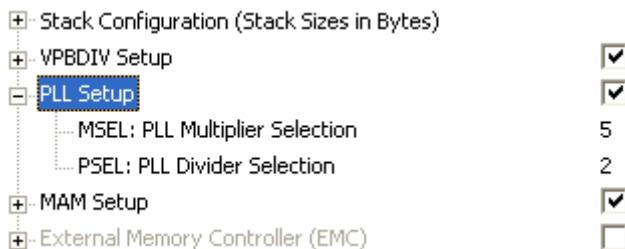
```
VPBDIV = 0x00000002;
```

Compile the code and load it into the board.

Run the initPLL routine and observe the results in the Peripherals\PLL window. This will show the true bus frequency.



Configuration of the PLL can be done via the Keil startup code but for this exercise this code is disabled. The equivalent settings for the startup code are shown below.



Exercise 11: Fast Interrupt

In this exercise we will configure an external interrupt to be handled as an FIQ. This code was used in exercise 5 to show the C handling of an interrupt function. This time we will see how to configure the hardware for an FIQ interrupt.

Open the project in EX12 Interrupt vectored\work

Configure the external interrupt as an FIQ by programming the IntSelect register

```
VICIntSelect = 0x00008000;
```

Compile the code and start the debugger and check the interrupt works as in exercise 5.

The name of the FIQ ISR must match the name in the vector constant table in startup.s. The default name given by Keil is FIQ_Handler.

If you are using the simulator the interrupt can be triggered by setting pin 0.14 low in the peripherals/GPIO window or by using the “Generate EINT1” button in the toolbox

Exercise 12: Vectored Interrupt

In this exercise we will configure a IRQ source to be handled as a vectored interrupt by the VIC. We will use the same external interrupt as exercise 9 but this time it will have its own dedicated ISR to give faster servicing of the interrupt request.

Open the project in C:\work\EX11 Interrupt vectored

Configure slot 0 in the VIC to service the EINT1 interrupt

```
VICVectCntl0 = 0x0000002F;
```

Place the address of the dedicated external interrupt routine in the correct vector address register.

```
VICVectAddr0 = (unsigned)EXTINTVectoredIRQ;
```

Complete the code in the EXTINTVectored to cancel the interrupt

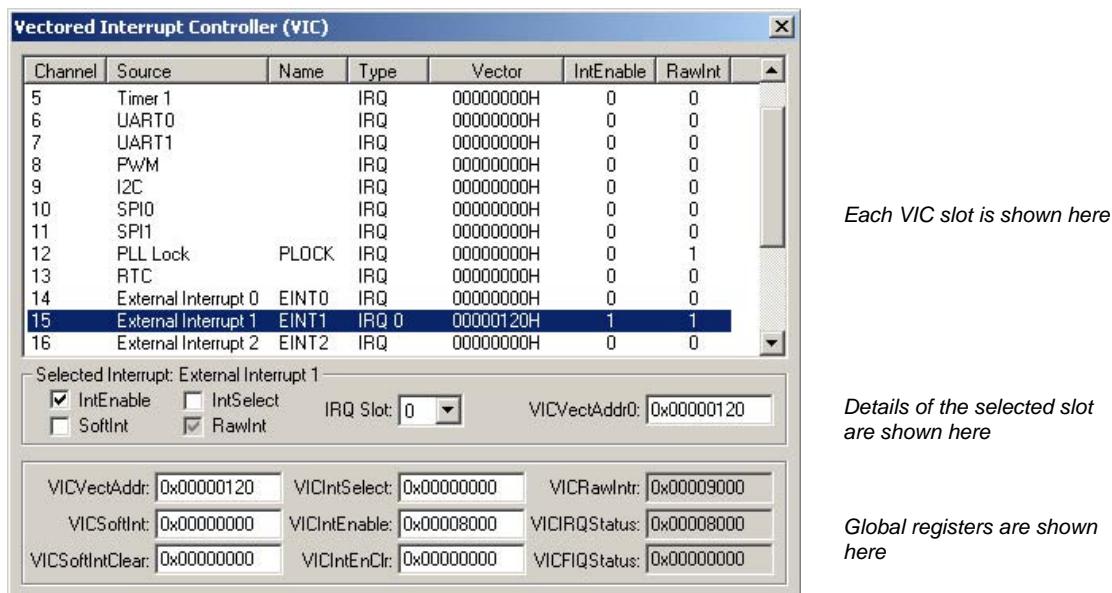
```
EXTINT      = 0x00000002;
VICVectAddr = 0x00000000;
```

Compile the code and start download it into the debugger.

Run the code and check that the interrupt is entered correctly and that it only runs once for each press of the EINT1 button pin 0.14.

If you are using the simulator open the GPIO Port 0 peripheral window. The interrupt can be triggered by bringing pin 0.14 low or use the “Generate EINT1” button in the toolbox.

Open the VIC peripheral window in the debugger and get familiar with its layout and how the VIC is configured.



In the simulator use the cycle counter in the register window calculate how long the device takes to enter to the first line of your code in the interrupt routine

Exercise 13 : Non Vectored Interrupt

In this exercise we will use the Vector interrupt unit to generate the slowest form of interrupt response a non vectored interrupt . We will use External interrupt one EINT1 to generate an interrupt. The VIC will respond to this interrupt event by providing the address of a general purpose ISR for the processor to jump to. On entry to this routine the ISR must calculate the source of the interrupt, take appropriate action and then correctly exit the ISR and resume normal processing.

Open the project in EX9-Interrupt Non Vectored\work

In Startup.s add the correct assembly code to the IRQ interrupt vector

```
LDR PC, [PC, #-0xFF0]
```

In main.c configure the Pin Connect Block to enable P0.14 as an External interrupt.

```
PINSEL0 = 0x20000000;
```

Place the address of the NonVectored ISR into the Default vector address register. Note in C this can be done as follows:

```
VICDefVectAddr = (unsigned)<Name of ISR routine>;
```

Enable the External interrupt channel in the VIC

```
VICIntEnable = 0x8000;
```

In the interrupt routine check the IRQ status register to find the source of the interrupt

```
if(VICIRQStatus&0x00008000)
```

At the end of the ISR clear the interrupt flag in EINT0 register and perform dummy write to the correct register in the VIC to clear the interrupt source

```
EXTINT = 0x00000002;
```

```
VICVectAddr = 0x00000000;
```

Compile the code and start the debugger

Run the code and check that the interrupt is entered correctly and that it only runs once for each press of the EINT1 button pin 0.14.

If you are using the simulator open the GPIO Port 0 peripheral window. The interrupt can be triggered by bringing pin 0.14 low. Alternatively open the toolbox and use the “Generate EINT1” button.

In the simulator use the cycle counter in the register window calculate how long the device takes to enter to the first line of your code in the interrupt routine

Exercise 14: Nested Interrupts

In this exercise we will setup two interrupt sources. First timer0 which sets a port pin high for the duration of the interrupt and secondly external interrupt one which also sets a second port pin high for the duration of the interrupt. Under normal operation both interrupts would block each other while they are running. However by adding the appropriate macros to the External interrupt routine we can allow the timer interrupt to interrupt the external interrupt routine so it will be guaranteed to run every 10ms. This can be observed on the LEDS on the MCB2100 or in the logic analyser of the simulator.

Open the project in EX14-Interrupt Non Vectored\work

In Main.c complete the two IENABLE and IDISABLE macros

```
#define IENABLE
    __asm { MRS      LR, SPSR      }          // Nested Interrupts Entry
    __asm { STMFD   SP!, {LR}      }          // Copy SPSR_irq to LR
    __asm { MSR     CPSR_c, #0x1F }          // Save SPSR_irq
    __asm { STMFD   SP!, {LR}      }          // Enable IRQ (Sys Mode)
                                            // Save LR

#define IDISABLE
    __asm { LDMFD   SP!, {LR}      }          // Nested Interrupts Exit
    __asm { MSR     CPSR_c, #0x92 }          // Restore LR
    __asm { LDMFD   SP!, {LR}      }          // Disable IRQ (IRQ Mode)
    __asm { MSR     SPSR_cxsf, LR }          // Restore SPSR_irq to LR
                                            // Copy LR to SPSR_irq
```

Add the two macros to the External interrupt service routine

```
void EXINT1_ISR (void) __irq
{
    EXTINT      = 2;                         // Clear EINT1 interrupt flag
    IENABLE;                                // allow nested interrupts
    IOSET1 = 0x00010000;                     // Switch on an LED
    delay (0x500000);                       // wait a long time
    IOCLR1 = 0x00010000;                     // Switch off the LED
    IDISABLE;                                // disable interrupt nesting
    VICVectAddr = 0;                         // Acknowledge Interrupt
}
```

Build the project and download it to the debugger

If you are using the MCB2100 run the code at full speed. One LED will flash for the period it is in the timer interrupt. Press the INT1 button and the second led will illuminate for the longer period it is in the External interrupt routine. However the timer led will continue to flash because the interrupt is not blocked.

If you are using the simulator the same behaviour can be observed using the logic analyser.

Remove the IDISABLE and IENABLE macros and observe the new behaviour.

Exercise 15: General purpose IO pins

In this exercise we will use the GPIO pins. A group of pins will be set as outputs and then each will be flashed sequentially.

Open the project in EX12 GPIO\work

Add the include file for the LPC21xx SFR's

```
#include <LPC21xx.H>
```

Program the data direction register to enable pins 1.16 – 1.23 as output

```
IODIR1 = 0x00FF0000;
```

Complete the ChangeGPIOPinState function to clear and set the relevant pins

```
IOCLR1 = ~state;  
IOSET1 = state;
```

Compile the code and download it into the debugger

Run the code on the MCB2100 to see the LED chaser.

If you are using the simulator step through the code and check it works by examining the state of the IO pins via the GPIO peripheral window

Exercise 16: Timer Capture

In this exercise we will use timer 0 in a simple compare mode to measure the time between starting the timer and getting a rising edge on pin 0.02. Because the MCB2100 does not have a button on a capture channel to generate the pulse this exercise is based on the Simulator. A target version is included but you would need to modify the MCB2100 hardware and observe the output on an oscilloscope.

Pclk = 30 MHz Pclock tick = 0.033MHz

Open the project in EX18 Timer capture\work

In main.c complete the code as follows

Enable pin 0.2 as capture channel 0:

PINSEL0 = 0x00000020;

Load prescaler with 1 micro second tick value:

TOPR = 0x0000001E;

Reset timer counter and prescaler counter registers:

TOTCR = 0x00000002;

Configure capture channel 0 to capture on the rising edge: T0CCR = 0x00000005;

Enable the timer:

TOTCR = 0x00000001;

In the interrupt routine copy the capture value into a dummy variable:

Value = T0CR0;

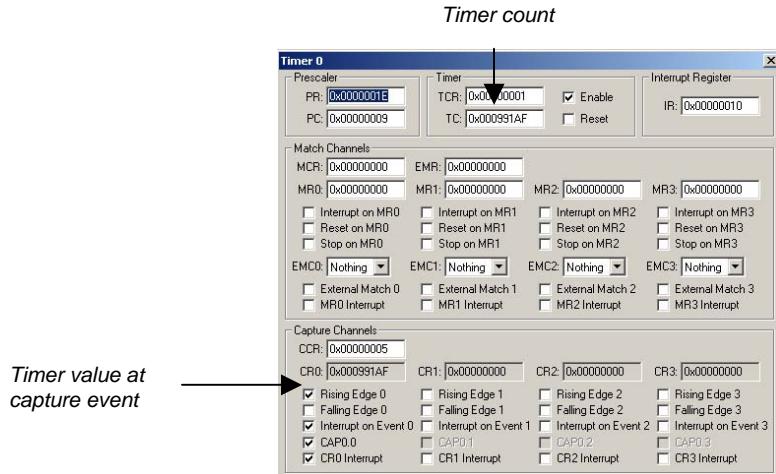
Compile and load the code into the debugger

Run the program and check the following

Test the capture interrupt is working by setting a breakpoint on the timer ISR, and running the code.

In the simulator a script has been added to the toolbox to generate a pulse on 0.02. If you are using the MCB2100 you need to pull the port pin up to Vcc via a 10K resistor.

In the simulator check that the timer is incrementing at the required rate by single stepping until the timer increments and measuring the time interval with the clock in the register window.



Exercise 17: Timer Match

In this exercise we will configure the timer to generate a single edge PWM signal using two match channels. Match zero will be used to generate the total PWM period. On match it will be used to reset the timer and generate an interrupt. Match one will be used to create the duty cycle. On match it will clear the external match one pin. At the beginning of the cycle the interrupt will set the external match one pin high.

Open the project in EX18 Timer Match\work

In Main.c complete the timer initialising code as follows

Configure the pin connect block with P0.5 as MAT1

PINSEL0 |= 0x00000800;

Set match zero to reset the counter and generate an interrupt

T0MCR = 0x00000003;

Set the PWM period to 16 msec

T0MR0 = 0x00000010;

Set the Match 1 to give a 50% duty cycle

T0MR1 = 0x00000008;

Configure match 1 pin to clear on Match. It is also set high for the first cycle

T0EMR = 0x00000042;

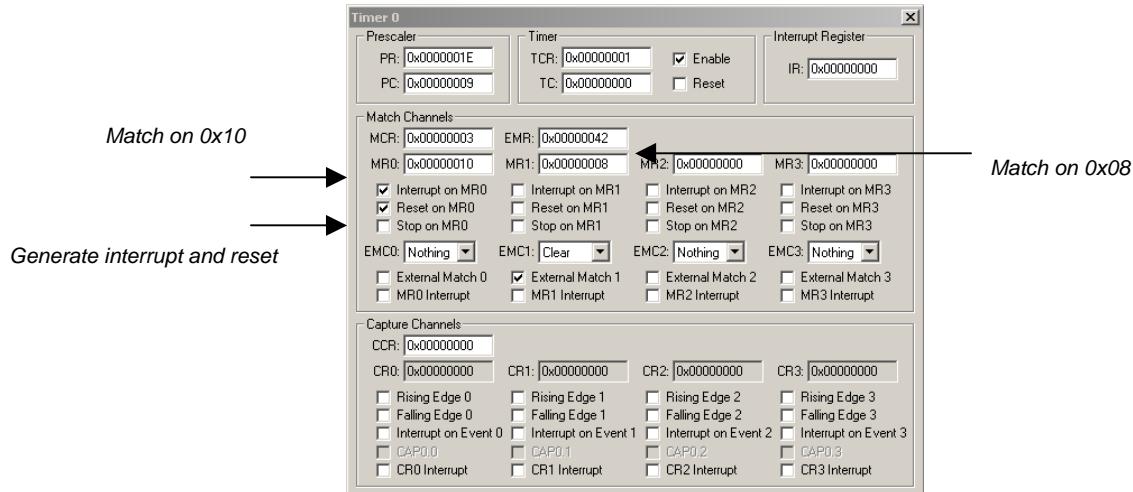
In the interrupt set Match 1 output pin High

T0EMR |= 0x00000002;

Compile and download the code into the debugger

Run the code and observe the activity on GPIO pin 0.5 with an oscilloscope.

Again we can see the configuration of the timer in the peripherals window

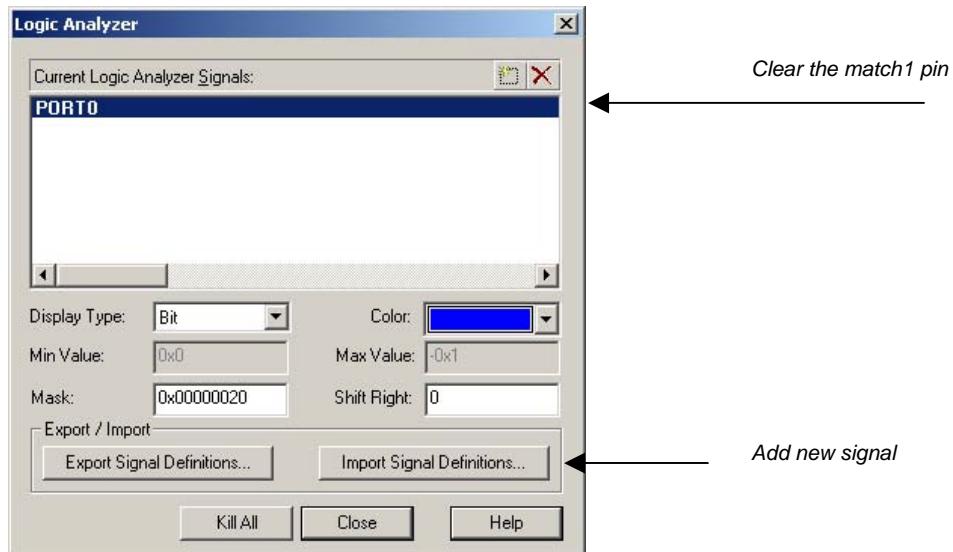


The simulator also includes a logic analyser window that can display a graphical trace of activity on a pin over time. The logic analyser is used as follows

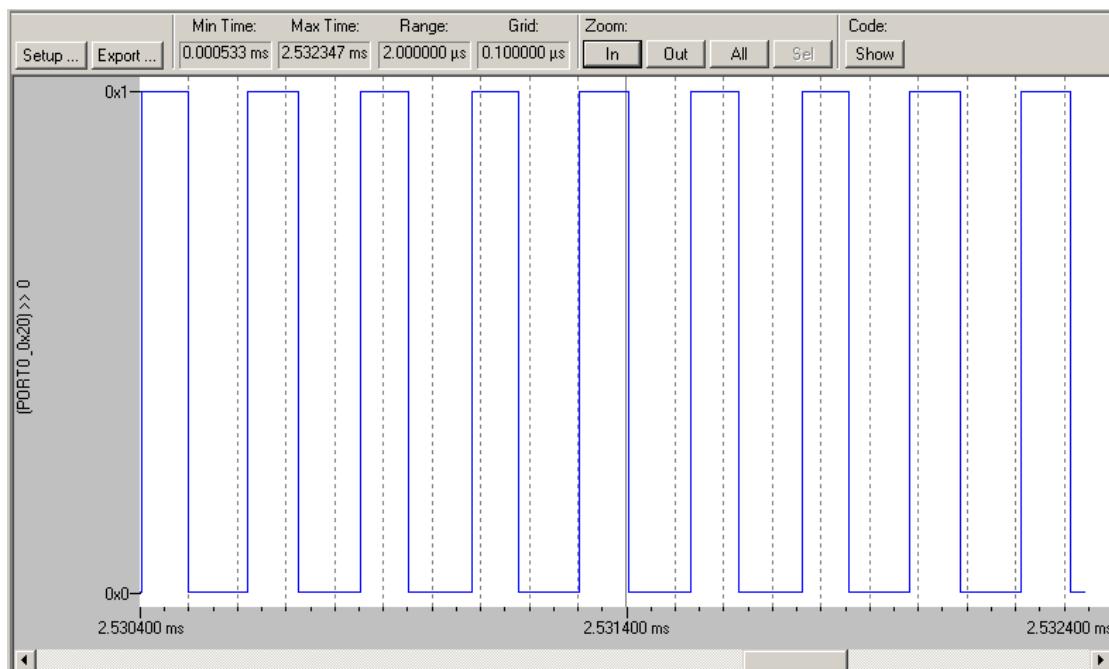
Open the logic analyser window with view\logic analyser

In the top left hand corner press the setup button

Add an new signal called PORT0, set the mask to 0x00000020 and set the display type to bit.



Now as you run the code any activity on the match1 pin will be recorder into the logic analyser window.



Exercise 18: Dual-Edge (Symmetrical) PWM Generation.

In this exercise we will use the PWM module to generate a single channel of symmetrically modulated PWM signal. This will use the MAT0 channel to generate the total signal period and will be configured to reset the timer. MAT1 will be used to generate the turn on off and MAT2 will generate the turn off point. The PWM duty cycle can then be modulated by the +- keys in serial window 2.

Open the project in C:\work\ EX20 PWMModule

In main.c complete the code as follows

Enable pin 0.7 as PWM2

```
PINSEL0 |= 0x00008000;
```

Configure PWM channel two to double edge control, output enabled

```
PWMPCR = 0x0000404;
```

Configure Timer 0 reset to the counter

```
PWMMCR = 0x00000003;
```

Set the reload values of mat1 and mat2 to give an initial “spike” which is gradually broadened in the main loop as the code runs.

```
PWMMR0 = 0x000000FF;
```

```
PWMMR1 = 0x00000080;
```

```
PWMMR2 = 0x00000080;
```

enable shadow latch for match 0 - 2

```
PWMLER = 0x00000007;
```

Reset counter and prescaler and enable PWM mode

```
PWMTCR = 0x00000002;
```

enable timer with PWM enabled

```
PWMTCR = 0x00000009;
```

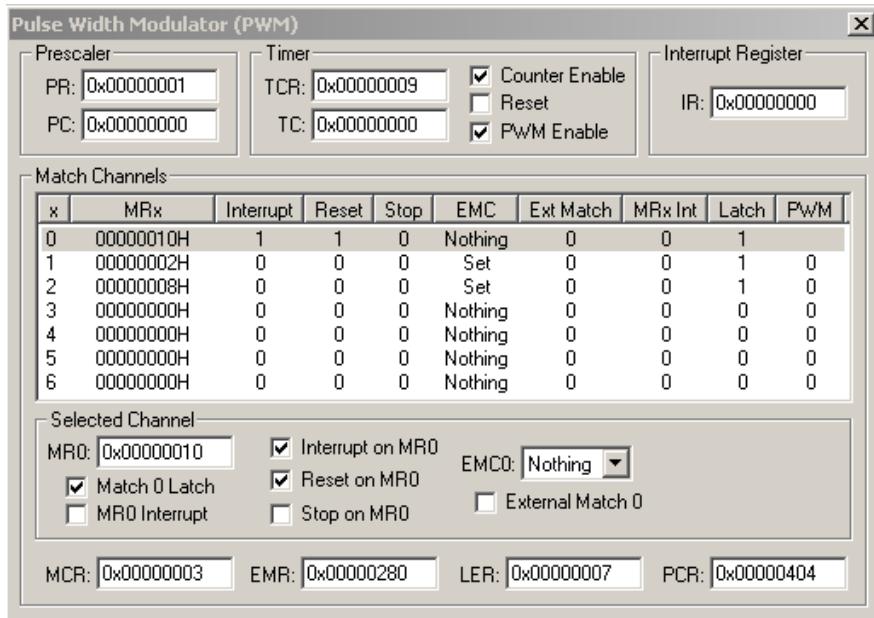
In the main while loop enable the shadow latch to update the match registers when a new values are available.

```
PWMLER = 0x00000006;
```

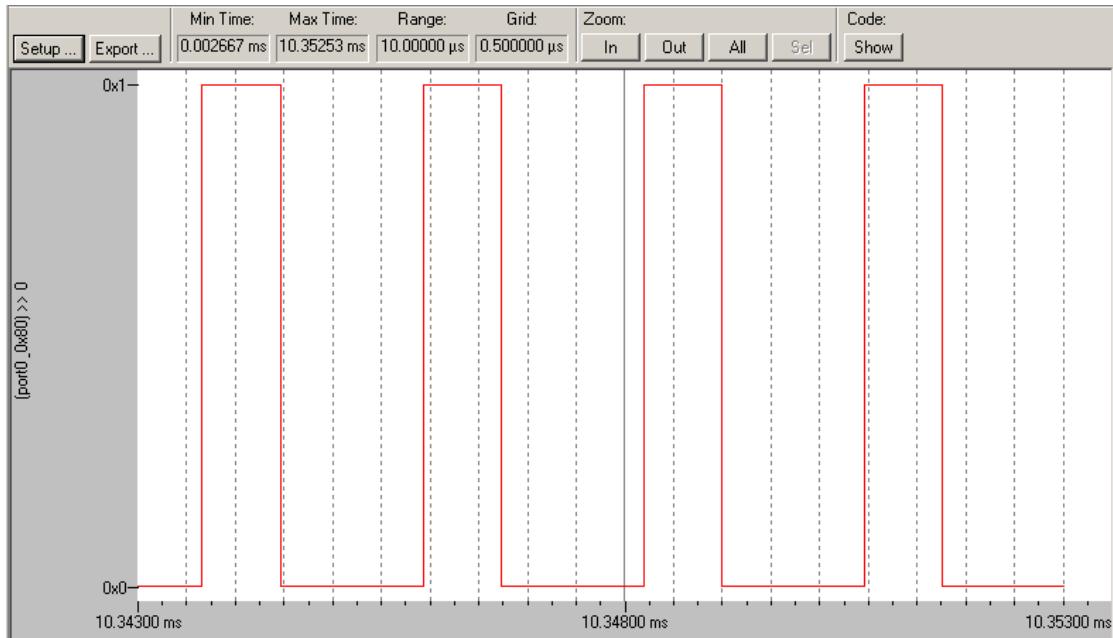
Compile and download the code into the debugger

Run the code and check that it is correct with the simulator

Again the peripheral window can be used to view the configuration of the PWM module



The logic analyser can also be used to examine the activity on the PWM 1 pin. This time the mask should be set to 0x00000080.



Exercise 19: Real Time Clock

In this exercise we will use all the major features of the real time clock. The first step is to configure the reference clock to give 32.768 KHz. We can then configure the counter increment registers to give an interrupt every second and the alarm registers to give an interrupt at 10 seconds.

External oscillator = 12.00MHz, Cclk = 60.000MHz, Pclk = 30.000MHz

Calculate the value of PREINT using $\text{PREINT} = (\text{int})(\text{pclk}/32768)-1$

(Answer 0x392)

Calculate the value of PREFAC using $\text{PREFAC} = \text{Pclk} - ((\text{PREINT}+1) \times 32768)$

(Answer 0x4380)

Open the project in ex16 RTC\work

In main.c enter the values for PREINT and PREFAC

Next enable the Seconds increment interrupt

CIIR = 0x00000001;

Set the Seconds alarm register for 3 seconds and enable the seconds alarm in the alarm mask register

ALSEC = 0x00000003;

Program the Clock control register to start the RTC running

CCR = 0x00000001;

In the interrupt routine test the interrupt location register to determine the cause of the RTC interrupt (seconds increment or alarm)

if(ILR&0x00000001) increment
if(ILR & 0x00000002) Alarm

In either case make sure the interrupt flag is cleared before exiting the interrupt

ILR = 0x00000001; increment
ILR = 0x00000002; Alarm

Using either the simulator or the MCB2100 hardware prove the code works correctly. Remember the simulator is not realtime and you have to use the Internal.Sec counter in the register window to get timing information.

Exercise 20: UART

In this exercise we will configure UART1 to transmit at 19200 baud 8 bits no parity on stop bit. We will use the UART to echo characters sent to it from the simulator terminal window.

Open the project in EX13 UART\work

Configure the UART for 9600 Baud 8/N/1 with pclk = 30MHz (The actual achievable baud rate is 9664)

```
UART1_LCR = 0x00000083  
UART1_DLL = 0x000000C2;  
UART1_LCR = 0x00000003;
```

In getchar and putchar add the code to monitor the line status register

- a For Putchar while (!(UART1_LSR & 0x20)); (This line is used twice)
- b For Getchar while (!(UART1_LSR & 0x01));

Compile and download the code into the debugger

If you are using the MCB2100 hardware connect UART1 to the serial port of a PC and start Hyperterminal. When you run the code enter characters in Hyperterminal and see them echoed back to the screen.

If you are using the simulator use the built in serial window to simulate a terminal (serial window #2).

Exercise 21: I2C interface

This exercise will demonstrate using the I2C interface to write and read back from a serial EEPROM of the type (see the CD for the full datasheet). Because the MCB2100 evaluation board is not fitted with such a memory a script file is used to simulate the EEPROM so we can test our routines and have them ready as soon as the hardware becomes ready.

Open the project in C:\work\I2C

The script file is in I2C.ini and this is added to the project in “Options for Target/Debug”

Add the following lines to the C code

Build the code and start the simulator

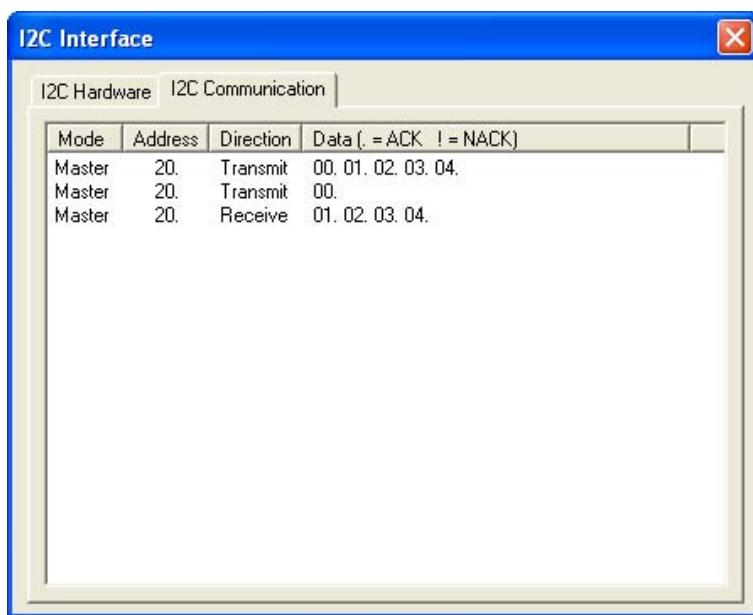
Open the peripherals/I2C window.

The I2C Hardware tab shows the configuration of the peripheral and the I2C communication tab will show all the bus transactions.

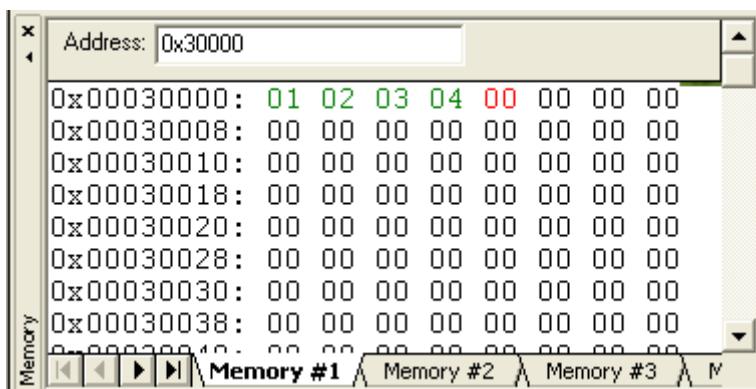
In this code the I2Ctransfer() function starts the I2C bus transaction which is then handled by the I2C interrupt function. The interrupt function is a state machine in the form of a case statement. By setting a breakpoint on the initial switch statement we can run the code and observe the activity on the I2C bus

```
092
093 void I2CISR (void) __irq
094 {
095
096 switch (I2STAT)
097 {
```

After running the code the transactions on the I2C bus will look as follows.



The Simulated EEPROM memory is mapped into an unused section of the processor address range. In this case 0x00030000 to 0x000300FF. As the code is run you can view data as it is written and read from the memory, this can be a great aid to debugging your drivers.



Exercise 22: SPI

Like the I2C example this example uses the SPI interface to communicate to a serial EEPROM and demonstrates how to write and read data to such a device. Since the MCB2100 is not fitted with such a memory this example uses the simulator with a script file to simulate the EEPROM.

Open the project in C:\work\EX16-SPI

The script file is in SPI.ini and this is added to the project in “Options for Target/Debug”

The script maps some memory for the SPI EEPROM at 0x700000-0x7000FF

In Main.c add the following lines of Code

Configure the operation mode and bit timing.

```
S0SPCCR      =      0x000000FF;  
S0SPCR       =      0x000000A0;
```

Kick off an SPI transfer

```
SPI_write(output_buffer,8);
```

Use a case statement for each state of the transaction

```
case (0x01): //Send Write opcode  
    S0SPDR = 0x00000005;  
    status = 0x02;           //set next state  
break;
```

Compile the code and start the simulator

Set a breakpoint on the switch statement in the SPI ISR

Run the code and examine the state of the virtual EEPROM memory and the debug output in the message window.

Exercise 23: Analog To Digital Converter

In this exercise we will configure the A/D in hardware mode and do an 8 bit conversion on channel 0. The results are modulated onto the LED pins.

Open the A to D project in C:\work\EX20-AtoD

Program the ADCR to give hardware conversion on channel zero, 8 bits resolution, Adjust the A/D clock for a Pclk of 30 MHz

ADCR = 0x00270601;

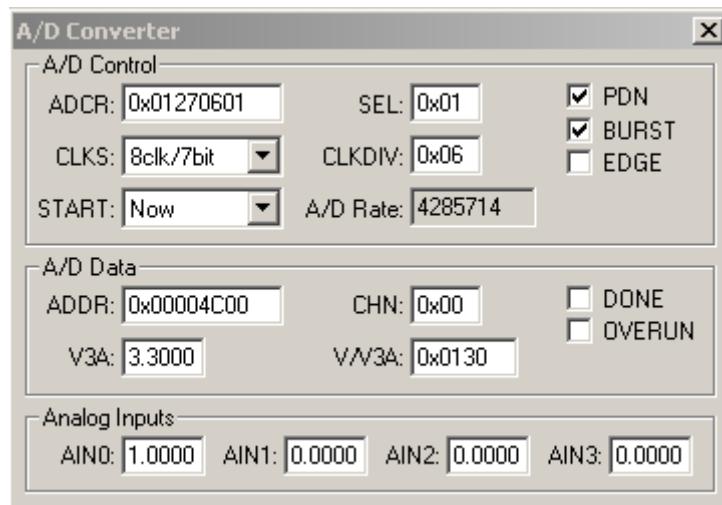
Test the done bit in ADDR to find a finished conversion

```
while ((val & 0x80000000) == 0);
```

Compile the code and load it into the debugger

If you are using the target board turning the potentiometer will change the analogue value analogue channel 0.

If you are using the simulator the peripheral window will allow you to simulate external voltages.



A new simulation script has been added that allows you to change the A-D voltage via buttons in the toolbox.

Exercise 24: Digital to Analogue Converter

In this exercise we will setup the A-D to make a conversion every 50uSec (20 KHz) by using a timer match channel to trigger the conversion. As soon as the conversion is made the converted value is fed into the Digital to Analogue converter. If you have access to a signal generator and a scope you can compare the input and output waveforms or use the simulator and the Logic analyzer.

Open the AtoD project in C:\work\EX20-AtoD

Configure the A/D so it is triggered by timer0 match channel

Copy the A/D result into the DAC register

Build the project and start the debugger

If you are using the MCB2100 connect a signal generator to the A/D input channel and an oscilloscope to the output.

Run the code and compare the two waveforms

Vary the input frequency and find the frequency at which aliasing occurs

If you are using the simulator an input script is provided in the toolbox and the input /output signals can be compared in the logic analyser window.

Exercise 25: Transmitting CAN Data

In the next two exercises we are going to look at transmitting and receiving CAN messages. In order for a CAN network to work we need a minimum of two nodes. Fortunately the LPC2194 has two independent CAN controllers which are brought out to two D-Type connectors on the evaluation board. For our examples we will connect the two channels together and send data from one channel to the other.

Setup the evaluation board with the JTAG debugger and connect Pin 2 (CAN low) of socket P3 to pin 2 of socket P4.

Do the same with pin7 (CAN High)

Open the project in c:\examples\work\EX22- TXCAN

Calculate the bit timing values for 125 Mbit/sec withPclk = 60 MHz

Complete the lines of code in the example

Program the bit timing register

```
C2BTR = 0x001C001D
```

To transmit the data set the Data length code to four bytes

```
C2TFI1 = 0x00040000;
```

Set the address to message 22 and a standard 11 bit identifier

```
C2TID1 = 0x00000022;
```

Copy the A/D result into the first byte of the TX buffer

```
C2TDA1 = val;
```

Schedule the message for transmission

```
C2CMR = 0x00000001;
```

Build the code and start the debugger

Once the code is running the A/D conversion will be transmitted from the CAN2 module and received by CAN1. The received data is then written to the GPIO to modulate the LED's

In the simulator the two CAN channels are looped back by a script file and the CAN traffic may be observed in the Peripherals/CAN/communication window.

Exercise 26: Receiving CAN Data

This is the same code as in the previous example but this time it is presented from the receive point of view.

Open the project in C:\work\EX24-CANRX

In main.c add the following lines to configure the acceptance filter to receive messages

Declare an array in the filter memory

```
unsigned int StandardFilter[2]           _at_ 0xE0038000;
```

Configure the standard message filter table

```
AFMR          =      0x00000001;
StandardFilter[0] =      0x20012002;
StandardFilter[1] =      0x20032004;
SFF_sa        =      0x00000000;
SFF_GRP_sa   =      0x00000008;
```

Release the receive buffer before leaving the ISR

```
C1CMR         =      0x00000004;
```

Build the code and start the debugger.

Run the code as far as initialising the CAN peripherals.

Examine the Acceptance filter memory in peripherals/CAN/Acceptance filter

This view shows you which messages may be received by the CAN peripherals.

Continue to run the code and receive a message.

Check that the message ID matches the Index assigned in the Acceptance filter.

Chapter 6: Tutorial With GNU Tools

Introduction

The following tutorial demonstrates how to setup a project in uVISION for the GNU compiler. Exercises 1 – 6 are repeated to show the non-ANSI aspects of the GNU compiler. Once you are familiar with these exercises, you can rejoin the main tutorial but use the exercise examples in the GCC directory.

GCC Startup Code

The startup code used in the GNU project is different in that the Keil Assembler has different directives and naming conventions. However, it is performing the same operations. It is up to the programmer to edit the vector table as discussed in the section on the Keil compiler startup code. The graphical editor allows you to configure the processor stacks and system peripherals in the same way as the Keil compiler startup code.

Interworking ARM/THUMB Code

The GCC compiler also supports the ARM procedure calling standard and allows interworking between the ARM and THUMB instruction sets. However, unlike the Keil compiler, it is not possible to select individual functions as ARM or THUMB. In the GCC compiler all ARM code must be in one module or modules and the THUMB code must be in separate modules. These modules are compiled as ARM or THUMB as required and then linked together. This process is described in example 3 in this section.

Accessing Peripherals

The Keil and GNU compilers can use the same include files to access the on-chip SFR registers.

Interrupt Service Routines

The GCC compiler has a set of non-ANSI extensions which allow functions to be declared as interrupt routines. The general form of the declaration is shown below

```
void IRQ_Routine (void) __attribute__ ((interrupt ("IRQ")));
```

The following keywords are available to define the exception source required:

FIQ, IRQ, SWI, UNDEF.

This function declaration is only required on the function prototype and should not be used on the main body of the function. An interrupt service routine is shown in example 5.

Software Interrupt

There is no real software interrupt support in the GCC compiler. To generate a software interrupt you must use inline Assembler as shown below:

```
#define SoftwareInterrupt2 asm (" swi #02" )
```

This will place a SWI instruction encoded with the value 2 in your code. Next it is possible to declare a pointer to a CPU register using the non-ANSI register keyword as shown below:

```
register unsigned * link_ptr asm ("r14");
```

This allows us to read the contents of the link register when we enter the ISR. When the SWI instruction is executed, the CPU will enter supervisor mode and jump to the SWI vector. The address of the SWI instruction plus four will be stored in the link register. On entry to the software interrupt ISR the following line of code is executed:

```
temp = *(link_ptr-1) & 0x00FFFFFF;
```

The address stored in the link register is rolled back by one instruction (word-wide pointer i.e. four bytes) so that it is pointing at the address of the SWI instruction which generated the exception. The top eight bits of the SWI instruction are masked off and bits 0-23 are copied into the temp variable. This in effect loads the number 2 into the temp variable. A switch statement can now be used to run the desired code. This method of handling software interrupts is shown in example 6.

Inline Functions

Within the GNU compiler functions may be declared as inline functions as follows:

```
inline int fast_function(char param1)
```

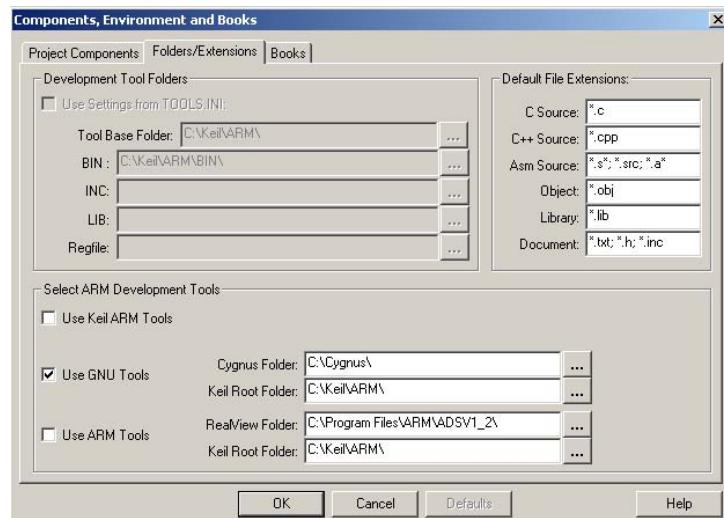
Exercise 1: Using The Keil Toolset With The GNU Compiler

This example is based on the source code which can be found in:

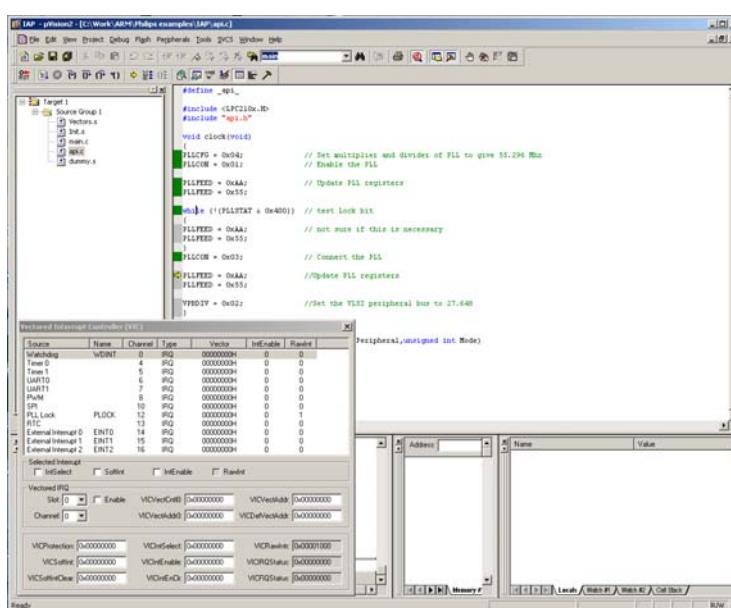
C:\Exercise\Work\EX1 first program

In this first exercise we will spend some time defining a first project, building the code and downloading it into the Simulator for debugging. We will then cover the basic debugging functions of the Keil simulator.

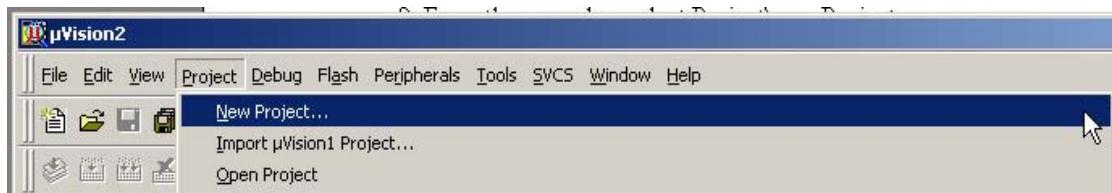
The Keil uVISION IDE is designed to support several compilers: the GNU C compiler, the ARM development suite and the Keil ARM compiler. Before compiling, make sure you have the GNU compiler selected. This is done by activating the project workspace, right-clicking and selecting ‘manage components’. In this dialog, select the Folders/extensions tab and make sure the GNU tools box is selected.



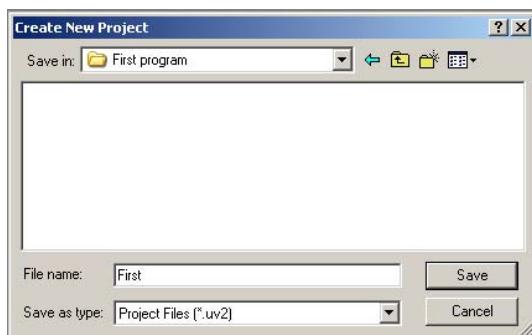
Double-click on the Keil Uvision3 icon to start the IDE.



From the menu bar select Project\New Project.

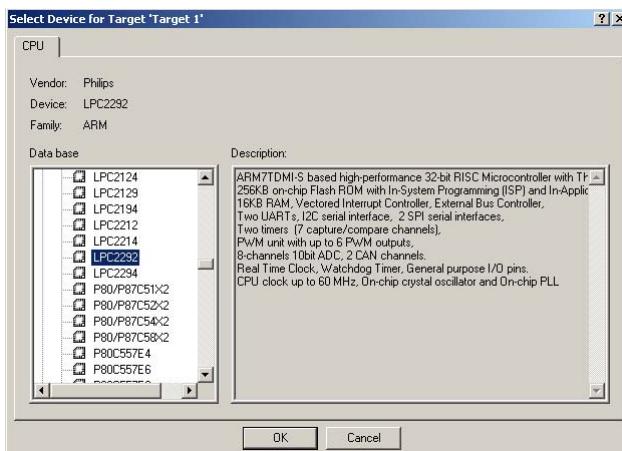


In the New Project dialog navigate to your desired project directory.

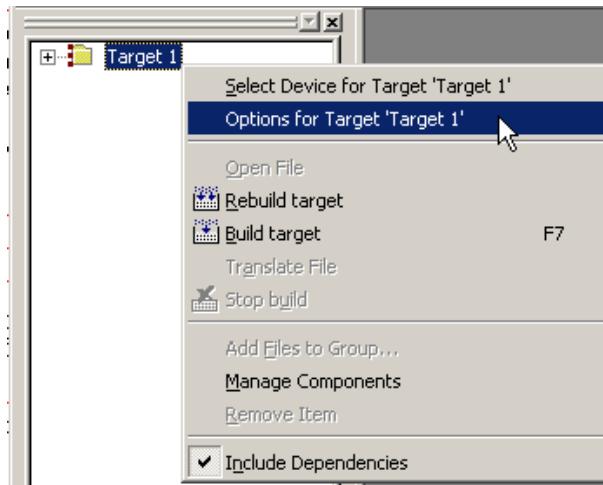


In the New Project dialog name the project first.uv2 and select Save.

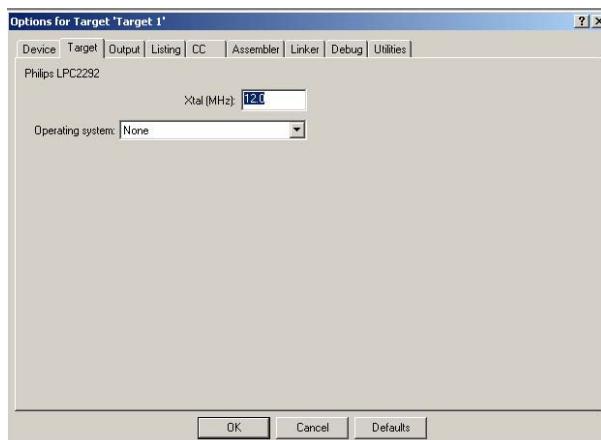
A ‘select new device for target’ dialog will appear. Navigate through the device data base and select the Philips\LPC2129 folder and select OK.



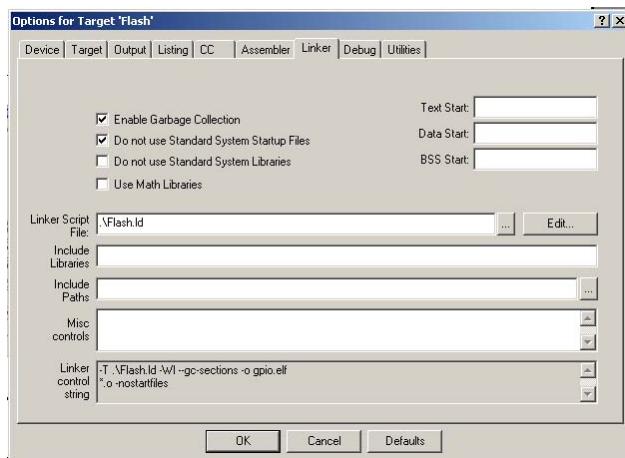
In the project browser highlight the ‘Target1’ root folder and select the local menu by pressing the right mouse button. In this menu select ‘Options for Target’.



In the ‘Target’ tab set the simulation frequency to 12.000 MHz.

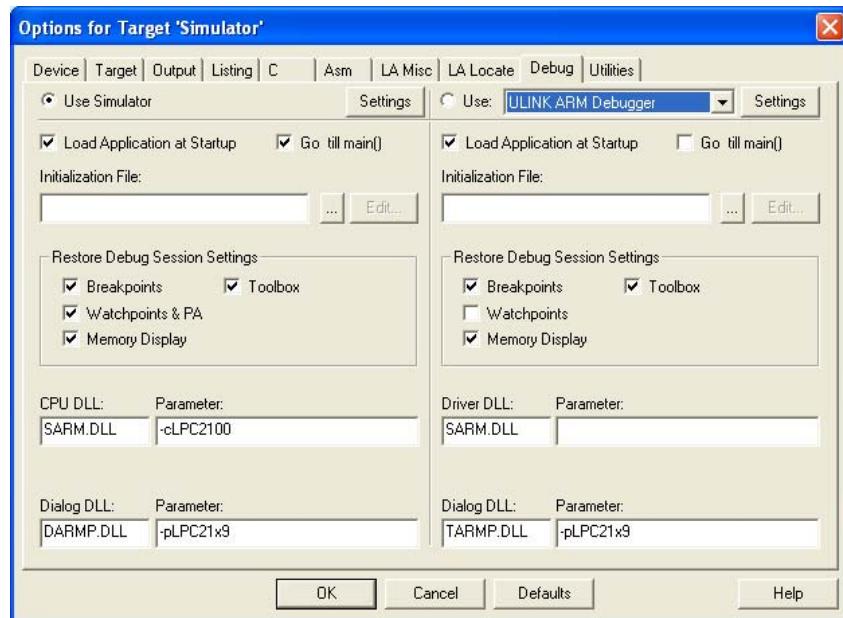


In the Linker tab select the linker file flash.ld and tick the “Garbage collection” and do not use “standard startup files” boxes



Note: To build the project so it will run within the on-chip RAM of the LPC2100 device, configure the Text start as select the linker file RAM.ld

In the debug tab make sure the “Use Simulator” radio button is active. Also make sure “Load Application at Startup” and “Go till main()” are checked.

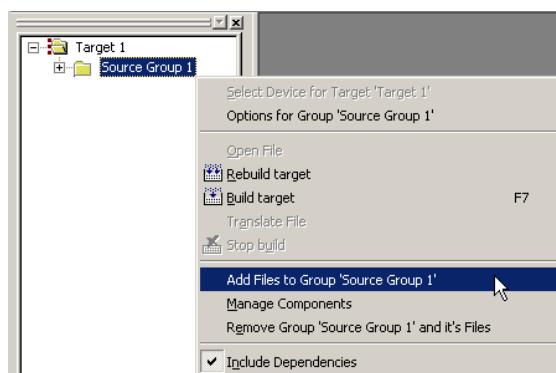


Select OK to complete the target options.

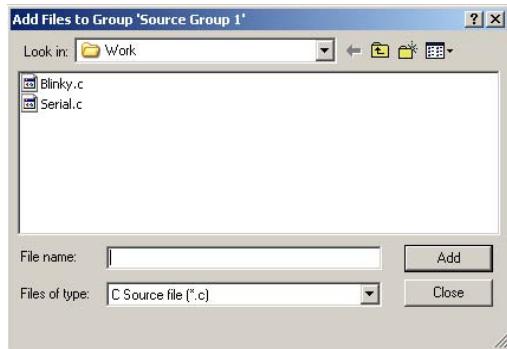
In the project browser expand the ‘Target1’ root node to show the Source group 1 folder.



Highlight the ‘Source Group 1’ folder, open the local menu with a right click and select ‘Add Files to group Source Group1’.



In the ‘Add files to Group’ dialog add the file blinky.c and serial.c.



Change the ‘Type of file’ filter to ASM and add the file startup.s

These are all the source files necessary for the project so select close.

You can view the source code contained in a file by double-clicking on the file name in the project browser window.

Once you have added all the source files the project can be built via the program menu or by the build button on the toolbar.



Once the code is built, you can start the simulator by pressing the debugger button. The use of the simulator and JTAG debugger are detailed in Exercise One in the Tutorial and are the same for the GNU compiler.

Exercise 2: Startup Code

In this exercise we will configure the compiler startup code to configure the stack for each operating mode of the ARM7. We will also ensure that the interrupts are switched on and that our program is correctly located on the interrupt vector.

Open the project in EX2 Startup\work

Open the file Startup.s and using the graphical editor configure the operating mode stacks as follows:

Stack Configuration	
Top of Stack Address	0x4000 4000
Stack Sizes (in Bytes)	
Undefined Mode	0x0000 0080
Supervisor Mode	0x0000 0080
Abort Mode	0x0000 0080
Fast Interrupt Mode	0x0000 0080
Interrupt Mode	0x0000 0080
User/System Mode	0x0000 0400
PLL Setup	<input checked="" type="checkbox"/>

Compile the code

Start the simulator and when the PC reaches main, examine the contents of each R13 register.

User/Syst...	R8	0x00000000
	R9	0x00000000
	R10	0x40003980
	R11	0x00000000
	R12	0x00000000
	R13 (SP)	0x40003d80
	R14 (LR)	0x00000000
Fast Interrupt	R8	0x00000000
	R9	0x00000000
	R10	0x00000000
	R11	0x00000000
	R12	0x00000000
	R13 (SP)	0x40003f00
	R14 (LR)	0x00000000
	SPSR	0x00000000
Interrupt	R13 (SP)	0x40003e80
	R14 (LR)	0x00000000
	SPSR	0x00000000
Supervisor	R13 (SP)	0x40003e00
	R14 (LR)	0x00000000
	SPSR	0x00000000
Abort	R13 (SP)	0x40003f80
	R14 (LR)	0x00000000
	SPSR	0x00000000
Undefined	R13 (SP)	0x40004000
	R14 (LR)	0x00000000
	SPSR	0x00000000

Each stack is allocated a space of 0x80. The user stack is 0x400 bytes so user data will start at 0x40003d80 – 0x400

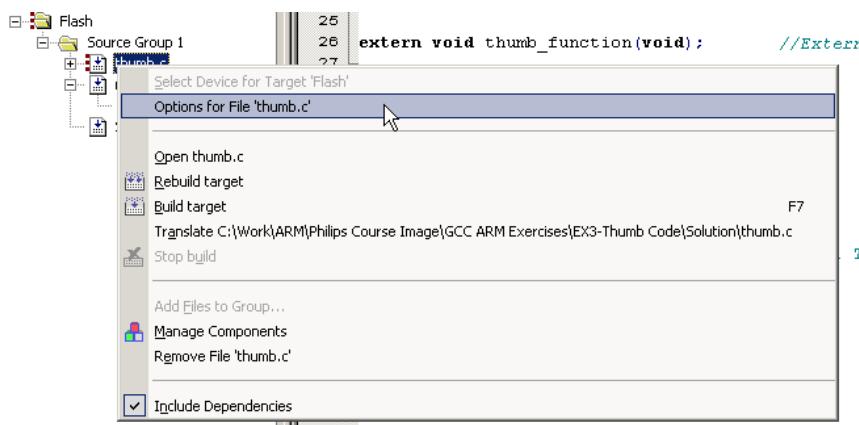
Start of stack space at the top of on-chip memory

Exercise 3: Using THUMB Code

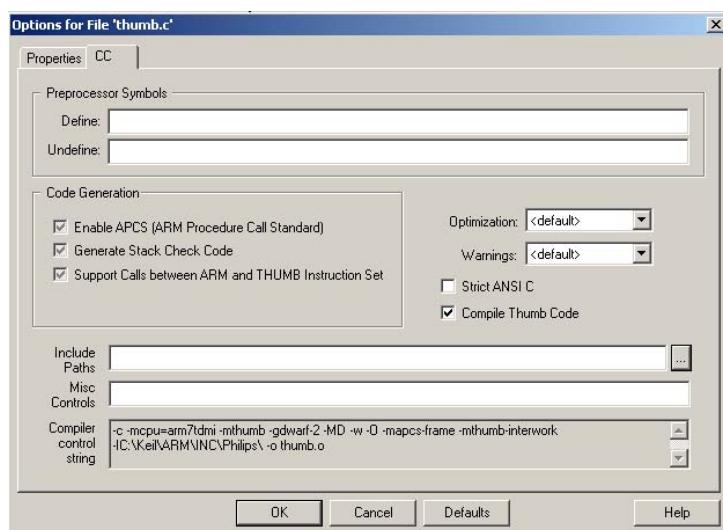
In this example we will build a very simple program to run in the ARM 32-bit instruction set and call a 16-bit THUMB function and then return to the 32-bit ARM mode.

Open the project in EX3 THUMB code\work

In the files browser select thumb.c open the local menu (right-click) and select “options for thumb.c”

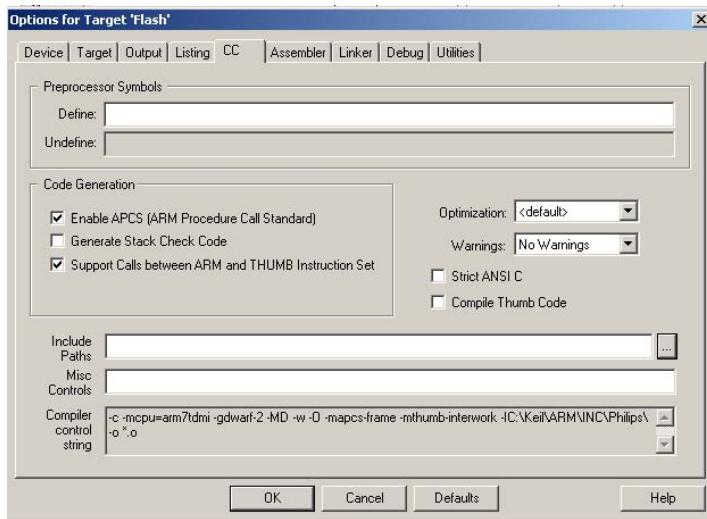


Select the CC tab and in the misc controls add –mthumb or tick the “compile thumb code” box and click OK



Again in the file browser select the root target (FLASH) and in the local menu “options for target”

In the CC tab tick the “enable APCS option and the “support calls between THUMB and ARM”



Compile and download the code into the debugger

Open the disassembly window and single step through the code using the F11 key

Observe the switch from 32-bit to 16-bit code and the THUMB flag in the CPSR

CPSR		0x60000010
N	0	
Z	1	36: thumb_function();
C	1	37: }
V	0	0x00000198 EB000000 BL 0x000001A0
I	0	0x0000019C EAFFFFFC B 0x00000194
F	0	0x000001A0 E59FC000 LDR R12,[PC]
T	0	0x000001A4 E12FFF1C BX R12
M	0x10	

The processor is running in ARM (32-bit) mode, the T-bit is clear and the instructions are 4 bytes long. A call to the THUMB function is made which executes a BX instruction forcing the processor into THUMB mode (16-bit).

The THUMB bit is set and on entry to the THUMB function a PUSH instruction is used to preserve registers on to the stack.

Exercise 4: Using The GNU Libraries

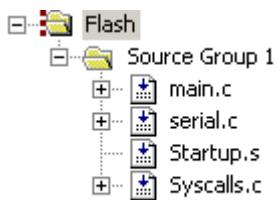
In this exercise we will look at tailoring the GNU Printf function to work with the LPC2100 UART. We will look at the registers of the UARTs in more detail later.

Open the project in EX4 printf\work

In main.c add a message for transmission to the printf statement

```
while(1)
{
    printf("Your Message Here \n"); //Call the printf function
}
```

Add the file syscalls.c in the work directory to the project.



In syscalls.c add modify the write function as follows:

Complete the for loop statement so it runs for the length of the printf string (len)
Inside the for loop add the putchar statement to write a single character to the stdio channel (putchar (*ptr))

Increment the pointer to the character string ptr++

```
int write (int file, char * ptr, int len)
{
    int i;

    for (i = 0; i < len; i++) putchar (*ptr++);
    return len;
}
```

Compile the code and download it to the development board

Run the code and observe the output within hyper terminal

If you are using the simulator, select view/serial window #1. This opens a terminal window within the simulator which displays the UART0 output.

Exercise 5: Simple Interrupt

1. In this exercise we will setup a basic FIQ interrupt and see it serviced.
2. Open the project in EX5-Interrupt\work
3. In main.c complete the definition of the EXTintFIQ function prototype to define it as the FIQ interrupt service routine

```
void EXTintFIQ (void) __attribute__ ((interrupt("FIQ")));
```

In startup.s complete the vector constants table to define EXTintFIQ as the FIQ ISR.

```
.global EXTintFIQ           Declare the name of the C ISR function as a
global
    .global _startup
    .func   _startup
_startup:

Vectors: LDR      PC, Reset_Addr
          LDR      PC, Undef_Addr
          LDR      PC, SWI_Addr
          LDR      PC, PAbt_Addr           Vector Table
          LDR      PC, DAbt_Addr
          .long 0xB8A06F58
          LDR      PC, [PC, #-0xFF0]
          LDR      PC, FIQ_Addr

Reset_Addr: .word   Reset_Handler
Undef_Addr:  .word   Undef_Handler
SWI_Addr:   .word   SWI_Handler   Constants table
PAbt_Addr:  .word   PAbt_Handler
DAbt_Addr:  .word   DAbt_Handler
            .word   0
IRQ_Addr:   .word   IRQ_Handler
FIQ_Addr:   .word   EXTintFIQ
```

Insert the name of the C ISR function in the constants table

5. Compile the code and download it onto the board.

6. Step through the code and observe the following using the disassembly window and the registers window.

Step through the code until you reach the while loop

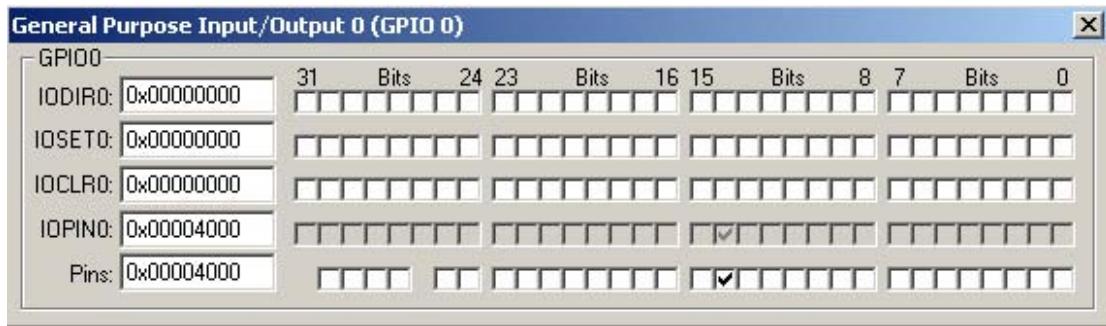
c. Set a breakpoint in the EXTintFIQ function

Press F5 to set the program running

On the MCB2100 board press the INT button to generate the interrupt.

If you want to see the entry and exit mechanisms to the exception, it is best to use the simulator and single step in the disassembly window. This way you can watch the program flow and the actions on the CPU registers.

To control the interrupt in the simulator, open the peripherals/GPIO port 0 window. Pin 0.14 is set high by the map.ini startup script. If you set the program running unchecking, the Pin1.4 box will generate the interrupt. You must raise the pin high again to stop interrupts.



Alternatively in the toolbox there is a “Generate EINT1” button. This button will generate a simulated pulse on to the interrupt pin.

Toolbox button



Toolbox with user configurable scripts



Within uVISION there is a full scripting language which allows you to simulate external events. These scripts are based on the C language and are stored in text files. The script used to simulate the pulse is shown below:

```
signal void Toggle(void)
{
    PORT0 = (PORT0 ^ 0x4000);
    twatch (200);
    PORT0 = (PORT0 ^ 0x4000);
}

KILL BUTTON *
DEFINE BUTTON "GenerateEINT1", "Toggle()"
```

This script is stored in the file signal.ini and is added to the project in the debug window. For more details on the scripting language see the uVISION documentation.

Exercise 6: Software Interrupt

In this exercise we will define an inline Assembler function to call a software interrupt and place the value 0x02 in the calling instruction. In the software interrupt SWI we will decode the instruction to see which SWI function has been called and then use a case statement to run the appropriate code.

Open the project in EX6 SWI\work

In main.c add the following code

As the first instruction in main add the assembler define which calls the swi instruction

```
#define SoftwareInterrupt2 asm (" swi #02")
```

In the SWI ISR complete the register definition to access R14

```
register unsigned * link_ptr asm ("r14");
```

Complete the code to pass value of the SWI ordinal into the temp variable

```
temp = *(link_ptr-1) & 0x00FFFFFF;
```

Compile and download the code into the debugger

Step the code and observe the SWI being serviced

In the disassembly window the first SWI instruction has been encoded with the value 1 at location 0x00000015C

```
42: SoftwareInterrupt1;
0x00000015C EF000001 SWI          0x00000001
```

On entry to the ISR the supervisor link register contains the value 0x000000160

```
R14 [LR]      0x000000160
```

The calculation for temp is $\text{temp} = \ast(\text{link_ptr}-1) \& 0x00FFFFFF$ or $0x164 - 4$ (word-wide pointer remember) which is 0x15C which points to the instruction which generated the SWI. The top 8 bits are masked off which yields a value of 1. This is used in the case statement to run the required code.

Appendices

Appendix A

Bibliography

ARM7TDMI datasheet	ARM
LPC2119/2129/2194/2292/2294 User Manual	Philips
ARM System on chip architecture	Steve Furber
Architecture Reference Manual	David Seal
ARM System developers guide	Andrew N. Sloss, Domonic Symes, Chris Wright Jean J. Labrosse Arthur Griffith
MicroC/OS-II	
GCC The complete reference	

Webliography

Reference Sites

<http://www.arm.com>

<http://www.philips.com>

<http://www.lpc2000.com>

Tools and Software Development

<http://ww.hitex.co.uk>

<http://www.keil.co.uk>

<http://www.uicos-ii.com>

<http://www.ristancase.com>

<http://gcc.gnu.org/onlinedocs/gcc/>

Evaluation Boards And Modules

<http://www.phytec.co.uk>

<http://www.embeddedartists.com>

The LPC2000 family from Philips semiconductors is the first of a new generation of microcontrollers based on the ARM7-TDMI 16/32-bit RISC processor.

This book is written as both an introduction to the ARM7-TDMI processor and the LPC2000 microcontroller architecture. The content is based on a series of one day seminars held for professional engineers interested in learning how to use the LPC2000 family as quickly as possible.

Topics covered in this book include:

- Introduction to the ARM7 processor
- Software development tools
- LPC2000 system architecture
- LPC2000 peripherals

In addition a comprehensive tutorial is included that takes you through practical exercises to reinforce the topics discussed in the main text. By reading this book and performing the accompanying exercises, you will quickly become well versed in the ARM7 processor and the LPC2000 microcontroller.

The CD accompanying the book includes an evaluation version of the popular uVISION3 IDE and compiler from Keil Elektronik plus all the example exercises for both the Keil CARM and the GCC ARM compilers.

www.hitex.co.uk/arm

